



KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

Break Through Cluster Boundaries to Autoscale Workloads Across Them on a Large Scale

XingYan Jiang, DaoCloud & Ying Zhang, Zendesk

Agenda



KubeCon



CloudNativeCon

North America 2023

- HPA overview and benefits of autoscaling across clusters
- Introduction to Karmada
- Karmada feature - FederatedHPA
- Karmada new feature - CronFederatedHPA
- Q&A

HPA overview

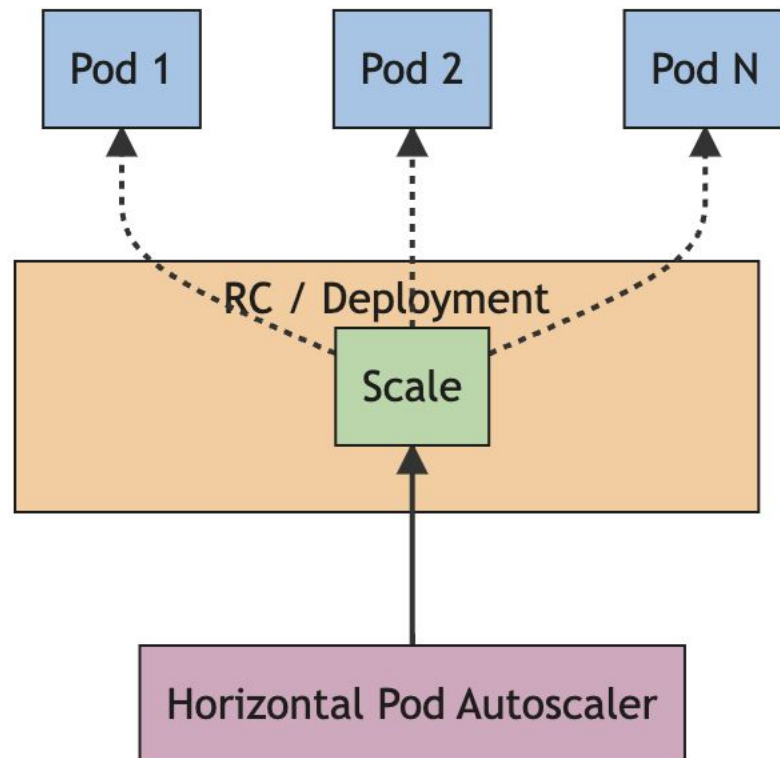


KubeCon



CloudNativeCon

North America 2023



Horizontal Pod Autoscaling (HPA) is a mechanism for automatically scaling the number of Pod replicas in a Kubernetes cluster to dynamically scale an application based on current load metrics.

HPA overview



KubeCon



CloudNativeCon

North America 2023

HPA example:

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```

The HPA controller will increase or decrease the number of Deployment replicas (by updating the Deployment) to maintain an average CPU utilization of 50% for all Pods.

Benefits of autoscaling across clusters



KubeCon



CloudNativeCon

North America 2023

- **Unified** management of autoscaling operation across clusters
- Break through the **resource limitations** of a single cluster
- With **rich strategies** in API definition, scale workloads across multiple clusters, and meet different scenarios
- Cluster-level autoscaling for disaster recovery.

Introduction to Karmada



KubeCon



CloudNativeCon

North America 2023

What is Karmada?

Karmada (Kubernetes Armada) is a Kubernetes management system that enables you to run your cloud-native applications across **multiple Kubernetes clusters** and **clouds**, with no changes to your applications. By speaking **Kubernetes-native APIs** and providing **advanced scheduling capabilities**, Karmada enables truly open, multi-cloud Kubernetes.

Karmada aims to provide turnkey automation for multi-cluster application management in multi-cloud and hybrid cloud scenarios, with key features such as **centralized multi-cloud management**, **high availability**, **failure recovery**, and **traffic scheduling**.

Introduction to Karmada



KubeCon



CloudNativeCon

North America 2023



Build an infinitely scalable container resource pool using Karmada.

Enable developers to use multi-cloud as easily as using a K8s cluster.

**Kubernetes Native API
Compatible**

Open and Neutrality

Avoid Vendor Lock-in

Out of the Box

**Fruitful Scheduling
Policies**

Centralized Management

Karmada Deep Dive

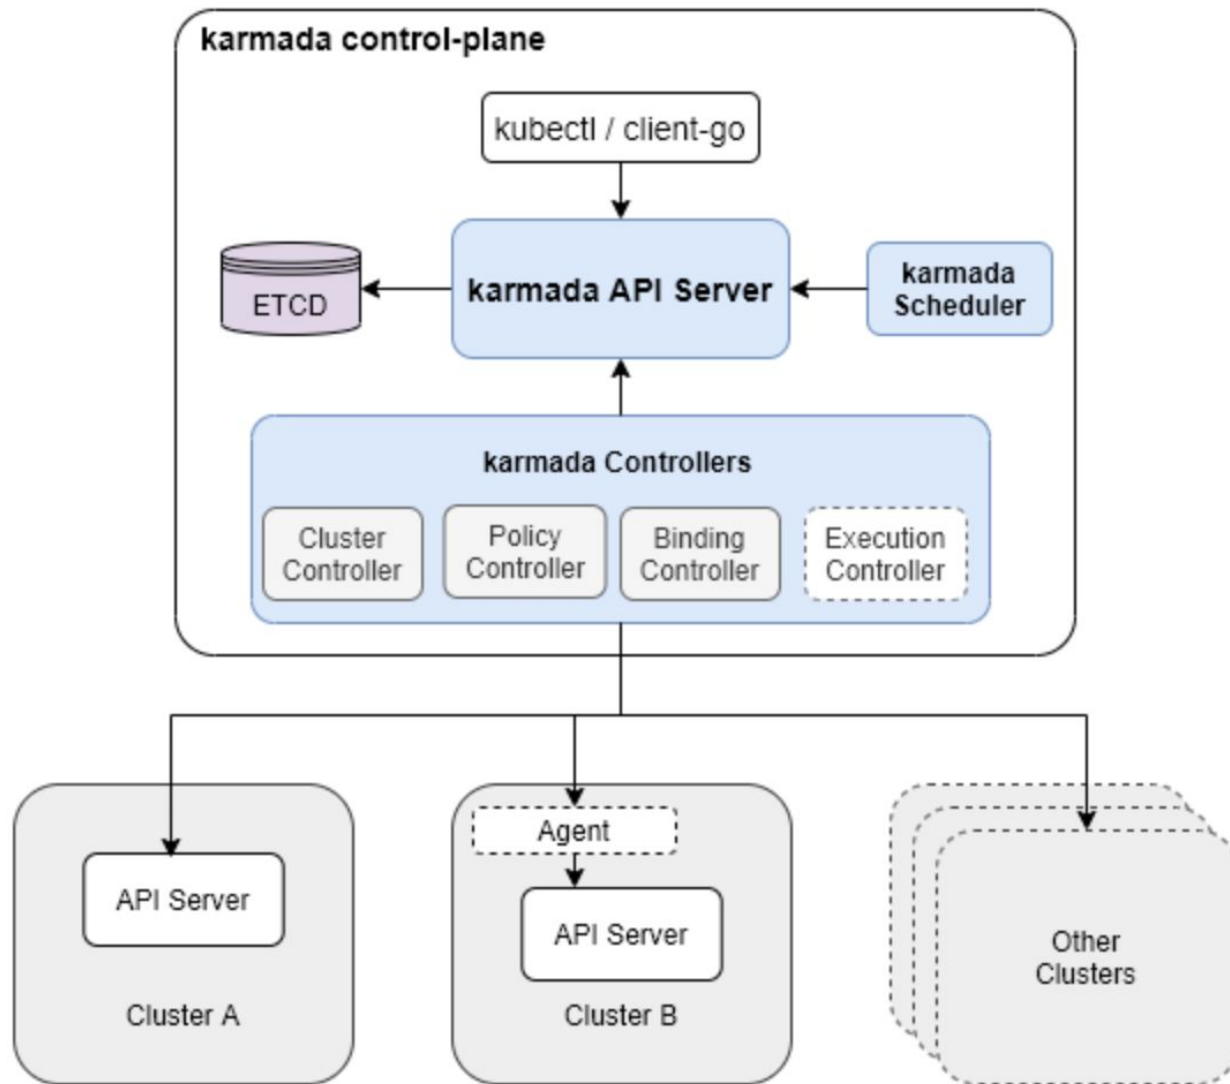


KubeCon

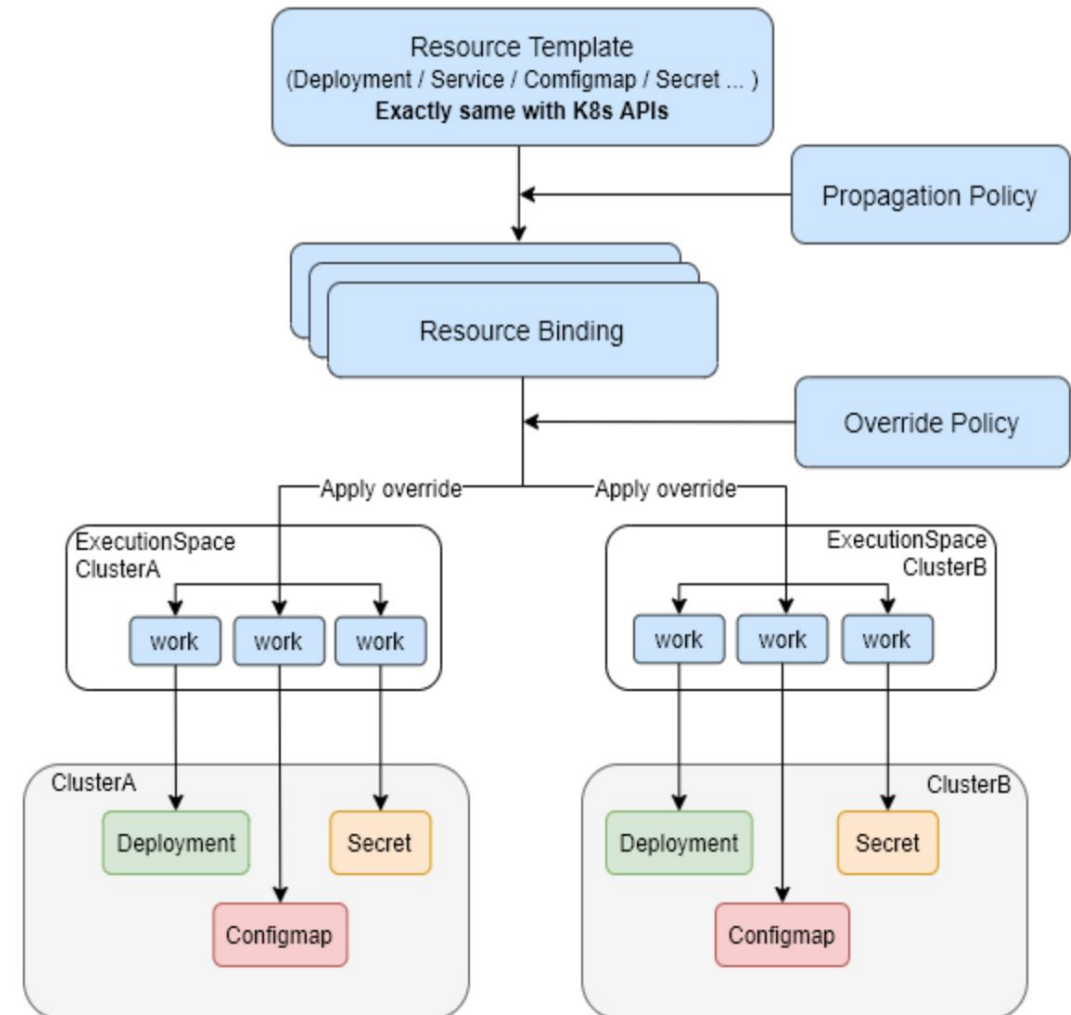


CloudNativeCon

North America 2023



Karmada Concepts



Multi-Cluster Management

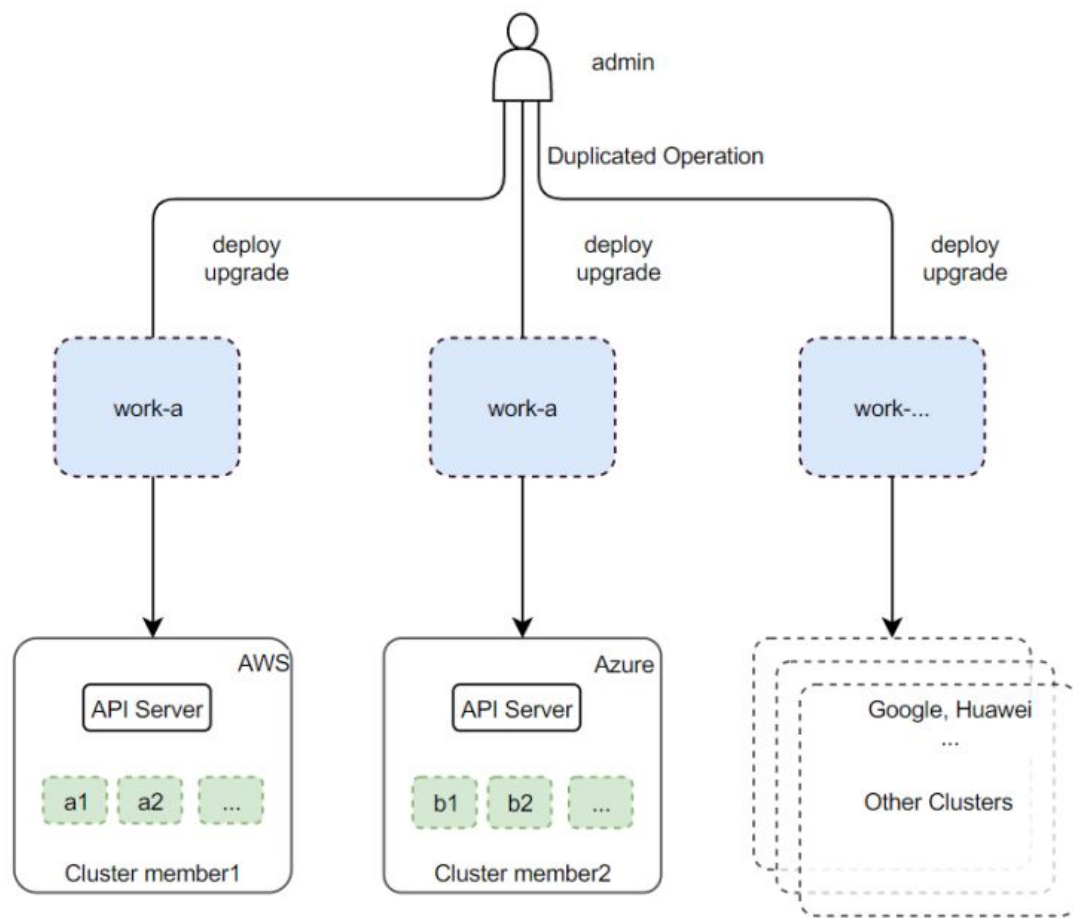


KubeCon

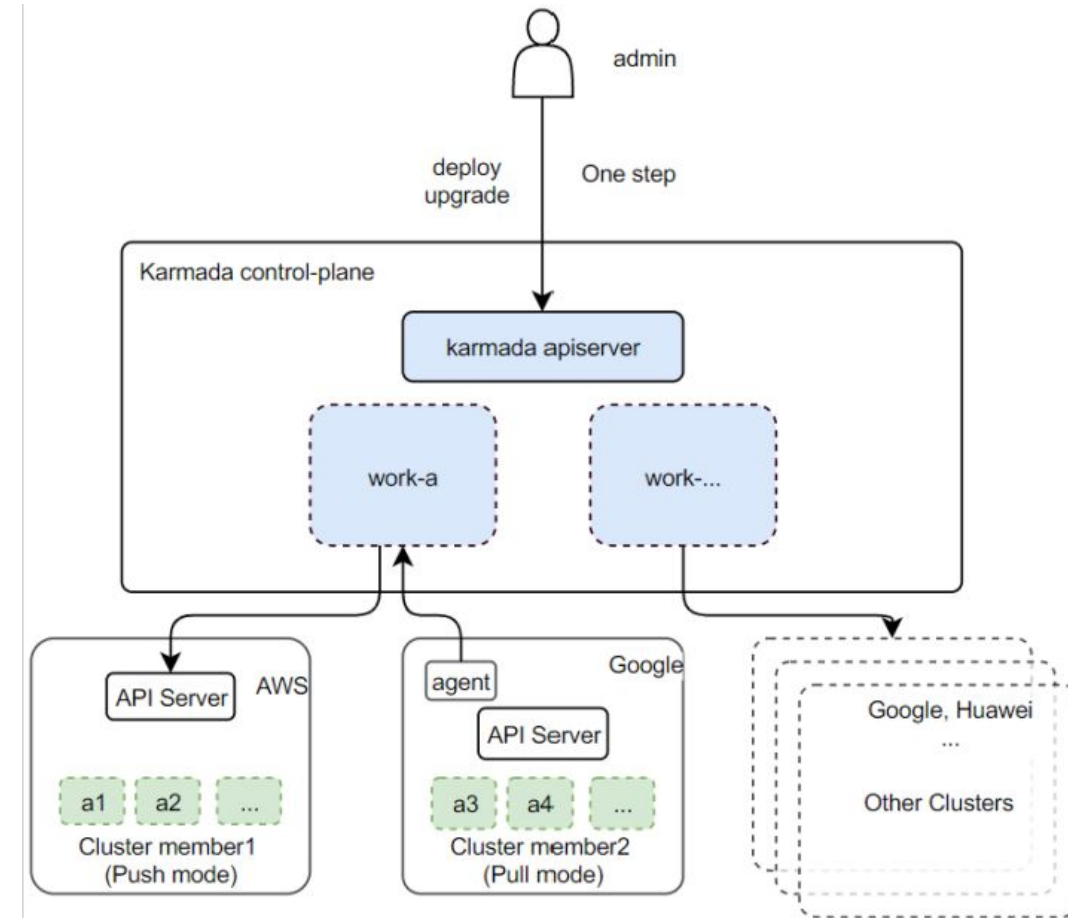


CloudNativeCon

North America 2023



Repeated management operations



Manage multi-cluster workloads in one step

Workload Propagation Across Clusters

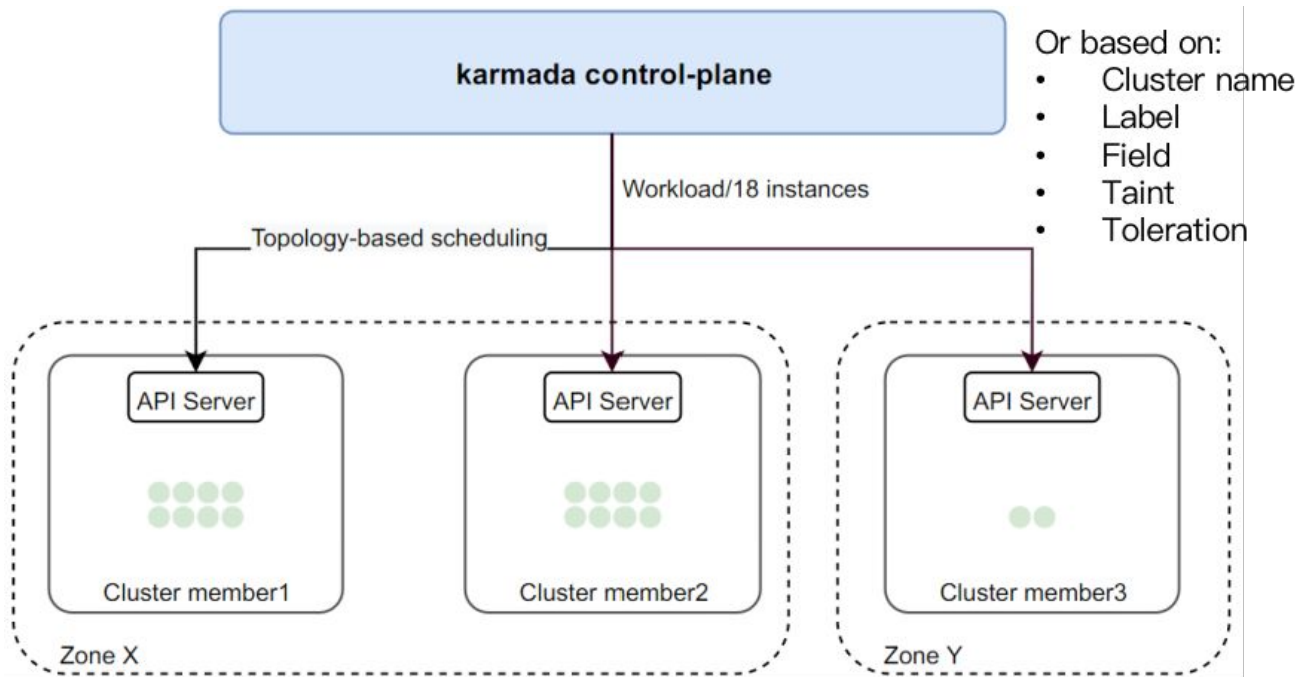


KubeCon

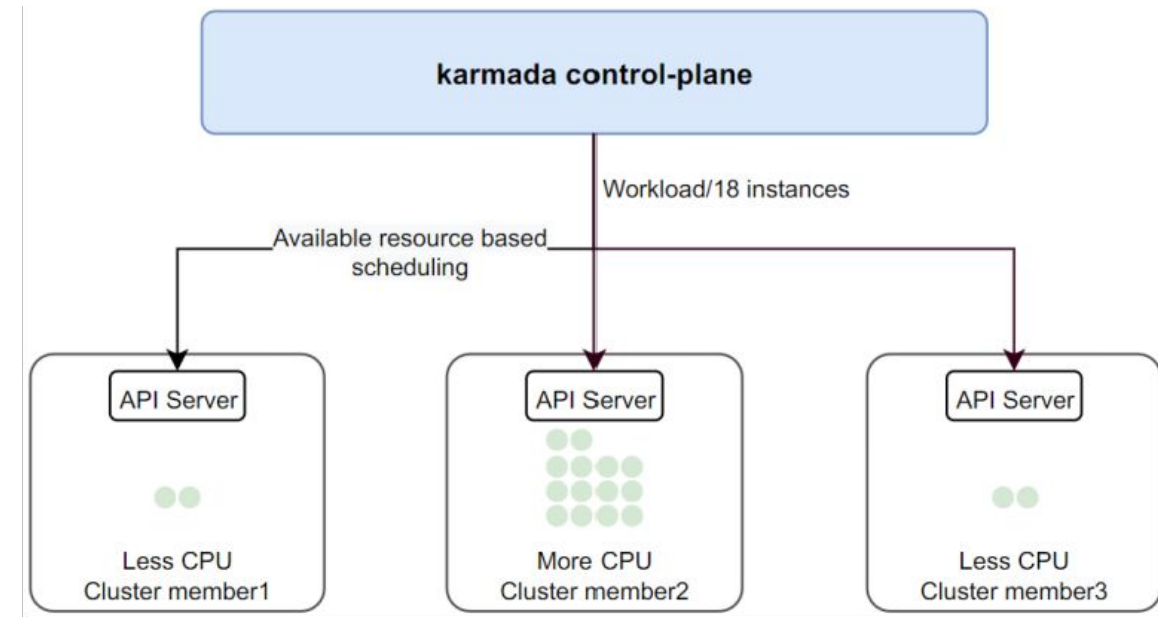


CloudNativeCon

North America 2023



Schedule based on topology



Schedule based on available resources

Advanced policies meet various scheduling requirements

PropagationPolicy



KubeCon



CloudNativeCon

North America 2023

```
# propagationpolicy.yaml
apiVersion: policy.karmada.io/v1alpha1
kind: PropagationPolicy
metadata:
  name: example-policy # The default namespace is `default`.
spec:
  resourceSelectors:
    - apiVersion: apps/v1
      kind: Deployment
      name: nginx
  placement:
    clusterAffinity:
      clusterNames:
        - member1
        - member2
  replicaScheduling:
    replicaDivisionPreference: Weighted
    replicaSchedulingType: Divided
```

Two replicaSchedulingTypes

- Duplicated
- Divided

Refer to: <https://karmada.io/docs/userguide/scheduling/resource-propagating/#multiple-strategies-of-replica-scheduling>

Cross-Cluster Application Failover

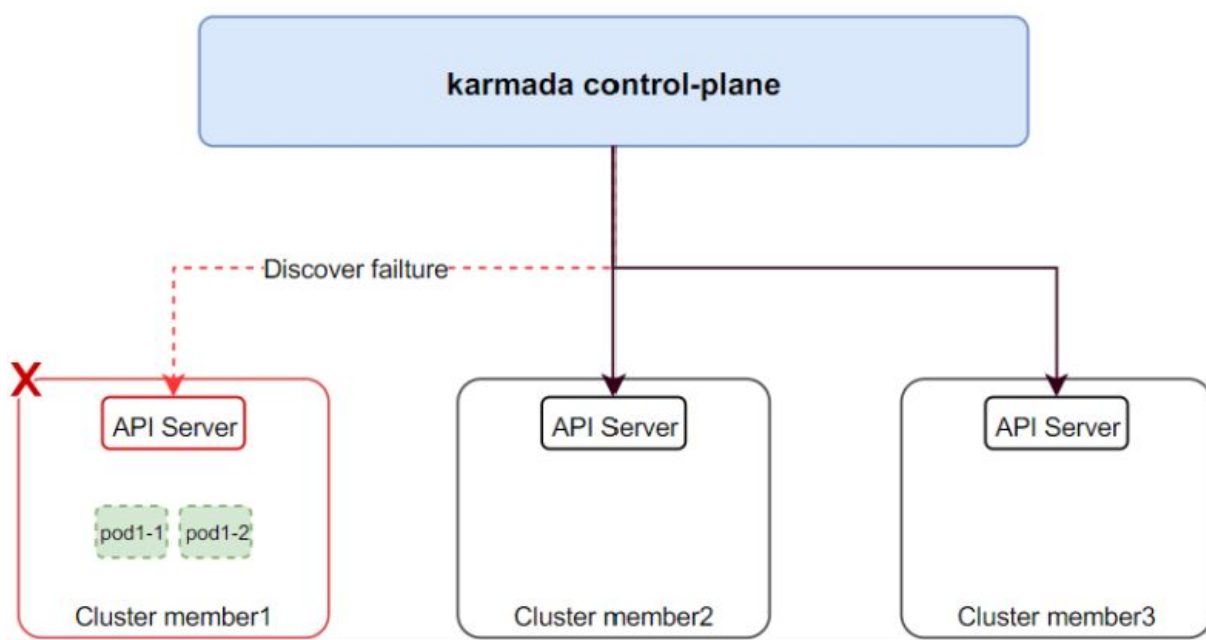


KubeCon

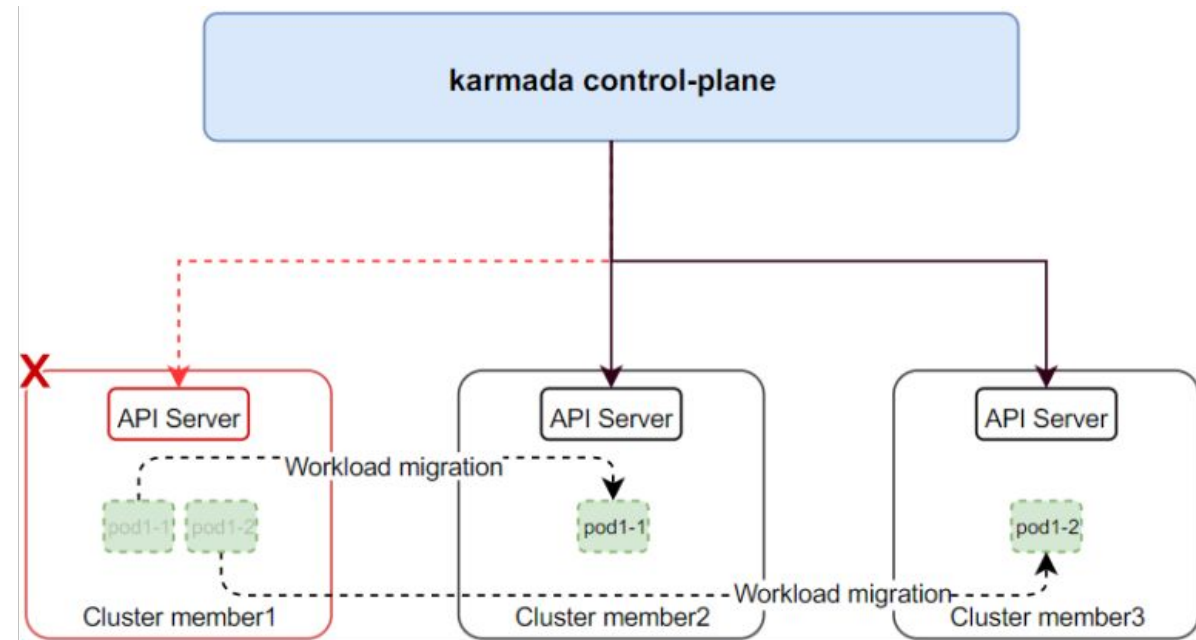


CloudNativeCon

North America 2023



Faulty cluster loses connection



Gracefully migrate workloads

Graceful migration ensures uninterrupted services

Agenda



KubeCon



CloudNativeCon

North America 2023

- HPA overview and benefits of autoscaling across clusters
- Introduction to Karmada
- **Karmada feature - FederatedHPA**
- Karmada new feature - CronFederatedHPA
- Q&A

FederatedHPA



KubeCon



CloudNativeCon

North America 2023

```
type FederatedHPA struct {
    metav1.TypeMeta    `json:",inline"`
    metav1.ObjectMeta  `json:"metadata,omitempty"`

    Spec FederatedHPASpec `json:"spec"`

    Status autoscalingv2.HorizontalPodAutoscalerStatus `json:"status"`
}

type FederatedHPASpec struct {
    ScaleTargetRef autoscalingv2.CrossVersionObjectReference `json:"scaleTargetRef"`

    MinReplicas *int32 `json:"minReplicas,omitempty"`

    MaxReplicas int32 `json:"maxReplicas"`

    Metrics []autoscalingv2.MetricSpec `json:"metrics,omitempty"`

    Behavior *autoscalingv2.HorizontalPodAutoscalerBehavior `json:"behavior,omitempty"`
}
```

Keep **consistent** with the core API of K8s native HPA. It can be **migrated** to FederatedHPA at a very low cost.

FederatedHPA



KubeCon



CloudNativeCon

North America 2023

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```

Minimal cost
migration

```
apiVersion: autoscaling.karmada.io/v1alpha1
kind: FederatedHPA
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```


FederatedHPA workflow

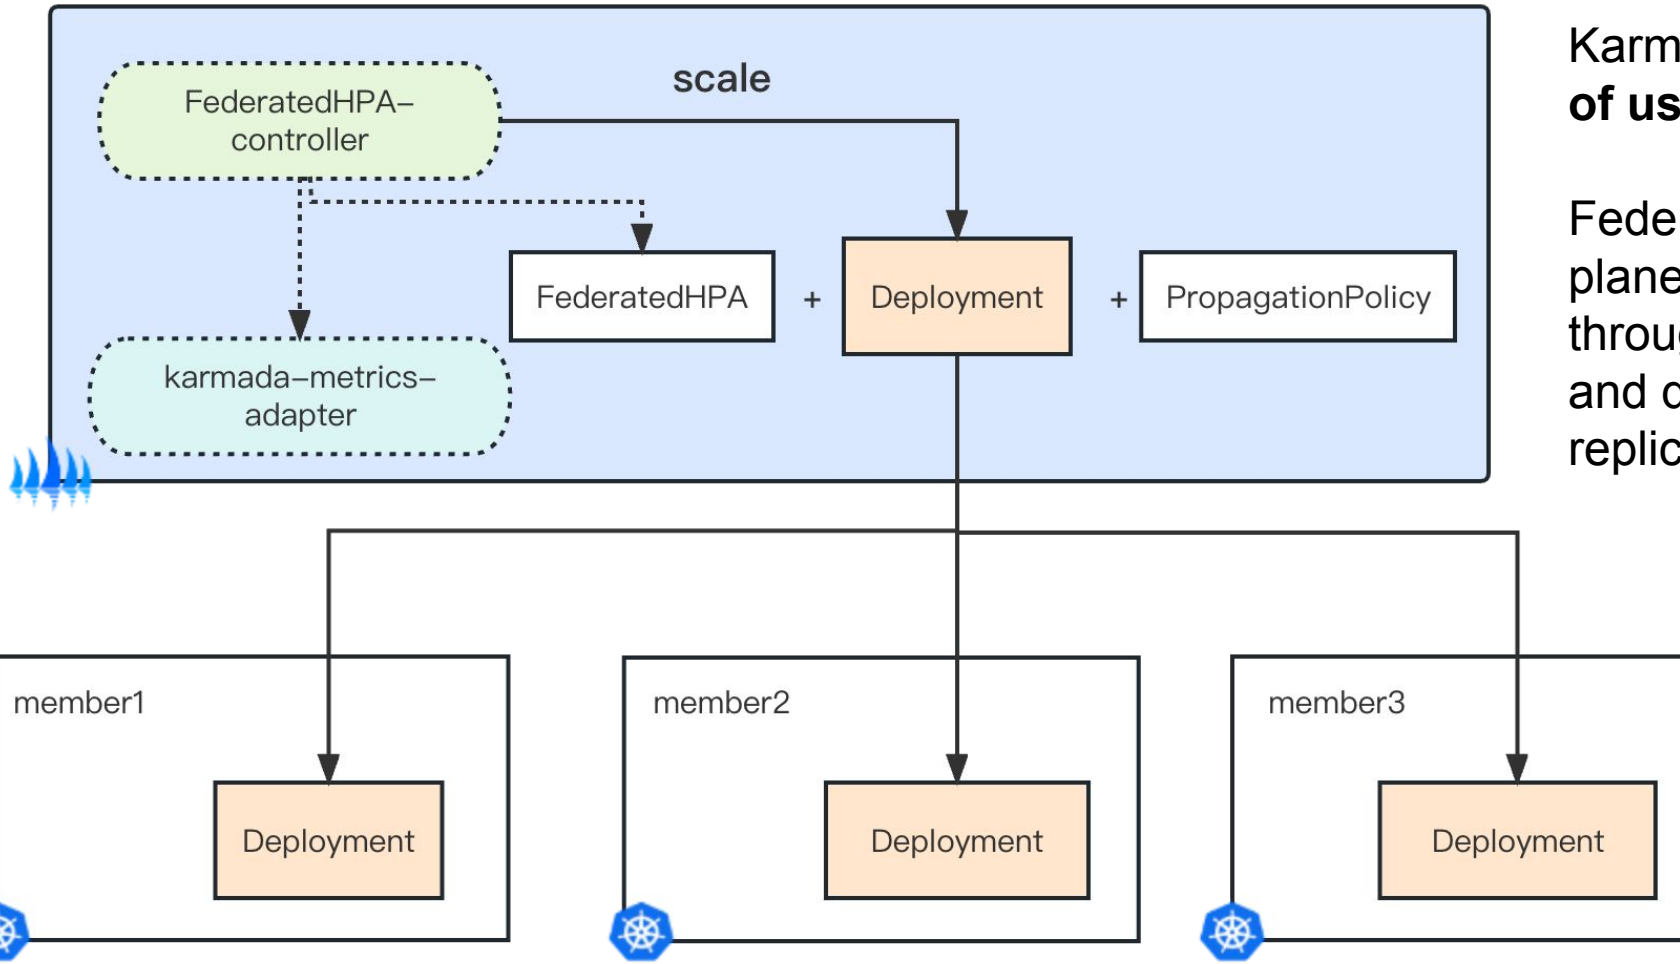


KubeCon



CloudNativeCon

North America 2023



Karmada FederatedHPA - **Experience of using single-cluster HPA.**

FederatedHPA controller on the control plane obtains metrics of Deployment through karmada-metrics-adapter, and dynamically scales the number of replicas for the Deployment.

FederatedHPA workflow

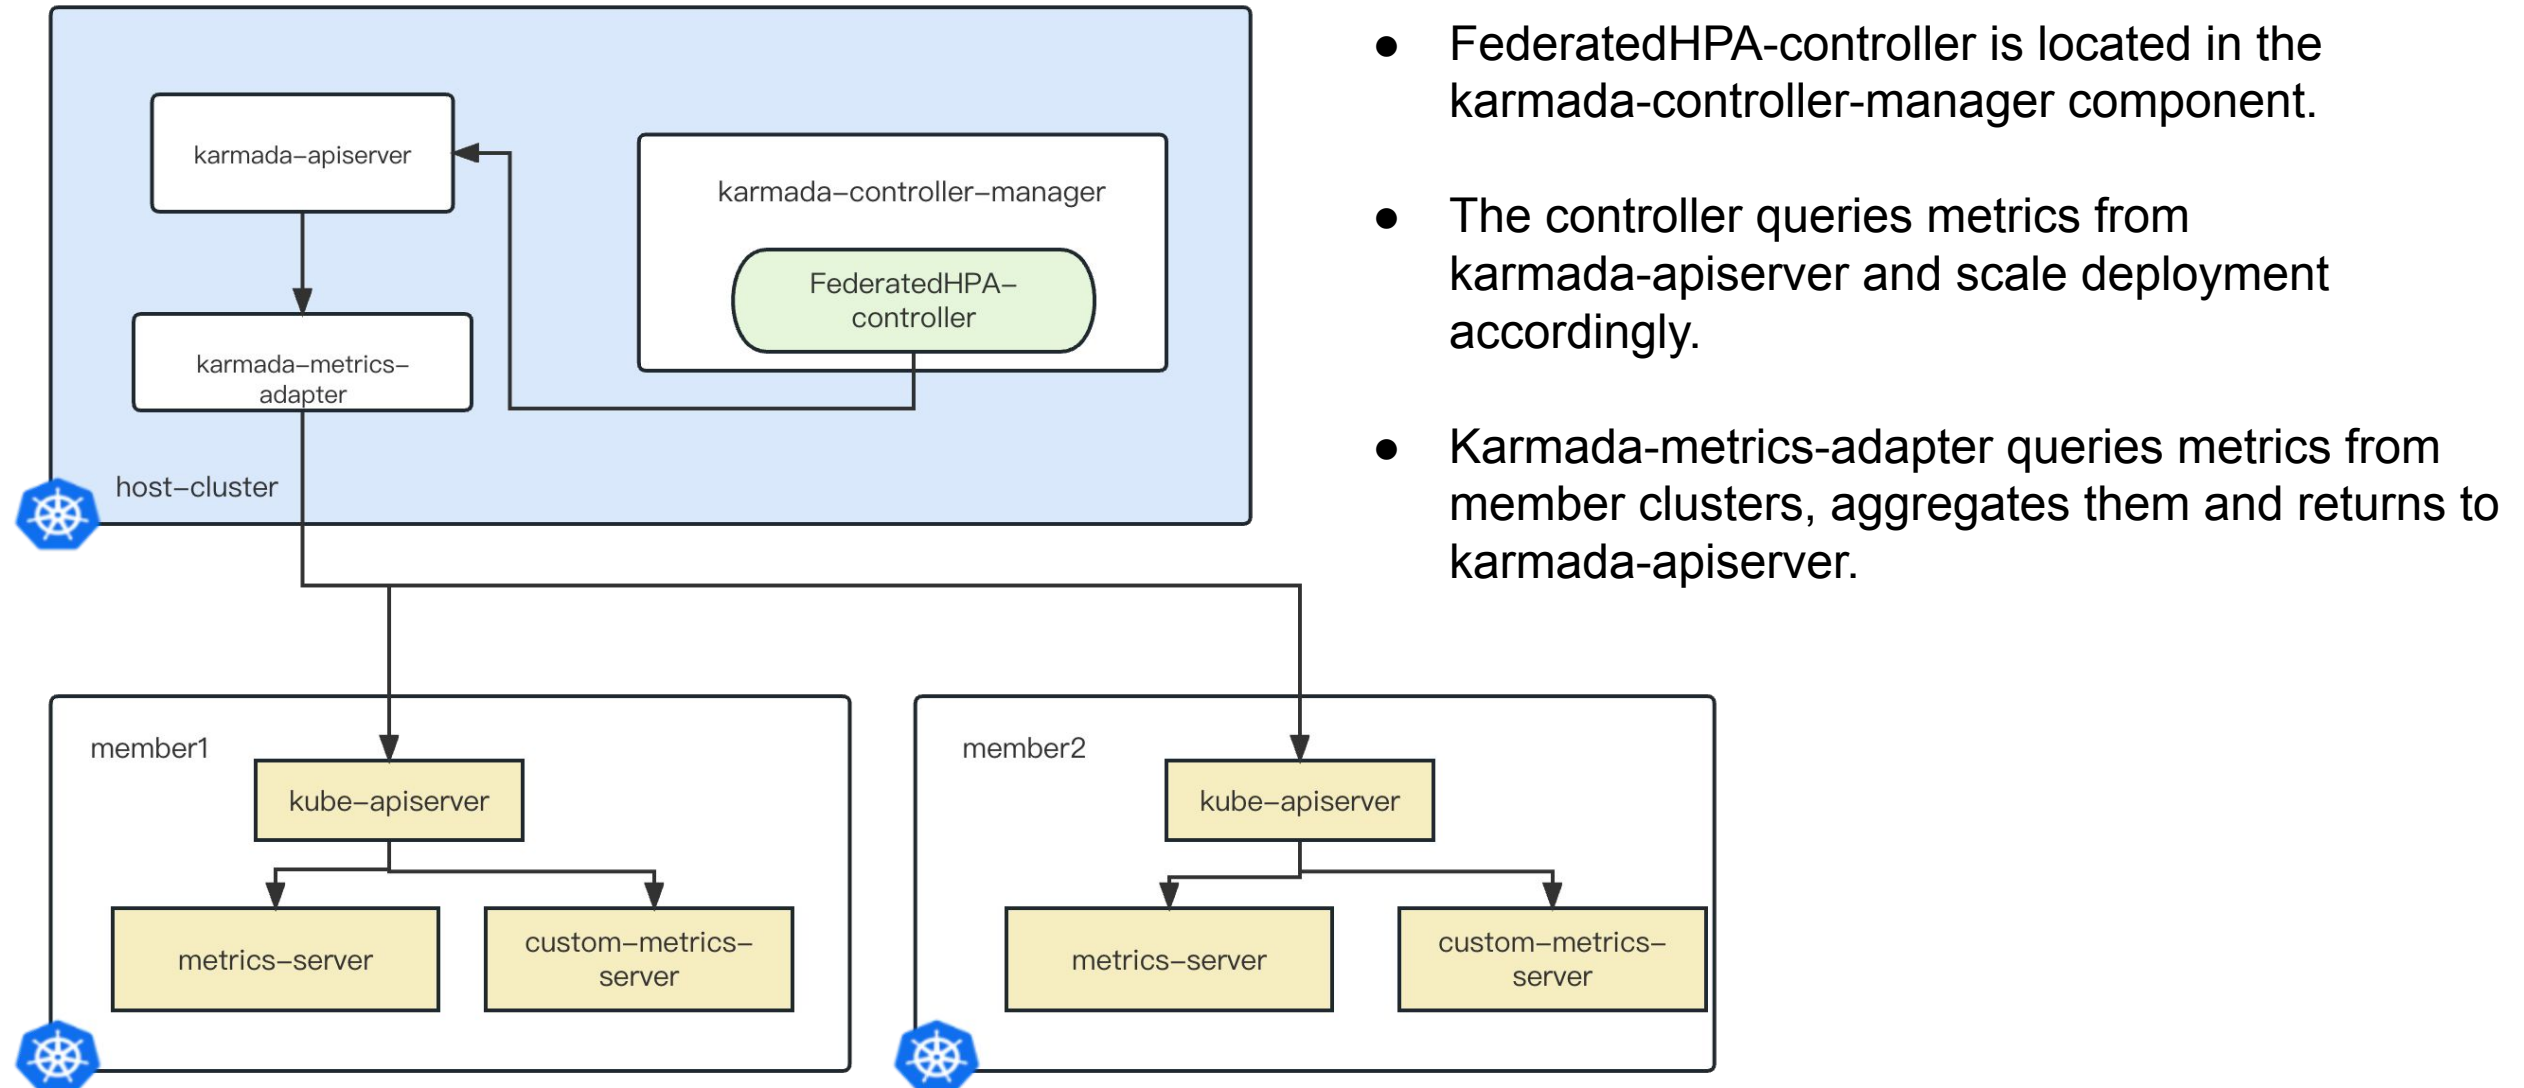


KubeCon



CloudNativeCon

North America 2023



Extension - CronFederatedHPA

CronFederatedHPA is used for regular autoscaling actions. It can scale workloads that have a scale subresource or FederatedHPA.

Scenarios

- Proactively scale up a workload for predictable spikes in traffic
- Schedule the non time-sensitive workloads on weekends
- And more

CronFederatedHPA



KubeCon



CloudNativeCon

North America 2023

```
apiVersion: autoscaling.karmada.io/v1alpha1
kind: CronFederatedHPA
metadata:
  name: cron-federated-hpa
spec:
  scaleTargetRef:
    apiVersion: autoscaling.karmada.io/v1alpha1
    kind: FederatedHPA
    name: shop-fhpa
  rules:
    - name: "Scale-Up"
      schedule: "30 08 * * *"
      targetMinReplicas: 1000
    - name: "Scale-Down"
      schedule: "0 11 * * *"
      targetMinReplicas: 1
```

Scale FederatedHPA

```
apiVersion: autoscaling.karmada.io/v1alpha1
kind: CronFederatedHPA
metadata:
  name: nginx-cronfhpa
  namespace: default
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx
  rules:
    - name: "Scale-Up"
      schedule: "30 08 * * *"
      targetReplicas: 1000
    - name: "Scale-Down"
      schedule: "0 11 * * *"
      targetReplicas: 1
```

Scale Deployment,etc

CronFederatedHPA workflow

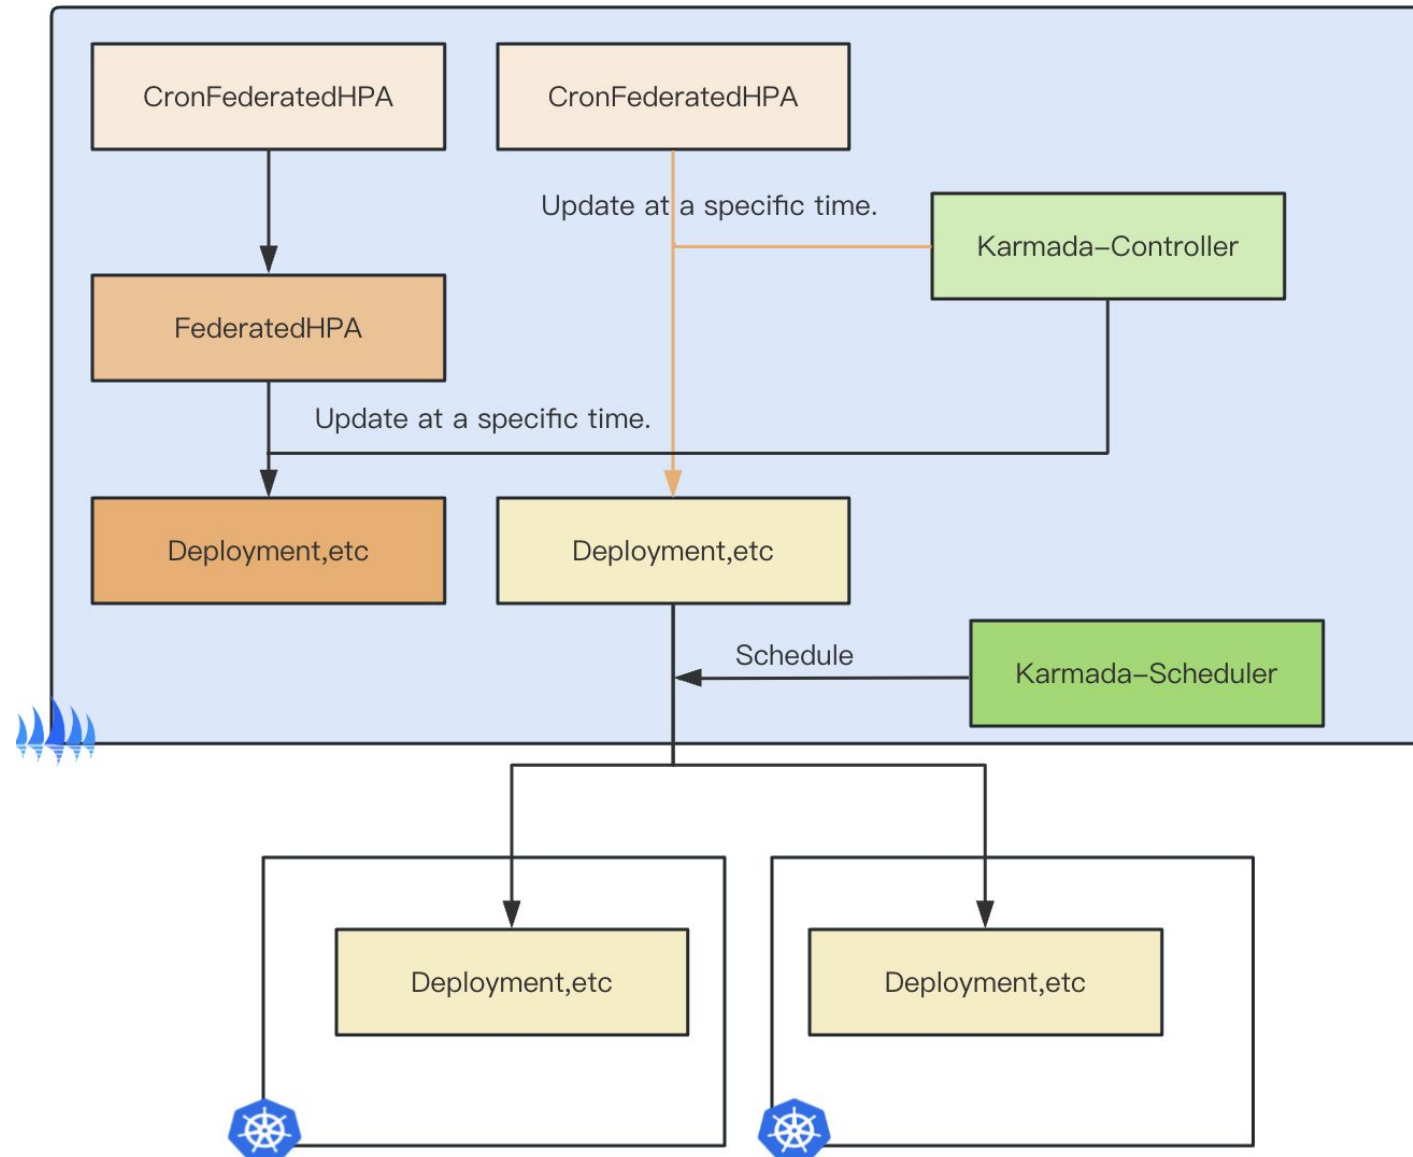


KubeCon



CloudNativeCon

North America 2023



CronFederatedHPA with FederatedHPA



Scaling workloads with scale subresource

CronFederatedHPA is capable of scaling workloads that have a **scale subresource**, such as Deployment and StatefulSet.

Scaling FederatedHPA

CronFederatedHPA scales resources at a specified time. When a workload is scaled only by CronFederatedHPA, it can't handle more requests until that time. To **proactively** meet peak loads and real-time demands, we recommend starting with **CronFederatedHPA to scale FederatedHPA**, and then, **FederatedHPA can scale workloads** based on their metrics.

Advantages and Notice



KubeCon



CloudNativeCon

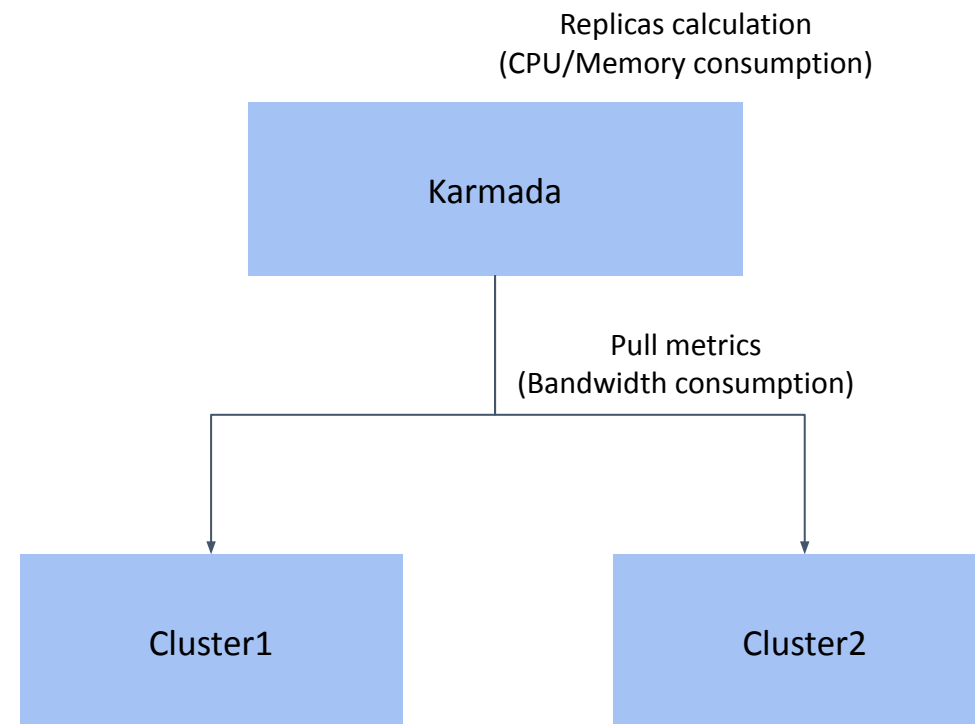
North America 2023

Advantages of FederatedHPA:

- The API and K8s native HPA are almost identical, with the same user experience and low migration cost.

Notice for using FederatedHPA:

- FederatedHPA is a type of centralized multi-cluster HPA.
- When scaling concurrently on a large scale, a larger bandwidth is required (maximum concurrent scaling metrics pull: 500M bandwidth, 200,000 Pods).
- Storing and computing corresponding data requires much CPU/Memory.



Explore distributed autoscaling



KubeCon

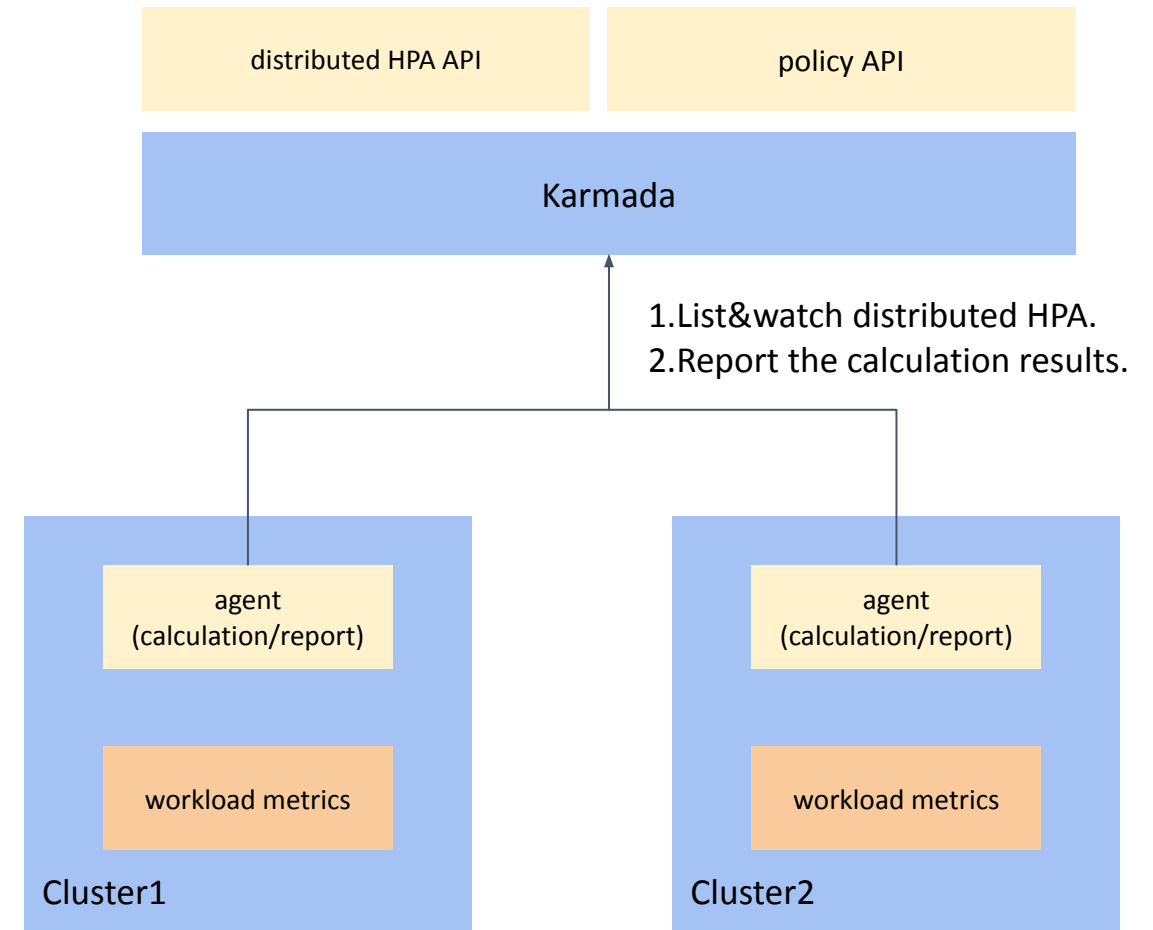


CloudNativeCon

North America 2023

Design :

- Define the metrics for scaling in DistributedHPA API.
- Member cluster **agents will list&watch the distributed HPA** and calculate the intermediate calculation state.
- Update the calculated intermediate state to the "status" of distributed HPA.
- Karmada processes data based on reports from member clusters, **calculates final results, and adjusts replicas.**
- Based on policy configuration and **real-time states** of clusters, Karmada schedules scaled instances to member clusters to achieve cross-cluster autoscaling.



Key Takeaways



KubeCon



CloudNativeCon

North America 2023

- Regular HPA cannot break through cluster boundaries to scale workloads
- FederatedHPA and CronFederatedHPA can autoscale workloads across multiple clusters
- Karmada enables you to run your cloud-native applications across multiple Kubernetes clusters and clouds

Join us



KubeCon



CloudNativeCon

North America 2023



<https://karmada.io>



<https://github.com/karmada-io/karmada>



<https://slack.cncf.io> (#karmada)



Thank you!