



KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

Building Better Controllers

John Howard, Google

Overview



KubeCon



CloudNativeCon

North America 2023

- What *are* controllers
- Writing controllers is really hard
- Challenges writing controllers for Istio
- A ~~better~~ different approach to writing controllers

What are controllers?




KubeCon



CloudNativeCon

North America 2023

"Controllers are control loops that watch the state of your cluster, then make or request changes where needed. Each controller tries to move the current cluster state closer to the desired state."



Writing controllers is *hard*



KubeCon



CloudNativeCon

North America 2023

Most controllers are *very* low level.

Main primitive is the Informer, allowing us to subscribe to events:

```
type EventHandler struct {  
    AddFunc    func(obj any)  
    UpdateFunc func(oldObj, newObj any)  
    DeleteFunc  func(obj any)  
}
```

Writing controllers is *hard*



KubeCon



CloudNativeCon

North America 2023

Real world code is [400 LOC](#) for a sample

```
func controller() {
    queue := NewQueue(handleDeployment)
    deployments := NewDeploymentInformer()
    deployments.AddEventHandler(EventHandler{
        AddFunc: func(a any) { queue.Enqueue(a)
    })
}

func handleDeployment(object types.NamespacedName) {
    deploy := deployments.Get(object)
    if deploy == nil {
        // ...handle deletion...
    }
    // ...handle deployment...
}
```

Writing controllers is *hard*



KubeCon



CloudNativeCon

North America 2023

istio / istio Public

Use queue

Merged

istio-te

istio / istio Public

crd watcher: fix edge case causing missed events #471

Merged

kubernetes / kubernetes Public

Admission configuration managers incorrectly

kubernetes / kubernetes Public

761

[Scheduler] Make sure handlers have synced before scheduling #116729

ments

Merged

k8s-ci-robot merged 1 commit into `kubernetes:master` from `AxeZhan:handlers_sync` on Jun 28

Example: Istio Service Mesh

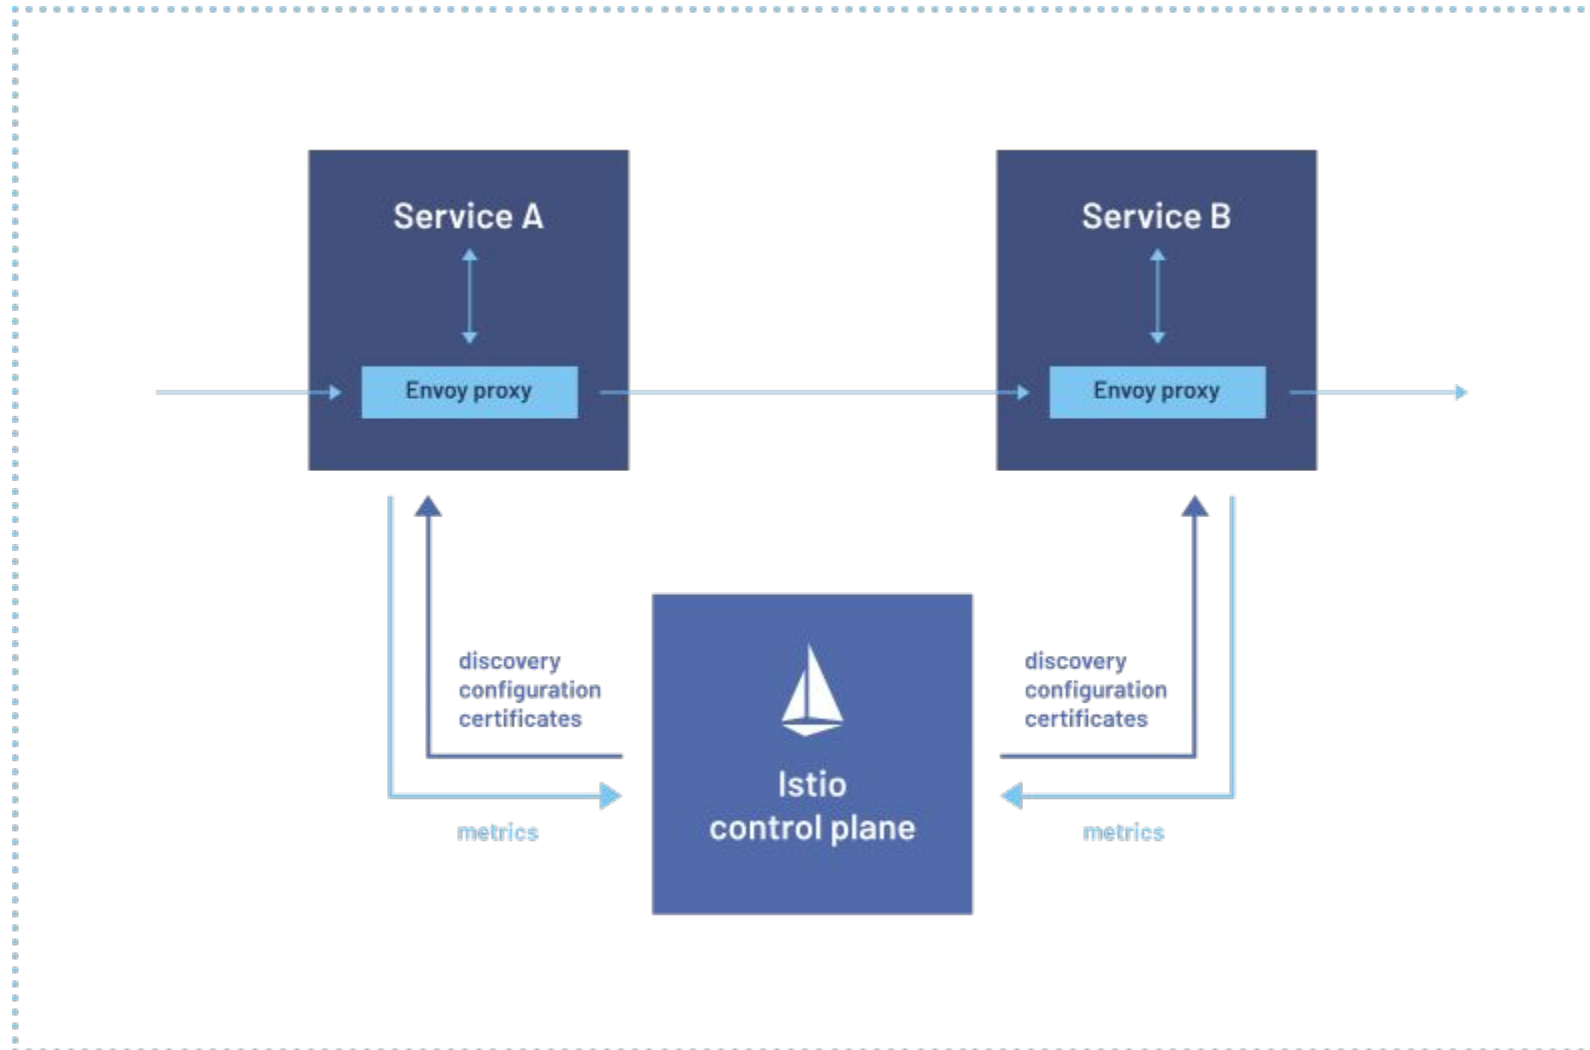


KubeCon



CloudNativeCon

North America 2023



Istio Challenges




KubeCon



CloudNativeCon

North America 2023

Istio is primarily one giant controller, but unique.



- Lots of inputs/outputs
- Intermediate state not persisted in cluster
 - Or memory!
- Outputs large (megabytes), and written to many places

Istio Implementation

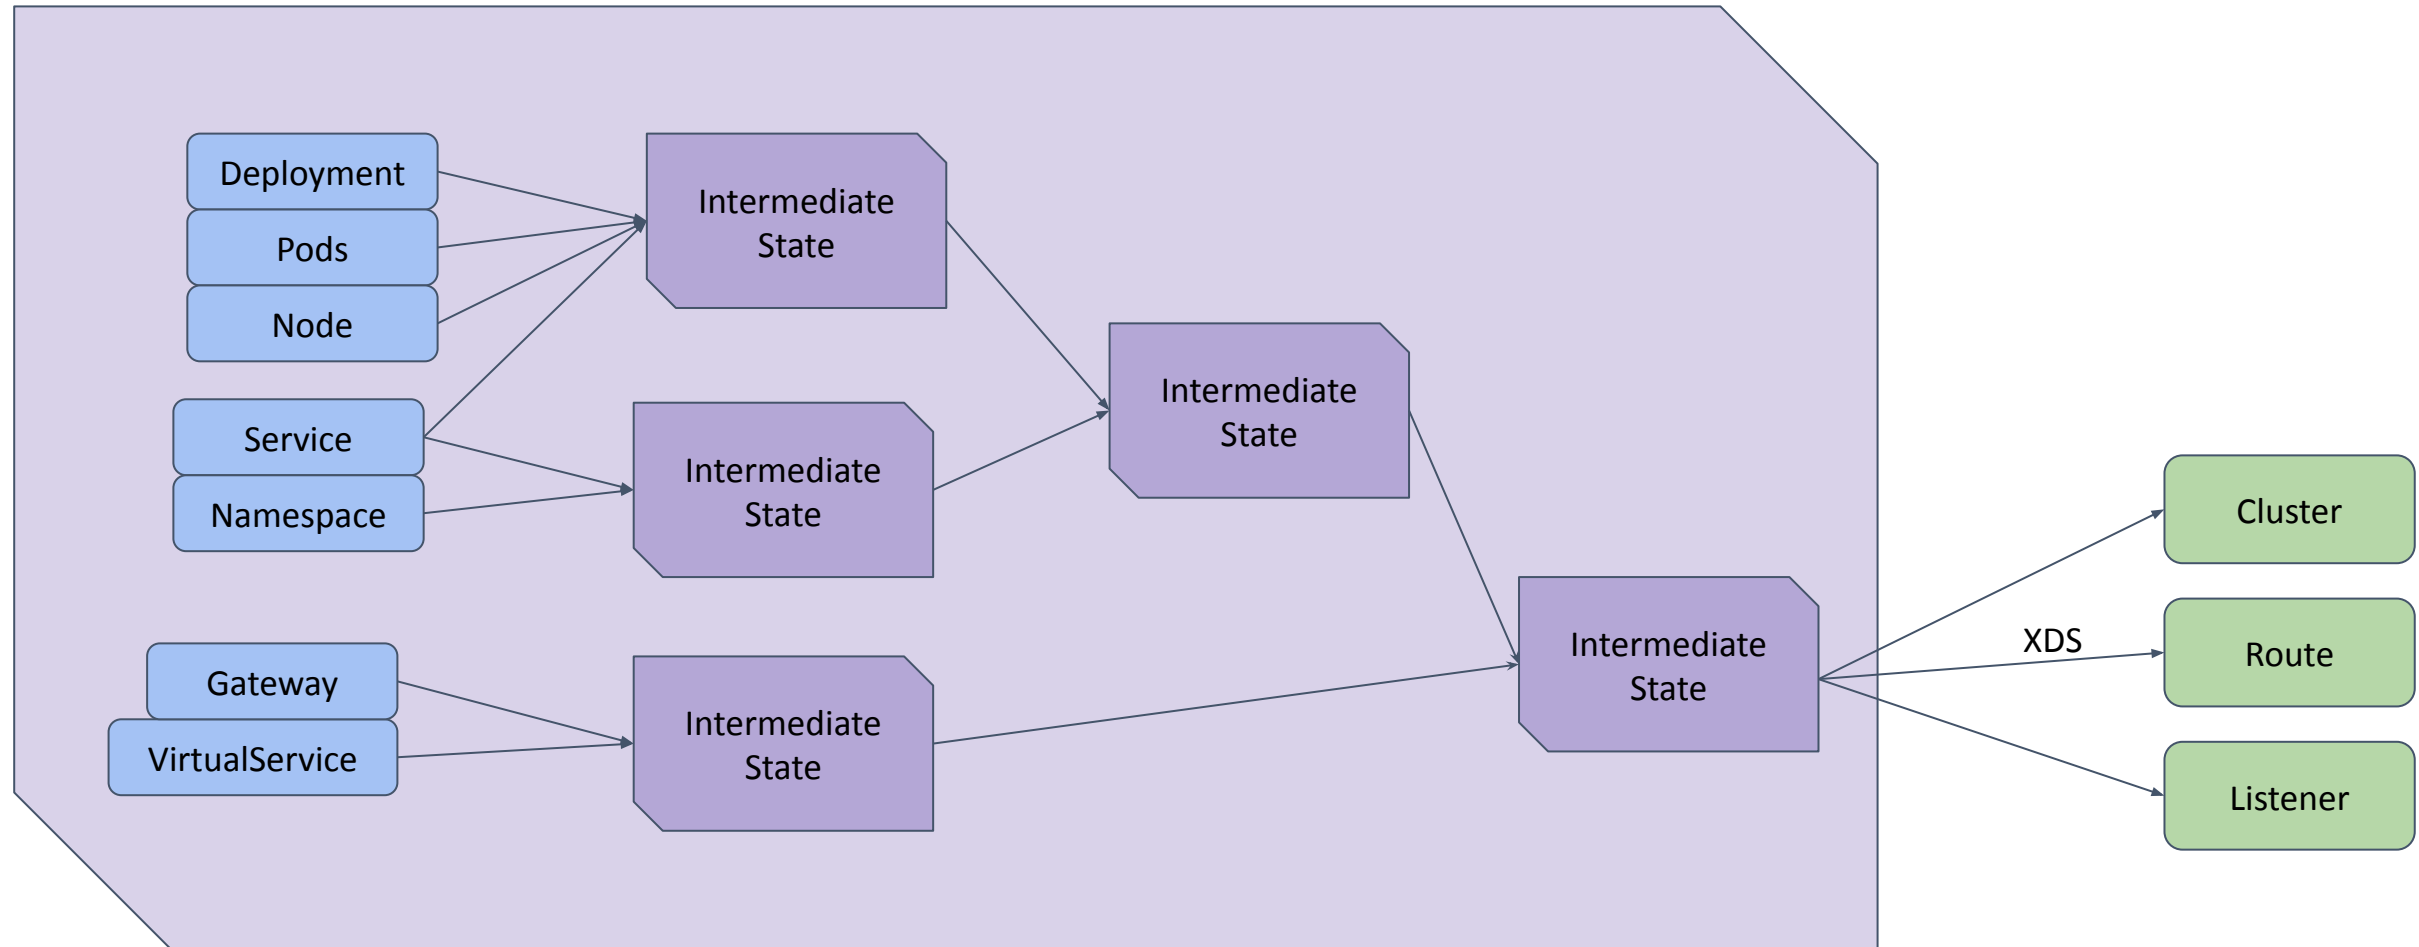


KubeCon



CloudNativeCon

North America 2023



State management



KubeCon



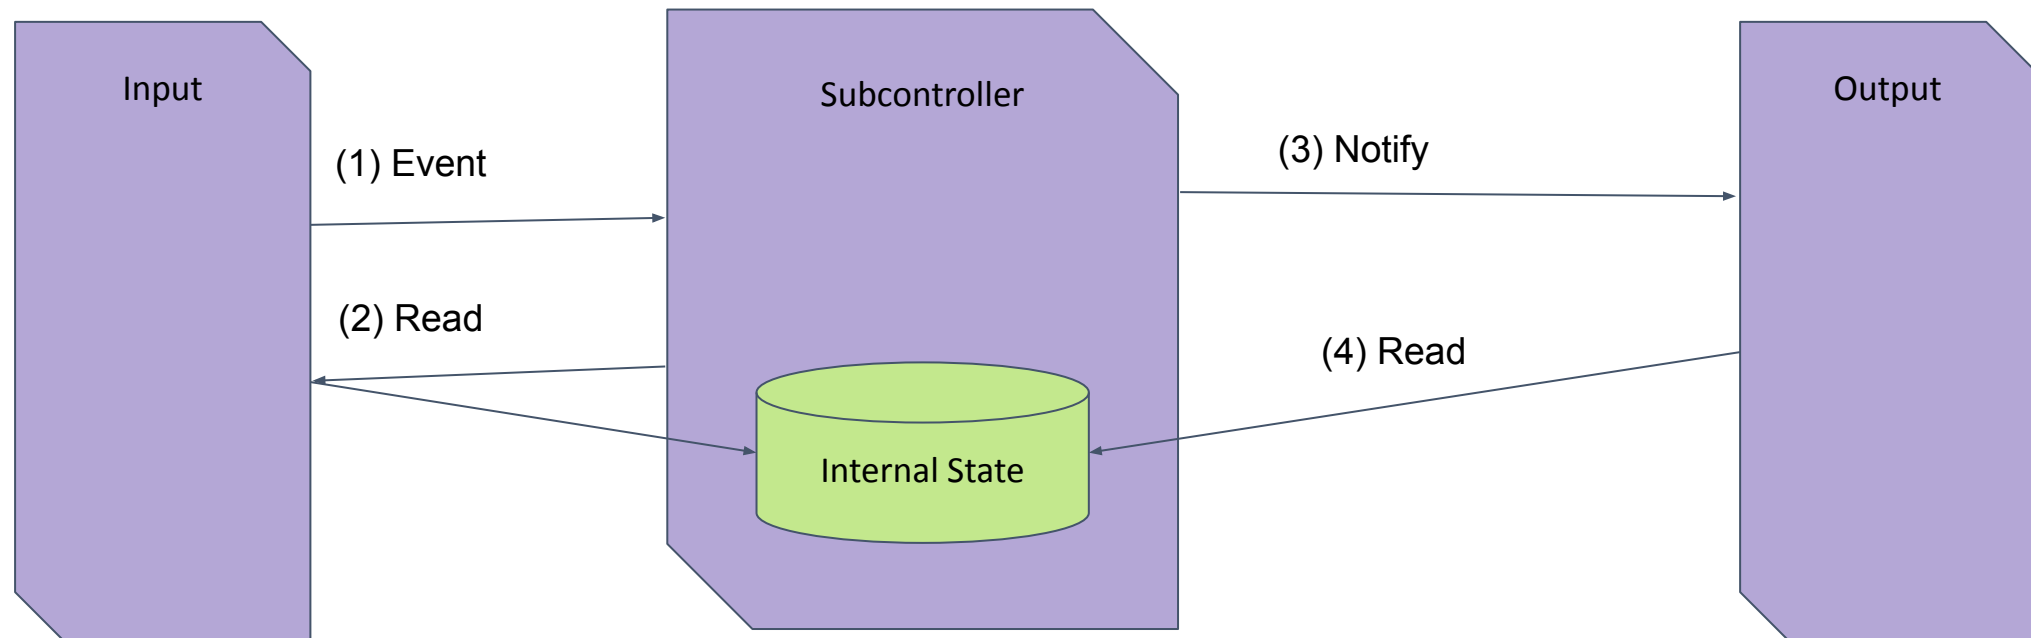
CloudNativeCon

North America 2023

Majority of Istio codebase consists of:

- Receive event that something has changed.
- Reconcile internal state with the change.
- Notify dependencies something changed
- Dependency queries our new internal state

Most controllers utilize Kubernetes as intermediate state to get this for "free"



Event detection




KubeCon



CloudNativeCon

North America 2023



Istio Implementation



KubeCon




CloudNativeCon

North America 2023

```
type AmbientIndex struct {
    func (a *AmbientIndex) handleService(...){
        pods := c.getPodsInService(svc)
        for _, p := range pods {
            wl := a.generateWorkload(p, c)
            for _, networkAddr := range networkAddressFromWorkload(wl) {
                a.byPod[networkAddr] = wl
            }
            a.byUID[wl.Uid] = wl
            wls[wl.Uid] = wl
        }
    }
}

a.serviceByAddr[networkAddr] = si
}
a.byService[namespacedName] = wls
a.serviceByNamespacedHostname[namespacedName] = si
}
```





- Easy to write **correctly** and **efficiently**
 - The obvious implementation should be the correct and efficient one.
- **High level**
 - The controller should describe business logic, and shouldn't be concerned with low level state management.
- **Composable**
 - Controllers should be able to build upon each other, even if they are not persisted to Kubernetes.

Goals



KubeCon



CloudNativeCon

North America 2023



Extremely simple core interface

```
type Collection[T any] interface {  
    Get(k Key[T]) *T  
    List(labels Labels) []T  
    RegisterWatcher(handler func(o Event[T]))  
}
```

Ecosystem of composable components built around this interface

Sources

- Collection from informer
- Collection from files
- Collection from in-memory objects
- Collection fetch from external state

Transformations

- Index
- Transformations
- Complex compositions

Outputs

- Write to Kubernetes
- Send over XDS
- Write to Cloud APIs
- Arbitrary event handler

Collection Creation



KubeCon



CloudNativeCon

North America 2023

```
ConfigMaps := FromInformer[ConfigMap]()  
Config := NewCollection(func() Config {  
    cfg := DefaultConfig()  
    cms := []ConfigMap{  
        // Fetch a single item from the ConfigMaps collection.  
        FetchOne(ConfigMaps, filter.Name("config")),  
        FetchOne(ConfigMaps, filter.Name("config-overrides"))  
    }  
    for _, c := range cms {  
        cfg = cfg.Merge(c.data.Config)  
    }  
    return cfg  
})
```

From Kubernetes Informer

Derived from various inputs

Config is automatically updated when the input ConfigMaps change

Config watchers notified any time Config changes

Collection Creation




KubeCon



CloudNativeCon

North America 2023



Event detection

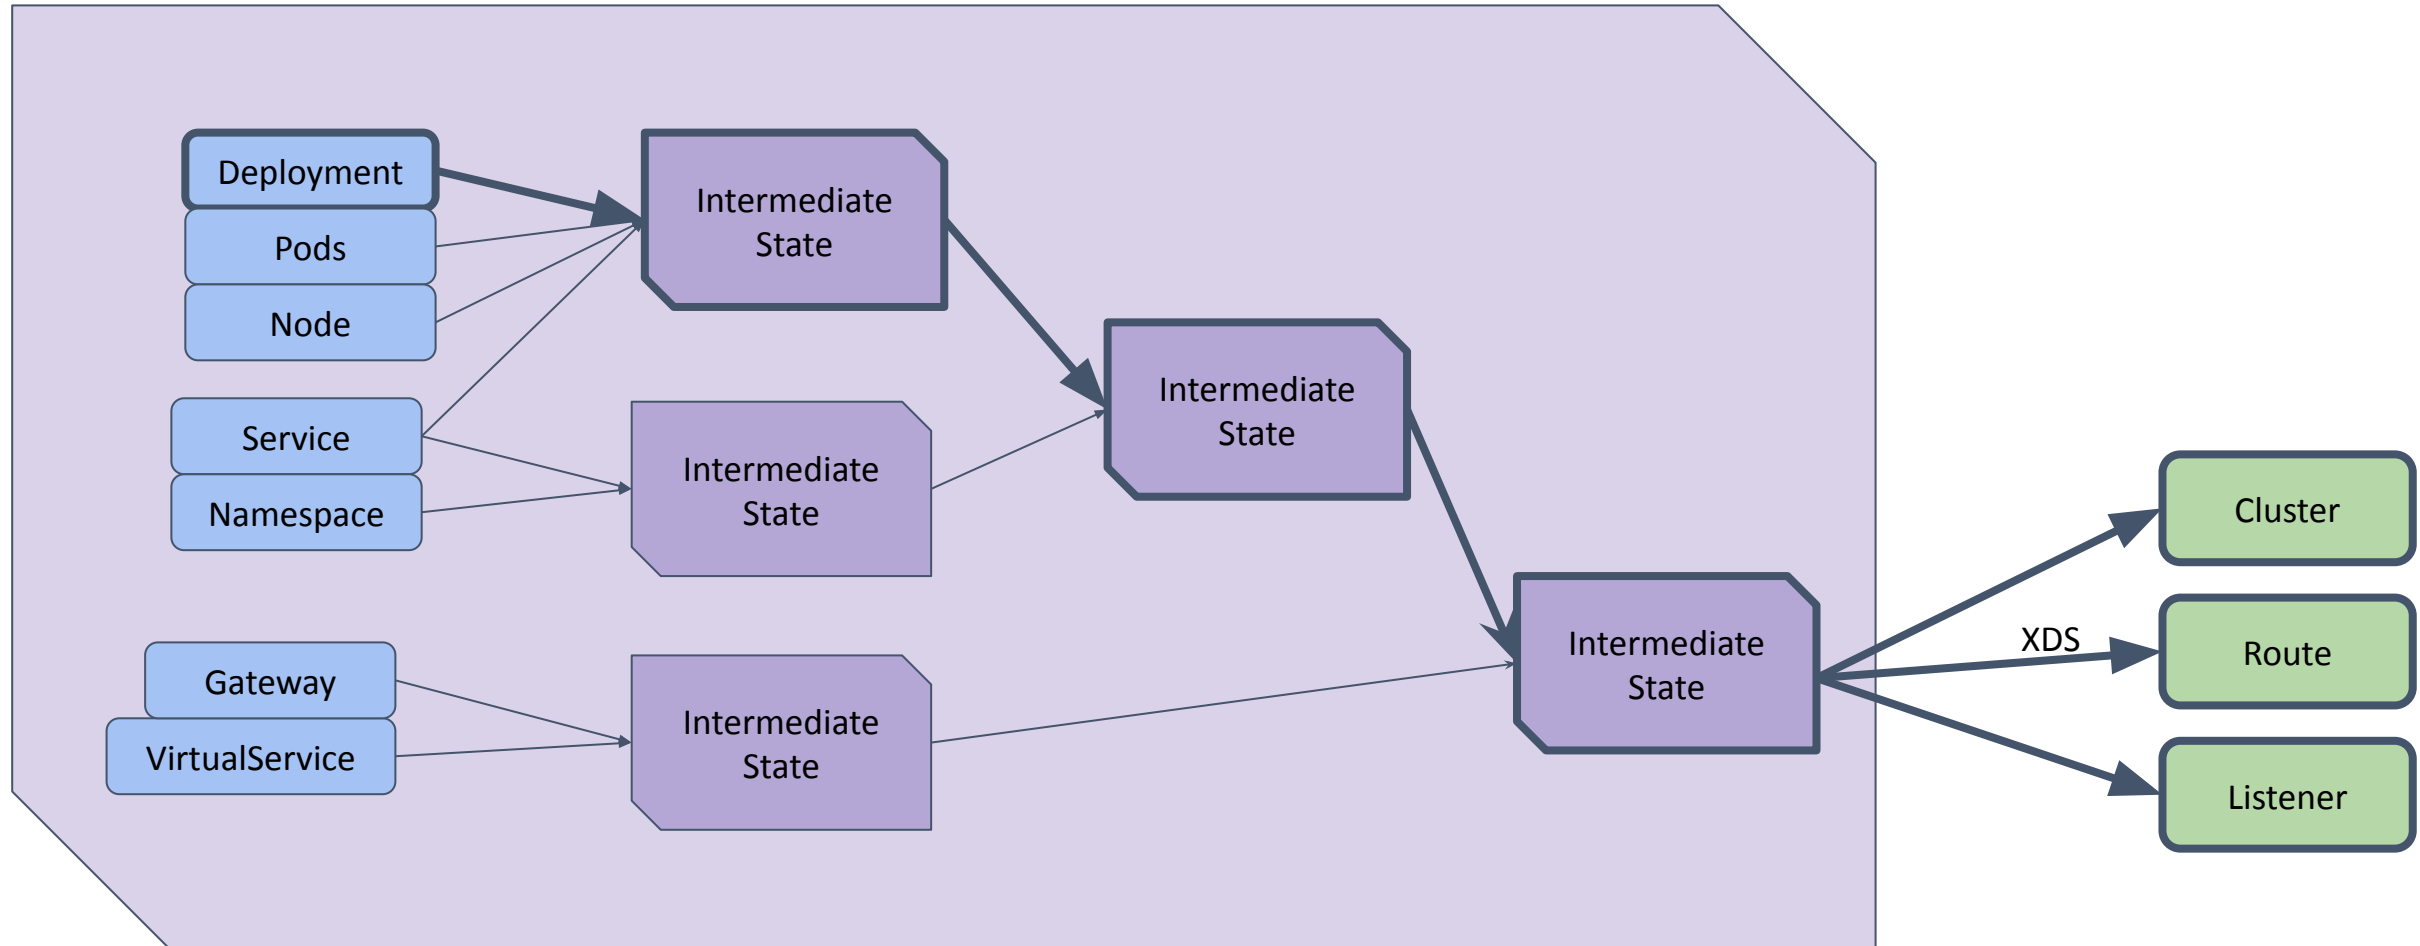


KubeCon



CloudNativeCon

North America 2023



Index Creation



KubeCon



CloudNativeCon

North America 2023

```
IpIndex := CreateIndex[Pod, string](Pods, func(p Pod) []string {  
    return pod.Spec.PodIPs  
})  
  
// Now we can do IpIndex.Lookup(podIP)
```

Outputs: Kubernetes



KubeCon

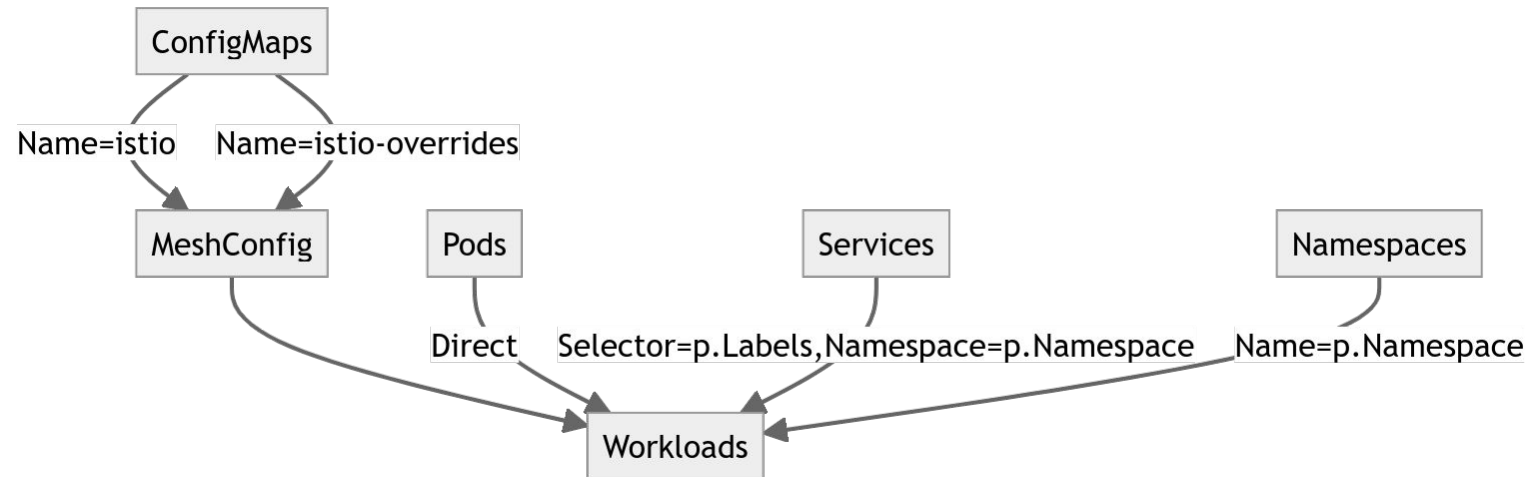


CloudNativeCon

North America 2023

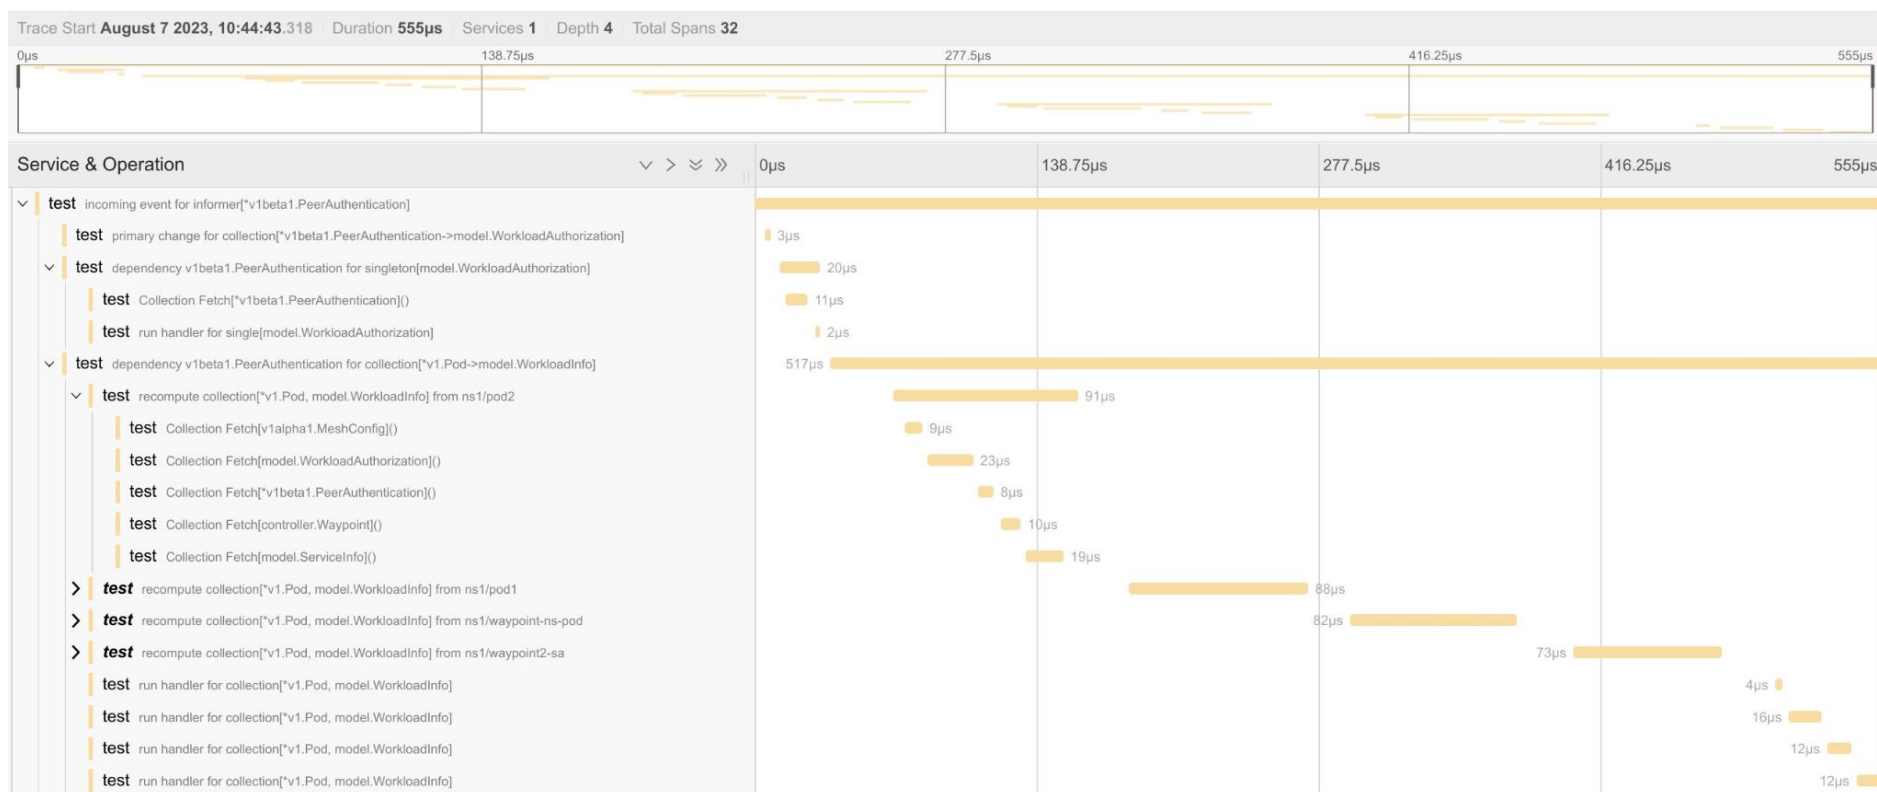
```
// Determine our desired Pods  
DesiredPods := NewCollection(...)  
// Watch Pods actually in the cluster  
LivePods := InformerWatch[pod]()  
// Apply the generated Pod to the cluster, and keep them in sync if they  
change by watching the actual Workloads in the cluster.  
SyncWithApply(DesiredPods, LivePods)
```

Automatically generated architectural diagrams





Auto-tracing instrumentation





Lots of opportunities to improve testing

- Testing pure functions is much easier than controllers
- Unified framework gives more possibilities
 - Automated fuzz testing?

Where are we?



KubeCon



CloudNativeCon

North America 2023

Demoware!

- Everything here is actually implemented as a prototype!
- I have a branch replacing ~50% of Istio with this model
- ... but deploying to production is another matter



PromCon
North America 2021



**Please scan the QR Code above
to leave feedback on this session**

Questions?

**Catch me at "Service Mesh Battle Scars"
In W176 at 5:25pm**