
KubeCon
CloudNativeCon
North America 2023


KubeCon | CloudNativeCon

North America 2023

Cloud Native Storage: Storage TAG Intro, Projects, Landscape & Technology

Alex Chircop, Akamai
Xing Yang, VMware
Raffaele Spazzoli, Red Hat

Agenda

- Overview of the TAG
  - How to join and how to help
  - Overview of storage projects in the CNCF
- What is Cloud Native Storage
  - Why it is important
  - What's New
- Overview of the CNCF Storage Landscape document
- Overview of Data on Kubernetes document
- Overview of the Performance and Benchmarking document
- Overview of the Cloud Native DR document
- Community

Meetings

TAG Storage

CNCF SIGs were renamed TAGs (Technical Advisory Groups)

Meetings are on the 2nd and 4th Wednesday of every month at 8AM PT (USA Pacific)

- **Home:** https://github.com/cncf/tag-storage
- **Conf call:** http://bit.ly/cncf-storage-tag-call
- **Agenda:** http://bit.ly/cncf-storage-tag-minutes
- **Mail list:** https://lists.cncf.io/g/cncf-tag-storage

Our calls and membership are open!

Who we are

- We are a diverse set of users & developers of Cloud Native technologies with a storage focus
- We are leaders & early adopters

Co-Chairs	Tech Leads	TOC Liaisons
■Alex Chircop ■Xing Yang ■Raffaele Spazzoli	■Luis Pabón ■Sheng Yang ■Nick Connolly	■Nikhita Raghunat ■Matt Farina

Questions? Reach out and feel free to connect on our mailing list, and CNCF Slack!

What we do

“Scale contributions by the CNCF technical and user community, while retaining integrity and increasing quality in support of the CNCF mission (to make cloud native computing ubiquitous).”

What we do

“Scale contributions by the CNCF technical and user community, while retaining integrity and increasing quality in support of the CNCF mission (to make cloud native computing ubiquitous).”

...this means we

What we do

“Scale contributions by the CNCF technical and user community, while retaining integrity and increasing quality in support of the CNCF mission (to make cloud native computing ubiquitous).”

...this means we

- **Educate**
- **Review Projects**
- **Engage with the user community**
- **Provide subject matter expertise**

CNCF Storage Projects

Graduated

etcd

Incubating

LONGHORN

CNCF Projects: https://www.cncf.io/projects/
Sandbox Projects: https://www.cncf.io/sandbox-projects/

CNCF Storage Projects

https://www.cncf.io/projects/

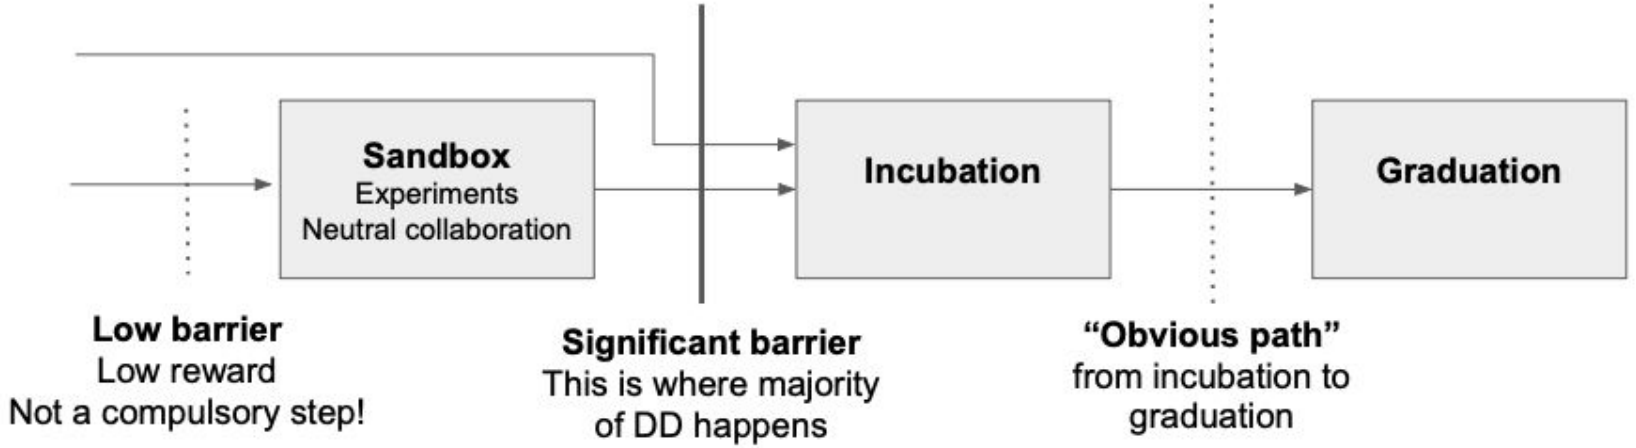

Sandbox	Incubation	Graduation
■Experiments ■IP Policy ■Build Community	■Used successfully in production ■Healthy number of committers ■Project metrics	■Mainstream production use ■Security audits ■Committers from multiple organisations

Why should you think about this?

Why should you think about this?

Why should you think about this?

There is no such thing as a stateless application …

… all applications store state somewhere!

Cloud Native Storage is Here!

Move Stateful Workloads to K8s

- Automation
- Scale
- Performance
- Failover

➪ **Broad ecosystem and CSI support**

➪ **Operators for databases, message queues, *and many more!***

CNCF Storage Whitepaper

Whitepaper: https://bit.ly/cncf-storage-whitepaperV2

1. Definition of the attributes of a storage system
2. Definition of the layers in a storage solution with a focus on terminology and how they impact the attributes
3. Definition of the data access interfaces in terms of volumes and application APIs
4. Definition of the management interfaces

Storage Attributes

Availability	Scalability	Performance	Consistency	Durability
Failover	Clients	Latency	Delay to access correct data after a commit	Data protection
Moving access between nodes	Operations	Operations		Redundancy
Redundancy	Throughput	Throughput	Delay between commit and data being committed to non-volatile store	Bit-Rot
Data Protection	Components			

Storage Layers

Orchestrator, Host and Operating System
Storage Topology (centralized, distributed, sharded, hyperconverged)
Data Protection (RAID, Erasure coding, Replicas)
Data Services (Replication, Snapshots, Clones, etc.)
Physical, Non-Volatile Layer

Why this matters …

Let’s take a look at a couple of different use cases and deployments:

- Hyperconverged
- Block Volumes
- Shared Filesystems
- Object Stores

Why this matters …

Let’s take a look at a couple of different use cases and deployments:

- **Hyperconverged**
  - availability: converged fault and change management domains
  - performance: shared network and compute

Why this matters …

Let’s take a look at a couple of different use cases and deployments:

- **Block Volumes**
  - useful to disaggregate compute and storage
  - availability: ability to move volumes between nodes
  - performance: typically lower latency, but needs good connectivity between compute and storage nodes

Why this matters …

Let's take a look at a couple of different use cases and deployments:

- **Shared Filesystems**
  - can be used by multiple nodes at the same time
  - consistency: distributed locks, cache coherency is hard
  - layers: could be built on block, object stores etc … and that determines many attributes

Why this matters …

Let’s take a look at a couple of different use cases and deployments:

- **Object Stores**
  - scale: almost infinite for capacity, and throughput
  - latency: higher than data on a volume
  - performance: RPS is often the determining factor

Data Workloads on Kubernetes

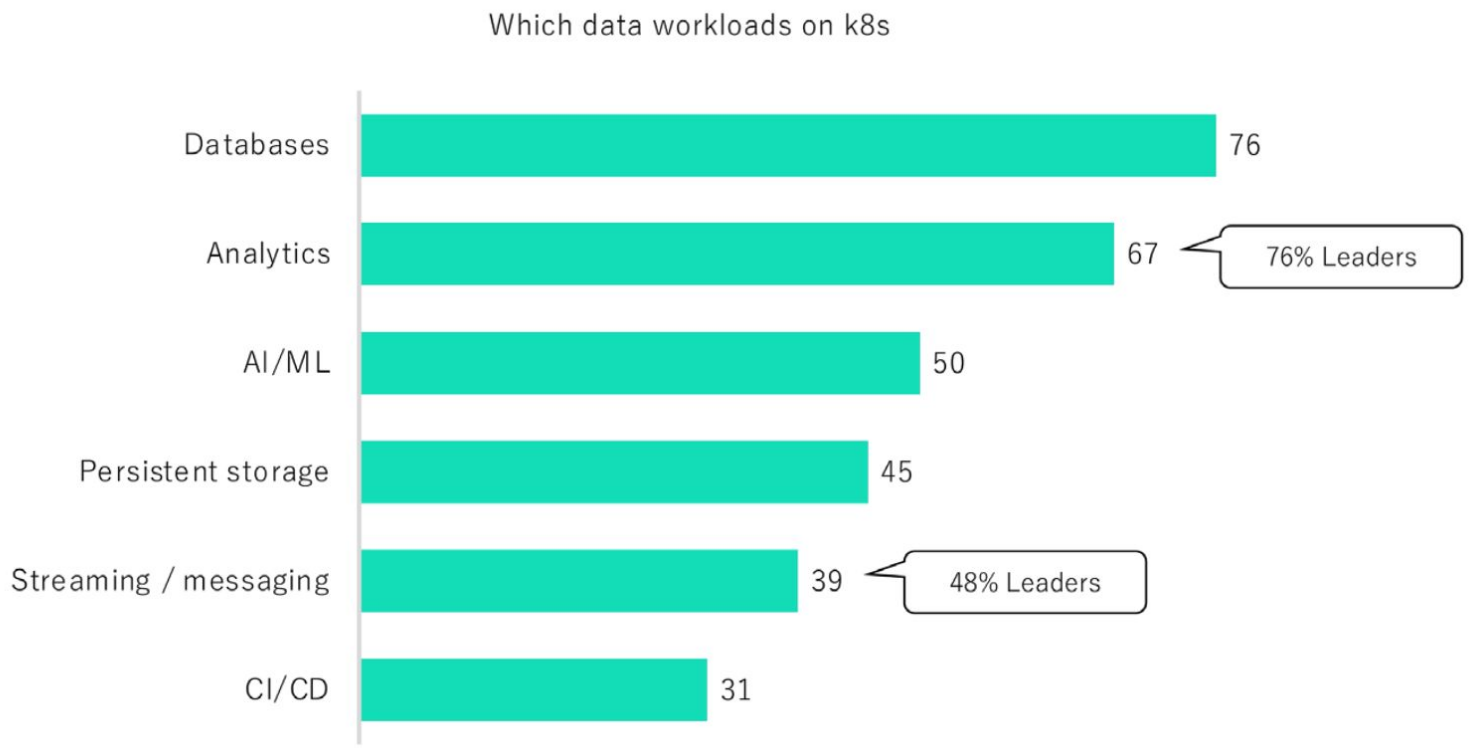

Data on Kubernetes Community 2022 Survey

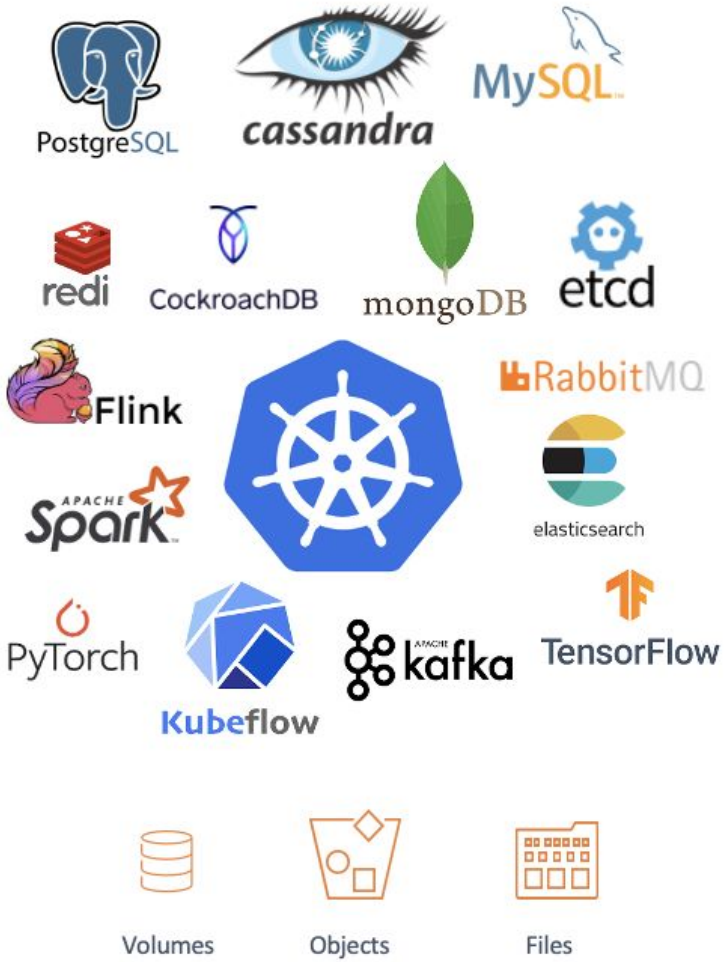

Data on Kubernetes Whitepaper

KubeCon CloudNativeCon North America 2023

https://bit.ly/cncf-storage-dok-whitepaper

- Describe patterns of running data on Kubernetes
- Collaborating with Data on Kubernetes Community (DoKC)
- Focusing on databases in v1
- Paper layout
  - Attributes of a storage system and how they affect running data in Kubernetes
  - Running data inside vs outside of Kubernetes
  - Common Kubernetes patterns and features used when running data on Kubernetes
  - Observability
  - Security
  - Day 2 operations

Storage Attributes and Running Data in K8s

- Storage System Attributes
  - Attributes
    - Availability
    - Consistency
    - Durability
    - Scalability
    - Performance
    - Observability
    - Elasticity
  - Storage Stacks
    - Stacks/Layers
    - Disaster Recovery

Running Data inside vs outside of K8s

- Managed database services, provided by most cloud providers
- Running data inside K8s with operators
  - Declarative approach
  - Automate “Day 2 Operations”
  - Externalize database functionalities such as monitoring, cert management to third parties

Kubernetes Operators

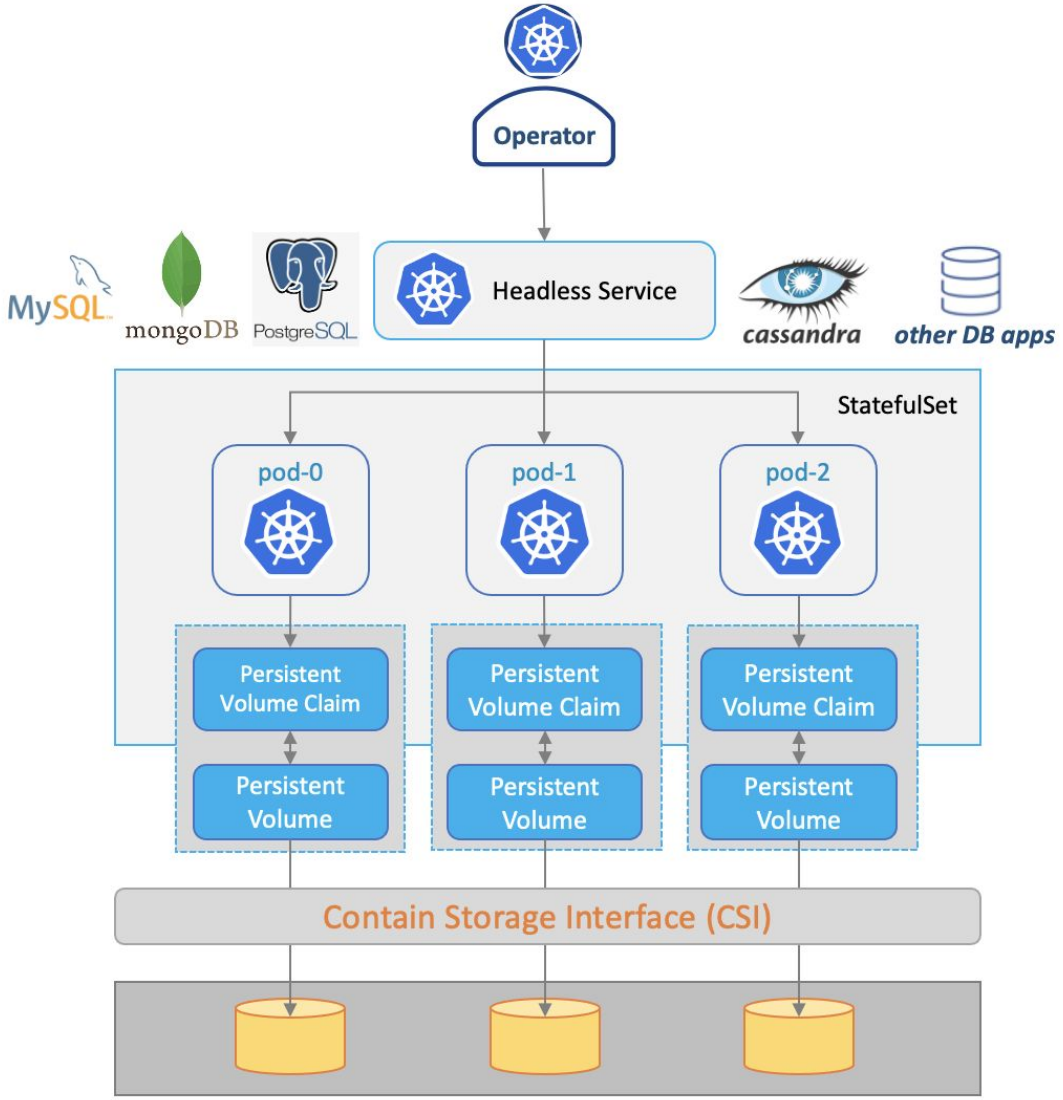

Kubernetes Operators (cont.)

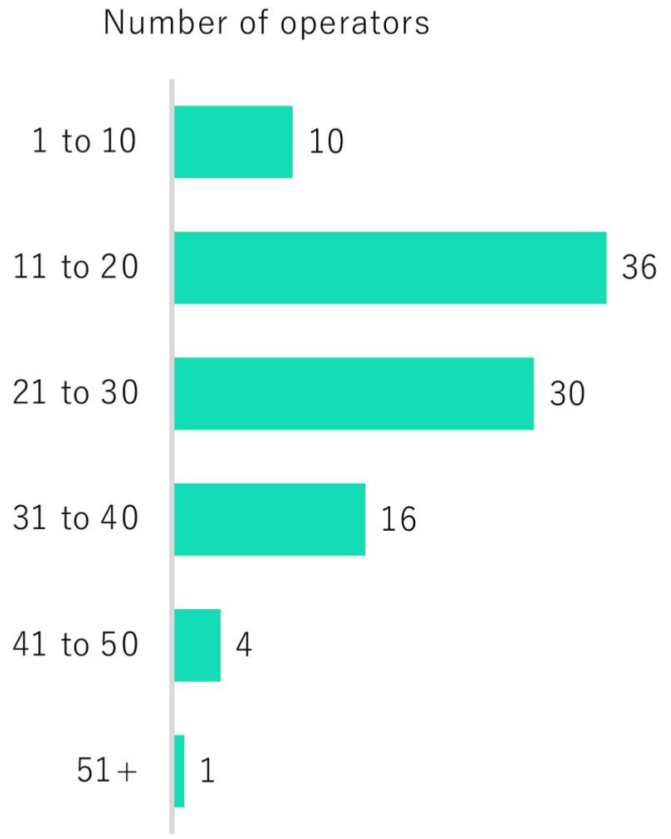

Data on Kubernetes Community 2022 Survey

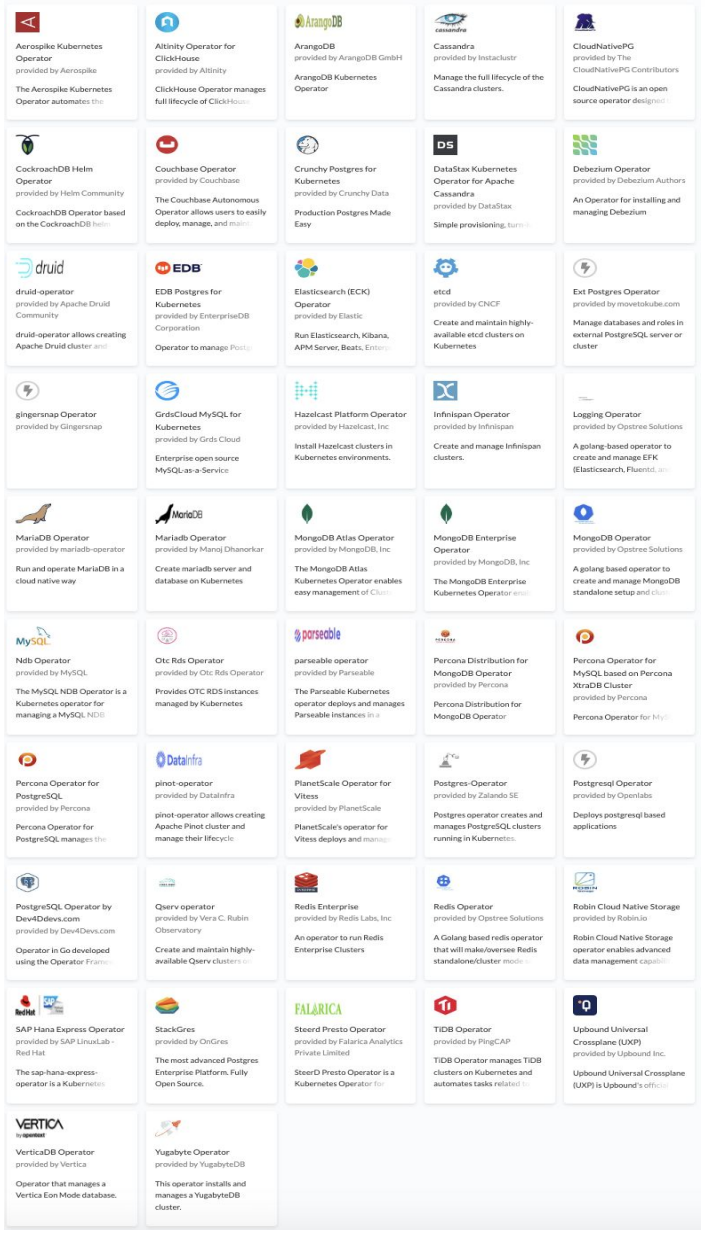

- https://operatorhub.io/
  - 331 operators
  - 47 database operators, including etcd and Vitess
  - 9 PostgreSQL operators, including CloudNativePG
- Other operators not listed in operatorhub

- https://github.com/dokc/operator-feature-matrix

Common K8s Patterns and Features

- Kubernetes Operators
- Container Storage Interface (CSI)
- Kubernetes Workload APIs
- Topology Aware Scheduling
- Pod Disruption Budget
- Resource Management
- Separation of Control Plane and Data Plane
- Default Secure

Performance Whitepaper

KubeCon CloudNativeCon North America 2023

Whitepaper: https://bit.ly/cncf-tag-storage-performance-benchmarking

- Definition of common concepts for measuring performance and benchmarking for volumes and databases

- Definition of common pitfalls and considerations
  - Operations vs Throughput
  - Topology, Data Protection, Data Reduction, Encryption matters …
  - But Latency often matters more …
  - Concurrency for queues, clients and backends
  - Caching happens at multiple layers
  - Be critical and beware of results that are too good to be true!

Performance Whitepaper

Whitepaper:
https://bit.ly/cncf-tag-storage-performance-benchmarking

TL;DR - important takeaway:

published results are not useful for making comparisons - it is hard to compare published results without a deep understanding of the test conditions, so it is always important to run your own test, on your own environment with your own applications

High-level DR Approaches

Failure Domain is either a data center or a cloud region

Active/Passive

Active/Active

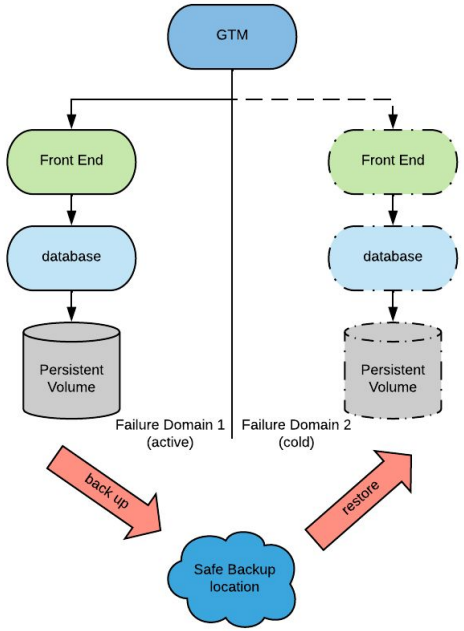

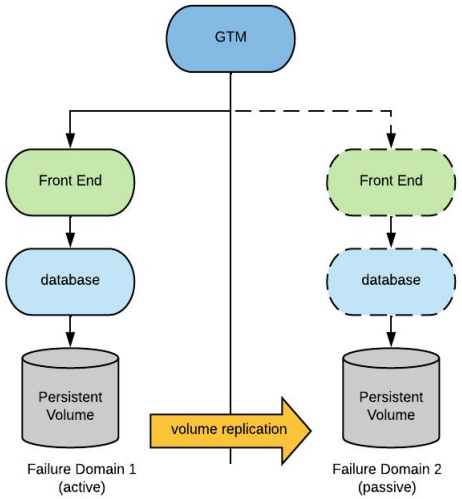

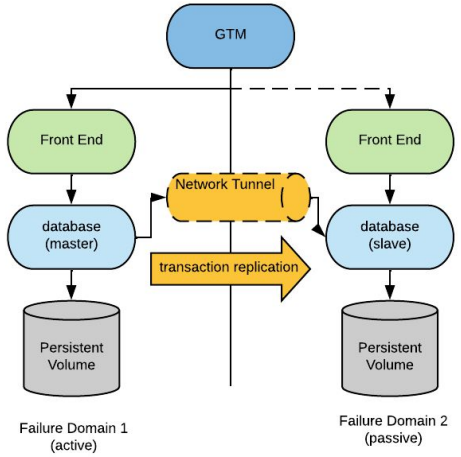

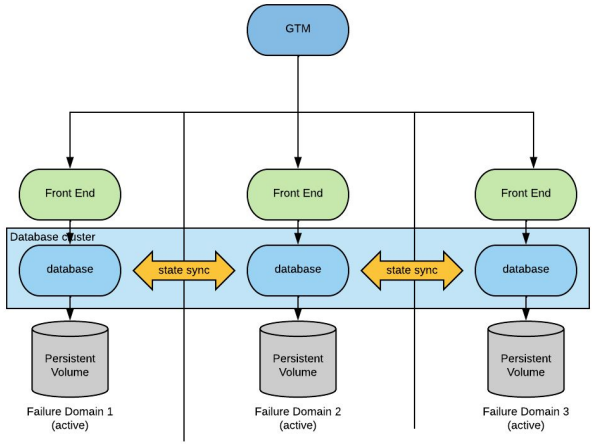

Backup / Restore	Volume Replication	Transaction Replication	Distributed Stateful Workloads
Storage Capabilities			
❑ backup/restore	❑ sync/async repl		
Network Capabilities			
❑ global load balancer	❑ global load balancer	❑ global load balancer	❑ global load balancer
		❑ east-west path	❑ east-west path
Workload Capabilities			
		❑ primary/secondary	❑ distributed stateful workload

Capabilities and Products

Backup & Restore

OADP

kasten by Veeam

TRILIOVAULT

COHESITY

portworx by Pure Storage

VELERO

k8up

Global Load Balancing

DNS

K8GB

Primary/ Secondary enabled middleware

MySQL

Cloud SQL

Azure SQL

Amazon Aurora

Volume Replication

Red Hat OpenShift Data Foundation

portworx by Pure Storage

LONGHORN

ROOK

East-West path

SUBMARINER

CALICO

cilium

Distributed middleware

Vault

yugabyteDB

mongoDB.

cassandra

CockroachDB

kafka

Azure Cosmos DB

Google Cloud Spanner

NATS

PULSAR

Cloud Native Disaster Recovery

Whitepaper: http://bit.ly/cncf-cloud-native-DR

Concern	Traditional DR	Cloud Native DR
Type of deployment	active/passive, rarely active/active	Active / active
Disaster Detection and Recovery Trigger	Human	Autonomous
Disaster Recovery Procedure execution	Mix of manual and automated tasks	Automated
Recovery Time Objective (RTO)	From close to zero to hours	Close to zero
Recovery Point Objective (RPO)	From zero to hours	Exactly zero for strongly consistent deployments. Theoretically unbounded, practically close to zero for eventual consistent deployments.
DR Process Owner	Often the Storage Team	Application Team
Capabilities needed for DR	From storage (backup/restore, volume replication)	From networking (east-west communication, global load balancer)

Anatomy of a Distributed Stateful Workload

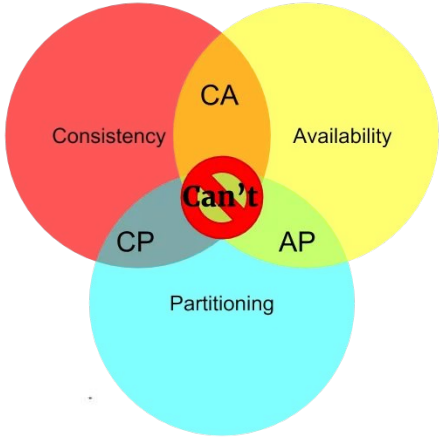

CAP Theorem

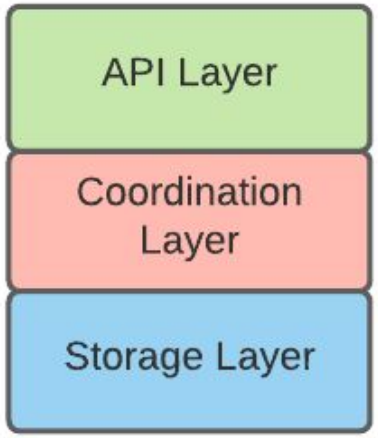

Stateful Workload Logical Tiers

Replicas & Partitions

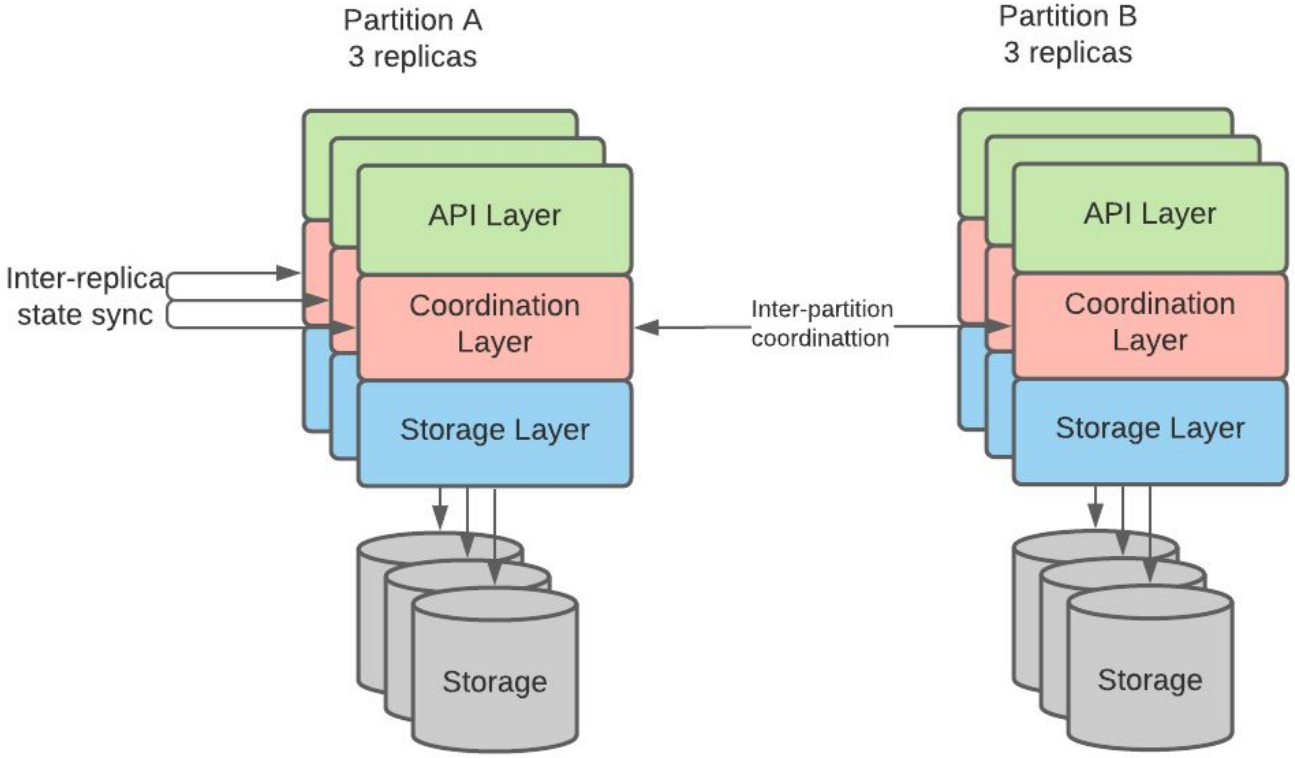

Examples of Consensus Protocol choices

Product	Replica consensus protocol	Shard consensus protocol
Etcd	Raft	N/A (no support for shards)
Consul	Raft	N/A (no support for shards)
Zookeeper	Atomic Broadcast (a derivative of Paxos)	N/A (no support for shards)
ElasticSearch	Paxos	N/A (No support for transactions)
Cassandra	Paxos	Supported, but details are not available.
MongoDB	Paxos	Homegrown protocol.
CockroachDB	Raft	2PC
YugabyteDB	Raft	2PC
TiKV	Raft	Percolator
Spanner	Raft	2PC+high-precision time service
Kafka	A custom derivative of PacificA	Custom Implementation of 2PC

CNDR -- Strong Consistency - Kubernetes Reference Architecture

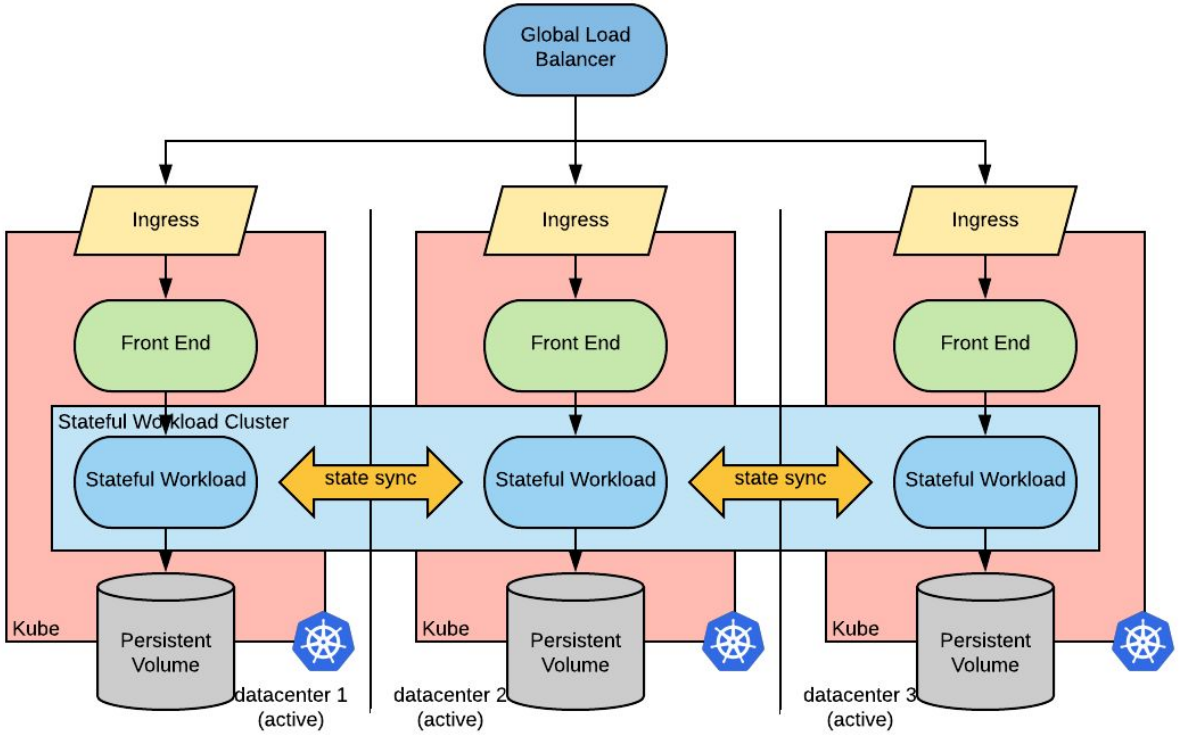

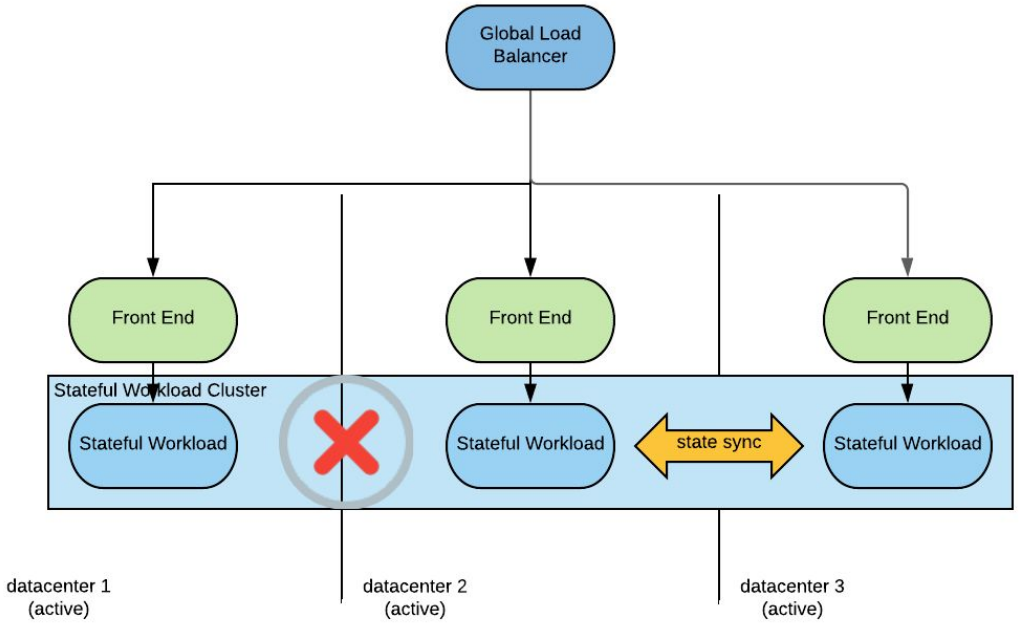

Community

- How you can get involved?
  - Join our meeting
    - 2nd & 4th Wednesday each month
  - Submit and help review projects for consideration

- We value community presentations of projects in the cloud native storage space including, but not limited to: *management frameworks, block stores, filesystems, object stores, key-value stores and databases* -> learn and work with the community

- Consider a role in the TAG!

- Contribute to TAG projects and help the community!

Please scan the QR Code above
to leave feedback on this session