



**KubeCon**



**CloudNativeCon**

**North America 2023**





KubeCon



CloudNativeCon

North America 2023

# Cluster API Deep Dive

## Improving performance up to 2k Clusters

*Fabrizio Pandini & Stefan Büringer, VMware*

# Who we are



KubeCon



CloudNativeCon

North America 2023



**Stefan Büringer**  
Staff Engineer  
*VMware*



**Fabrizio Pandini**  
Staff Engineer  
*VMware*

# Before we start...



KubeCon



CloudNativeCon

North America 2023

*All facts in this presentation are based on a true story*



This is a talk about the performance optimization work done in Cluster API v1.5.0, but the **theory** and the concepts behind this work **applies to all Kubernetes controllers**.

Same for the lessons learned and **experiences** we gathered.

## Agenda:

- How to get ready for performance optimization
- Scaling Cluster API up to 2K clusters
  - Controller optimization 101
  - Performance optimization in practice

# Get the right tools for the job



KubeCon



CloudNativeCon

North America 2023

## What you need

## Why do you need it

## TL;DR

Metrics

A software metric is a measurement of quantifiable or countable software characteristics.

Profiling

Profiling is a form of dynamic code analysis that measures, for example, the space (memory), the frequency and duration of function calls.

Tracing

A specialized use of logging to record information about a program's execution, for example, the duration of function calls.

Logs

Logging is the act of keeping a log of events that occur in a computer system, such as problems, errors or just information on current operations.

Automation

The act of making a process repeatable without human intervention.

Mocks

A mock, in software engineering, is a simulated object or module that acts as a stand-in for a real object or module.

*Performance  
Analysis  
&  
Investigation  
of bottlenecks*

*Inner dev loop*


*Understand the system at scale,  
across many reconcile operations*

*Dig into a single  
reconcile operation*

*Observe, optimize, repeat  
Prevent regressions*

*Increase speed  
Reduce costs*

Must have



Nice to have

# What metrics do you need ?



KubeCon



CloudNativeCon

North America 2023

## Metric categories

### “User-facing” Metrics

A user-facing metric is a measurement of quantifiable or countable software characteristics that **directly relates to the user perception of how the system works**.  
E.g. Machine provisioning time

### “Internal” Metrics

A “internal” metric is a measurement of quantifiable or countable software characteristics that **documents how the system works under the hood**.  
Not performing “internal” metrics are usually perceived by the users as slowness, unresponsiveness etc.  
E.g. Average reconcile time

## Why do you need it

*Define goals*  
*Measure success*

*Understand how the system performs at scale*

*Point to what needs to be optimized next*

# How can I get all of this?



KubeCon



CloudNativeCon

North America 2023

You get most of it for free! (not everything)



“User-facing”  
Metrics

kube-state-metrics provides the basic tooling to get metrics from CRDs  
**You need to provide a config defining how metrics should be derived from your CRDs**

“Internal”  
Metrics

Go SDK, client-go, and controller-runtime have you covered  
Consider adding custom metrics only if strictly necessary

Profiling

Go SDK has you covered



Tracing

**You need to instrument your code**  
Start small, add more when you need it

Logs

You probably already have them  
Consider adding more logs only if necessary



Automation

**You need to write your own automation**, but you probably already have something



Mocks

**Mocking complex systems like a cloud machine can be time-consuming**  
The investment usually pays back

## CAPI observability stack

[Config for CAPI CRDs](#)  
[kube-state-metrics deployment](#)  
[Prometheus deployment](#)  
[Grafana deployment](#)



**Dashboards  
included!**

[Parca deployment](#)

[Tempo deployment](#)

[Promtail deployment](#)  
[Loki deployment](#)

[Scale E2E test](#)

[Cluster API In Memory provider](#)

# 2k Clusters!



KubeCon




CloudNativeCon

North America 2023

After getting the right tools for the job, we defined our performance goal:

**Cluster provisioning time (a business metric) must remain ~constant from Cluster n°1 to n° 2k**

Spoiler, we did it! The delta between the average time to provision the first 100 clusters and the last 100 clusters show a neglectable increase of +2.3%.



This is great!

But the most interesting part is to look back at how we managed to **keep the system responsive** while scaling up (and while working at scale).



# Controller optimization 101

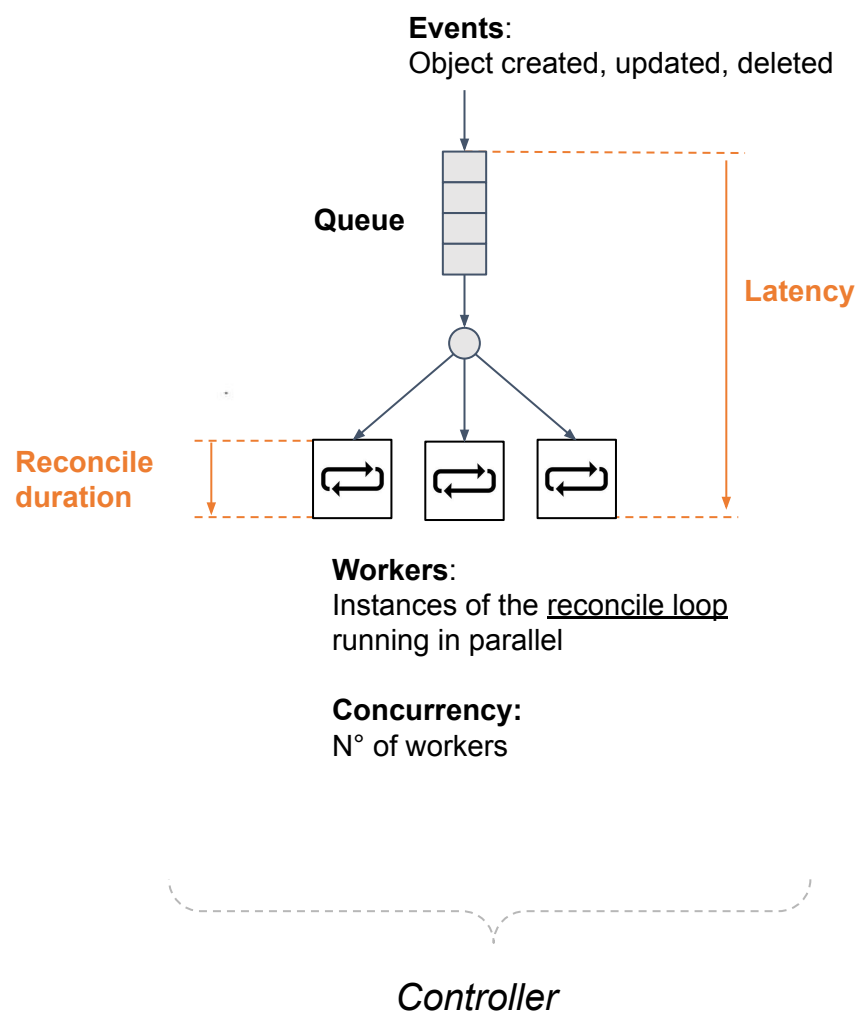


KubeCon



CloudNativeCon

North America 2023



The main performance characteristic of a controller is the **latency** between an **object create, update or delete** and its **successful reconciliation**

The “*worst case latency*” is the time required until the last object in the queue can be processed, and it can be determined via:

$$\frac{\text{\# objects in the queue}}{\text{concurrency}} * \text{reconcile duration}$$

An example:

$$\frac{2k \text{ clusters}}{10 \text{ workers}} * 3s = 10m \text{ 😞}$$

# Controller optimization 101



KubeCon



CloudNativeCon


North America 2023

Is it possible that all 2k Clusters are in the reconcile queue at the same time?

Yes, because on top of events generated when an object is created, updated or deleted, controller-runtime does a **periodic resync**, i.e. it periodically adds all objects to the queue.

Additionally controllers can **“requeue” an object back into the queue**.

To stay responsive after the resync/requeue **controllers have to process events as fast as possible**.



# Controller optimization 101



KubeCon



CloudNativeCon

North America 2023

How to improve:  $\frac{2k \text{ clusters}}{10 \text{ workers}} * 3s = 10m ?$

**Option 3**  
Reduce # objects in the queue



$$\frac{\text{\# objects in the queue}}{\text{concurrency}} * \text{reconcile duration}$$



**Option 1**  
Increase concurrency



**Option 2**  
Reduce reconcile duration

# Option 1: Increase concurrency

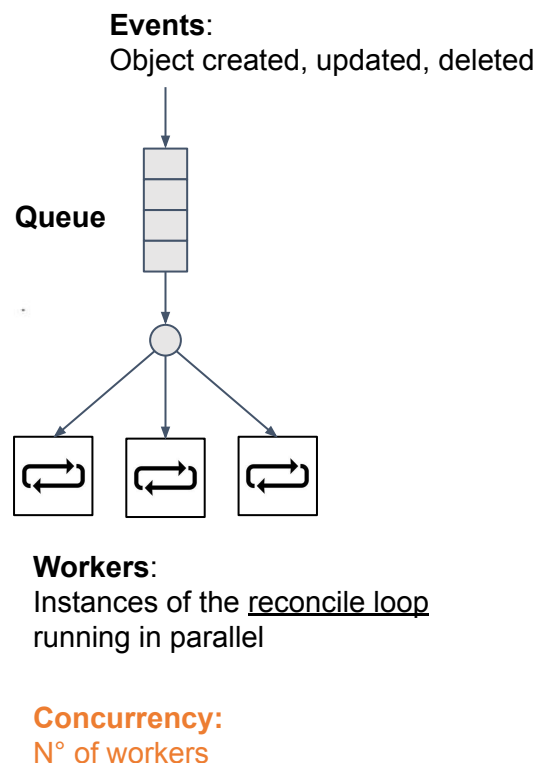


KubeCon



CloudNativeCon

North America 2023



If we use only one worker (i.e. concurrency 1) the system will reconcile one object at a time.

This not performant enough for high scale.

On the other side, **if we use too many workers, this might not help the system to scale**, because it will put too much pressure on the API server.

In Cluster API the default concurrency is 10 workers:

- Based on our observation 10 is good trade off when controllers have an average reconcile duration lower than 200-250ms
- We only run the controller with the highest reconcile duration (KCP) with more workers, but to make it sustainable, **we invested a lot in improving KCP**.

**TL;DR increasing concurrency is not a silver bullet**

# Option 2: Reduce reconcile duration



KubeCon



CloudNativeCon

North America 2023

Reducing reconcile durations is more complex than tuning the number of workers.

However, **if you follow the lead of metrics, profiles and traces** you can **improve performance by doing small, surgical changes**.

In Cluster API we **iteratively improved** the performance by:

- Removing noise

- Identifying slow reconcilers

- Removing bottlenecks

- Repeat

# Remove noise: client-side rate limiting



KubeCon




CloudNativeCon

North America 2023

Kubernetes has client- and server-side rate-limiting<sup>1</sup> of requests to the API server, which are mechanisms that protect the system from clients doing excessive API calls.

Cluster API's default rate-limiting is 20 QPS and 30 burst, **which is enough for up to 100 Clusters**. When you cross this limit, rate-limiting “randomly” delays your API calls, and this makes reconcile durations non-deterministic<sup>2</sup>.



**TL;DR; when running at scale increase QPS and burst (we used 100 QPS, 200 Burst)**

<sup>1</sup> [API Priority and Fairness](#)

<sup>2</sup> Increasing concurrency doesn't help as all reconcilers share the same client and thus also the same rate-limiting.

# How to identify slow reconcilers




KubeCon



CloudNativeCon

North America 2023

This step is easy! Just run the scale test and look at the “reconcile duration” controller-runtime metric



Next: focus on slow reconcilers one-by-one and dig deeper via metrics, profiling and tracing.

The following slides show what we found and how we improved performance.

# Avoid duplicate expensive operations



KubeCon

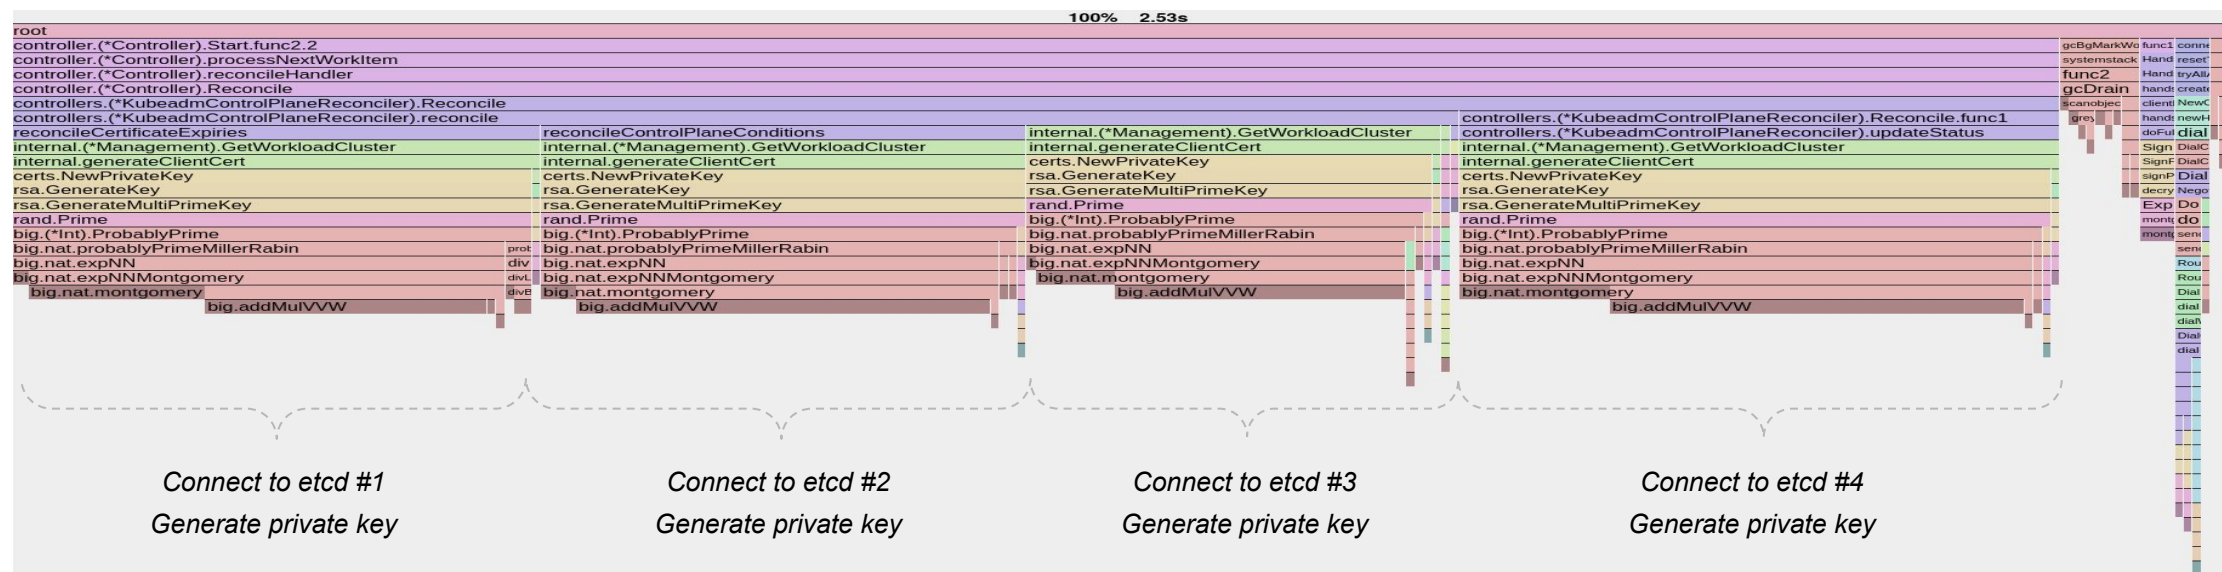


CloudNativeCon

North America 2023

We started looking at profiling to check where controllers were spending most of the time. The first thing we noticed was that some controllers were **doing the same expensive operation multiple times**. This is obviously not good for performance.

For instance, KCP, when connecting to etcd, was generating a private key and a certificate multiple times; this is an expensive operation. Another expensive operation is e.g. creating Kubernetes clients.



This issue was addressed by **caching** private keys/clients for each cluster, and **reusing them across many reconciles**. As a result reconcile duration went down by ~75%



# API calls



KubeCon

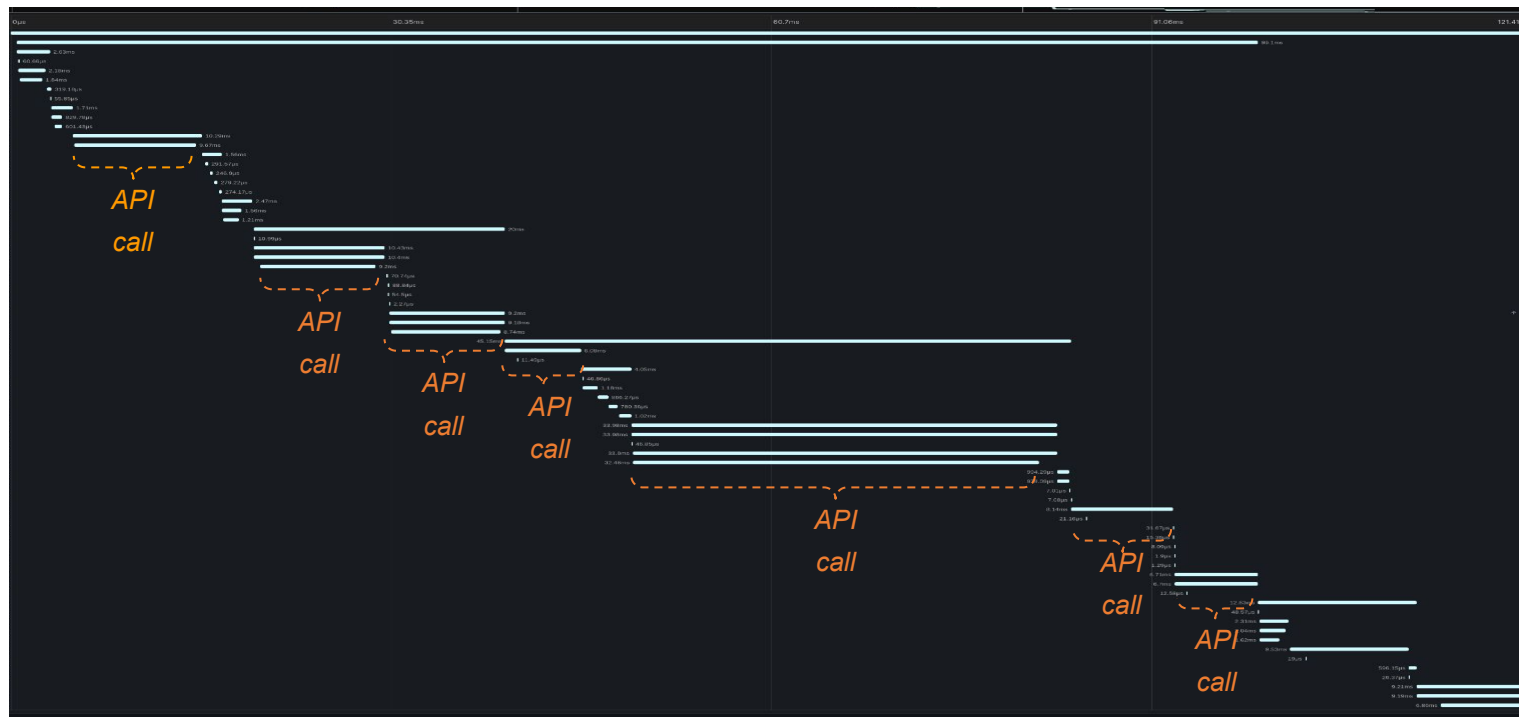


CloudNativeCon

North America 2023

Controllers use clients to read/write data from/to the Kubernetes API server.

As the complexity of your controller grows, **the number of API calls quickly increases, and this can impact performance because of network latency**. This becomes evident as soon as you start looking at traces.



How can we **improve performance and avoid pressure on the API server with many API calls?**

# Caching API calls



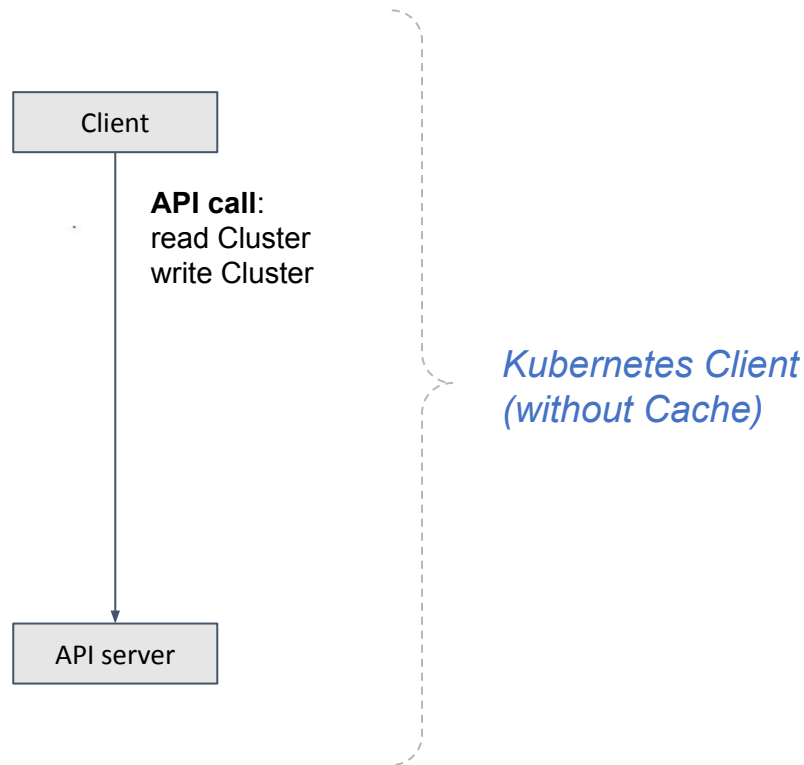
KubeCon



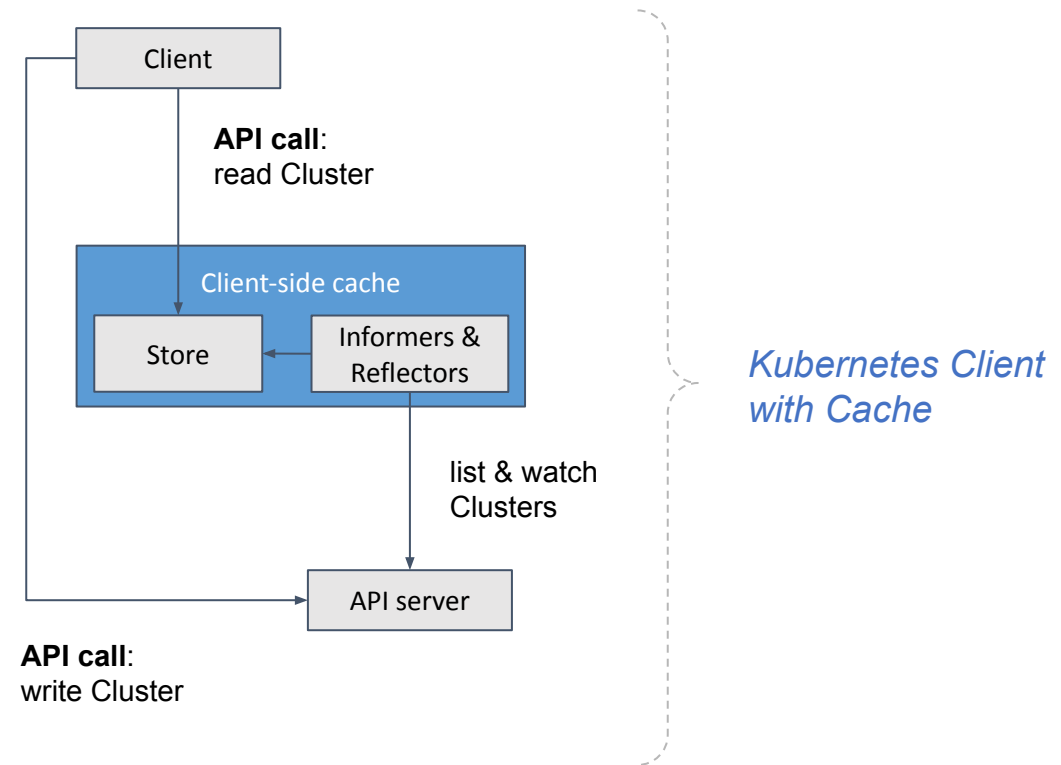
CloudNativeCon

North America 2023

Kubernetes clients in their simplest form, forward all API calls to the API server.



Controller-runtime's default client is an enhanced client, which includes a cache implemented using client-go primitives.



# Caching caveats

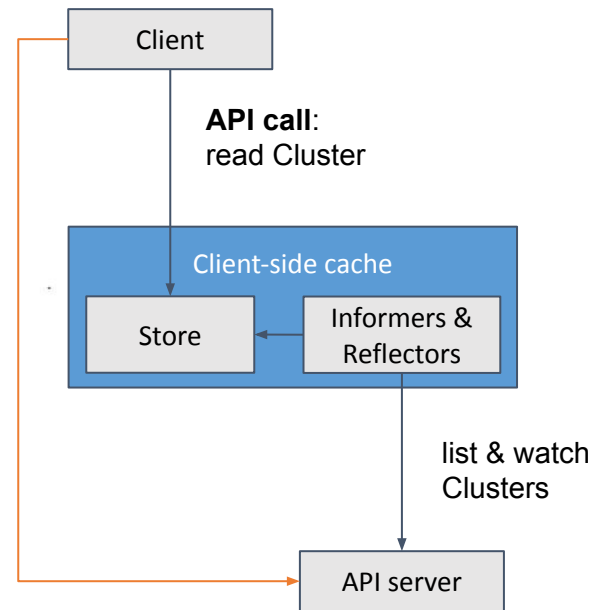


KubeCon



CloudNativeCon

North America 2023



API call:  
write Cluster  
Unstructured read and write (default)

In Cluster API we are using the controller-runtime's default client almost everywhere, so **most of the API calls are already cached** ( $\mu$ s vs ms).

However, not all API calls are cached by default:

- **Write calls are not cached**, so we made sure to have as few write calls as possible.
- **Read calls are not cached by default for Unstructured types**. But it is possible to cache them with a custom cache configuration, which we did.

But caching comes with other caveats too:

- The cache **increases the memory** used by controllers (with 2k clusters, CAPI core controller required 2-4GB depending on the topology).
- You can have **stale reads** from the cache. This *usually* is not a problem, because as soon as the new data arrives a new reconcile will be triggered.

# Option 3: Reduce # objects in the queue




KubeCon



CloudNativeCon

North America 2023



We can't control events for objects being created, updated or deleted.

We can't control periodic resync\*.

But we can **control how our controller requeues objects.**

At the end of reconcile we can tell controller-runtime to requeue the current object, if we want to reconcile the object later again.

\* Theoretically we could reduce the load by increasing the resync period, but this would be just a workaround as eventually the resync will happen

# Requeueing from controllers

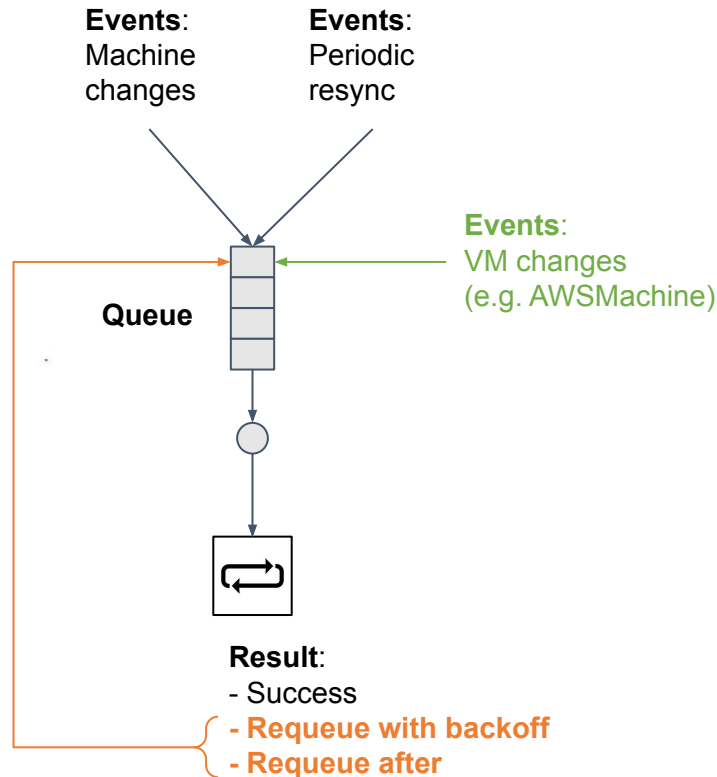


KubeCon



CloudNativeCon

North America 2023



We use *requeue with backoff* for errors. That's ok.

On the other hand sometimes we use *requeue after* to wait for something to happen in the system, e.g. while reconciling the Machine we requeue to wait for the corresponding VM to be ready.

But **reconciling an object every few seconds puts unnecessary load on the system**, and this doesn't help at scale.

Instead we can **use an additional watch** for VM objects, so a reconcile will be only triggered when necessary.

\* Theoretically we could reduce the load by increasing the resync period, but this would be just a workaround as eventually the resync will happen

# Reconciler best practices



KubeCon

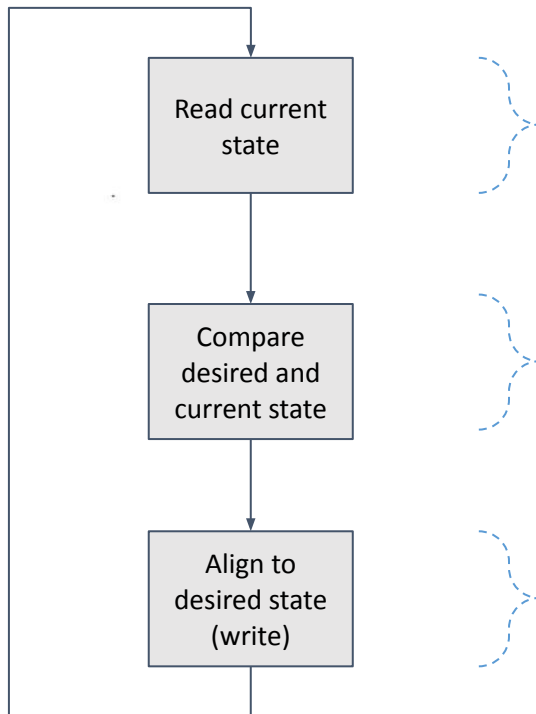


CloudNativeCon

North America 2023



## “Ideal” Reconcile implementation



Read everything at the beginning of the reconcile to avoid duplicate read API calls later

Pay attention to client caching caveats

Avoid duplicate expensive operations  
e.g. generating private keys, creating Kubernetes clients, etc.

Write only once per reconcile (in Cluster API we use defer patch)  
Avoid reconcile after, use watches instead if possible

# Optimization Do's and Don'ts



KubeCon



CloudNativeCon

North America 2023



- Get the right tools for the job
- Define measurable goal(s)
- Invest time in learning how controllers and client-go works
- Iterate fast and with small improvements



- Don't optimize without data that proves there is a problem and where (don't drive blind!)

# Thank you!



KubeCon



CloudNativeCon

North America 2023

Kudos to:

- Christian Schlotter for driving the work for generating metrics from CRDs
- Lennart Jern for triggering initial discussions about reconciler optimizations
- Killian Muldoon and Yuvaraj Kakaraparthi for helping to develop Mocks & Automation in Cluster API

**Please scan the QR Code above  
to leave feedback on this session**

Session QR  
Codes will be  
sent via email  
before the  
event