




KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

eBPF + Wasm: Lightweight Observability on Steroids

Michael Yuan & Yusheng Zheng

Introduction to eBPF and Wasm

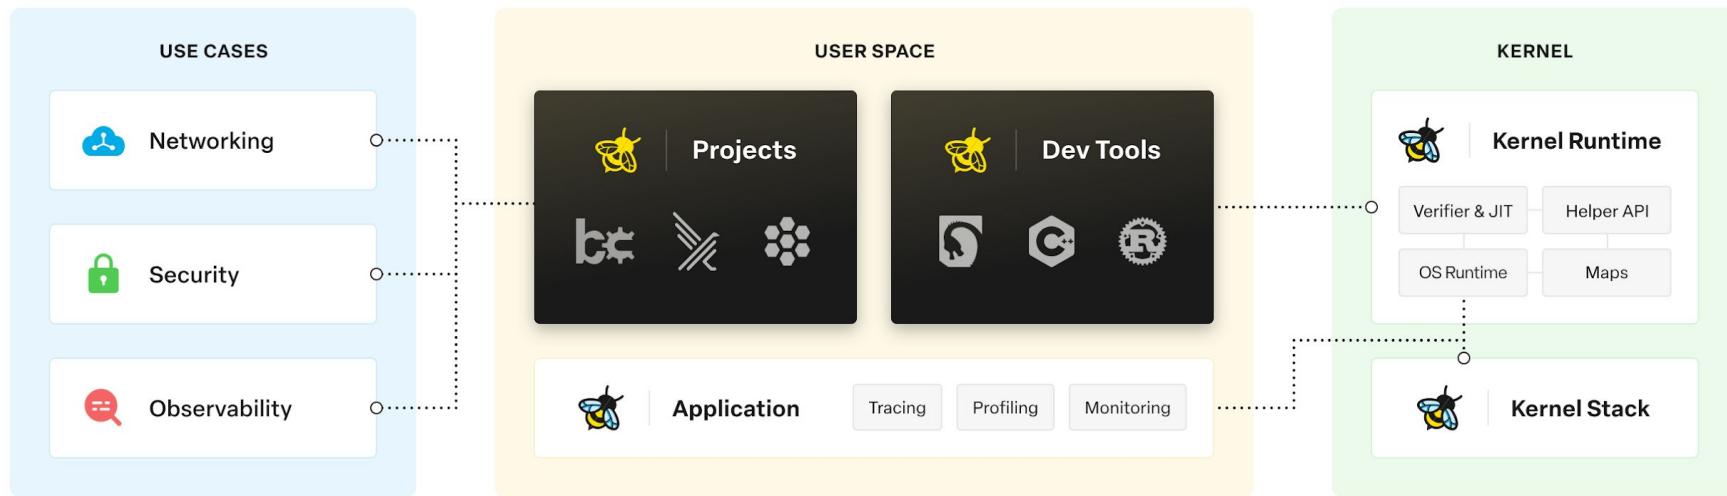
How Wasm improves eBPF developer experience:

- Non-intrusive deployment into k8s pods
- Decoupled from application workloads
- Declarative security checks at deployment time
- Supports user-space eBPF
- Downstream analytics of eBPF data

How eBPF improves Wasm developer experience:

- Better security sandboxing for WASI
- Better observability for host calls
- Better debugging

eBPF (Extended Berkeley Packet Filter): Dynamically and safely program the Linux kernel for efficient networking, observability, tracing, and security



WebAssembly (*Wasm*):

- A binary format for user space security sandbox
- Support multiple programming languages
- Very lightweight and fast
- Access to host resources through capability-based security declarations
- Cross platform portability



WasmEdgeRuntime

<https://github.com/WasmEdge/WasmEdge>




KubeCon



CloudNativeCon

North America 2023



How Wasm improves eBPF DevEx

eBPF Deployment models



KubeCon



CloudNativeCon

North America 2023

- Integrated control plane: bundled in the workload app
cilium, pixie, tetragon, falco...



- Decoupled / sidecar: a eBPF daemon
bumblebee, inspektor-gadget, bpf-d...



bpf-d



Integrated eBPF deployment

Closely integrate model and deploy with containers in a k8s cluster

- Embedding in the agent to release the eBPF probe
- Allows direct control and efficient management

Challenges

- Security and Permissions: CAP_BPF
- Heavy weight for small eBPF programs
- More than one eBPF User: potential for conflicts in hooks

Integrated eBPF deployment



KubeCon



CloudNativeCon

North America 2023

Closely integrate model and deploy with containers in a k8s cluster

Customs
Protocols

HTTP/2

SSL/TLS

Challenges

- New eBPF probes needs to be released with the agent
- Hard to trace private Protocols or Application logic with only pre-defined probes
- Hard to perform user-defined complex data processing

Decoupled sidecar eBPF deployment



KubeCon



CloudNativeCon

North America 2023

Daemon-Based Approach

- BPF daemon manages lifecycle and permissions via RPCs
- Decouples BPF functionality for modularity

Challenges

- Support and Maintenance
- Critical Component Risk
- Consistency in Updates
- Kernel Feature Integration

Wasm + eBPF: best of both worlds

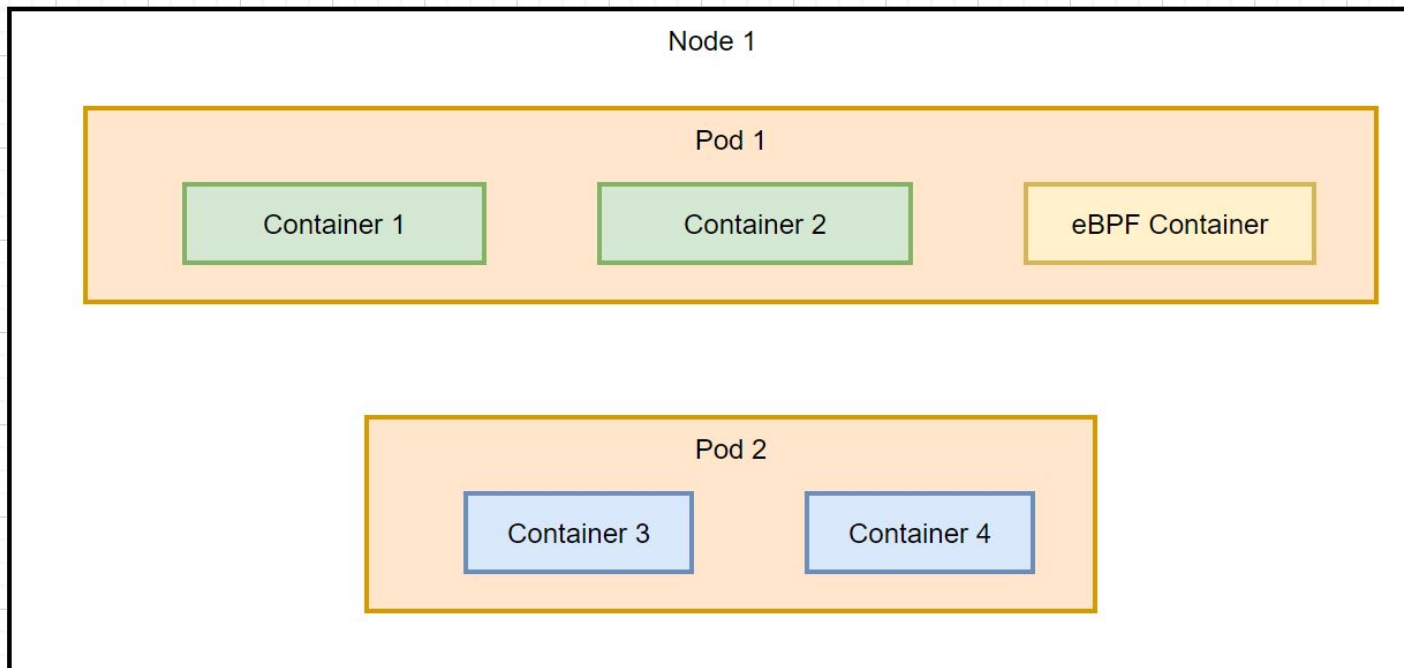


KubeCon



CloudNativeCon

North America 2023



Wasm + eBPF: best of both worlds

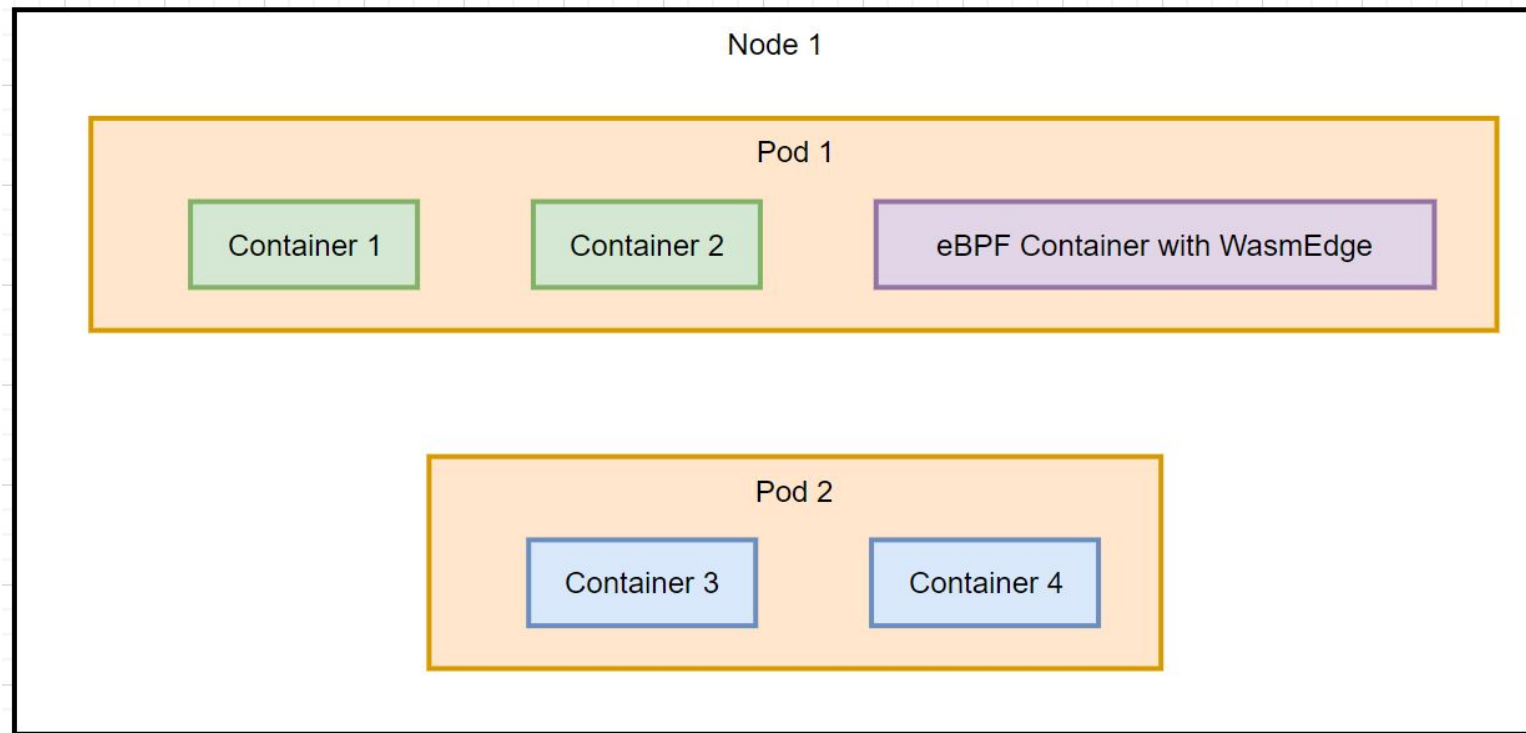


KubeCon



CloudNativeCon

North America 2023



WasmEdge eBPF plugin: wasm-bpf



KubeCon



CloudNativeCon

North America 2023

- Library and tool chains:
<https://github.com/eunomia-bpf/wasm-bpf>
- Runtime support:
https://github.com/WasmEdge/WasmEdge/tree/master/plugins/wasm_bpf

Faster, easier & safer eBPF deployment



KubeCon



CloudNativeCon

North America 2023

- [Faster] WasmEdge app is only 1/100 the size of containers
- [Easier] Portable across platforms and kernel versions with CO-RE
 - with user space eBPF runtime, no kernel eBPF support needed
- [Easier] Manage eBPF packages as OCI images in Wasm module
- [Safer] Configurable and limited eBPF WASI behavior
- [Safer] Improved security with access control to lower kernel resources
- [Safer] Sandbox the user space and enable eBPF plugin in tools



Too good to be true? There is no free lunch.

- Libraries and toolchains need to be ported,
 - e.g. libbpf and libbpf-rs -> libbpf-wasm
- eBPF features are limited in Wasm
 - e.g., sockets

Use container tools to run Wasm + eBPF



KubeCon



CloudNativeCon

North America 2023

The screenshot shows a terminal window with a dark background. The prompt is `root@b6768d3348a3:/WasmEdge#` and the command `clear` is entered. On the right side, a list of recent commands is visible: `bash yunwei`, `bash crun`, `bash WasmE...`, `bash crun`, `docker was...`, and `docker crun`. The `docker was...` command is highlighted. At the bottom of the terminal, there is a status bar with various icons and text: `192.168.174.128`, `master*`, `01 21`, `3 3 2`, `5`, `You, 2分钟前`, `行 145, 列 1`, `空格: 4`, `UTF-8`, `LF`, `Markdown`, and some icons.

Use podman to deploy eBPF with WebAssembly



Developer experience



KubeCon



CloudNativeCon

North America 2023

Similar developing experience as the libbpf-bootstrap

```
/* Load and verify BPF application */
skel = bootstrap_bpf__open();
if (!skel) {
    fprintf(stderr, "Failed to open and load BPF skeleton\n");
    return 1;
}

/* Parameterize BPF code with minimum duration parameter */
skel->rodata->min_duration_ns = env.min_duration_ms * 1000000ULL;

/* Load & verify BPF programs */
err = bootstrap_bpf__load(skel);
if (err) {
    fprintf(stderr, "Failed to load and verify BPF skeleton\n");
    goto cleanup;
}

/* Attach tracepoints */
err = bootstrap_bpf__attach(skel);
if (err) {
    fprintf(stderr, "Failed to attach BPF skeleton\n");
    goto cleanup;
}

/* Set up ring buffer polling */
rb = bpf_buffer__open(skel->maps.rb, handle_event, NULL);
if (!rb) {
    err = -1;
    fprintf(stderr, "Failed to create ring buffer\n");
    goto cleanup;
}
```

```
/* SPDX-License-Identifier: (LGPL-2.1 OR BSD-2-Clause) */

/* THIS FILE IS AUTOGENERATED! */
#ifndef __MINIMAL_BPF_SKEL_H__
#define __MINIMAL_BPF_SKEL_H__

#include <stdlib.h>
#include <bpf/libbpf.h>

struct minimal_bpf {
    struct bpf_object_skeleton *skeleton;
    struct bpf_object *obj;
    struct {
        struct bpf_map *bss;
    } maps;
    struct {
        struct bpf_program *handle_tp;
    } progs;
    struct {
        struct bpf_link *handle_tp;
    } links;
    struct minimal_bpf__bss {
        int my_pid;
    } *bss;
};

static inline void minimal_bpf__destroy(struct minimal_bpf *obj) { ... }
.
static inline int minimal_bpf__attach(struct minimal_bpf *obj) { ... }
static inline void minimal_bpf__detach(struct minimal_bpf *obj) { ... }

#endif /* __MINIMAL_BPF_SKEL_H__ */
```

Examples



KubeCon



CloudNativeCon

North America 2023

Support eBPF use cases from Observability,
Networking to Security



```
#include <linux/bpf.h>
#include <linux/ptrace.h>
#include <bpf/bpf_helpers.h>
#include <bpf/bpf_tracing.h>

char LICENSE[] SEC("license") = "GPL";

SEC("uprobe././target:probe_add")
int BPF_KPROBE(probe_add, int a, int b) {
    bpf_printk("uprobe_add ENTRY: a = %d, b = %d", a, b);
    return 0;
}
```

Uprobe

```
#include "vmlinux.h"
#include <bpf/bpf_helpers.h>

SEC("xdp")
int xdp_pass(struct xdp_md* ctx) {
    void* data = (void*)(long)ctx->data;
    void* data_end = (void*)(long)ctx->data_end;
    int pkt_sz = data_end - data;

    bpf_printk("packet size: %d\n", pkt_sz);
    return XDP_PASS;
}

char __license[] SEC("license") = "GPL";
```

XDP

```
#include "vmlinux.h"
#include <bpf/bpf_helpers.h>

char _license[] SEC("license") = "GPL";

// all lsm the hook point refer
// https://www.kernel.org/doc/html/v5.2/security/LSM.html
SEC("lsm/path_rmdir")
int path_rmdir(const struct path* dir, struct dentry* dentry)
{
    char comm[16];
    bpf_get_current_comm(comm, sizeof(comm));

    unsigned char dir_name[] = "can_not_rm";
    unsigned char d_iname[32];
    bpf_probe_read_kernel(&d_iname[0], sizeof(d_iname),
        &(dir->dentry->d_iname[0]));

    bpf_printk("comm %s try to rmdir %s", comm, d_iname);
    for (int i = 0; i < sizeof(dir_name); i++) {
        if (d_iname[i] != dir_name[i]) {
            return 0;
        }
    }

    return -1;
}
```

LSM

<https://github.com/eunomia-bpf/wasm-bpf/blob/main/examples>

How it works

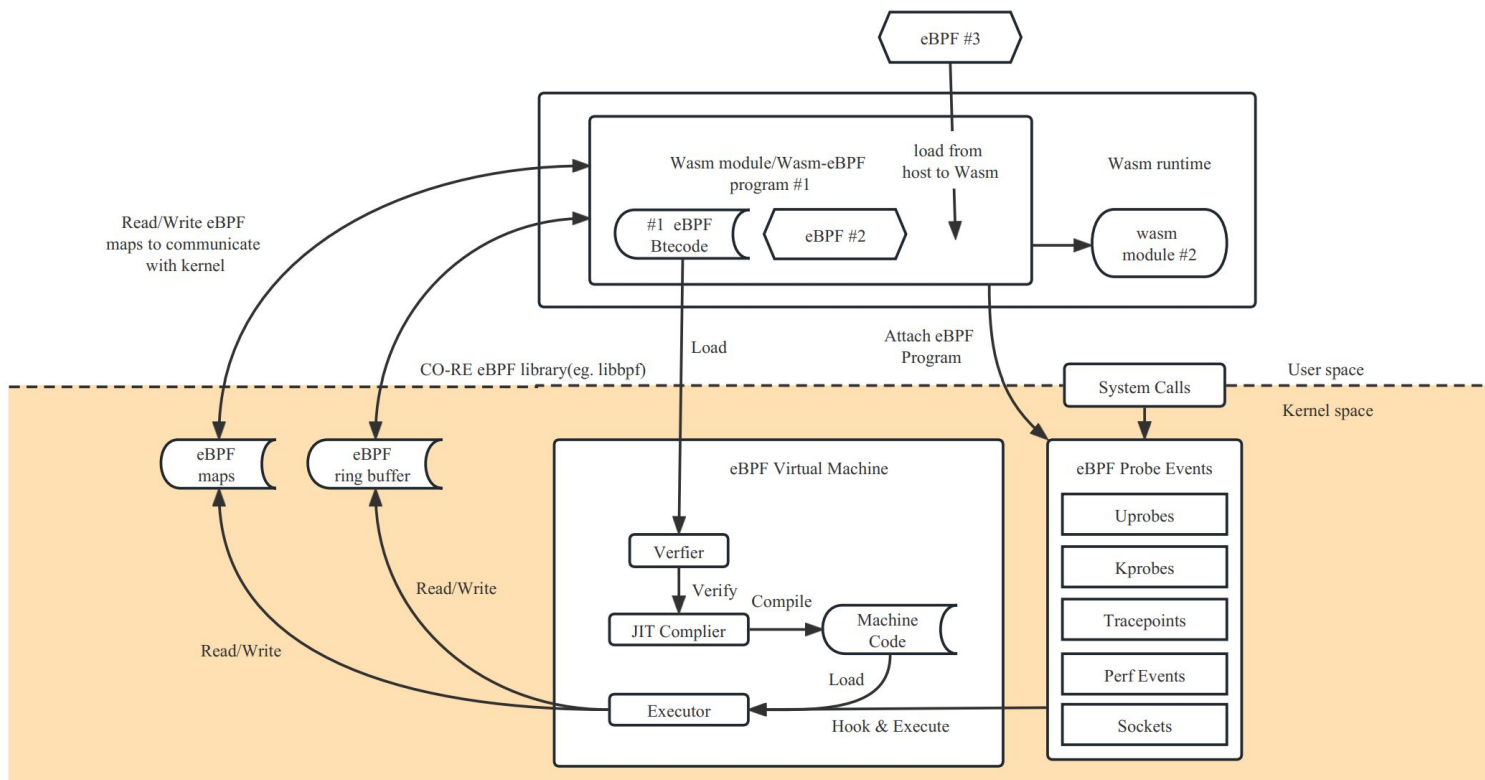


KubeCon



CloudNativeCon

North America 2023



How it works



KubeCon

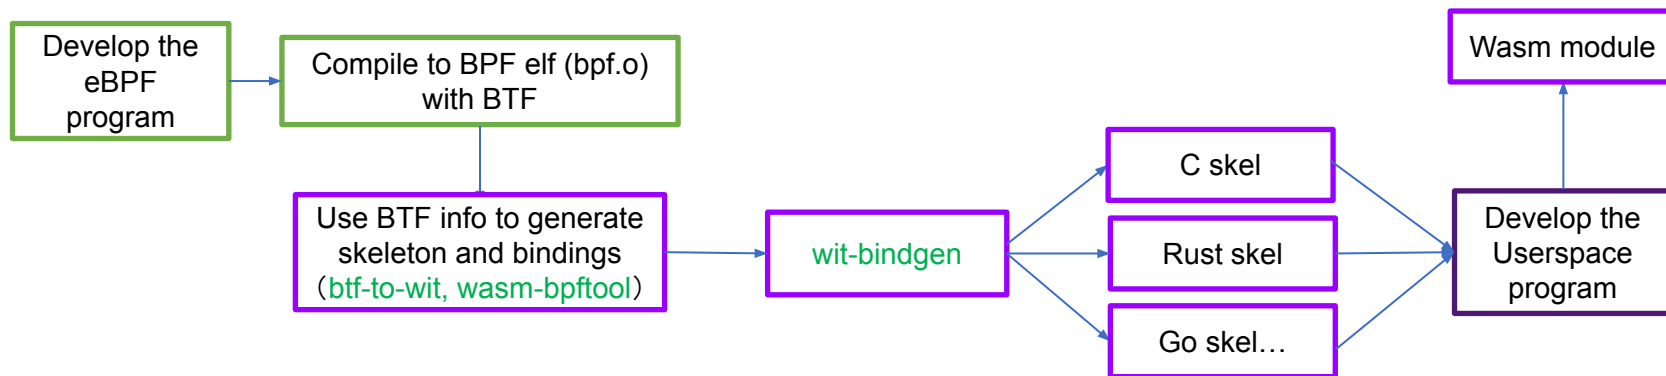


CloudNativeCon

North America 2023

The toolchains:

- Auto generate the Wasm-eBPF skeleton headers and type definitions for bindings
- No serialization overhead for complex data type
- Share memory maps between Wasm and kernel



Challenges



KubeCon



CloudNativeCon

North America 2023

- Port libraries and prepare toolchains to Wasm:
 - C/C++: libbpf
 - Rust: libbpf-rs
 - Go: cilium/go
- Wasm is 32bit and eBPF is 64 bit: data layout is different
 - Use toolchains to generated bindings and avoid serialization
- eBPF may require kernel version support:
 - Use CO-RE to enable portable between different kernel
 - Use user space eBPF runtime for optional
 - A compatible layer

Wasm with user space eBPF



KubeCon

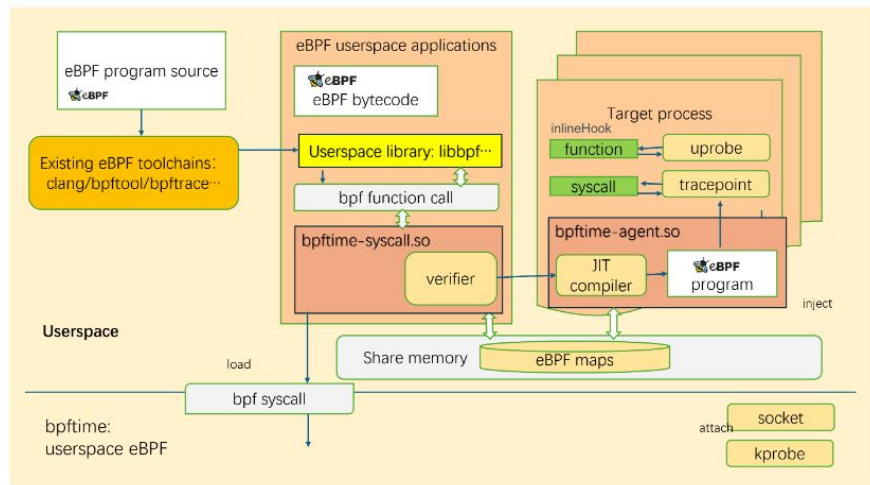


CloudNativeCon

North America 2023

Wasm with kernel eBPF enable more possibilities, but require kernel version support and privilege.

- ❑ Completely Userspace eBPF runtime
- ❑ Uprobe and Syscall hooks, Interprocess maps
- ❑ No manual integration or restart required
- ❑ Compatible with kernel eBPF toolchains
- ❑ Can run bcc tools and bpfftrace in userspace
- ❑ Speed up 10x than kernel uprobe



<https://github.com/eunomia-bpf/bpftime>




KubeCon



CloudNativeCon

North America 2023



How eBPF improves Wasm DevEx

WASI and Debugging needs



KubeCon



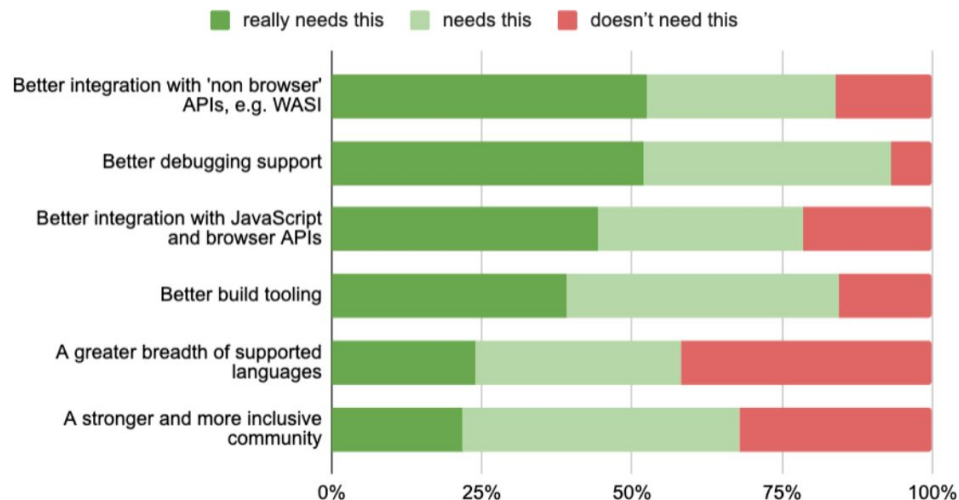
CloudNativeCon

North America 2023

WebAssembly needs:

- A better WASI (WebAssembly System Interface)
- A better debugging toolchain

What WebAssembly needs (n=303)



eBPF makes WASI more secure



KubeCon



CloudNativeCon

North America 2023

WASI checks if an accessed path is on an "allowed" list before granting access:

- Non-trivial task to implement this works
- Error prone: reliance on code-review
- Less configurable: directory-level granularity, not file-level.

Solution:

- Use eBPF LSM and seccomp to provide programmable access-control in Linux kernel for Wasi

Example: hook in dir remove and check the permission to remove a spec directory

```
SEC("lsm/path_rmdir")
int path_rmdir(const struct path* dir, struct dentry* dentry) {
    char comm[16];
    bpf_get_current_comm(comm, sizeof(comm));

    unsigned char dir_name[] = "can_not_rm";
    unsigned char d_iname[32];
    bpf_probe_read_kernel(&d_iname[0], sizeof(d_iname),
                          &(dir->dentry->d_iname[0]));

    bpf_printk("comm %s try to rmdir %s", comm, d_iname);
    for (int i = 0; i < sizeof(dir_name); i++) {
        if (d_iname[i] != dir_name[i]) {
            return 0;
        }
    }

    return -1;
}
return go(f, seed, [])
}
```

<https://github.com/eunomia-bpf/wasm-bpf/tree/main/examples/lsm>

Debug tracing for Wasm runtime



KubeCon



CloudNativeCon

North America 2023

Enabling Detailed Tracing of WebAssembly with eBPF

- Wasm provides limited tracing methods
- eBPF's Uprobe can trace any user-space function
- No additional instrumentation needed.
- userspace eBPF runtime like bpftrace, enable fast Uprobe without kernel support requirement and without root privilege.

Examples: trace memory allocation in WasmEdge

```
#include <vmlinux.h>
#include <bpf/bpf_helpers.h>
#include <bpf/bpf_core_read.h>
#include "allocnoop.h"

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
    __uint(key_size, sizeof(u32));
    __uint(value_size, sizeof(u32));
} events SEC(".maps");

SEC("uprobe/wasmedge:allocate")
int BPF_KPROBE(trace_alloc, uint32_t PageCount)
{
    u64 id;
    struct event event={0};

    uid_t uid = (u32)bpf_get_current_uid_gid();
    id = bpf_get_current_pid_tgid();

    event.pid = id >> 32;
    bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    return 0;
}

char LICENSE[] SEC("license") = "GPL";
```

Take away



KubeCon



CloudNativeCon

North America 2023

1. A “Wasm container” can deploy eBPF on hosts of k8s pods — it has the benefits of both the tight integration and RPC approaches.
2. Wasm can run user space eBPF apps for tasks like data analytics.
3. eBPF could become part of a WASI implementation on Linux to enhance security and support debugging.




KubeCon



CloudNativeCon

———— North America 2023 ————

Thanks





PromCon
North America 2021

Session QR Codes will be
sent via email before the event

**Please scan the QR Code above
to leave feedback on this session**

Agenda



KubeCon



CloudNativeCon

North America 2023

Introduction to eBPF and Wasm

How can WebAssembly enhance eBPF:

- Current eBPF deployment model and challenges
- What Wasm can bring

How can eBPF enhance WebAssembly:

- Extend the WASI
- Enable better observability for Wasm Runtime

Compare eBPF and Wasm



KubeCon



CloudNativeCon

North America 2023

WebAssembly (Wasm):

- Sandboxed
- Use run-time checks for any arbitrary code
- performance overhead
- Better language support and ecosystem

Extended Berkeley Packet Filter (eBPF):

- Static Verification
- minimize run-time checks, known before execution
- Performance-Oriented
- Small C programs

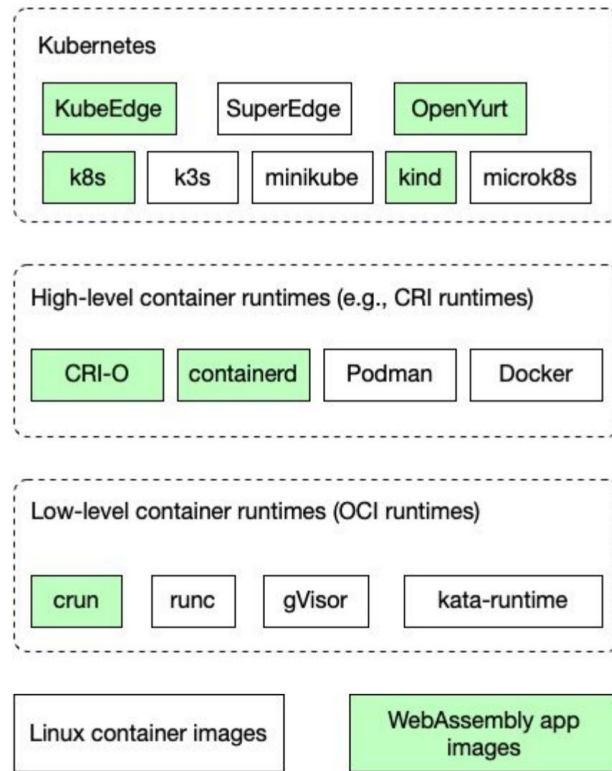
WasmEdge: A lightweight, secure, high-performance and extensible WebAssembly Runtime

- Integrated support for WasmEdge in Docker CLI
- Transparently supports containers and Wasm apps from the same tools and in the same cluster

<https://github.com/WasmEdge/WasmEdge>



WasmEdgeRuntime



The container ecosystem

eBPF Deployment Challenges



KubeCon



CloudNativeCon

North America 2023

In summary:

- Security Risks with Privileged Access
- Compatibility Issues with eBPF Hooks and Kernel Versions
- Complex Lifecycle and Deployment Management
- Challenges with Versioning and Plugability

What Wasm can bring



KubeCon



CloudNativeCon

North America 2023

WebAssembly's possible Role in Enhancing eBPF Deployments:

- Lightweight and portable
- Fine-Grained Permissions
- Help with Isolation
- Lifecycle Management
- Versioning and Update with plugin

WasmEdge eBPF plugin: wasm-bpf



KubeCon



CloudNativeCon

North America 2023

WebAssembly library, toolchain and runtime for eBPF programs

- Create Wasm-based eBPF control plane applications
- Enable managing eBPF programs in k8s pods with lightweight Wasm container
- Enables Wasm plugin in eBPF core applications

https://github.com/WasmEdge/WasmEdge/tree/master/plugins/wasm_bpf

<https://github.com/eunomia-bpf/wasm-bpf>

Make eBPF deployment safer



KubeCon



CloudNativeCon

North America 2023

Runtime optimized for eBPF in Wasm lightweight container

- Configurable and limited eBPF WASI behavior
- Improved security with access control to lower kernel resources
- Sandbox the user space and enable eBPF plugin in tools

Generate eBPF with natural language and run in WebAssembly:

- The AI generated code cannot be trusted.
- The generated eBPF program needs a isolation runtime to deploy.
- WebAssembly can be lightweight and cheaper than Docker or vm.

```
yunwei37@ebpf-plctlab:~/GPTtrace$ ./GPTtrace.py -e "Count page faults by process"
Sending query to ChatGPT: Count page faults by process
Press Ctrl+C to stop the program...
Attaching 1 probe...
^C

@[sudo]: 2
@[bash]: 60
@[tail]: 112
@[sleep]: 421
@[sshd]: 566
@[which]: 854
@[ps]: 989
@[sed]: 1380
@[cat]: 1685
@[python]: 2206
@[zsh]: 2461
@[sh]: 2464
@[git]: 2747
@[cpuUsage.sh]: 6788
@[node]: 36701

yunwei37@ebpf-plctlab:~/GPTtrace$
```

With our agent and GPT4, can have 80% rate to generate simple eBPF program successfully

<https://github.com/eunomia-bpf/GPTtrace>