



KubeCon



CloudNativeCon

North America 2023

GoTo Financial's Story: Towards 10k+ ArgoCD Apps to Support Billions of \$ Transactions

Giri Kuncoro & Yudi Phanama

goto financial



 @girikuncoro



 @yphanama

intro



KubeCon
North America 2023



CloudNativeCon
North America 2023

goto financial

Part of **goto** group,
the leading Indonesia's
digital ecosystem



ArgoCD Overview



@

goto financial



kubectl describe cluster



KubeCon



CloudNativeCon

North America 2023

- Distributed across Singapore and Indonesia region
 - **~50 Kubernetes clusters** in AWS, GCP, and Private Datacenter
 - **700+** compute nodes
 - **15,000+ CPU**
 - **120+ TB** memory
 - **30,000+** pods



argocd snapshot



KubeCon



CloudNativeCon

North America 2023

Overview



Uptime

1 day

Clusters

44

Applications

11167


Repositories

5979

Operations



Applications



argocd 11.2 K


Application Status (2 panels)

Sync Stats

Sync Activity



Reconciliation Performance



kubectl describe argocd



KubeCon



CloudNativeCon

North America 2023

- **11,000+** applications
- **6,000+** repositories
- **~60** projects
- **380,000+** total objects

on largest cluster

- **2,000** applications
- **40,000+** objects



centralized argocd (push model)

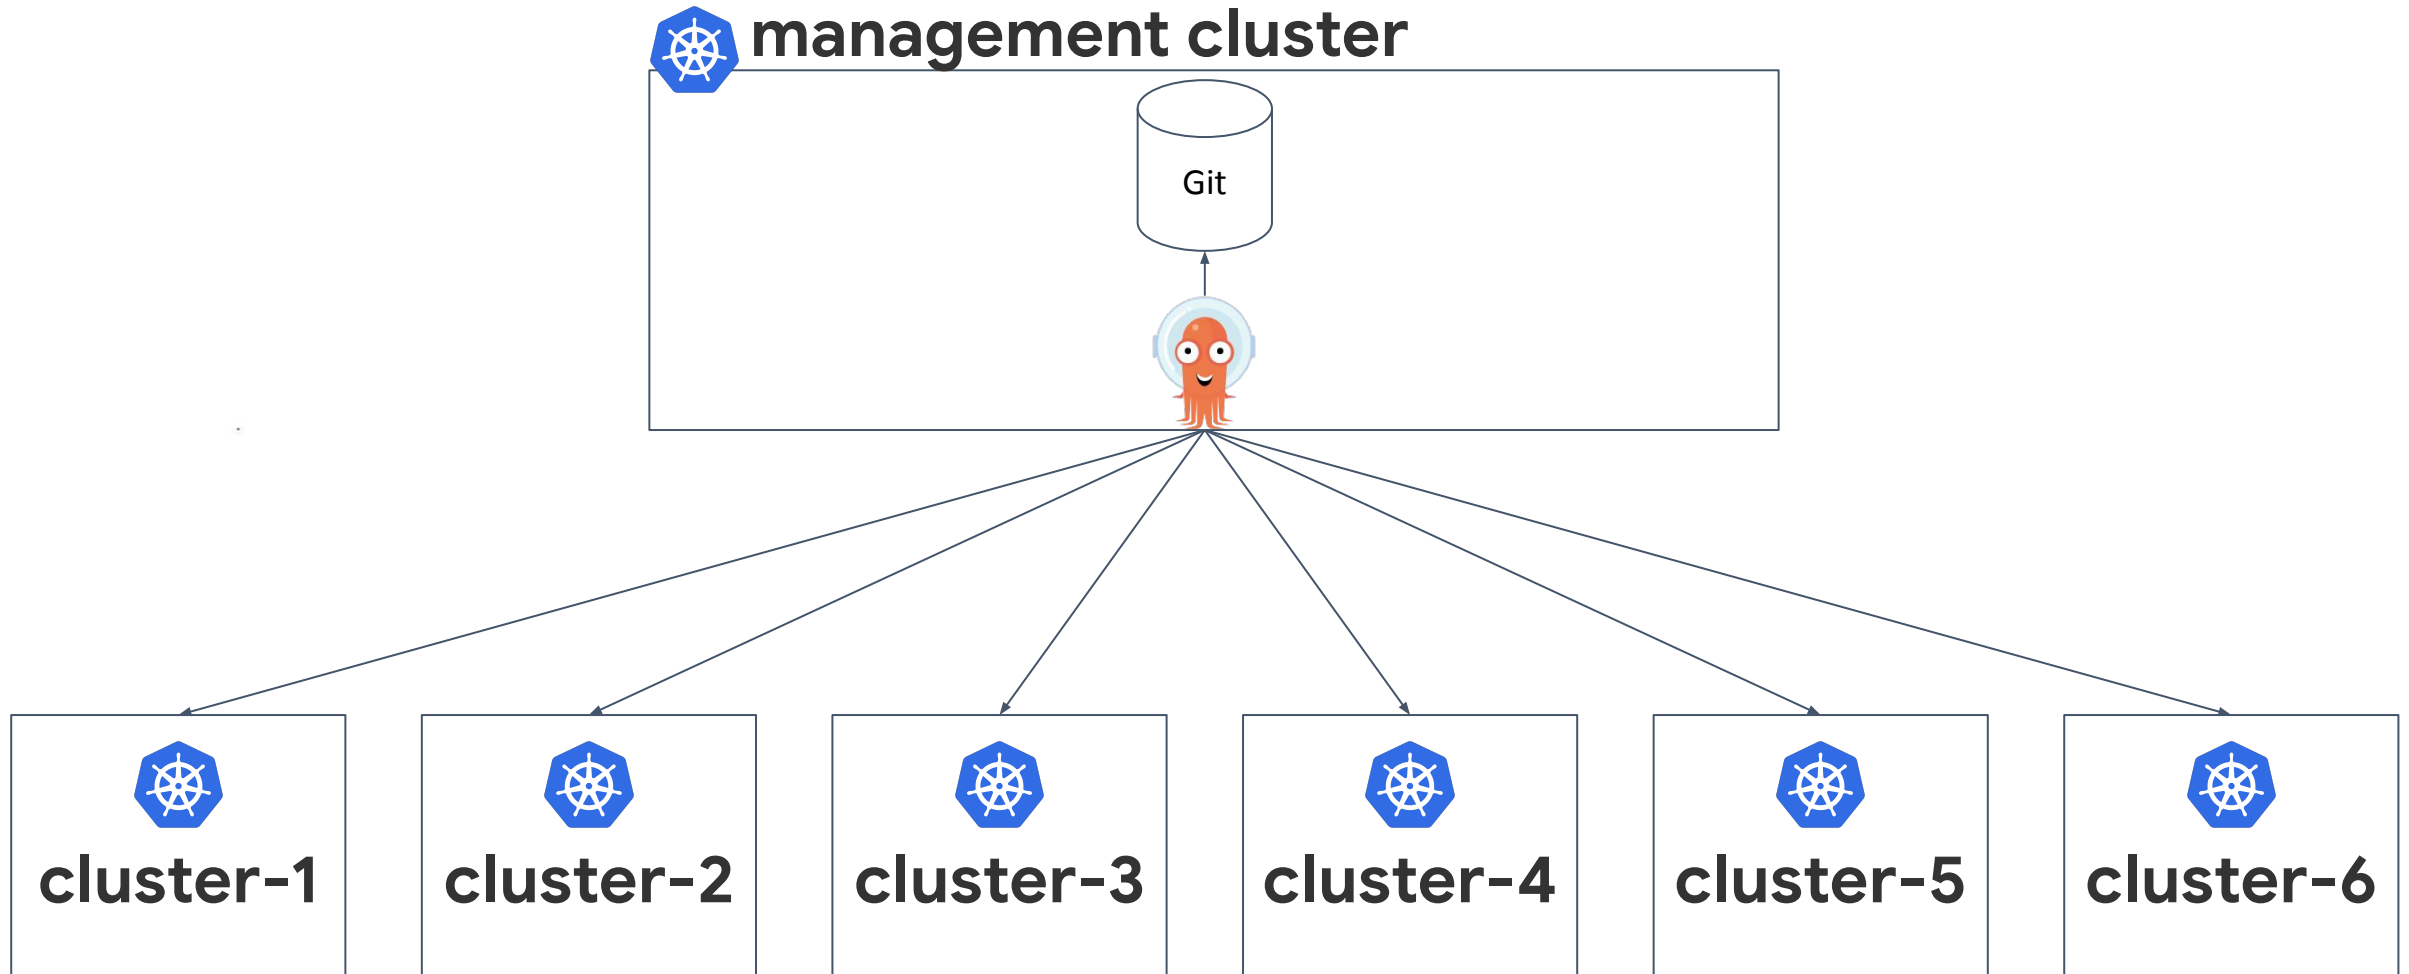


KubeCon



CloudNativeCon

North America 2023



centralized argocd



KubeCon



CloudNativeCon

North America 2023

Pros

- Easy to maintain and upgrade
- Easy to integrate with our automation and platform
- Easy to manage centralized RBAC
- Single dashboard to view and control all clusters

platform integration

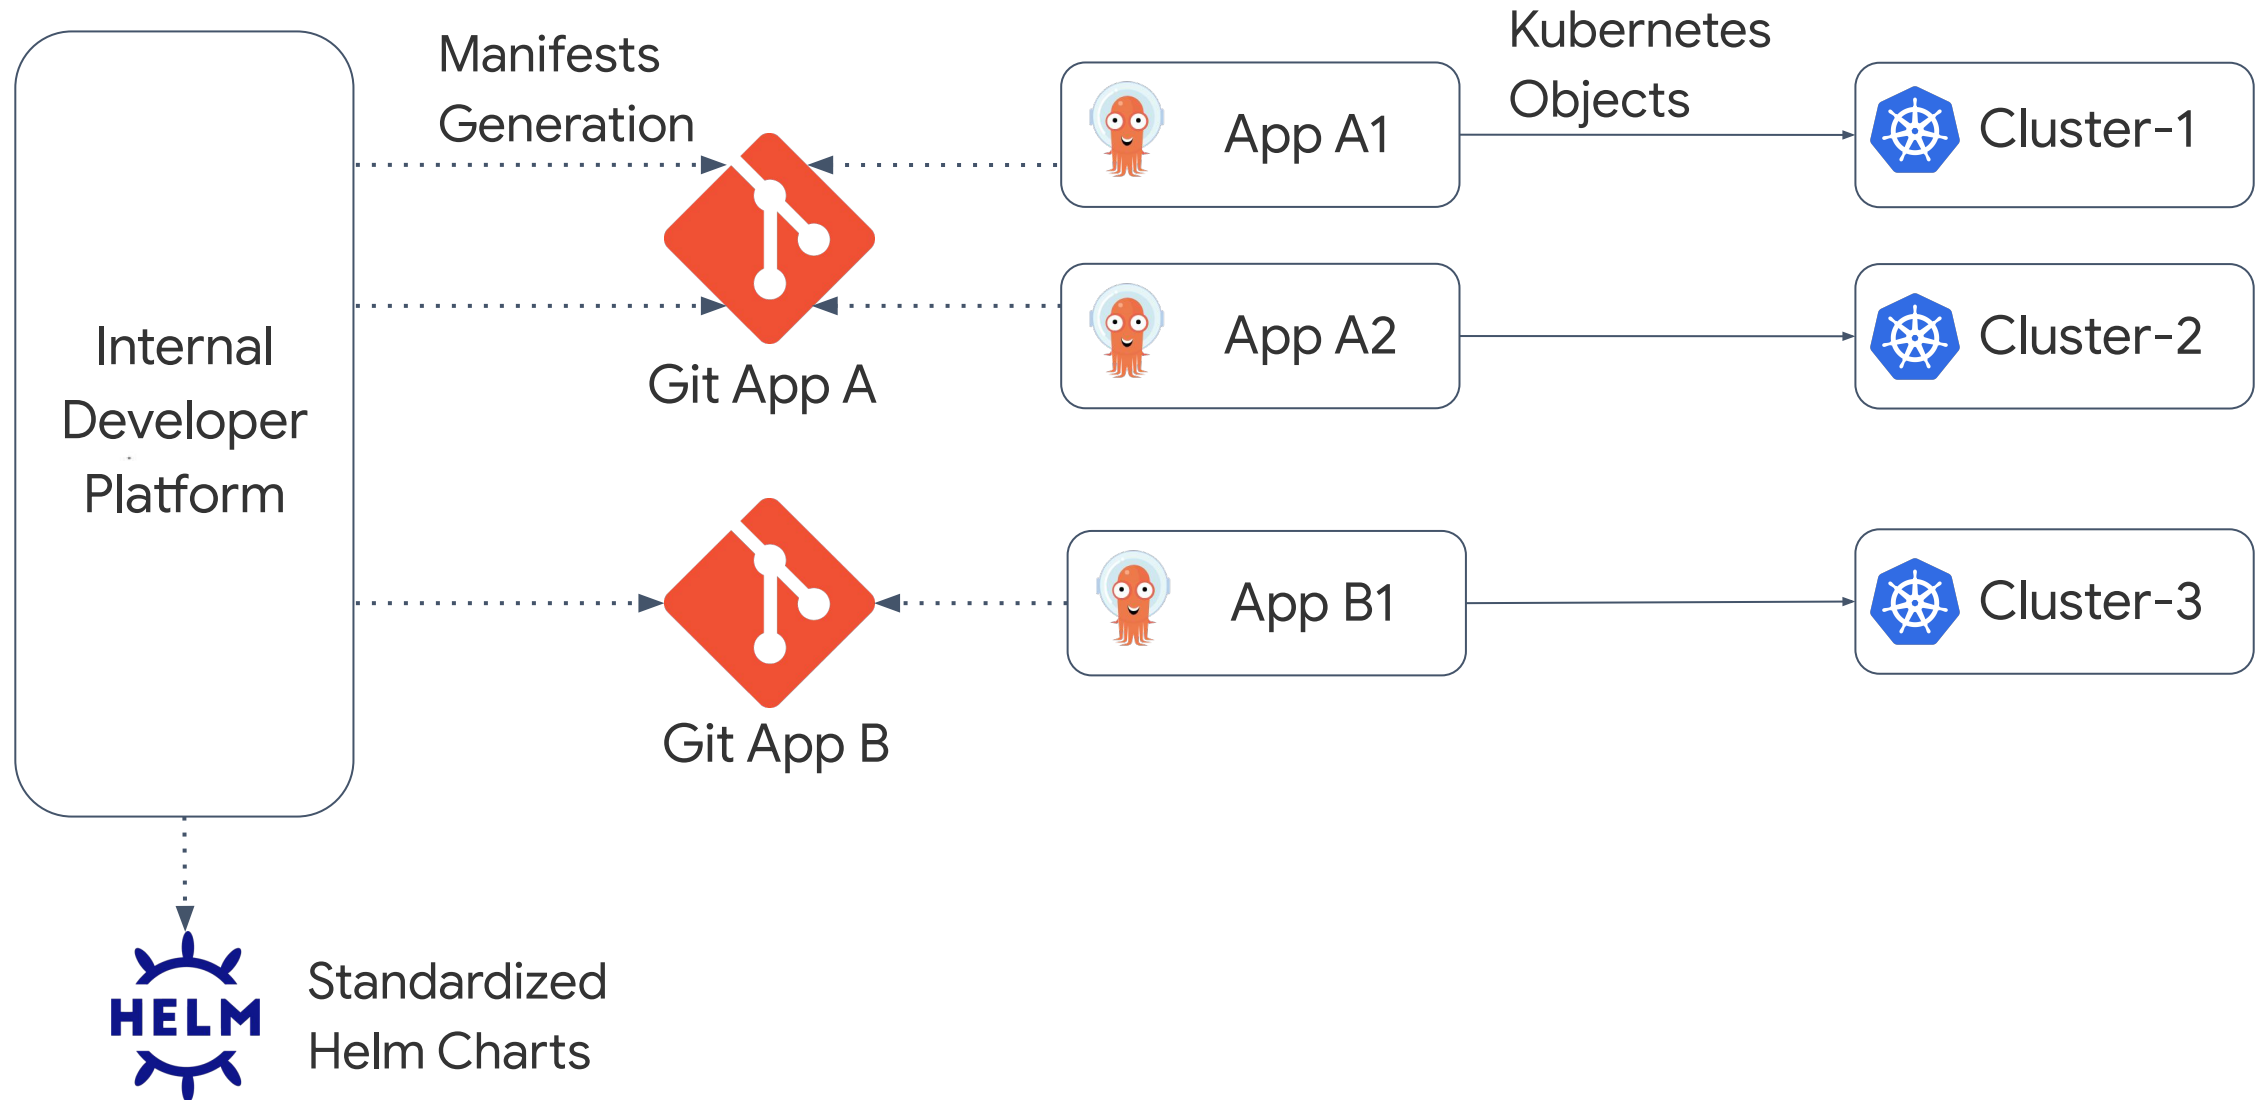


KubeCon



CloudNativeCon

North America 2023



service lifecycle



KubeCon



CloudNativeCon

North America 2023

1 service contains 3-5 ArgoCD apps with different lifecycle



cluster runtime components



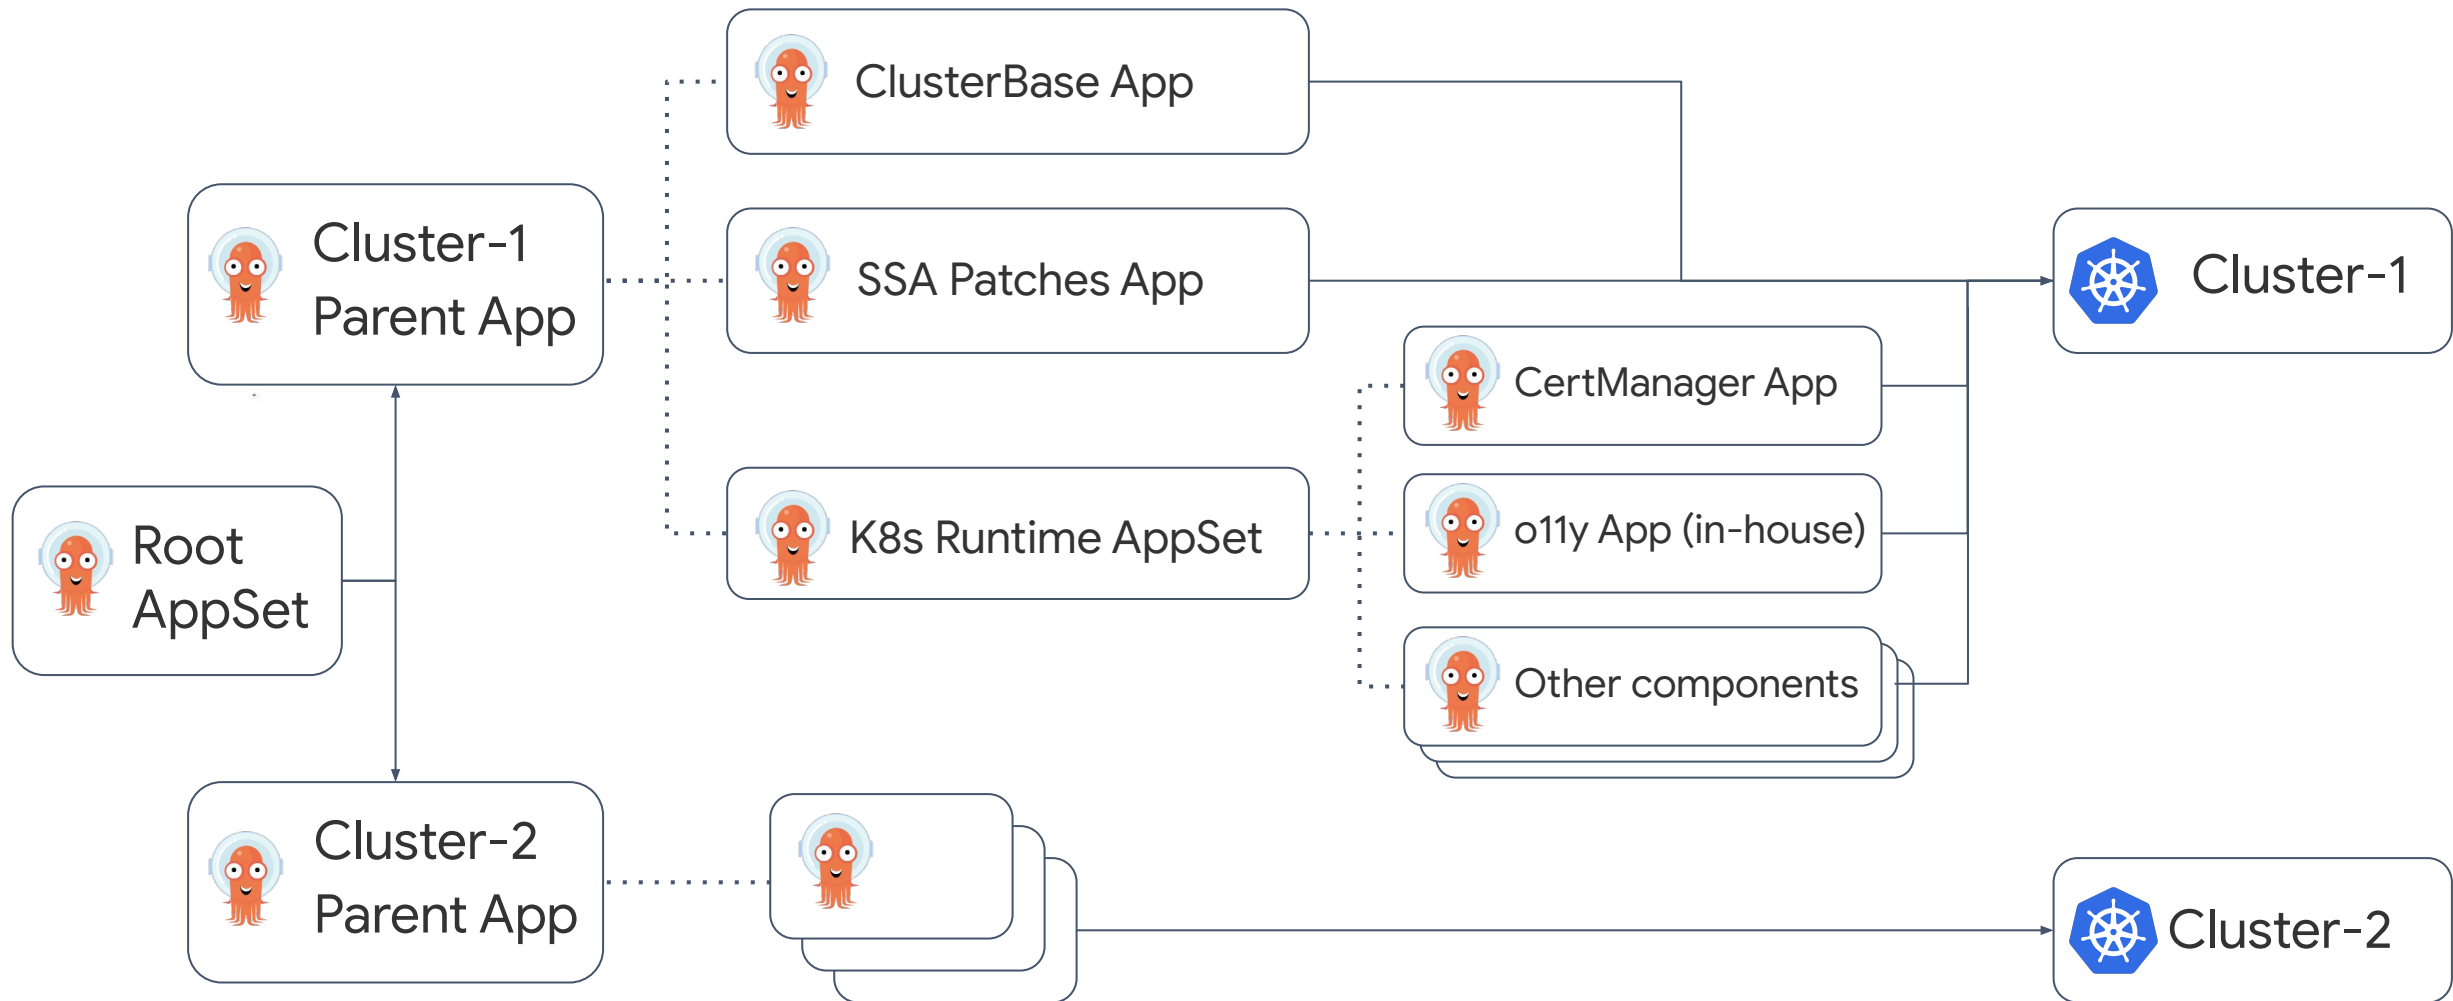
KubeCon



CloudNativeCon

North America 2023

appset & app of apps pattern - monorepo



argocd on argocd

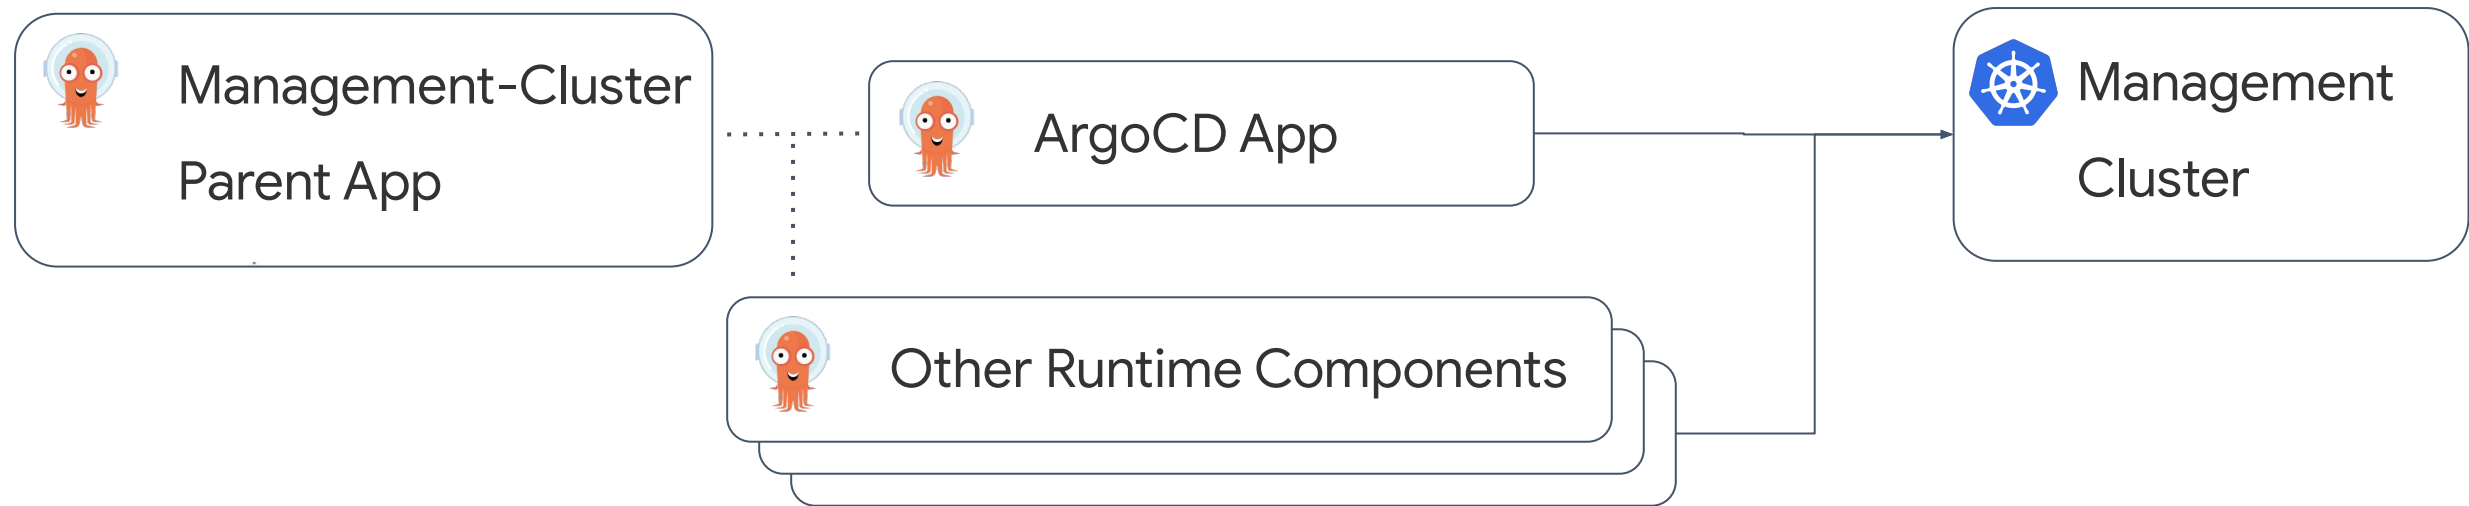


KubeCon



CloudNativeCon

North America 2023



challenges of centralized argocd



KubeCon



CloudNativeCon

North America 2023

Connectivity to all target clusters

- Tunnels / peering
- Public mTLS

Functionalities

- Must maintain unique application name globally (63 chars constraint)
- Single point of failure

challenges of centralized argocd



KubeCon



CloudNativeCon

North America 2023

Performance issues

- Slow reconciliation & sync
 - `workqueue_depth`
 - `argocd_app_reconcile_bucket`
- Slow UI loading (> 1 min)
 - very obvious + browser inspect
- Repo server OOM kills
 - kube events

challenges of centralized argocd



KubeCon



CloudNativeCon

North America 2023

- High rate of Git API calls (both ls-remote and fetch)
 - `argocd_git_request_total`
- High repo cache miss
 - `repo-server logs`
- Imbalanced shards & noisy cluster
 - `controller process_cpu_seconds_total`

Performance Tuning Journey



argocd components

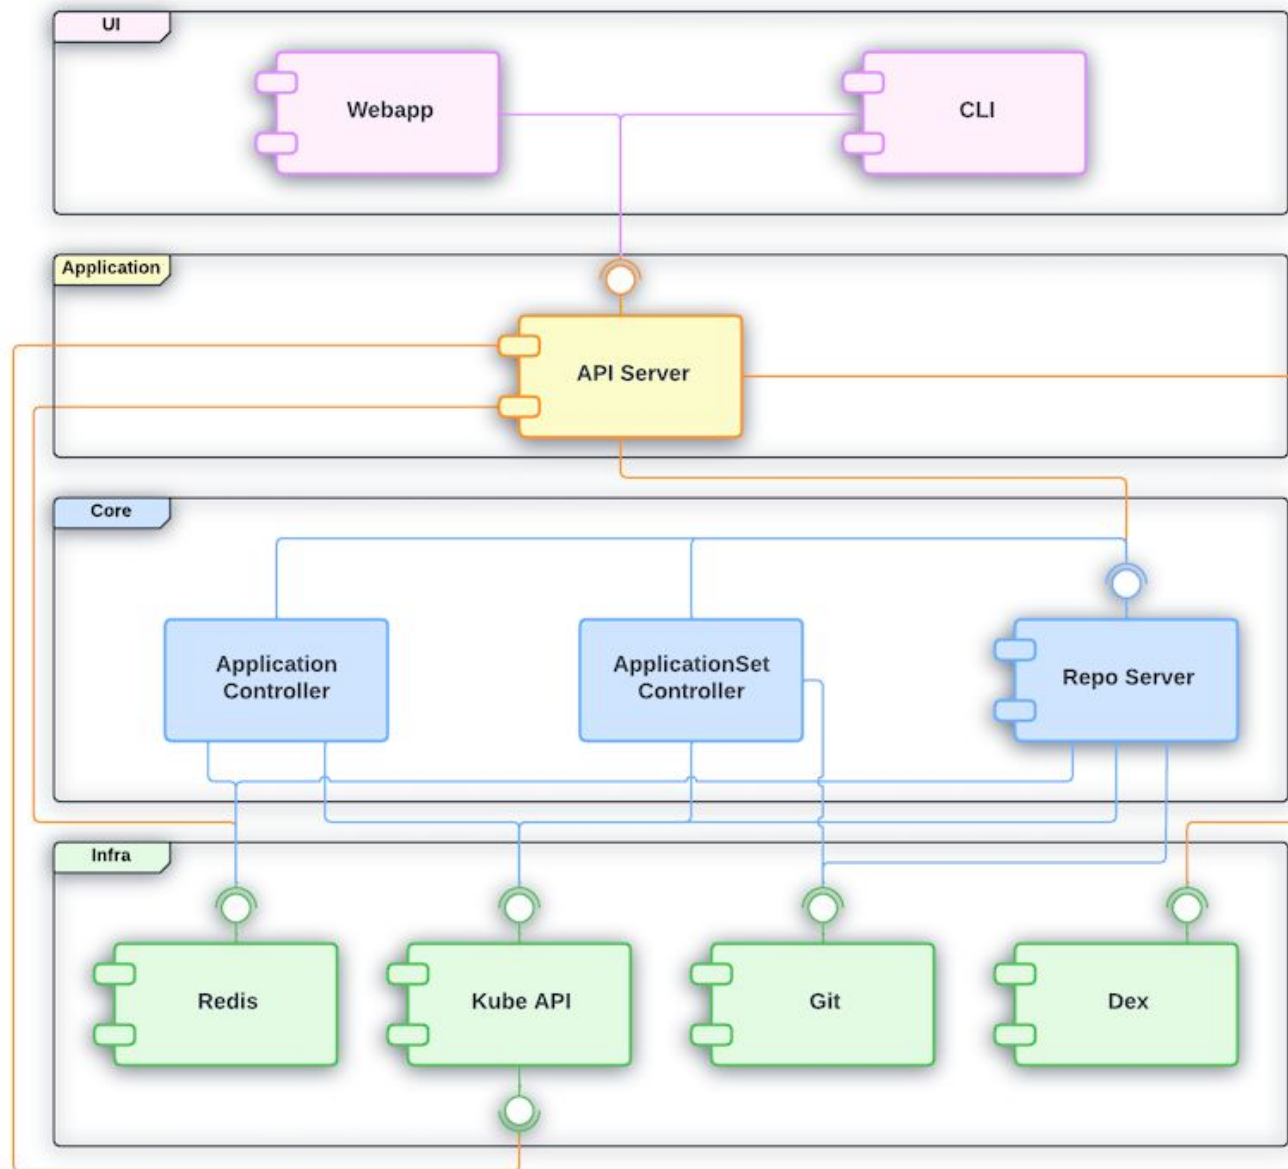


KubeCon



CloudNativeCon

North America 2023



Source:
<https://argo-cd.readthedocs.io/en/stable/developer-guide/architecture/components/>

tuning: argocd-server



KubeCon



CloudNativeCon

North America 2023

Problem: Slow UI load

🔧 enable gzip compression

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: "argocd-server"
spec:
  containers:
    - name: argocd-server
      env:
        - name: "ARGOCD_SERVER_ENABLE_GZIP"
          value: "true"
```

before:

<input type="checkbox"/> applications?fields=metadata.resou...	200	xhr	main.e135b04....js:2	37.8 MB	18.90 s	<div><div></div></div>
--	-----	-----	----------------------	---------	---------	------------------------

after (~5x faster load + ~7x smaller data)

<input type="checkbox"/> applications?fields=metadata.resou...	200	xhr	main.b6d0ded....js:2	5.1 MB	3.51 s	<div><div></div></div>
--	-----	-----	----------------------	--------	--------	------------------------

tuning: argocd-server



KubeCon



CloudNativeCon

North America 2023

argocd ui tip - use selectors (labels, projects, namespaces)

🍪 The last selectors are saved the next time we load the ArgoCD UI

The screenshot displays the ArgoCD web interface. At the top, a status bar shows the URL 'applications?projects...', 200 items, 'xhr' method, 'main.dac...' path, 655 kB size, and 570 ms response time. The left sidebar contains navigation links: Applications, Settings, User Info, and Documentation. Below these are filters for 'Favorites Only', 'SYNC STATUS', 'HEALTH STATUS', 'LABELS', 'PROJECTS' (with a 'CLEAR' button), 'CLUSTERS', 'NAMESPACES' (with a 'CLEAR' button), and 'AUTO SYNC'. The main content area shows a grid of application cards. Each card displays the project name, labels, status (Healthy and Synced), repository, target, path, destination, namespace, creation time, and last sync time. The 'internal-tooling' project is selected in the filters. The bottom of the interface shows pagination controls and sorting options.

tuning: k8s cpu limits



KubeCon




CloudNativeCon

North America 2023

Problem: CPU gets throttled across all components

🔧 remove k8s CPU limit, use k8s requests only

- k8s request & limit uses cgroup, which uses CFS [\[1\]](#)
- CFS guarantees or throttle CPU proportional to container shares in a Node
- ref [\[2\]](#)[\[3\]](#)



tuning: repo-server



KubeCon



CloudNativeCon

North America 2023

Problem: repo-server got OOMkilled frequently

🔧 Increase replicas and use HPA

```
apiVersion: apps/v1
kind: HorizontalPodAutoscaler
metadata:
  name: "argocd-repo-server"
spec:
  ...
  maxReplicas: 15
  minReplicas: 5
  metrics:
    - resource:
        name: memory
        target:
          averageUtilization: 75
        type: Resource
    - resource:
        name: cpu
        target:
          averageUtilization: 75
        type: Resource
  ...
```

- Automatically scales with memory usage
- Distributes requests to more pods
- Alternatively, you could use the `--parallelism-limit` flag to control how many manifest generation requests that can be served in parallel and help avoid OOM kills.

tuning: repo-server



KubeCon



CloudNativeCon

North America 2023

Problem: repo-server timeout errors from our logs

🔧 increase repo server client timeout

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: "argocd-server"
spec:
  ...
  containers:
    - name: argocd-server
      env:
        - name: "ARGOCD_SERVER_REPO_SERVER_TIMEOUT_SECONDS"
          value: "120" #default is 60
  ...

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: "argocd-application-controller"
spec:
  ...
  containers:
    - name: argocd-application-controller
      env:
        - name: "ARGOCD_APPLICATION_CONTROLLER_REPO_SERVER_TIMEOUT_SECONDS"
          value: "120" #default is 60
  ...
```

- argocd-server and app-controller talks to the repo-server for manifests generation.
- We started seeing these timeout errors from them when syncing or refreshing apps.
- Increase it in **both** the argocd-server and argocd-application-controller.

tuning: repo-server



KubeCon



CloudNativeCon

North America 2023

Problem: persistently high git fetch requests

🔧 extend repo cache expiration

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: "argocd-repo-server"
spec:
  ...
  containers:
    - name: argocd-repo-server
      env:
        - name: "ARGOCD_REPO_CACHE_EXPIRATION"
          value: "24h"
      ...
```



- ArgoCD caches generated manifests (24h).
- When **remote files change** often **even though** the repository **tag hasn't changed**, shorter expiry is desirable to pick up the updates.
- We can use higher expiry if the Helm/Kustomize/Git remote refs are already hermetic.
- Set ARGOCD_REPO_CACHE_EXPIRATION to extend repo cache expiration.

tuning: monorepo usage



KubeCon



CloudNativeCon

North America 2023

Problem: Multi Sources apps caused high git requests

```
sources:
- directory:
  exclude: "*"
  path: "my-app/"
  repoURL: https://REPO_URL
  targetRevision: HEAD
  ref: values
- chart: CHART_NAME
  repoURL: CHART_URL
  targetRevision: v0.0.1
  helm:
    version: v3
    releaseName: RELEASE_NAME
    valueFiles:
    - $values/my-app/values.yaml
```

Git Requests Total (ls-remote)



Git Requests Total (checkout)



- Very high *git-fetch* and *git-ls-remote* requests.
- Potentially a bug. We implemented an undocumented workaround [#14725](#).
- Our *git-fetch* requests dropped dramatically.
- We're still seeing high *ls-remote* requests. The bug issue is still open.

tuning: monorepo usage w/webhook



KubeCon



CloudNativeCon

North America 2023

Problem: all monorepo apps refreshed every time there's commit to unrelated apps
🔧 use manifest paths annotation

```
...  
kind: Application  
metadata:  
  name: my-kustomize-app  
  namespace: argocd  
  annotations:  
    argocd.argoproj.io/manifest-generate-paths: ".*;/bases/my-base;"  
spec:  
...
```

- (in monorepo) ArgoCD webhook server refresh all apps when it receive a webhook.
- In the refresh process, ArgoCD invalidates cache for all apps and calls k8s API to annotate all Application objects, slowing update process when having 1k+ apps.
- Using the annotation filters out unrelated apps, speeding up the update process.

tuning: argocd-application-controller



KubeCon



CloudNativeCon

North America 2023

Problem: workqueue depth started piling up

🔧 Increase # of operation processors and status processors

```
apiVersion: v1
kind: ConfigMap
metadata:
  labels:
    app.kubernetes.io/name: argocd-cmd-params-cm
    app.kubernetes.io/part-of: argocd
  name: argocd-cmd-params-cm
data:
  ...
  # ARGOCD_APPLICATION_CONTROLLER_STATUS_PROCESSORS
  controller.status.processors: '500'
  # ARGOCD_APPLICATION_CONTROLLER_OPERATION_PROCESSORS
  controller.operation.processors: '250'
  ...
```

For every 1000 application, use

`--status-processors=50` and

`--operation-processors=25`

tuning: argocd-application-controller



KubeCon



CloudNativeCon

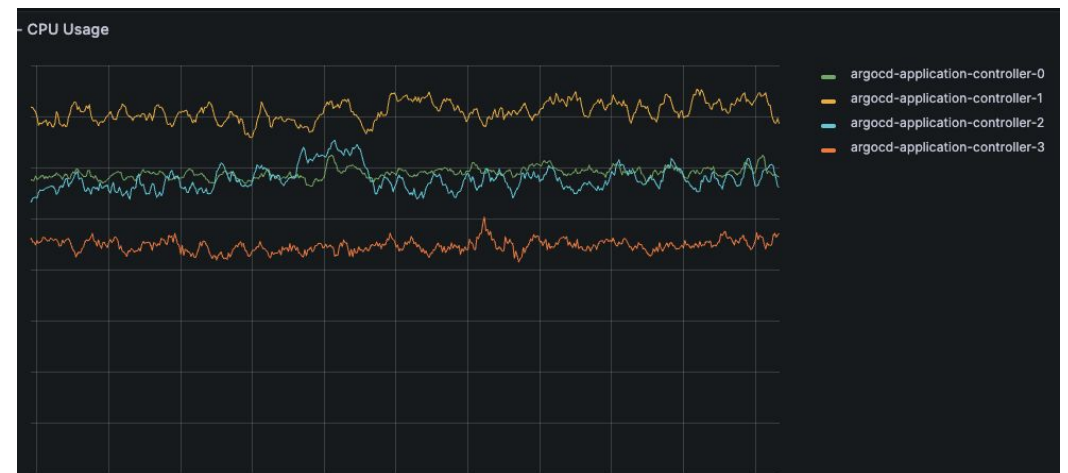
North America 2023

Problem: scaling argocd-application-controller

🔧 shard controllers to multiple pods - distribute load

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: "argocd-application-controller"
spec:
  replicas: 4
  template:
    spec:
      containers:
        - args:
          - "/usr/local/bin/argocd-application-controller"
          env:
            - name: "ARGOCD_CONTROLLER_REPLICAS"
              value: "4"
  ...
```

- argocd-app-controller is horizontally shardable
- sharding algorithm is on the cluster-level
- increase replicas and set **ARGOCD_CONTROLLER_REPLICAS**



tuning: argocd-application-controller



KubeCon




CloudNativeCon

North America 2023

Problem: uneven shards CPU usage, some shards are higher than others

🔧 manual shard assignment

```
...
apiVersion: v1
kind: Secret
metadata:
  annotations:
    argocd.argoproj.io/sync-options: ServerSideApply=true
  name: CLUSTER_A_SECRET_NAME
  namespace: argocd
type: Opaque
stringData:
  shard: "2" # the shard number
...
```



- ArgoCD shards per cluster, not per app. Large clusters could be hosted by the same shard.
- The round-robin sharding might not help much either. There's still chance large clusters could get into same shards.
- We did manual shard allocation instead to fine-tune shard resources.
- ArgoCD app-level sharding as a feature would be really great. See discussion [\[1\]](#).

tuning: argocd-application-controller



KubeCon



CloudNativeCon

North America 2023

Problem: we still see very high app-controller CPU usage, slow reconciles

🔧 optimize high-churn Reconciliations

```
resource.customizations.ignoreDifferences.all: |
  managedFieldsManagers:
  - kube-controller-manager
  jsonPointers:
  - /spec/replicas

resource.ignoreResourceUpdatesEnabled: "true"

resource.customizations.ignoreResourceUpdates.all: |
  jsonPointers:
  - /status

resource.customizations.ignoreResourceUpdates.autoscaling_HorizontalPodAutoscaler: |
  jsonPointers: #for autoscaling/v1 compatibility
  - /metadata/annotations/autoscaling.alpha.kubernetes.io~1behavior
  - /metadata/annotations/autoscaling.alpha.kubernetes.io~1conditions
  - /metadata/annotations/autoscaling.alpha.kubernetes.io~1metrics
  - /metadata/annotations/autoscaling.alpha.kubernetes.io~1current-metrics

resource.customizations.ignoreResourceUpdates.apps_ReplicaSet: |
  jsonPointers:
  - /metadata/annotations/deployment.kubernetes.io~desired-replicas
  - /metadata/annotations/deployment.kubernetes.io~max-replicas
  - /metadata/annotations/deployment.kubernetes.io~revision

resource.customizations.ignoreResourceUpdates.discovery.k8s.io_EndpointSlice: |
  jsonPointers:
  - /endpoints
  - /ports
  - /metadata/annotations

resource.compareoptions: |
  ignoreDifferencesOnResourceUpdates: true
```

- ArgoCD watches all field changes of tracked objects.
- K8s fields could get very concise or frequently update, even for fields that we don't really need to reconcile.
- Use the features: ignoreResourceUpdates (available v2.8) [1] and, ignoreDifferences [2].
- 🔍 might need to enable debug log to find org-specific high-churn objects



fix: argocd apiclient



KubeCon



CloudNativeCon

North America 2023

Problem: http2 GOAWAY / grpc ENHANCE_YOUR_CALM errors

  Fix missing grpc parameter when using `grpcWeb=true`

- Our in-house developer platform implements argocd's apiclient library.
- As we scale, we started seeing http2 GOAWAY errors from our platform.
- We investigated, found and fixed a bug in the apiclient library when using `grpcWeb=true` ([#15707](#)).
- Also, the argocd CLI may fallback to use `grpcWeb=true` even though `--grpc-web` flag is not specified.
- It may also implicitly use `grpcWeb=true`. Check your `.config/argocd/config` file!
- Alternatively, use native grpc.

Future Improvements



decentralized argocd (pull model)




KubeCon



CloudNativeCon

North America 2023





Pros

- Application controller workload distributed across clusters
- Access to Kubernetes API server is local only

Cons

- Maintenance and upgrade headache
- Automation headache: maintain multiple argocd client versions
- No more centralized dashboard
- Still require tuning for large clusters

agent-based argocd (hybrid model)



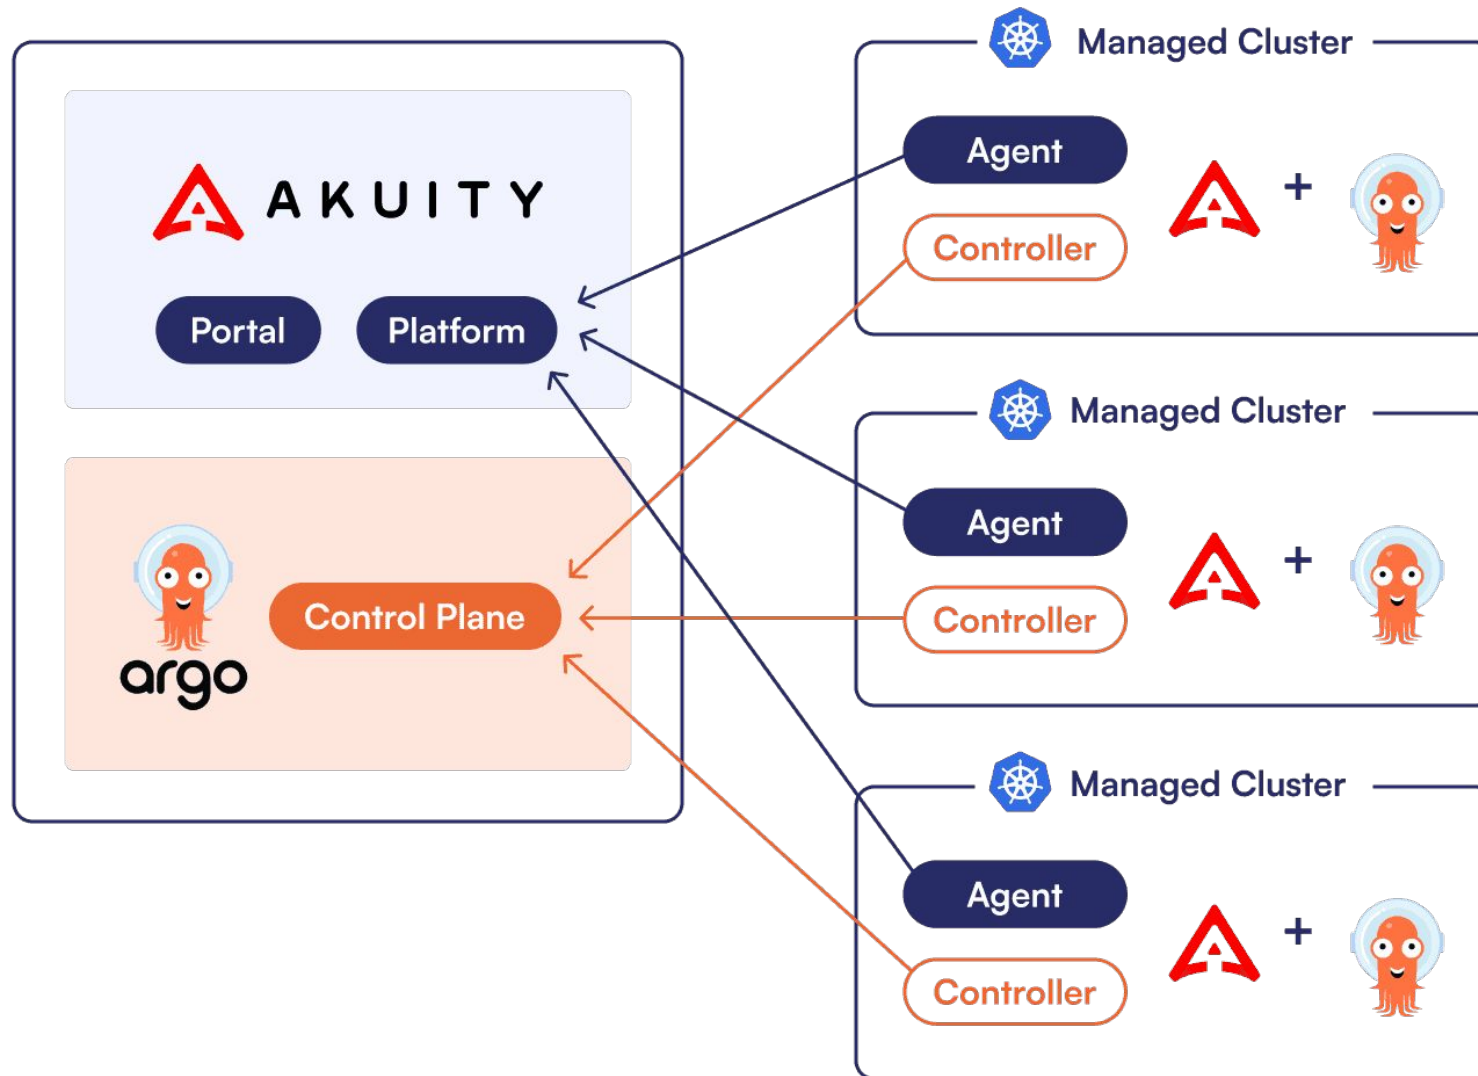
KubeCon



CloudNativeCon

North America 2023

popularized by Akuity platform



Source:
<https://akuity.io/blog/argo-cd-architectures-explained/>

agent-based argocd (hybrid model)



KubeCon



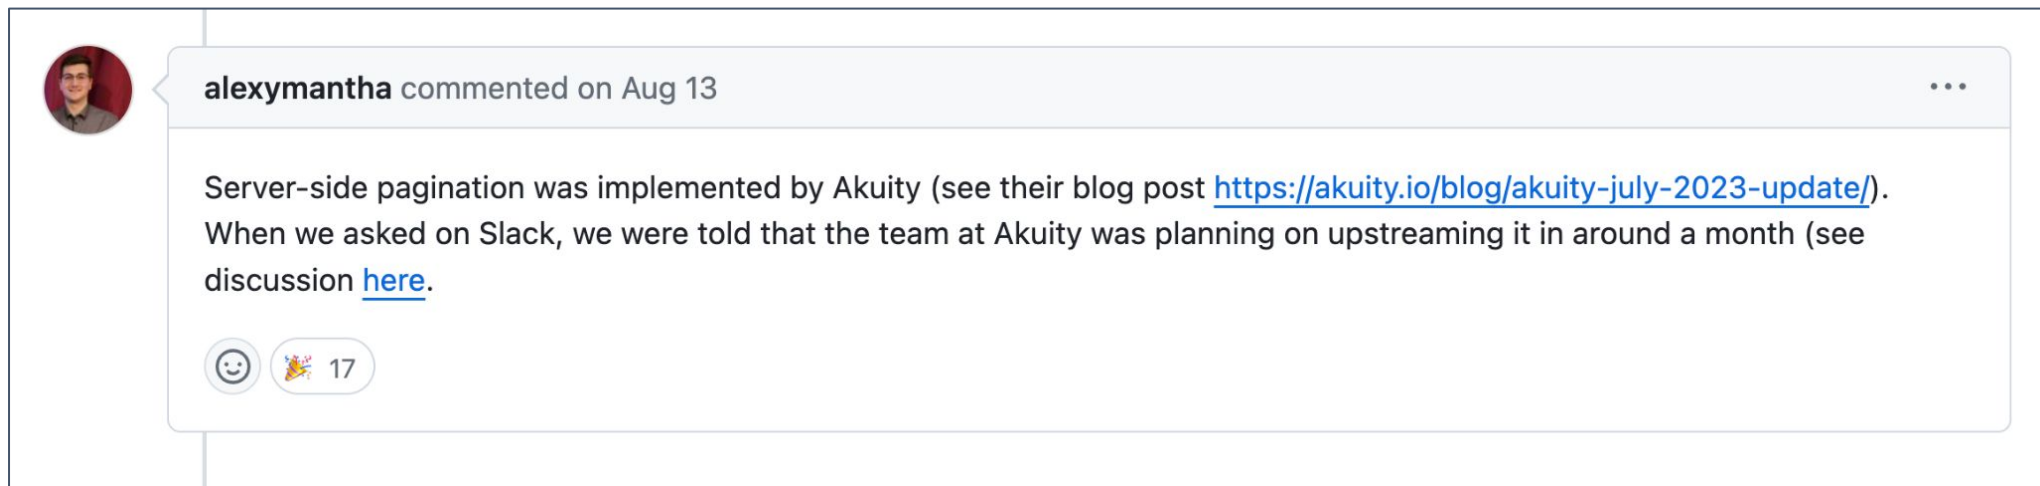
CloudNativeCon

North America 2023

WIP in the community

- Optional Pull Mechanism for ApplicationSet (merged: [#10908](#))
- Centralized UI for Multiple ArgoCD Instances (open: [#11498](#))

- Server-Side Pagination 🚀 (open: [#14947](#))
- Already [on Akuity](#), and they planned bring it to upstream 🎉



references



KubeCon



CloudNativeCon

North America 2023

- [TikTok, Managing Thousands of Apps with ArgoCD](#)
- [Adobe, Managing Hundreds of Clusters with ArgoCD](#)
- [Alexander Matyushentsev, ArgoCD Best Practices](#)
- [ArgoCD High Availability Documentation](#)



PromCon
North America 2021




**Please scan the QR Code above
to leave feedback on this session**

 @girikuncoro

 @yphanama



KubeCon



CloudNativeCon

North America 2023

