



Jaeger: The Future with OpenTelemetry and Metrics

Maintainer Talk - Kubecon NA 2023

Pavol Loffay [@ploffay]





Pavol Loffay?





Agenda

1. Why Distributed Tracing?
2. Jaeger demo
3. Service performance monitoring / Monitor tab
4. OpenTelemetry collector vs Jaeger
5. Jaeger V2
6. OpenTelemetry auto instrumentation
7. New Key Features + Roadmap for Jaeger
8. Q&A from the audience

Why Distributed Tracing?

Microservices Architectures

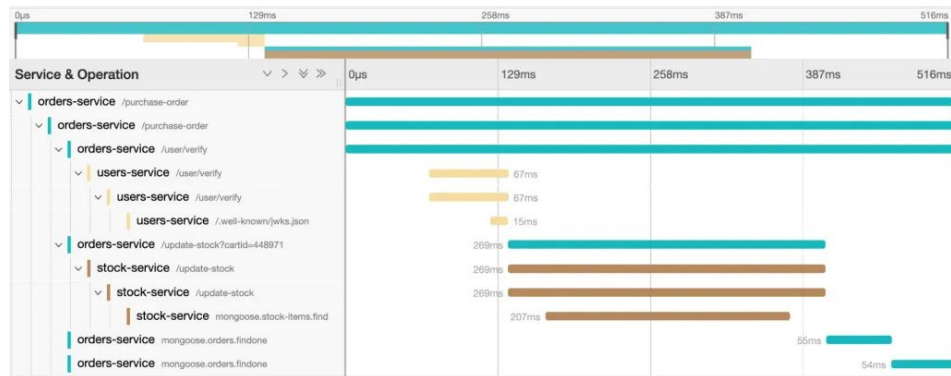
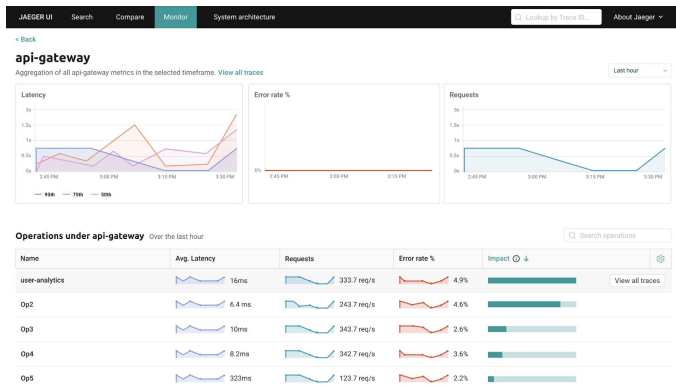
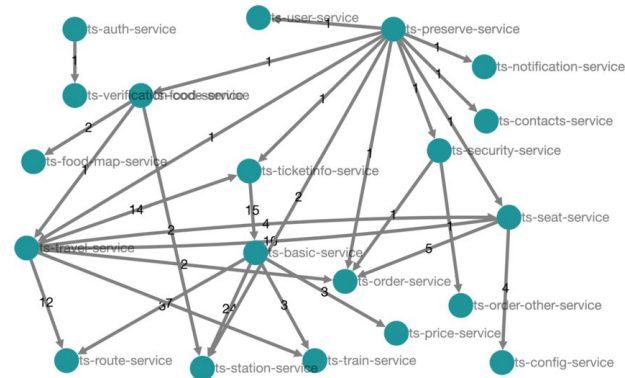
63% of enterprises are using Microservices most often managed by different teams.

When things **break or are slow**, how do you find the right team to help resolve the problem?



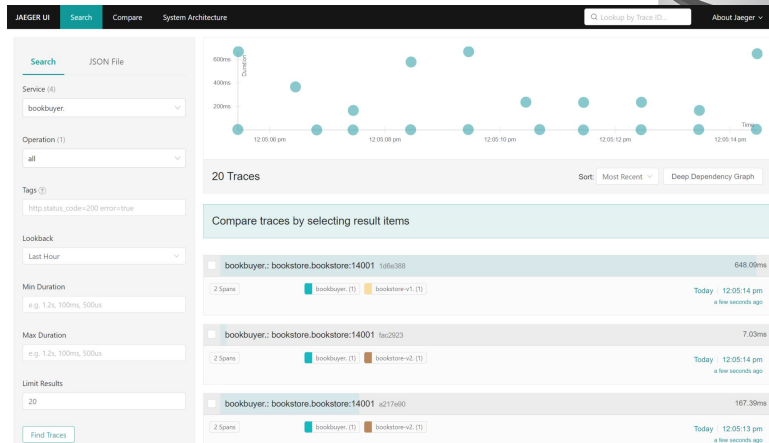
Benefits of Distributed Tracing

- Root cause analysis
- Build a service map of relationships and dependencies
- Measure and drill into specific user actions
- Collaborate across teams
- Monitor and maintain SLAs



Tracing Tools

- Instrumentation
- Data Collection
- Storage & Analysis & Visualization



Intro to Jaeger and demo



OpenTelemetry Semantics

Trace represents an end-to-end request (and response); made up of single or multiple **Spans**

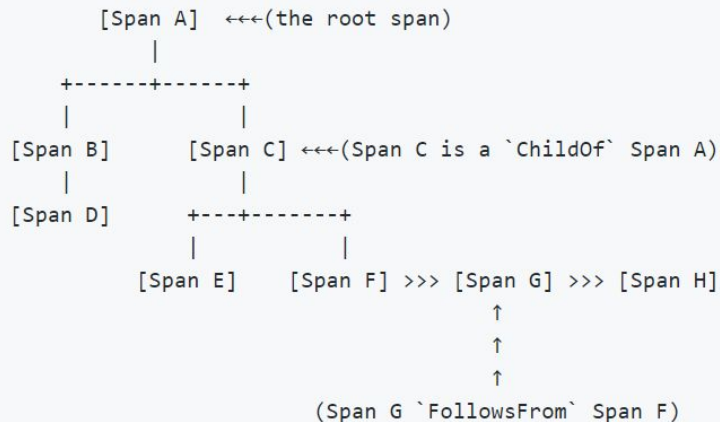
Span represents work done by a single-service or component with time intervals and associated metadata such as **Tags**

Tags/Attributes contain metadata to help contextualize a span

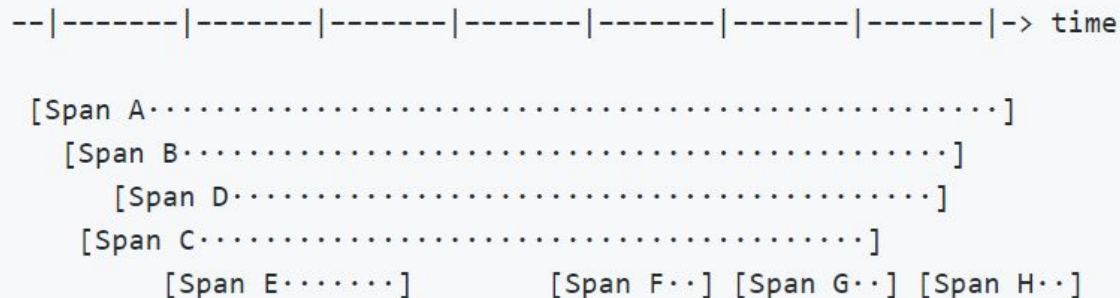


Relationships in tracing

Causal relationships between Spans in a single Trace



Temporal relationships between Spans in a single Trace



Measure errors, latency, and other indicators across each span

Trace Start **November 2 2023, 09:29:02.443** Duration **954.16ms** Services **6** Depth **4** Total Spans **40**

Service & Operation

```

  frontend /dispatch
  frontend HTTP GET
    customer /customer
      mysql → mysql SQL SELECT

```

- ▼ frontend driver.DriverService/FindNearest

▼ driver driver.DriverServiceFindNearest

```
redis-manual FindDriverD
● redis-manual GetDriver
redis-manual GetDriver
redis-manual GetDriver
redis-manual GetDriver
● redis-manual GetDriver
redis-manual GetDriver
redis-manual GetDriver
redis-manual GetDriver
● redis-manual GetDriver
redis-manual GetDriver
redis-manual GetDriver
```

frontend HTTP GET

route /route/

▼ frontend HTTP GET

route /route

▼ frontend HTTP GET

```
route route
```

frontend HTTP GET

```
route route
```

frontend HTTP GET

```
route /route
```

frontend HTTP GET

route /route

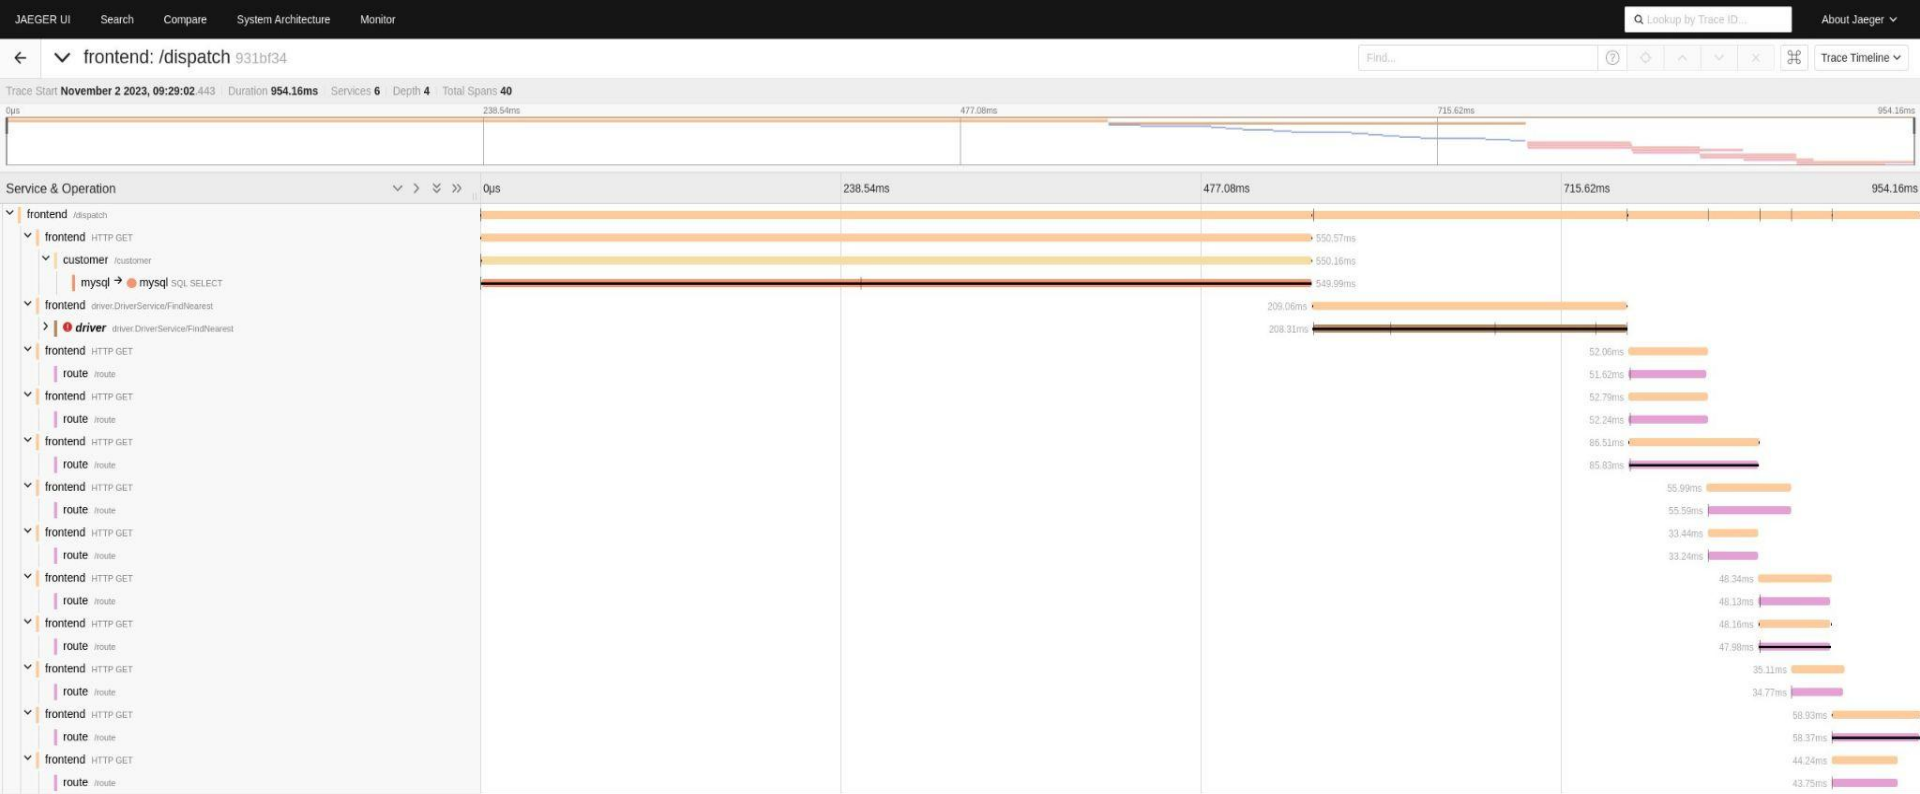
frontend HTTP GET

route /route/

frontend HTTP GET

route /route





New Monitoring tab and Prometheus support in Jaeger

Tracing and Monitoring



- Moving Jaeger from “distributed tracing” towards APM
 - Traces / Events
 - Metrics
- Additional use cases
 - Operational Monitoring
 - Operational Alerting
 - Change Planning



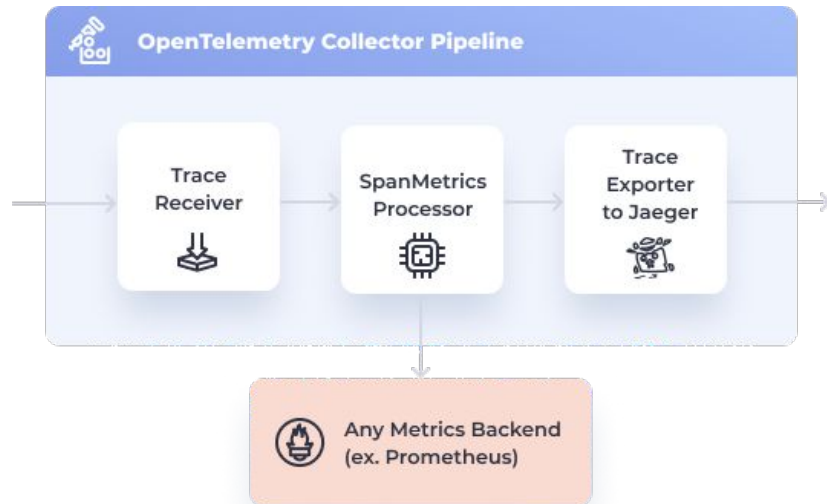
SpanMetrics Connector Flow

Prometheus handles all metrics use cases (monitor, alert, plan)
So can many other metrics platforms...



OpenTelemetry to derive
aggregated metrics from traces

SpanMetrics Connector



SpanMetrics Connector in OpenTelemetry



```
connectors:
  spanmetrics:
    exemplars:
      enabled: true

service:
  pipelines:
    traces:
      receivers: [otlp]
      exporters: [otlp/traces, spanmetrics]
  metrics:
    receivers: [spanmetrics]
    exporters: [otlphttp/metrics] #export data to Prometheus
```


Using Prometheus Metrics in Grafana

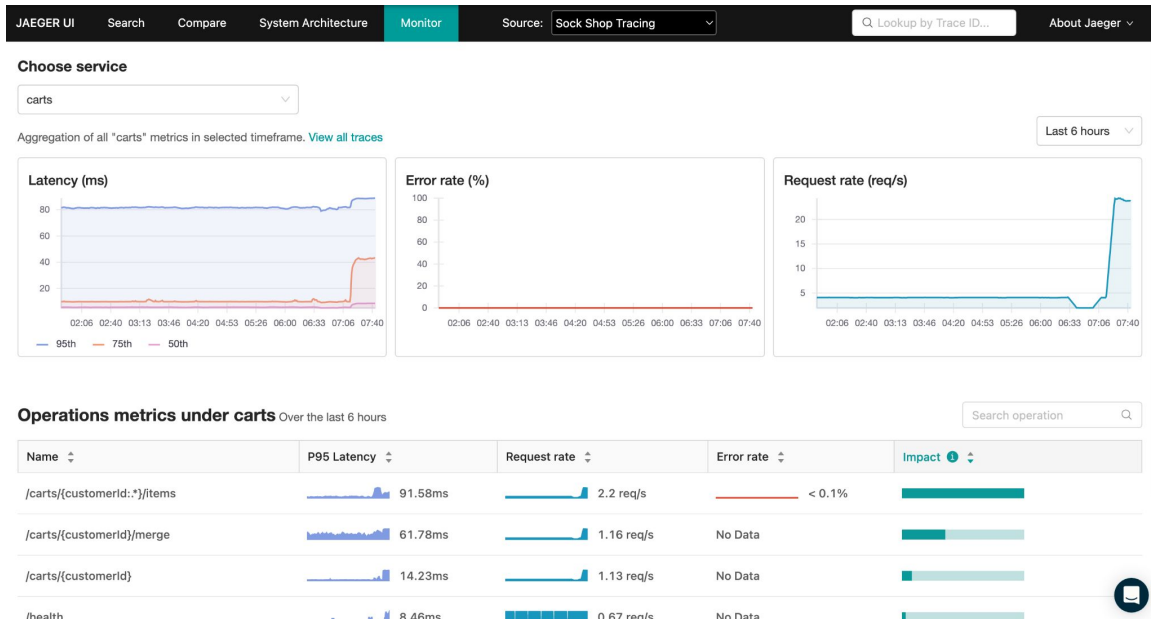
MetricQuery service in Jaeger to query metric backends.

- Any promql compatible backend (ex: Prometheus, Cortex, Thanos, M3DB, and Mimir)
- Community can add other metric systems as needed



Using Prometheus Metrics inside Jaeger UI

- New “monitor” homepage in Jaeger to provide status and health of transactions



Pipelines with Jaeger and OpenTelemetry

OTel Collector + OTel Clients + Jaeger



- OTel Supports Jaeger ingestion protocols
 - Agent (legacy)
 - thrift compact (6831)
 - thrift binary (6832)
 - Collector
 - protobuf/gRPC (14250)
 - thrift over HTTP (14268)
- Jaeger remote sampler extension (HTTP 5778, 14250 gRPC)
- Can also write to Kafka in supported Jaeger formats

OTel Processors

processors:

redaction:

```
# allow_all_keys is a flag which when set to true, which can disable the
# allowed_keys list. The list of blocked_values is applied regardless of
# you just want to block values, set this to true.
```

```
allow_all_keys: false
```

```
# allowed_keys is a list of span attribute keys that are allowed to pass
# through. The list is designed to fail closed. If allowed_keys is empty,
# no span attributes are allowed and all span attributes are removed. If
# allow all keys, set allow_all_keys to true. To allow the span attributes
# you know are good, add them to the list.
```

```
allowed_keys:
```

- description
- group
- id
- name

```
# blocked_values is a list of regular expressions for blocking values
# allowed span attributes. Values that match are masked
```

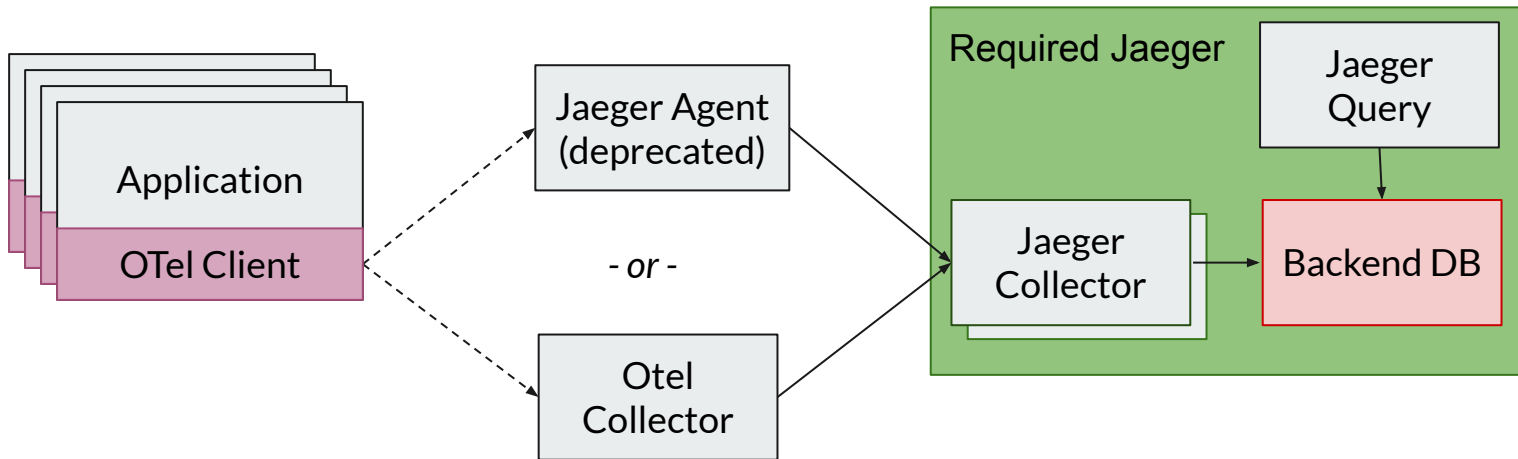
```
blocked_values:
```

- "4[0-9]{12}(?:[0-9]{3})?" ## Visa credit card number
- "(5[1-5][0-9]{14})" ## MasterCard number

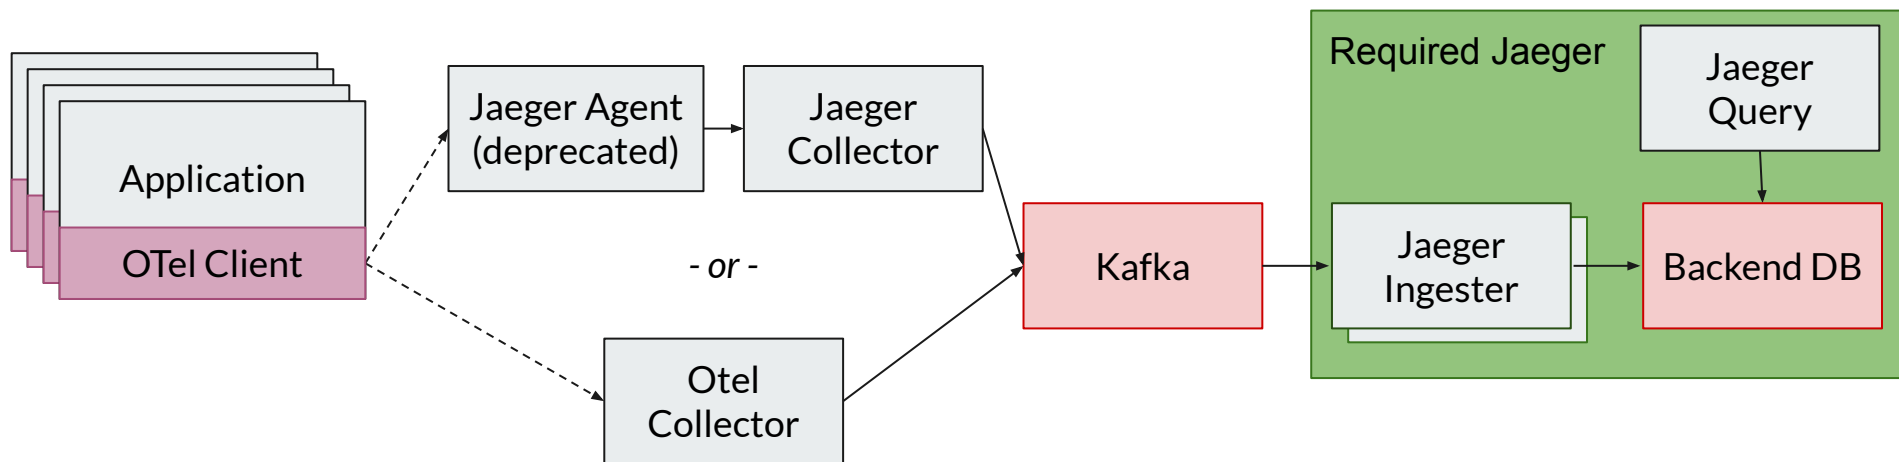
attributesprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
cumulativetodeltaprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
deltatorateprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
filterprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
groupbyattributesprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
groupbytraceprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
k8sattributesprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
logstransformprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
metricsgenerationprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
metricstransformprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
probabilisticsamplerprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
redactionprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
resourcedetectionprocessor	[chore] dependabot updates Mon Oct 24 22:43:13 UTC 2022 (#15650)
resourceprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
routingprocessor	[chore] dependabot updates Mon Oct 24 22:43:13 UTC 2022 (#15650)
schemaprocessor	[chore] dependabot updates Mon Oct 24 22:43:13 UTC 2022 (#15650)
servicegraphprocessor	[chore] dependabot updates Mon Oct 24 22:43:13 UTC 2022 (#15650)
spanmetricsprocessor	[chore] dependabot updates Mon Oct 24 22:43:13 UTC 2022 (#15650)
spanprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
tailsamplingprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)
transformprocessor	[chore] dependabot updates Mon Oct 24 18:28:10 UTC 2022 (#15536)

<https://github.com/open-telemetry/opentelemetry-collector-contrib/tree/main/processor>

Without Queueing



With Queueing



Jaeger operator

“operates” Jaeger on Kubernetes

Getting Started

- jaegertracing.io/docs/latest/operator
- Storage schema creation
- Jaeger upgrades
- Auto recognizes available APIs - OpenShift/Kubernetes, ES/Strimzi
- Can generate plain Kubernetes manifest files



Jaeger CRD



- CRD - Custom Resource Definition

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  strategy: allInOne | production | streaming

  storage:
    type: memory|elasticsearch|cassandra|badger|grpc-plugin|kafka
```

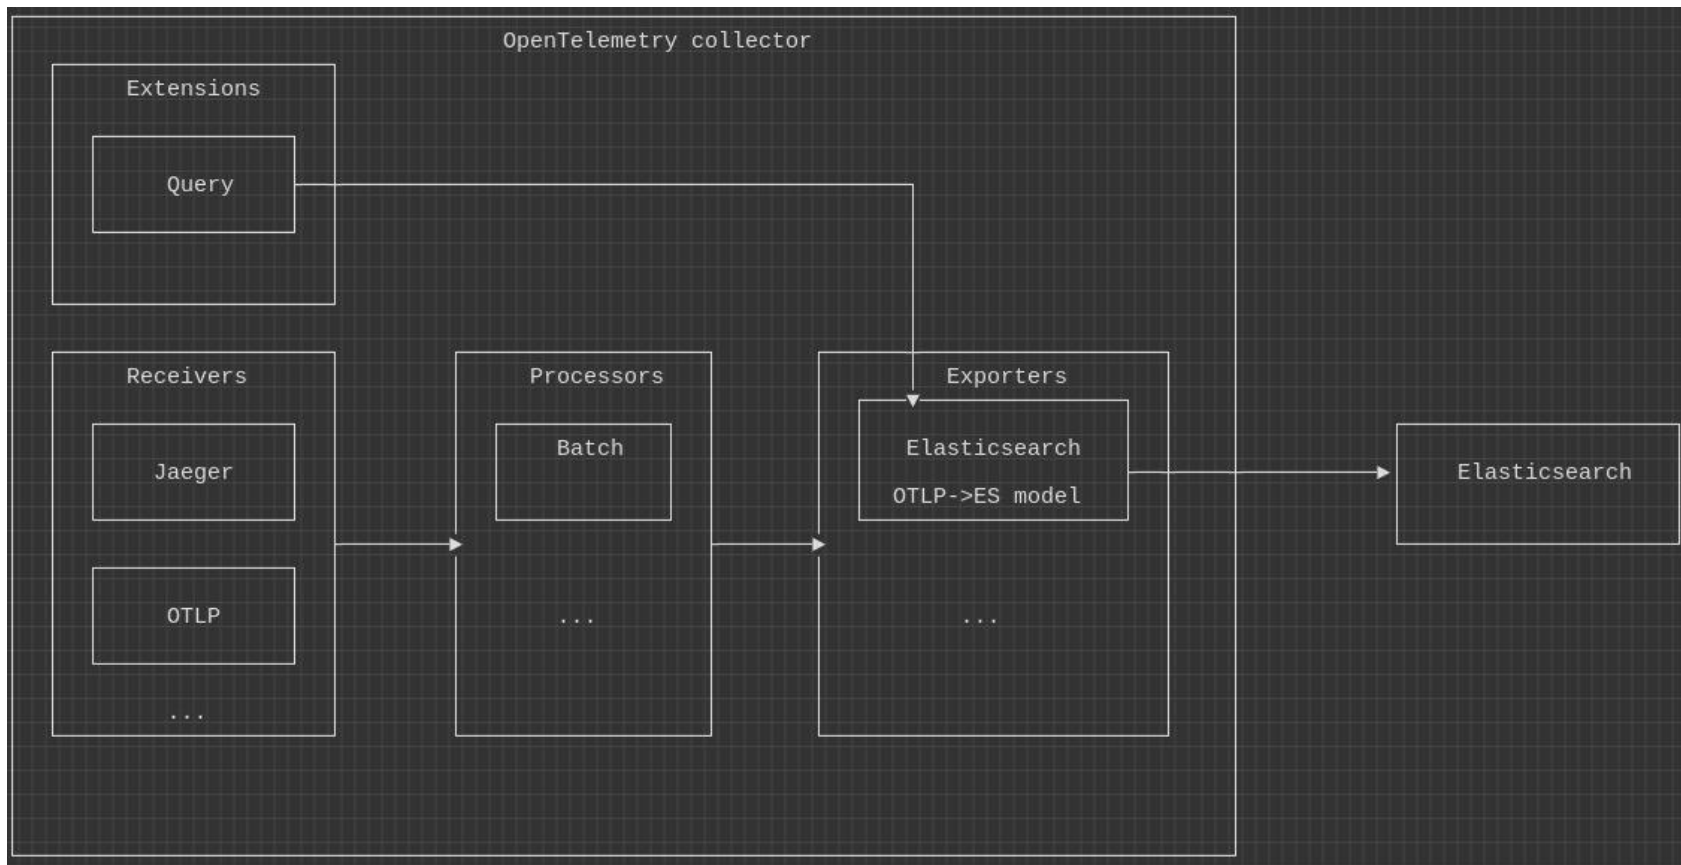
- `kubectl get jaegers`

Jaeger agent sidecar injection

- Jaeger agent has been deprecated in Jaeger 1.50.
 - it will be removed in ~April/May 2024
- Use OpenTelemetry operator!
 - Jaeger receiver *gRPC, Thrift binary/compact/HTTP)



Jaeger v2 based on OpenTelemetry collector



V2 Configuration



```
service:
  extensions: [jaeger_storage, jaeger_query]
  pipelines:
    traces:
      receivers: [otlp, jaeger, zipkin]
      processors: [batch]
      exporters: [jaeger_storage_exporter]

  extensions:
    jaeger_query:
      trace_storage: memstore

    jaeger_storage:
      memory:
        memstore:
          max_traces: 100000

  exporters:
    jaeger_storage_exporter:
      trace_storage: memstore
```

OpenTelemetry (auto) instrumentation

Deprecated Clients & Jaeger agent

Sunsetting Jaeger clients 🚩 #3362



Open



7 of 9 tasks

yurishkuro opened this issue on Oct 31, 2021 · 17 comments



yurishkuro commented on Oct 31, 2021 • edited ▾

The Jaeger clients have faithfully served our community for several years, providing a wide range of controlled samplers and per-operation / adaptive sampling which

Version 1.37

Latest

See also:

- [Client Features](#)



Jaeger clients are being retired.



This library is DEPRECATED!

No new pull requests are accepted except for security fixes.

We urge all users to migrate to [OpenTelemetry](#). Please refer to the [notice in the documentation](#) for details.

Jaeger Bindings for Go OpenTracing API

Wide Language Support (OpenTracing Shim)



<https://opentelemetry.io/docs/migration/opentracing/>

Environment Variable Based Config



JAeger_SERVICE_NAME
JAeger_TAGS

JAeger_AGENT_HOST

JAeger_SAMPLER_TYPE
JAeger_SAMPLER_PARAM
JAeger_PROPAGATION




OTEL_SERVICE_NAME
OTEL_RESOURCE_ATTRIBUTES

OTEL_TRACES_EXPORTER

OTEL_TRACES_SAMPLER
OTEL_TRACES_SAMPLER_ARG
OTEL_PROPAGATORS

<https://opentelemetry.io/docs/reference/specification/sdk-environment-variables/>

Trace Propagation



OTEL_PROPAGATORS

default: tracecontext, baggage (w3c)

supported: jaeger, b3, b3multi

<https://opentelemetry.io/docs/reference/specification/sdk-environment-variables/>

(Auto) Instrumentation



Auto!

- Java
- Python
- .NET
- Node.JS

Less Auto!

- Go
 - eBPF auto-inst
- Ruby
 - minor global code change
- C++, PHP, Rust...

<https://opentelemetry.io/docs/instrumentation/>

New Key Features Since Kubecon Europe + Roadmap

Other New Features



- OTLP ingestion is enabled by default
 - OpenTelemetry collector removed support for Jaeger exporter
- Support for spanmetrics connector
- Jaeger V2 based on OpenTelemetry collector
- Jaeger V3 query API exposes stable V1 OTLP
- UI: Critical path visualization
- UI: In-page search on key value pairs
- UI: batch download of traces



Roadmap

- Jaeger V2 based on OpenTelemetry collector
 - Kafka support
 - OpenTelemetry collector builder (OCB)
 - Testing & Docs & Release
- Native Clickhouse storage support in Jaeger V2
- Elasticsearch 8 storage support

And more interesting capabilities coming in the future

Q&A

Resources



jaegertracing.io/docs



[monthly community call and Notes](#)
[CNCF Slack #jaeger : https://slack.cncf.io](https://slack.cncf.io)



[@jaegertracing](https://twitter.com/jaegertracing)



medium.com/jaegertracing



Leave feedback on
this session