

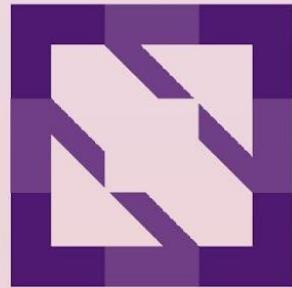


KubeCon

---

North America 2023

---



CloudNativeCon



KubeCon



CloudNativeCon

North America 2023

# Network Policy API: Intro and Project Update

*Andrew Stoycos, Dan Winship, Surya Seetharaman, Red Hat  
Yang Ding, VMware*



# Overview

- What is the network-policy-api subgroup? - Andrew
  - How did we get here
    - ([link](#) to last year's contributor summit talk so we can pull stuff from that...)
- NetworkPolicy API - Dan
- AdminNetworkPolicy API
  - A brief history (KEP), later work - Yang
  - Status - Yang
  - Implementations - Antrea, OVNK
    - (encourage more implementations)
- Future work
  - NPEPs, ANP v1beta1, etc - Surya (will start the format idea and delegate)
  - conformance tests + collaboration with GWAPI + conformance profiles- Surya
  - “NPv2” / DNP -Dan
- How and why you should get involved Whether you are a K8s user, developer, administrator, operator or CNI maintainer



KubeCon



CloudNativeCon

North America 2023

# sig-network-policy-api subgroup

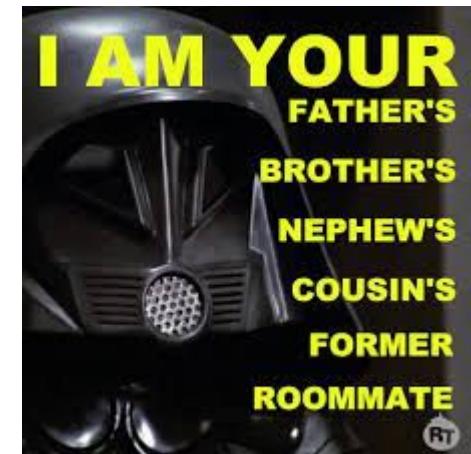
# sig-network-policy-api: Who are we?

- **Who are we?**

- A passionate group of developers who spawned out of the Kubernetes Networking Special interest group, focused explicitly on the maintenance and development of policy related APIs
- <https://github.com/kubernetes-sigs/network-policy-api/graphs/contributors> + many more :)

- **What do we do?**

- Help maintain the legacy NetworkPolicy v1 resource
- Rework of the upstream [network policy test suite](#)
- [The AdminNetworkPolicy KEP](#)
- Stay tuned...





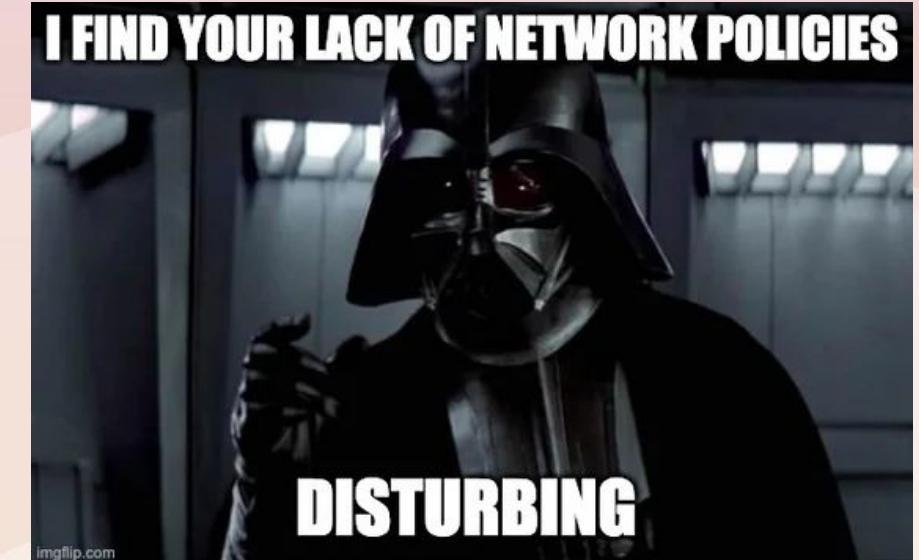
KubeCon



CloudNativeCon

North America 2023

# NetworkPolicy API



# NetworkPolicy v1: Designed for Developers



KubeCon



CloudNativeCon

North America 2023



As an **application developer**, I want to control which other users within the cluster can access my application



As an **application developer**, I want to have a multi-tiered architecture, where I can control access between pods in each tier

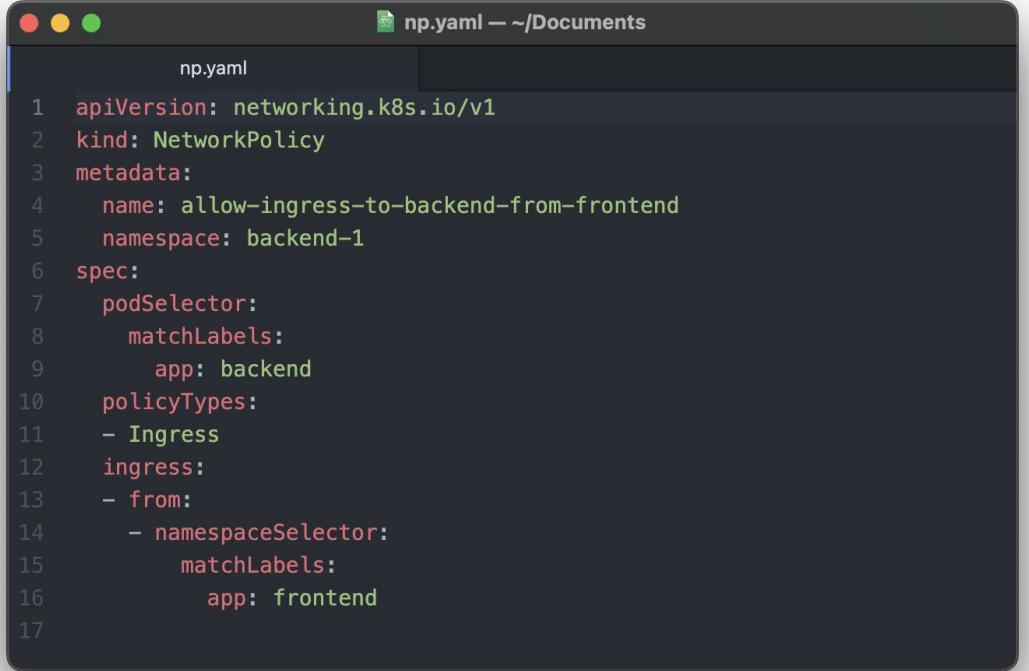
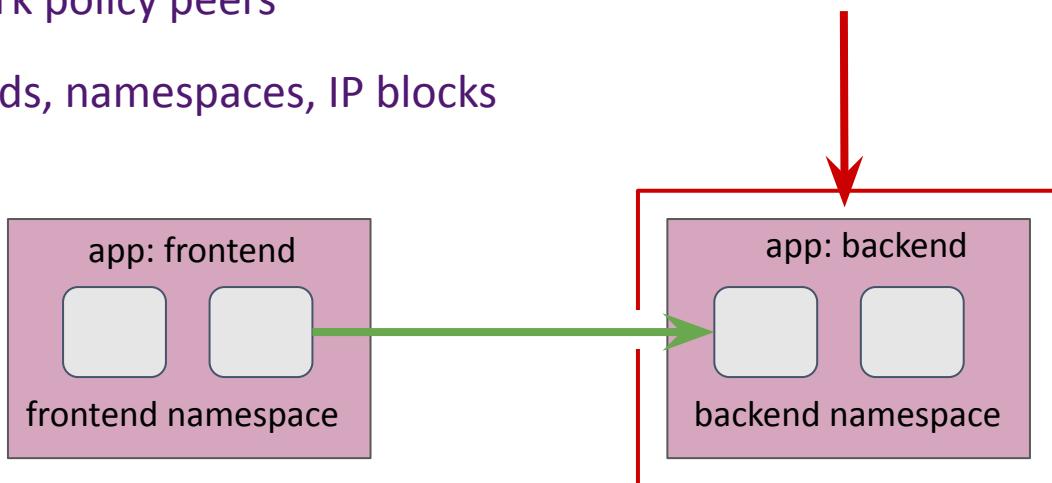
# NetworkPolicy v1: Designed for Developers



In 2016 we were all still drunk on DevOps, and Kubernetes mostly didn't even have clear “administrator” / “user” distinctions. (RBAC was added in the same Kubernetes release as NetworkPolicy.)

# NetworkPolicy v1

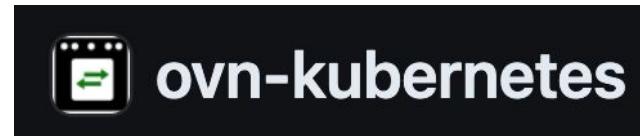
- How can app owners control traffic to/from their workloads?
  - example: backends can get traffic only from frontends, databases can only get traffic from backends etc..
- Namespace scoped Policy API that lets users define simple ingress/egress rules
- API design is implicit in nature
- Network policy peers
  - pods, namespaces, IP blocks



```
np.yaml — ~/Documents
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: allow-ingress-to-backend-from-frontend
5   namespace: backend-1
6 spec:
7   podSelector:
8     matchLabels:
9       app: backend
10  policyTypes:
11    - Ingress
12    ingress:
13      - from:
14        - namespaceSelector:
15          matchLabels:
16            app: frontend
17
```

# NetworkPolicy v1

- Has been a stable v1 API for over 6 years
- Numerous implementations exist, below are just a few



# NetworkPolicy v1: Issues



Policies only specify what is *allowed* and leave everything else implicitly denied in a way that is confusing and hard to work with

Implicit isolation model



The lack of deny rules, priority, and namespace-spanning policies, means it doesn't solve administrator use cases

Not for admin use cases



Behavior of policy as it applies to traffic entering/exiting a cluster is underspecified

N/S traffic selection



KubeCon



CloudNativeCon

North America 2023

# Time for a new API!

## The KEP process

- 10+ contributors
- 100+ individual commits
- 600+ comments resolved
- 1 year from PR open to merge



# AdminNetworkPolicy API - KEP

Contributors / reviewers <https://github.com/kubernetes/enhancements/pull/2522>



Casey



Gobind



Ricardo



Abhishek



Dan



Sanjeev



Thomas



Andrew



Jay



Yang



Satish



Tim



Vinay



Christopher

## Motivation

A cluster-scoped, admin-facing, CNI-agnostic policy API for complete control of cluster networking

## Challenges

- Consolidate on use cases (which were a lot)
- Need to be explicit in API design
- Priority modeling
- Provide rule actions that make sense to everyone
- Iterate design so that the API is minimal but solves all use cases agreed on
- Future extensibility



KubeCon



CloudNativeCon

North America 2023

# AdminNetworkPolicy API

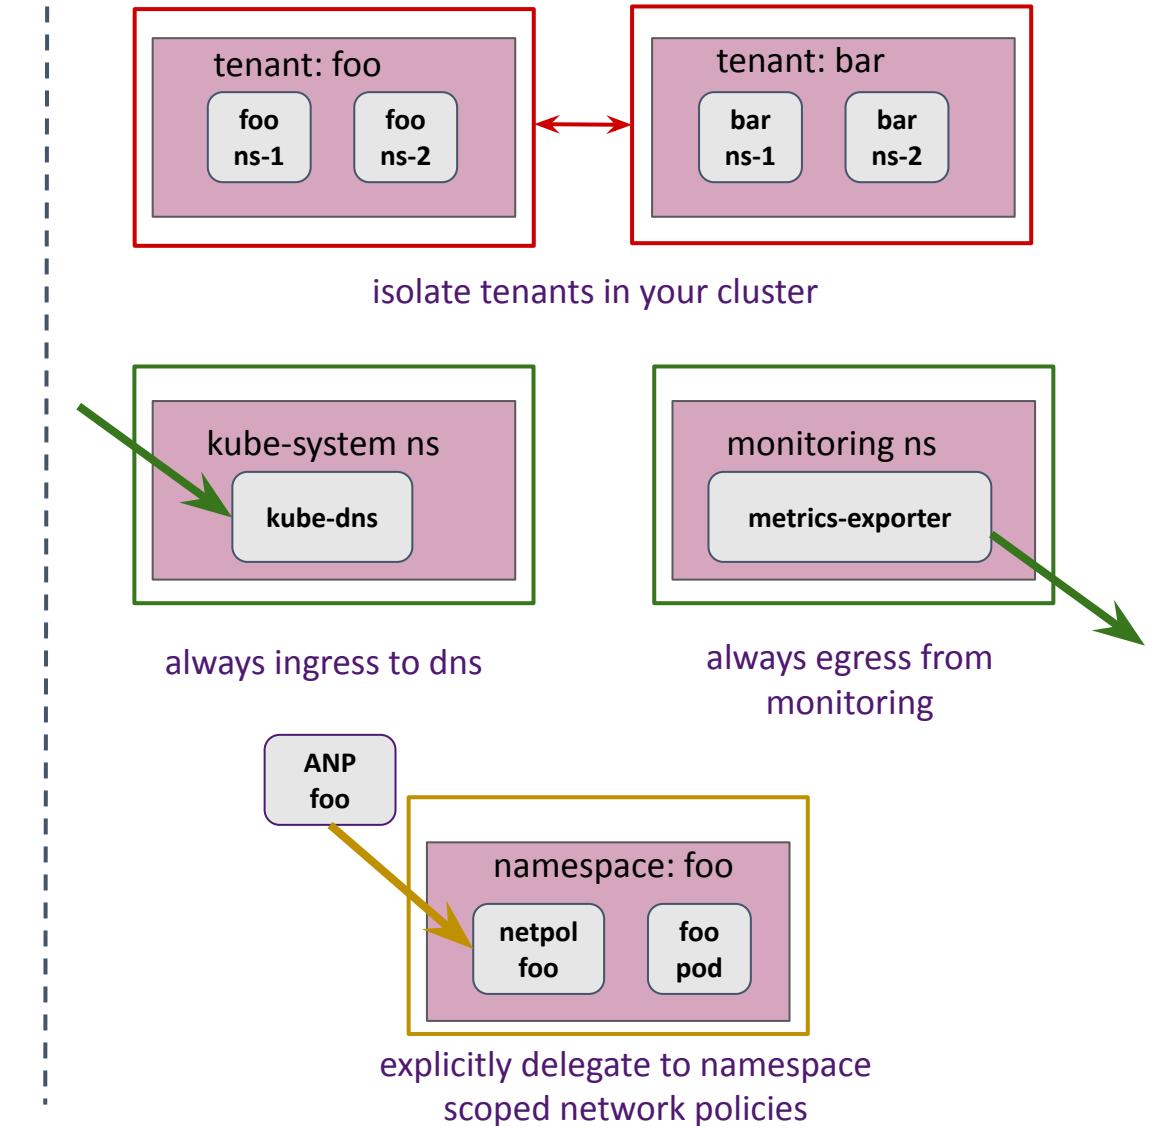
Started from the KEP now we're here

- Two objects: AdminNetworkPolicy and BaselineAdminNetworkPolicy
- out-of-tree
- v1alpha
- Supports intra-cluster controls



# AdminNetworkPolicy API

- How can cluster admins control traffic to/from their clusters' workloads in a non-overridable manner?
  - isolate tenants
  - always allow ingress to DNS namespace
  - always allow egress from monitoring namespace
  - delegate to namespace scoped network policies
  - always deny traffic to sensitive namespace from everywhere



# AdminNetworkPolicy API

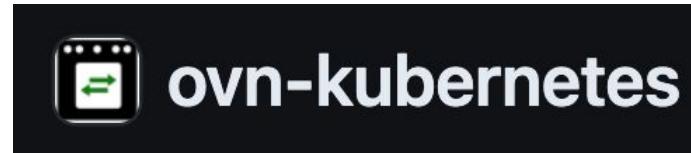
- Let admins define simple ingress/egress rules to secure their cluster that takes precedence over NetworkPolicy API
- There's also BaselineAdminNetworkPolicy to define the default security posture of a cluster in absence of NetworkPolicies
- API design is explicit in nature



```
np.yaml -- ~/Documents
anp.yaml
1 apiVersion: policy.networking.k8s.io/v1alpha1
2 kind: AdminNetworkPolicy
3 metadata:
4   name: cluster-wide-allow-example
5 spec:
6   priority: 30
7   subject:
8     namespaces: {}
9   ingress:
10    - action: Allow
11      from:
12        - namespaces:
13          namespaceSelector:
14            matchLabels:
15              kubernetes.io/metadata.name: monitoring-ns
16   egress:
17    - action: Allow
18      to:
19        - pods:
20          namespaces:
21            namespaceSelector:
22              matchLabels:
23                kubernetes.io/metadata.name: kube-system
24            podSelector:
25              matchLabels:
26                app: kube-dns
27
```

# AdminNetworkPolicy Current Status

- Has been a v1alpha1 API for over 1 year; on the journey towards beta....
- Two implementations exist, we welcome more implementations and call for feedback!!



- On the roadmap of:



1. Calico tracker: <https://github.com/projectcalico/calico/issues/7578>
2. Cilium tracker: <https://github.com/cilium/cilium/issues/23380>
3. KubeOVN Tracker: <https://github.com/kubeovn/kube-ovn/issues/3247>



KubeCon



CloudNativeCon

North America 2023

# Network Policy Enhancement Proposals + Feature development

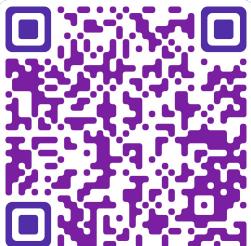
- Similar to KEPs but for the sig-network-policy-api project...
- Have a cool idea? Open a NPEP!

When I gotta go to work again even though I just went yesterday



# NPEP for Conformance Profiles

- How can API maintainers track implementations and establish a feedback loop?
- How can implementations be more involved in the API conformance design?
- NPEP proposes to add a new **ConformanceReport** CRD that captures the results of each *core* and *extended feature profile*



Reports Folder



Help Wanted!!



Contributor talk

```
profiles:  
- core:  
  failedTests:  
    - AdminNetworkPolicyIngressSCTP  
    - AdminNetworkPolicyEgressUDP  
  result: failure  
  statistics:  
    Failed: 2  
    Passed: 5  
    Skipped: 0  
    summary: ""  
    name: AdminNetworkPolicy  
- core:  
  result: success  
  statistics:  
    Failed: 0  
    Passed: 7  
    Skipped: 0  
    summary: ""  
    name: BaselineAdminNetworkPolicy
```



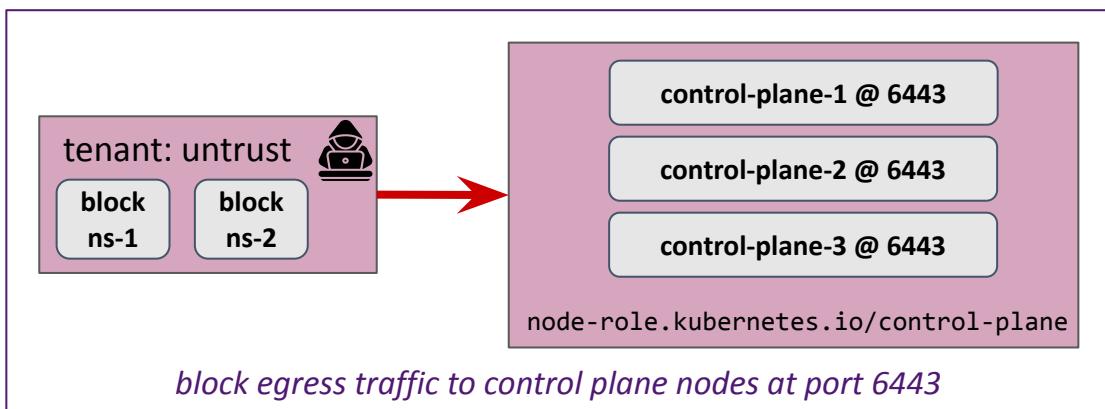
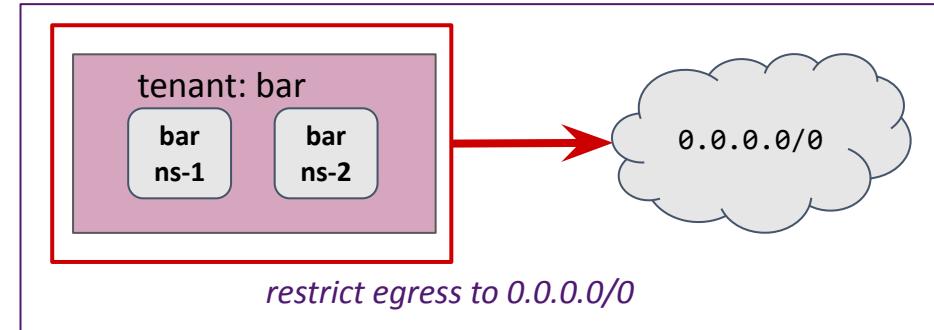
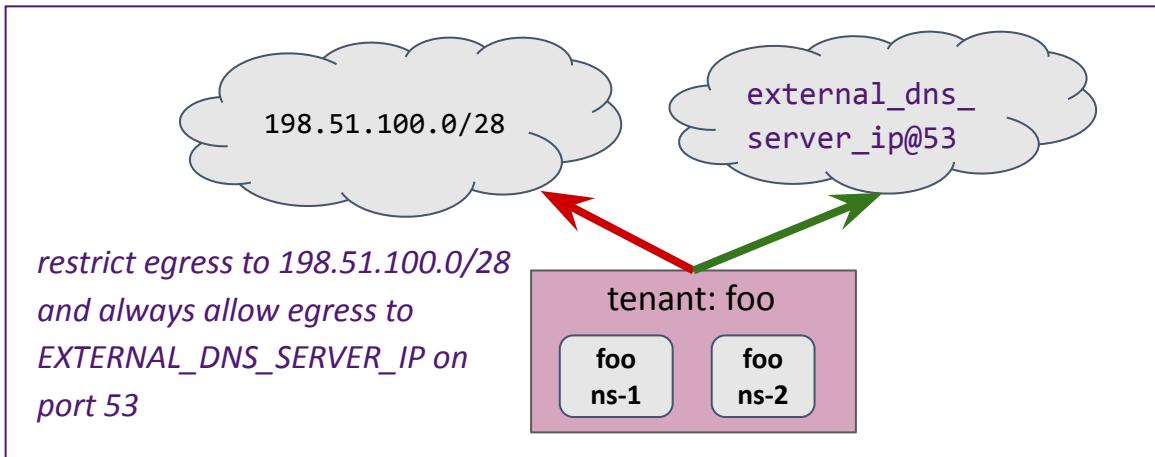
NPEP 137



Surya  
@tssurya

# NPEP for cluster egress traffic controls

- Current API has two egress peers: **namespaces** and **pods**
- How can cluster admins control northbound traffic from their clusters' workloads?



- NPEP proposes to add two new egress peers - **nodes** and **networkCIDRs**



Help Wanted!!  
Leave feedback!  
Comment on the Issue!



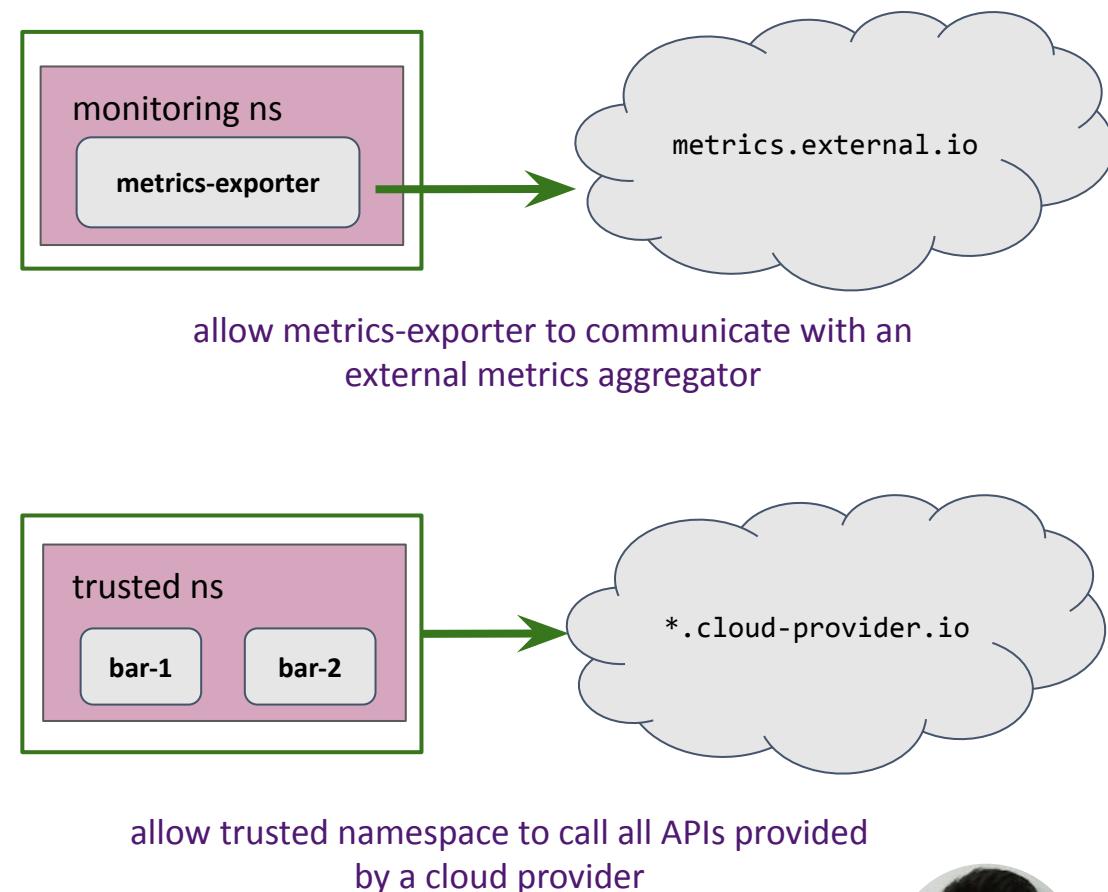
Surya  
@tssurya

# NPEP for cluster egress FQDN traffic controls

- Proposal: Allow specifying Fully-Qualified Domain Names (for example [www.kubernetes.io](http://www.kubernetes.io)) as egress peers
  - Also allows for wildcard selectors
- Caveats:
  - Only focused on egress
  - Only “ALLOW” type rules are permitted. The NPEP proposes prohibiting “DENY” type rules with a FQDN peer.
  - Not a true L7 policy, but rather uses DNS information for L3/4 enforcement

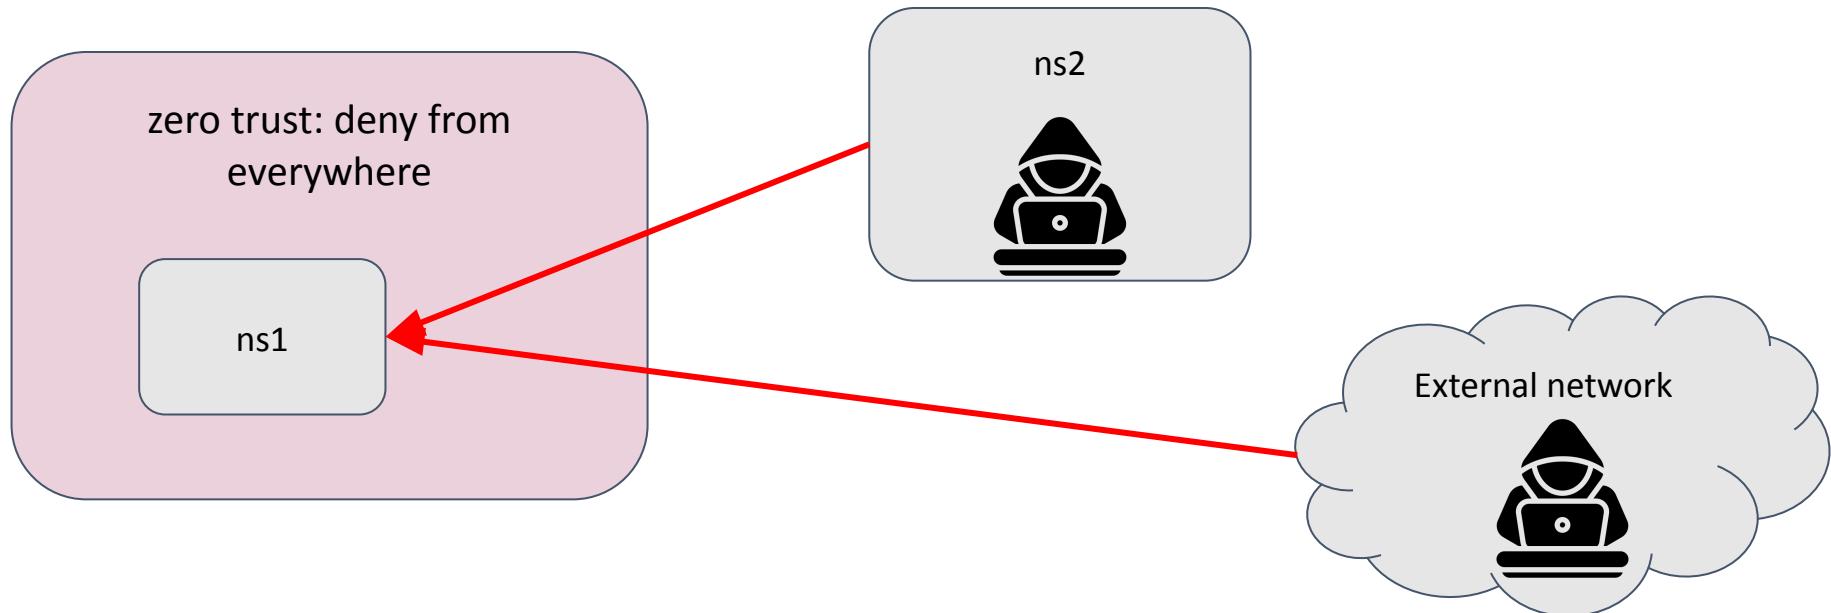


NPEP 133



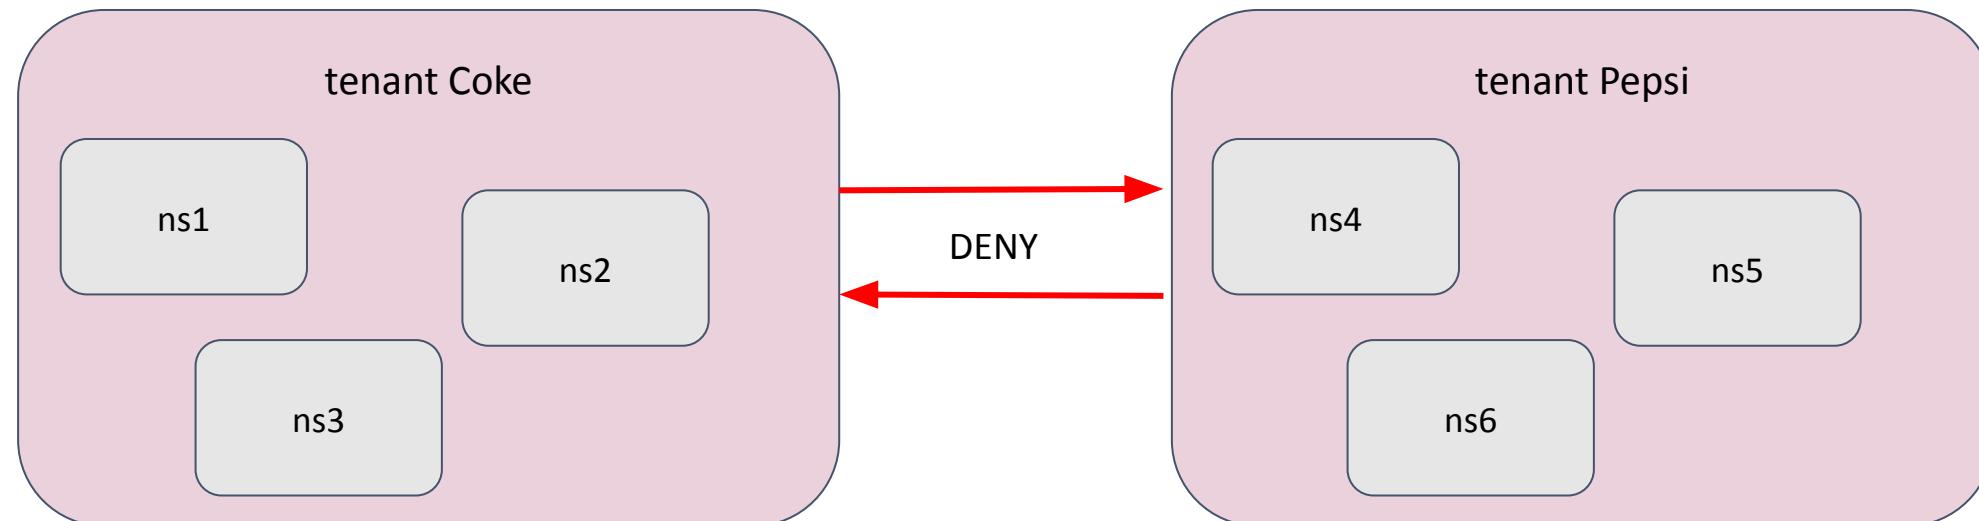
# NPEP for cluster Ingress traffic controls

- Problem: Existing API doesn't allow to specify external endpoints for ingress.
- Goal: Implement ingress traffic control from external destinations (outside the cluster)
- Issue: <https://github.com/kubernetes-sigs/network-policy-api/issues/127>



# NPEP for tenancy

- Problem: Existing API allows much more configurations that were initially required, which may be confusing for users.
- Goal: Simplify existing API to only allow configurations that are required by the outlined user stories.
- Issue: <https://github.com/kubernetes-sigs/network-policy-api/issues/122>



Nadia  
@npinaeva

# Policy Assistant Feature (#150) (built on Cyclonus)

## Goal 1: Develop/Understand your Policies

- Predict verdict from policies:

TRAFFIC	VERDICT	INGRESS FLOW	EGRESS FLOW
test/client → prod/server:80 (TCP)	Denied by ANP	[ANP] default/anp1: deny (pri=3)	[NPv1] prod/np1: deny
prod/client → prod/server:81 (UDP)	Allowed by BNP	[ANP] default/anp1: pass (pri=44) [BNP] default/bnp: allow (pri=55)	(No matching rules)

- Explain policies:

TYPE	TARGET	SOURCE RULES	PEER	ACTION	PORTS
Ingress	namespaces with same labels for "key1"	[ANP] default/anp1 [BNP] default/bnp	all namespaces pods w/ label "role=client"	ANP: allow (pri=3), deny (pri=4), pass (pri=7)  BNP: allow (pri=1)	80-99, TCP
			namespaces w/ label "user=3" all pods	ANP: deny (pri=9)	
	namespace: default pods w/ label "a=b"	[NPv1] default/np1 [NPv1] default/np2	namespace: default pods w/ label "role=client"	NPv1: peers allowed	80, UDP

## Goal 2: Extended Test Framework

- Run hundreds of test scenarios, filtering by tag (e.g. run ANP tests).
- Quick-to-code tests automatically calculate expected connectivity from policies.



Hunter  
@huntergregory



KubeCon



CloudNativeCon

North America 2023

# NetworkPolicy v2?

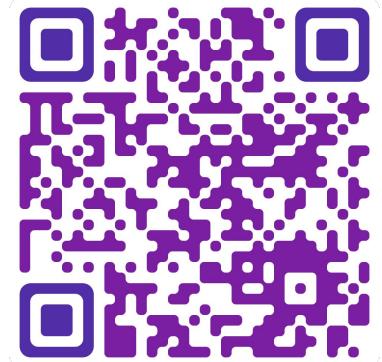


# NetworkPolicy v2?

- AdminNetworkPolicy solves the fact that NP doesn't implement administrator use cases, but it still leaves NP's other problems:
  - “Implicit deny on first policy in the namespace” is hard to work with.
  - Weird syntactic quirks make NetworkPolicies hard to write and understand.
  - Backward compatibility requires us to leave a hole in policies for kubelet probes.
  - It's hard for SIG Network to add new features to NetworkPolicy without causing old implementations to misinterpret new NetworkPolicy objects.
- A literal networking.k8s.io/v2 version of NetworkPolicy wouldn't help, because of backward-compatibility constraints.
- ANP is “better NetworkPolicy for administrators”. Could we also create “better NetworkPolicy for developers”?

# DeveloperNetworkPolicy

- (Still just an idea: [NPEP-136](#).)
- Let's take a step back from what NetworkPolicy became, and reinvent developer-focused networking policy from the user stories up.
  - Don't let administrator-focused use cases sneak in, as happened with NetworkPolicy; keep it solely focused on developer use cases.
  - Ensure that it fits perfectly between AdminNetworkPolicy and BaselineAdminNetworkPolicy, allowing use cases that are harder to implement with NetworkPolicy.
- AdminNetworkPolicy + DeveloperNetworkPolicy = "NetworkPolicy v2"
  - NetworkPolicy v1 would still keep working (at least, as long as implementors wanted to support it), but it would stop getting new features, and we would push people toward ANP and DNP.



NPEP 136



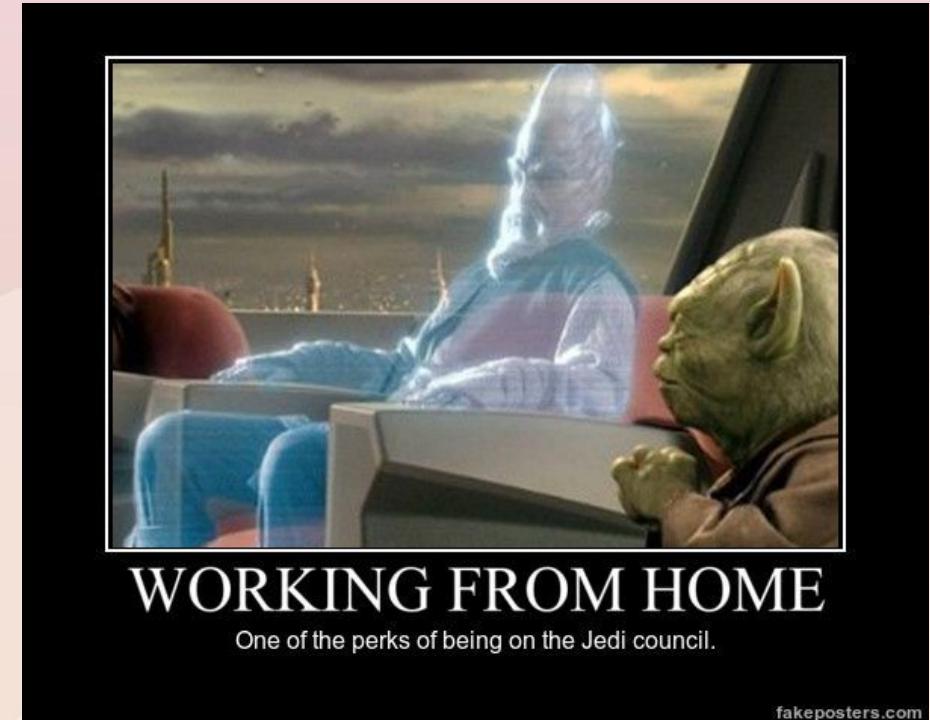
KubeCon



CloudNativeCon

North America 2023

# How To Get Involved



# Getting Involved in the project!

- **Level 1: Completely new to K8s and Don't know where to start!?**
  - Checkout the [K8s contributor guide!](#)
  - Read the [Network Policy V1 documentation](#)
  - Read the [AdminNetworkPolicy KEP](#)
- **Level 2: Familiar with the basics but still unsure where to dive in?**
  - Read though our documentation at <https://network-policy-api.sigs.k8s.io/> please open issues I am sure there are documentation bugs we've yet to find 😷
  - Checkout our [help wanted](#) and [good first](#) issues
  - Review any of the open PRs
- **Level 3: Want more features! Want more implementations!**
  - Read about our [NPEP \(Network Policy Enhancement Proposal\) process](#)
  - Propose a new NPEP with an issue.



[kubernetes-sigs/network-policy-api](https://github.com/kubernetes-sigs/network-policy-api)



[#sig-network-policy-api](#)



<https://network-policy-api.sigs.k8s.io>



PromCon  
North America 2021



**Please scan the QR Code above  
to leave feedback on this session**