



KubeCon



CloudNativeCon

North America 2023

Rapidly Scaling for Breaking News with Karpenter and KEDA



Mel Cone (they/she/he/ella), @melmaliacone
Senior Software Engineer, The New York Times



Deepak Goel (he/him), LinkedIn: deepak-goel12
Software Engineer, The New York Times

Breaking News Alert



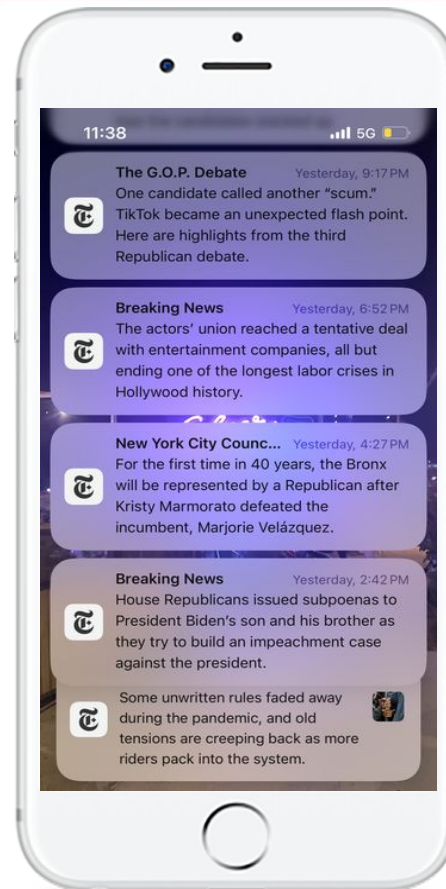
KubeCon



CloudNativeCon

North America 2023

- Breaking News Alerts (BNAs) are the push notifications notifying users of an article



The New York Times

The Problem




KubeCon



CloudNativeCon

North America 2023

- Sudden spiky traffic
- Reaches peak in 20 seconds for a duration of 2-3 minutes
- Goal: Use the fewest and cheapest node as possible



The New York Times

Predictable Traffic Spikes




KubeCon



CloudNativeCon

North America 2023

- Games daily release leads to up to 3x spike in traffic
- Spike lasts for a small time period



The New York Times



KubeCon



CloudNativeCon

North America 2023

1 What we're scaling

2 How we're scaling: Karpenter

3 Scaling Demo

4 How we're scaling: KEDA

5 Lessons Learned

What We're Scaling: Unified HTTP Ingress

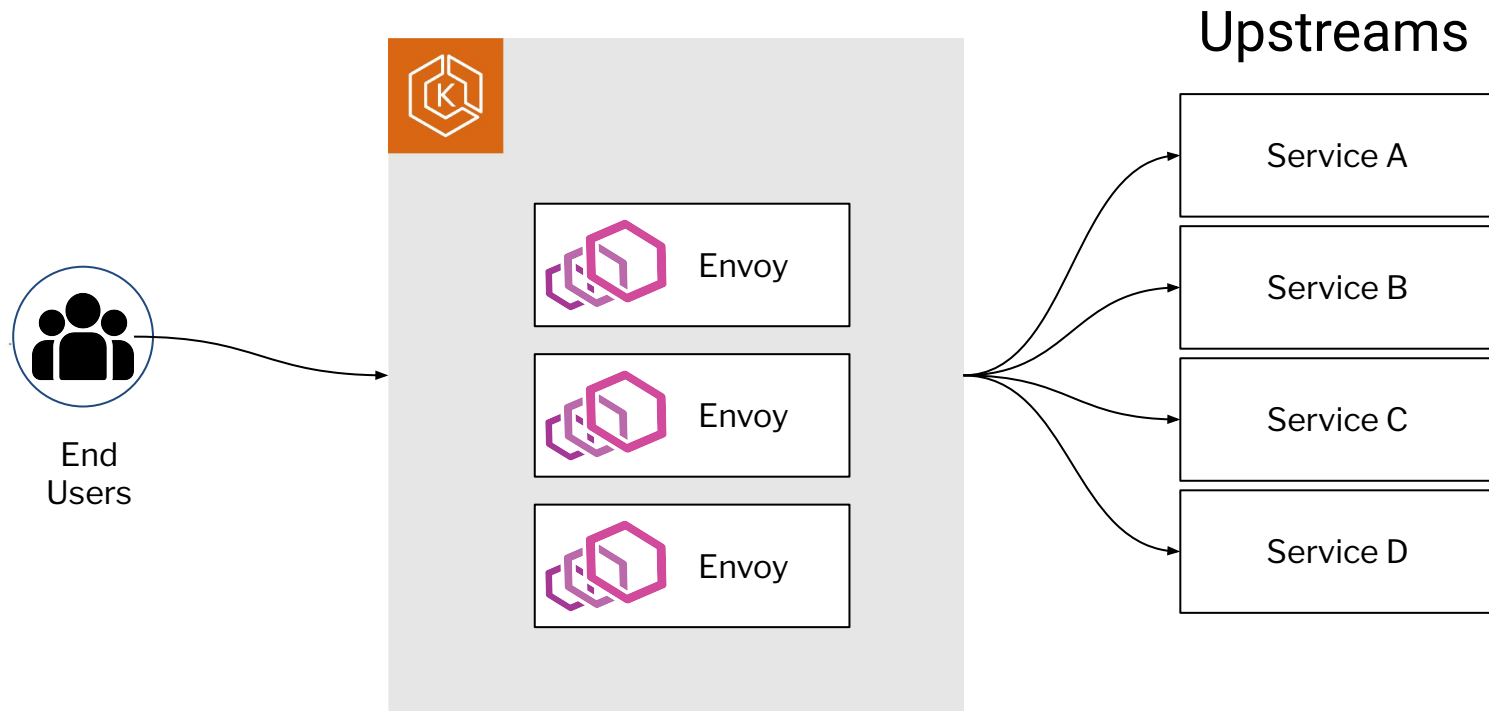


KubeCon



CloudNativeCon

North America 2023



Multi-Tenant Architecture




KubeCon



CloudNativeCon

North America 2023



The New York Times



KubeCon



CloudNativeCon

North America 2023

- 1 What we're scaling
- 2 How we're scaling: Karpenter**
- 3 Scaling Demo
- 4 How we're scaling: KEDA
- 5 Lessons Learned

Why not just Autoscaler?





KubeCon




CloudNativeCon

North America 2023



VS



HashiCorp

Terraform

The New York Times

Why not just Autoscaler?

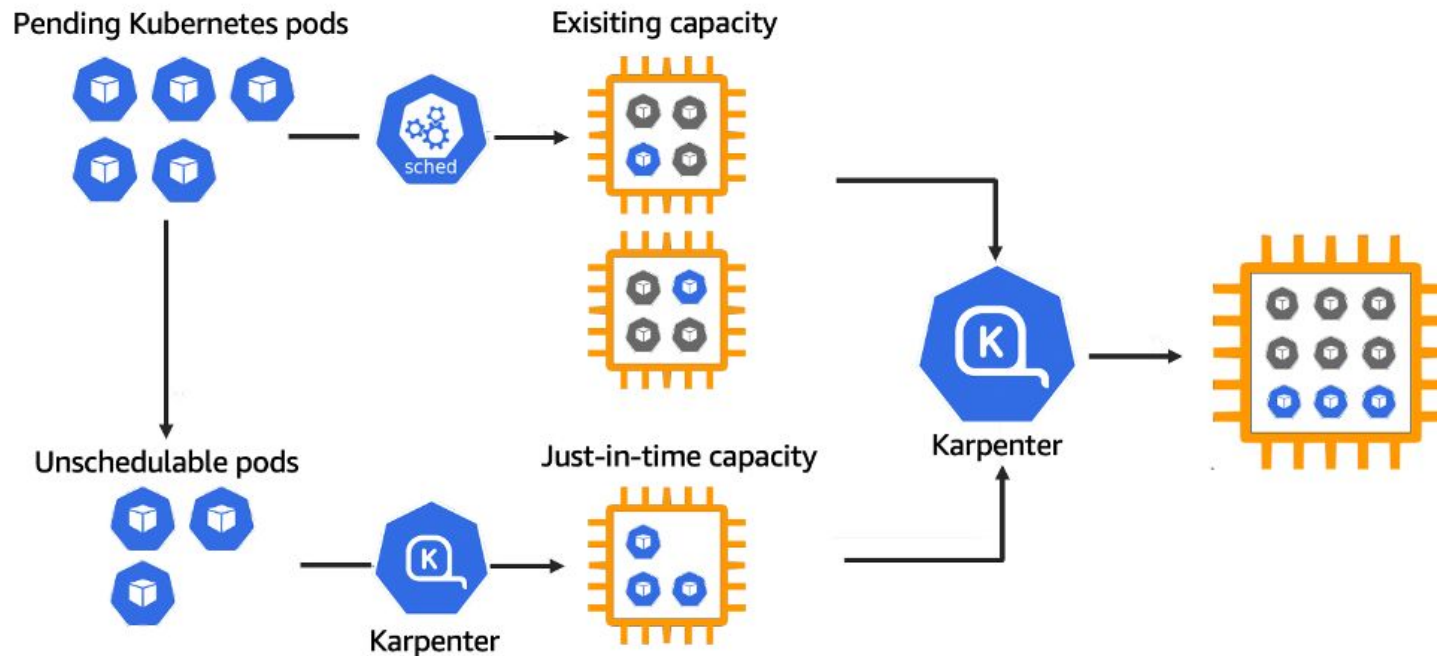


KubeCon



CloudNativeCon

North America 2023



Migrating to Karpenter



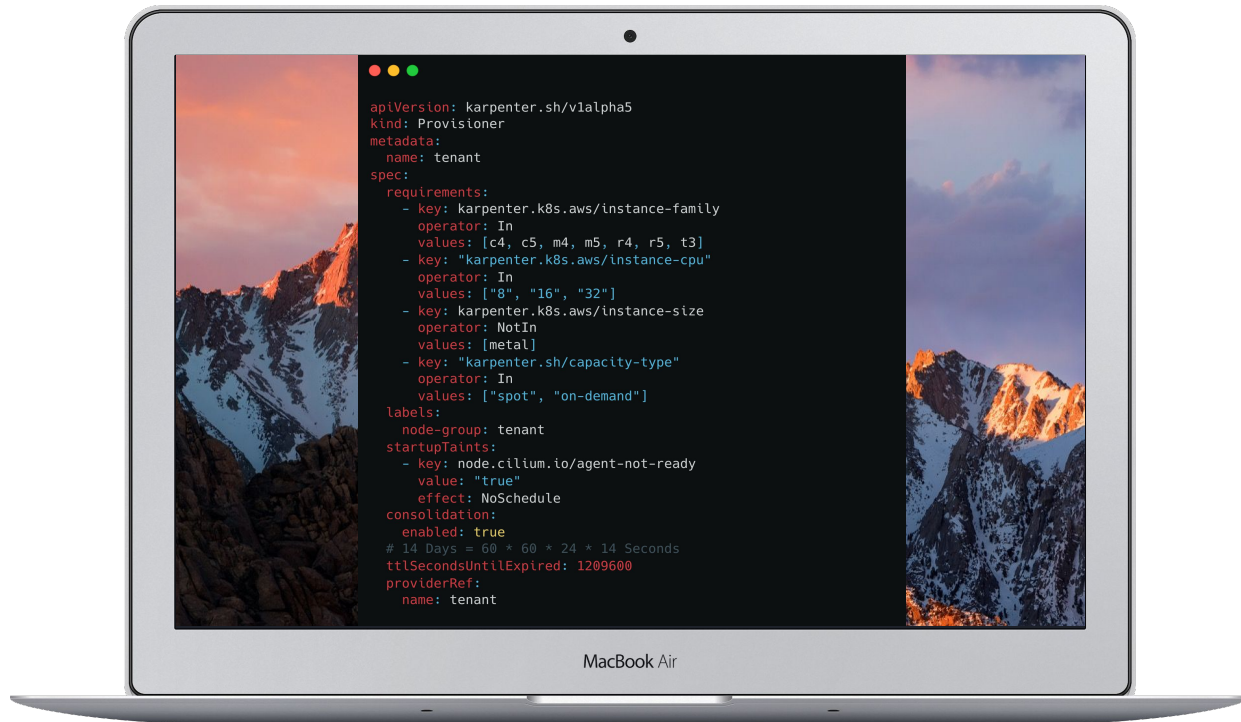
KubeCon



CloudNativeCon

North America 2023

- Groupless autoscaling
- Provisioners
 - NodePools
 - Spot instances vs. on-demand
 - Expiration
 - Consolidation



Reduces cluster cost by deleting or replacing nodes when:

- They're empty
- Their workloads can run on other nodes
- They can be replaced with cheaper nodes

```
apiVersion: karpenter.sh/v1beta1
kind: NodePool
metadata:
  name: default

# Describes the ways to disrupt and replace nodes
disruption:

# Types of nodes to consider for consolidation
consolidationPolicy: WhenUnderutilized | WhenEmpty

# Time to wait after discovering a consolidation decision
consolidateAfter: 30s

# Time a node can live on the cluster before it's removed
expireAfter: 720h
```

Karpenter and Spot Instances



KubeCon



CloudNativeCon

North America 2023

```
func filterUnwantedSpot(instanceTypes []*cloudprovider.InstanceType) []*cloudprovider.InstanceType {
    cheapestOnDemand := math.MaxFloat64

    //first, find the price of our cheapest available on-demand instance type that could support this node
    for _, it := range instanceTypes {
        for _, o := range it.Offerings.Available() {
            if o.CapacityType == v1alpha5.CapacityTypeOnDemand && o.Price < cheapestOnDemand {
                cheapestOnDemand = o.Price
            }
        }
    }

    //Filter out any types where the cheapest offering, which should be spot, is more expensive than the
    //cheapest on-demand instance type
    instanceTypes = lo.Filter(instanceTypes, func(item *cloudprovider.InstanceType, index int) bool {
        available := item.Offerings.Available()
        if len(available) == 0 {
            return false
        }
        return available.Cheapest().Price <= cheapestOnDemand
    })
    return instanceTypes
}

return go(f, seed, [])
}
```



KubeCon



CloudNativeCon

North America 2023

- 1 What we're scaling
- 2 How we're scaling: Karpenter
- 3 Scaling Demo**
- 4 How we're scaling: KEDA
- 5 Lessons Learned

Karpenter Scaling Demo




KubeCon



CloudNativeCon

North America 2023



The New York Times



KubeCon



CloudNativeCon

North America 2023

- 1 What we're scaling
- 2 How we're scaling: Karpenter
- 3 Scaling Demo
- 4 How we're scaling: KEDA**
- 5 Lessons Learned

Kubernetes Event-Driven Autoscaling (KEDA)

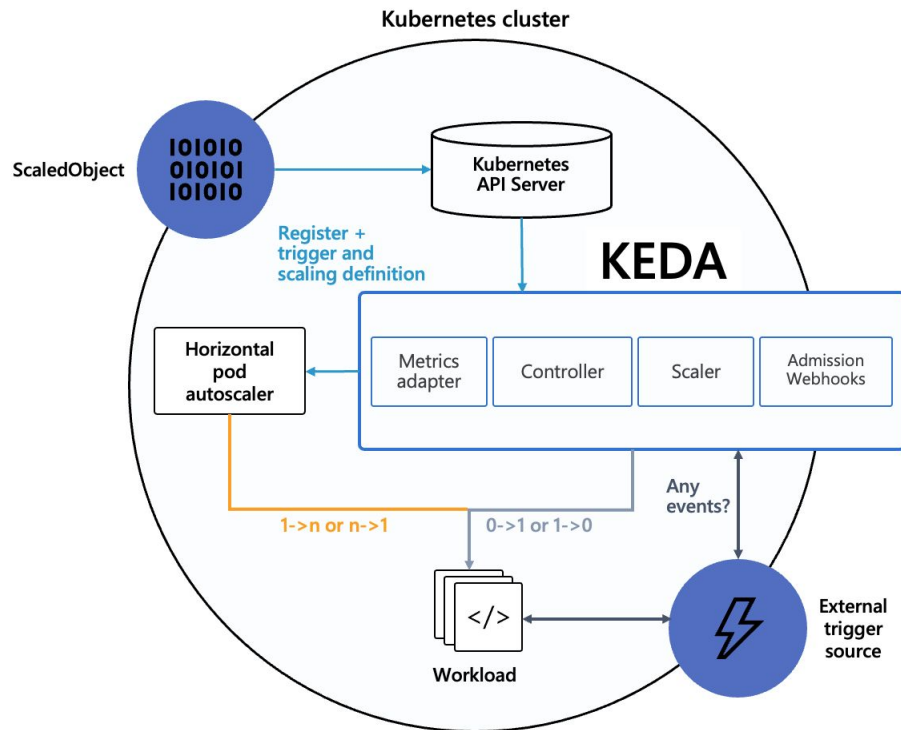


KubeCon



CloudNativeCon

North America 2023



ScaledObject

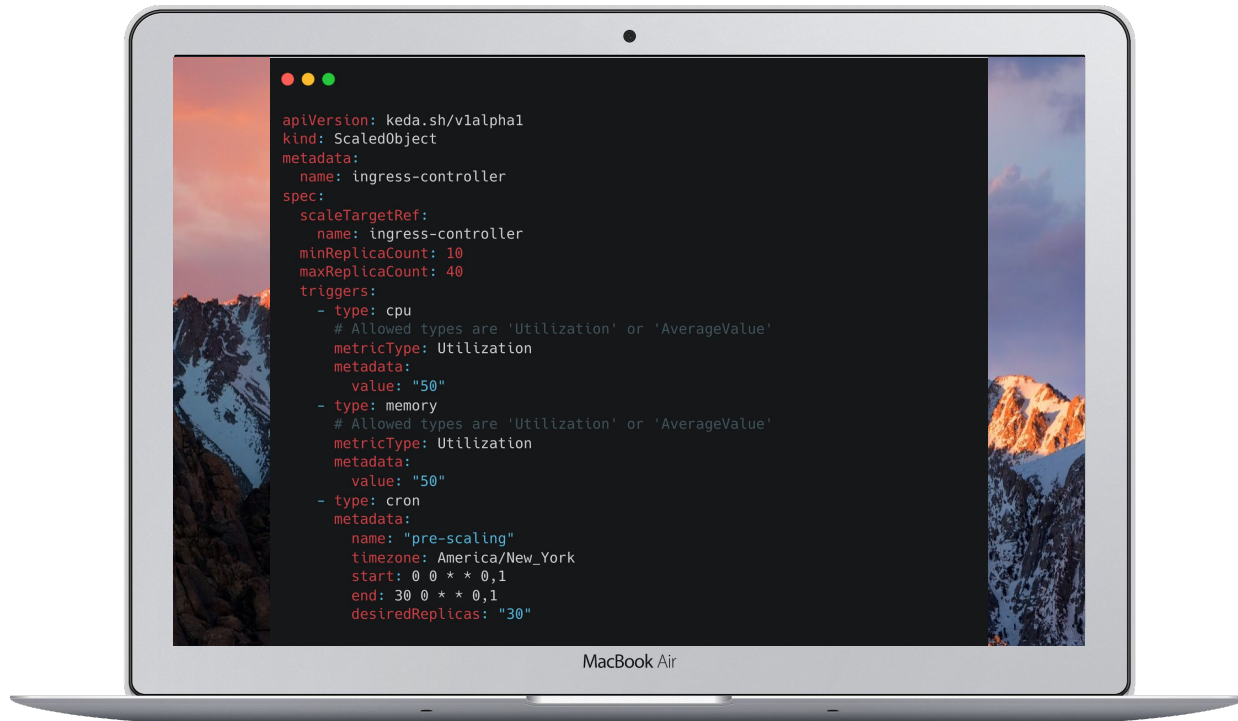



KubeCon



CloudNativeCon

North America 2023



The New York Times

Games Release




KubeCon



CloudNativeCon

North America 2023



The New York Times

Exploring Other Events




KubeCon



CloudNativeCon

North America 2023



The New York Times



```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: my-deployment-scaling
  namespace: my-team
spec:
  scaleTargetRef:
    name: my-deployment
  triggers:
    - type: external-push
      metadata:
        scalerAddress: bna-scaler.svc.local:9090
        supportedPodRPS: 2000
        eventWindow: 10m
```

External Scaler




KubeCon



CloudNativeCon

North America 2023



The New York Times



KubeCon




CloudNativeCon

North America 2023

- 1 What we're scaling
- 2 How we're scaling: Karpenter
- 3 Scaling Demo
- 4 How we're scaling: KEDA
- 5 Lessons Learned**

25% decrease in cost

- Less overprovisioning
- Node consolidation
- Spot instances
- More instance types



Further improvements

- More flexibility with unified HTTP ingress instances
- External BNA scaler
- KEDA as a service for tenants

Shared Platform: Centralization & Standardization



KubeCon



CloudNativeCon

North America 2023

- First place problems surface
- Establishes baseline
- Locate outliers
 - Why does one service have significantly higher RPS?
- Unified HTTP ingress traffic spikes
 - Some tenants making significantly more requests
 - As we onboard more tenants, will likely see similar thing
- We scale the architecture, tenants don't need to

THANK YOU



KubeCon



CloudNativeCon

North America 2023

New York Times Talks @ KubeCon NA 2023

ArgoCon - Monday, November 6

The ArgoCD AppProject - What Is a Project and How to Power Your Multi-Tenant Security

Luke Philips & Serhiy Martynenko

Multi-TenancyCon - Monday, November 6

What We Learned Designing & Securing a Multi-Tenant Developer Platform at The New York Times

Ahmed Bebars & David Grizzanti

The New York Times

Questions?



KubeCon



CloudNativeCon

North America 2023



**Please scan the QR Code above
to leave feedback on this session**

The New York Times