



# **Self-Service Stream Processing Platform on Kubernetes at Apple**

Chenya Zhang

KubeCon 2023 | Apple Inc. | November 7th

# Agenda

Moving to the Cloud and Kubernetes

Adopting Apache Flink for Real-Time Data Processing

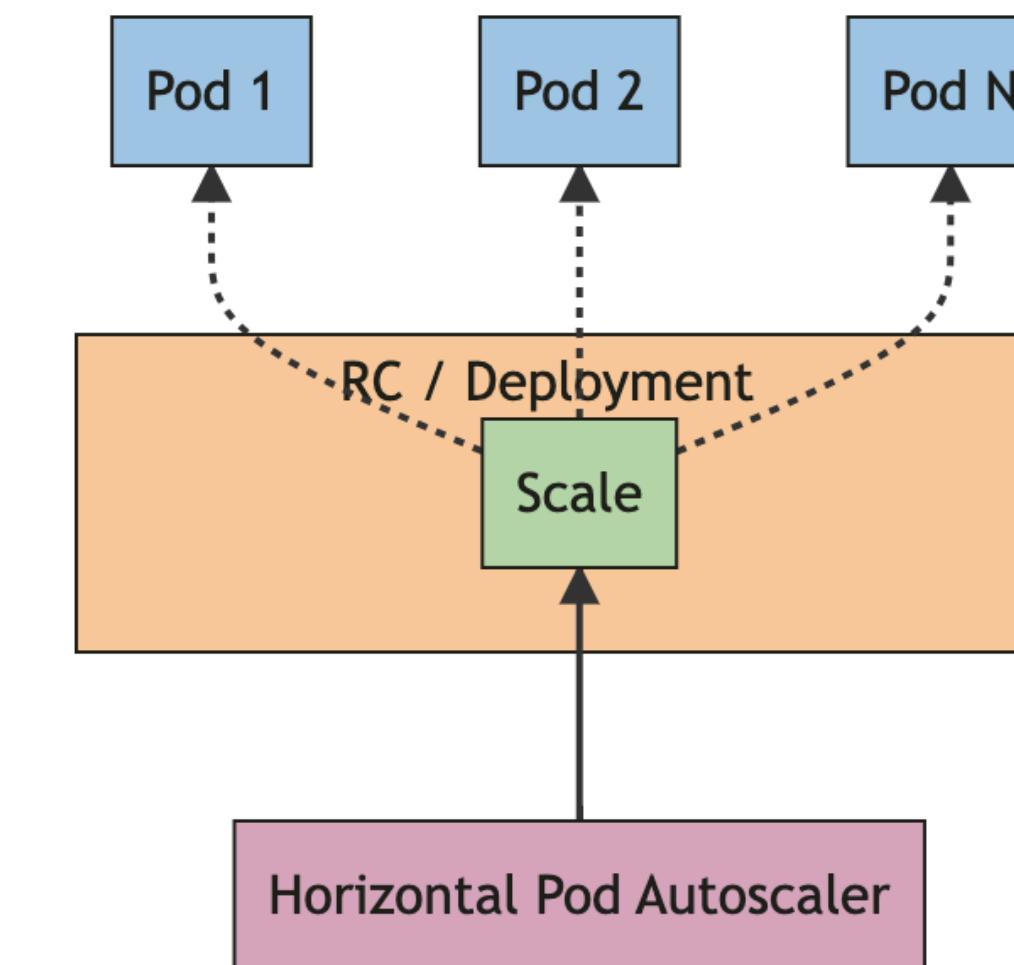
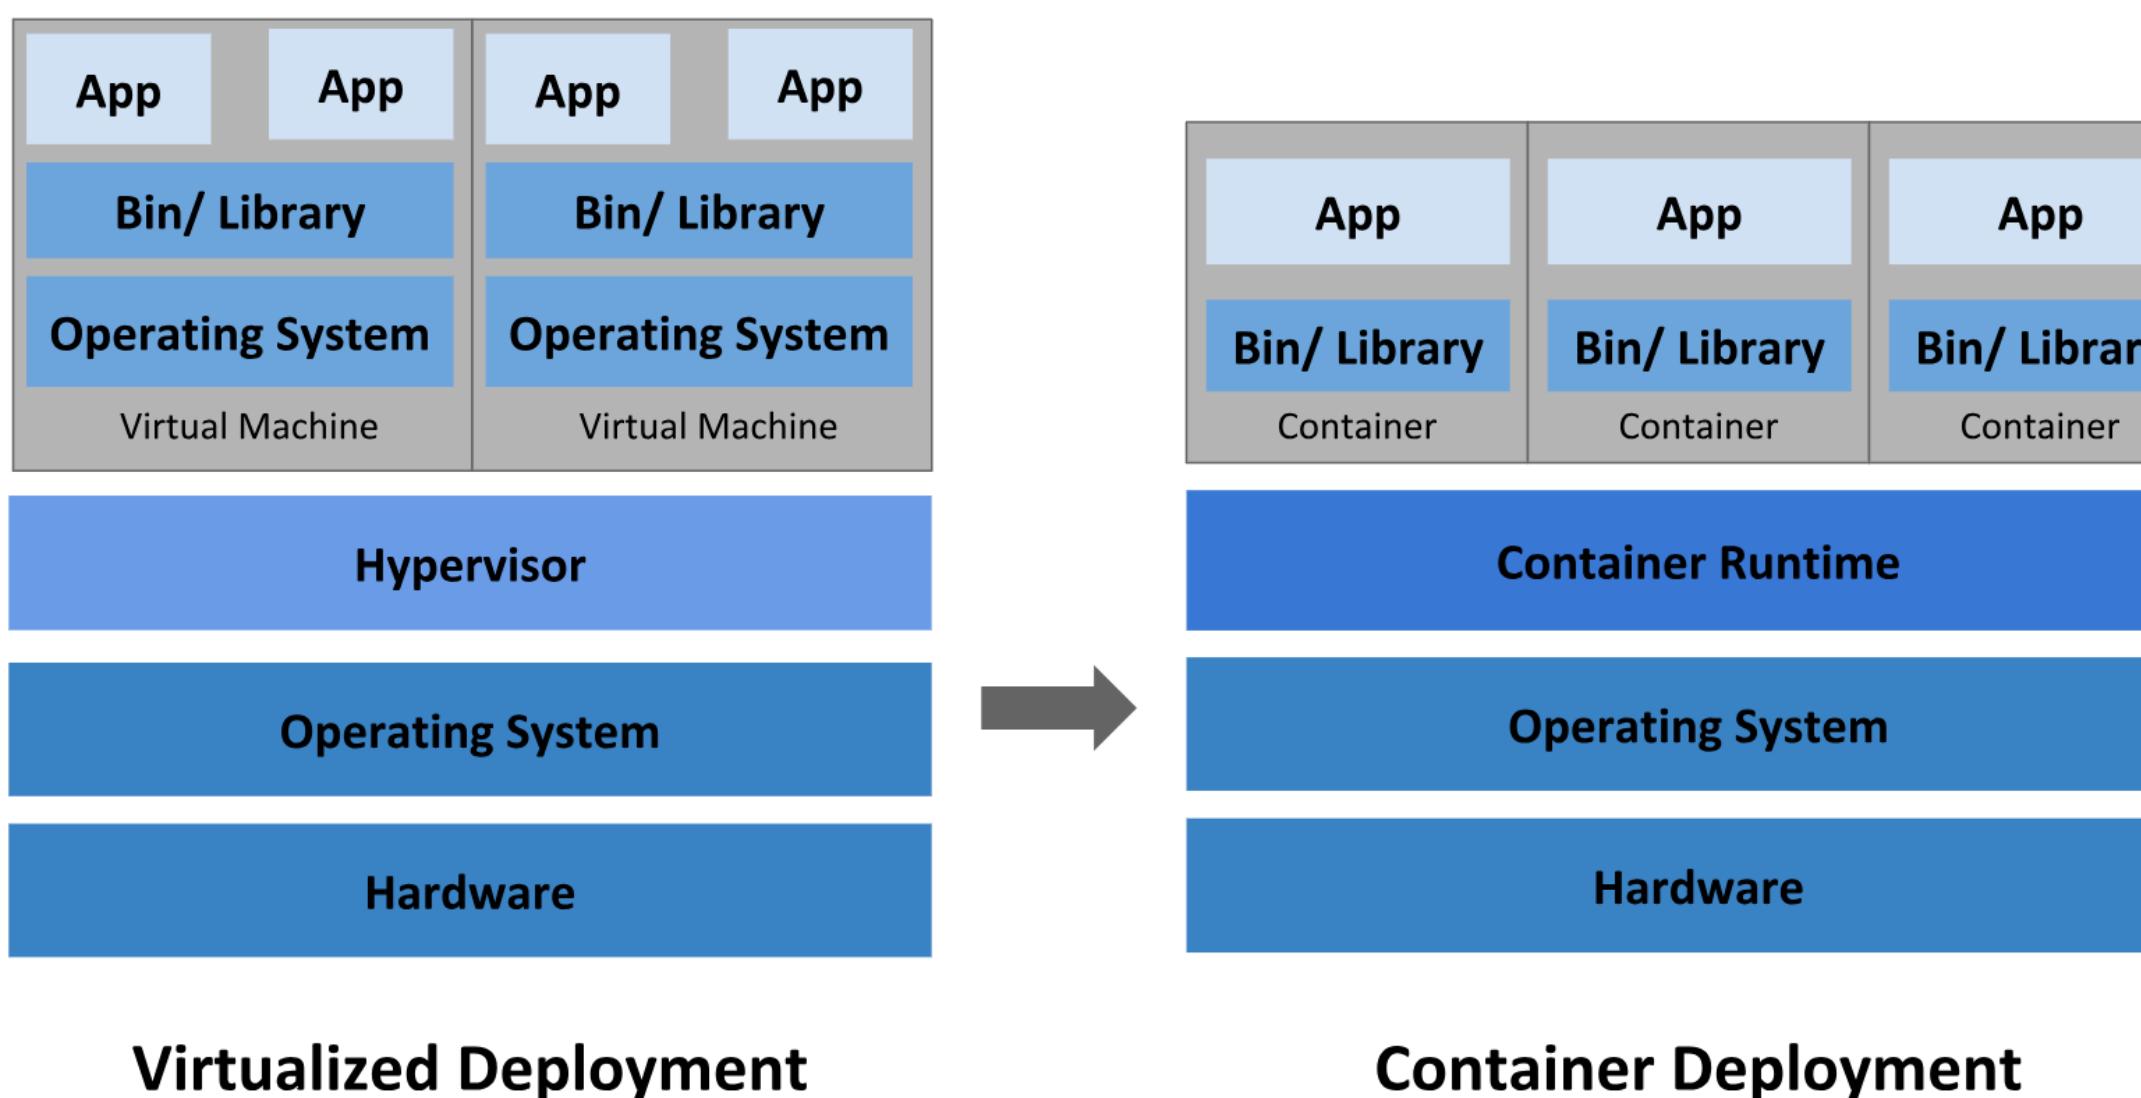
Challenges at Production Scale

Introducing Flink Control Plane on Kubernetes

Solutions on Cost Efficiency and Automation

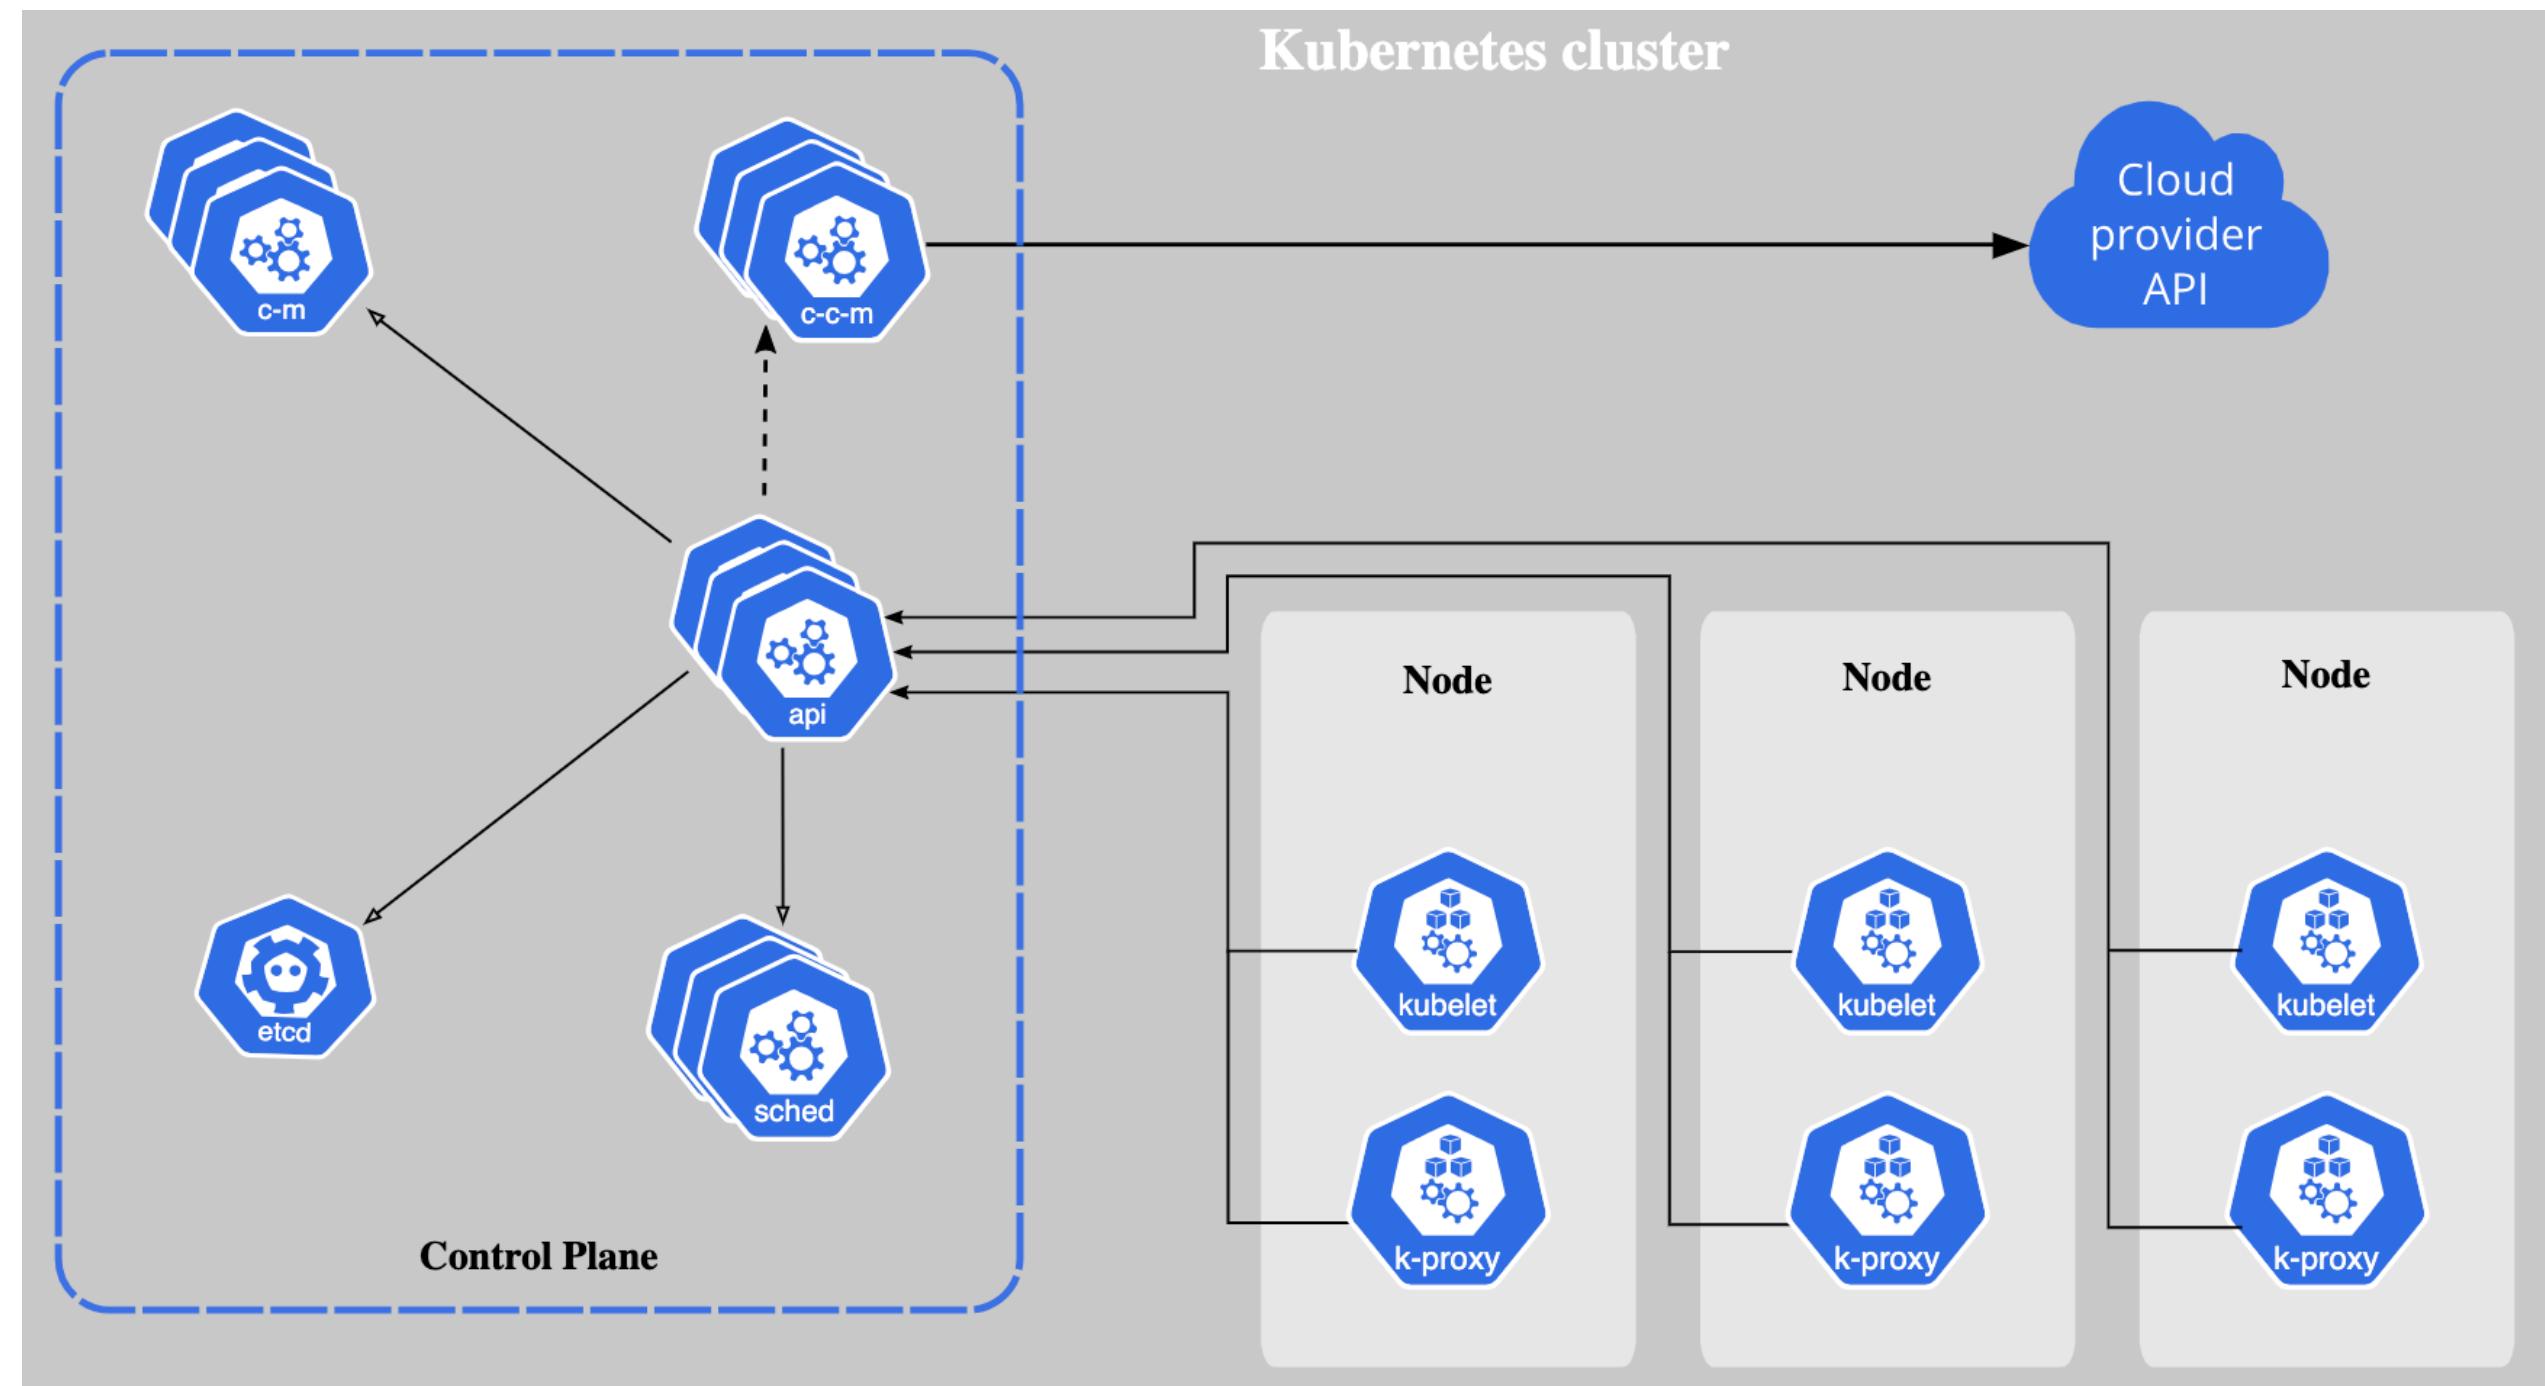
# Moving to the Cloud and Kubernetes

- Automatic provision based on real-time demands
- Physical abstraction via containerized infrastructure
- Flexible interoperation across multi-cloud environments



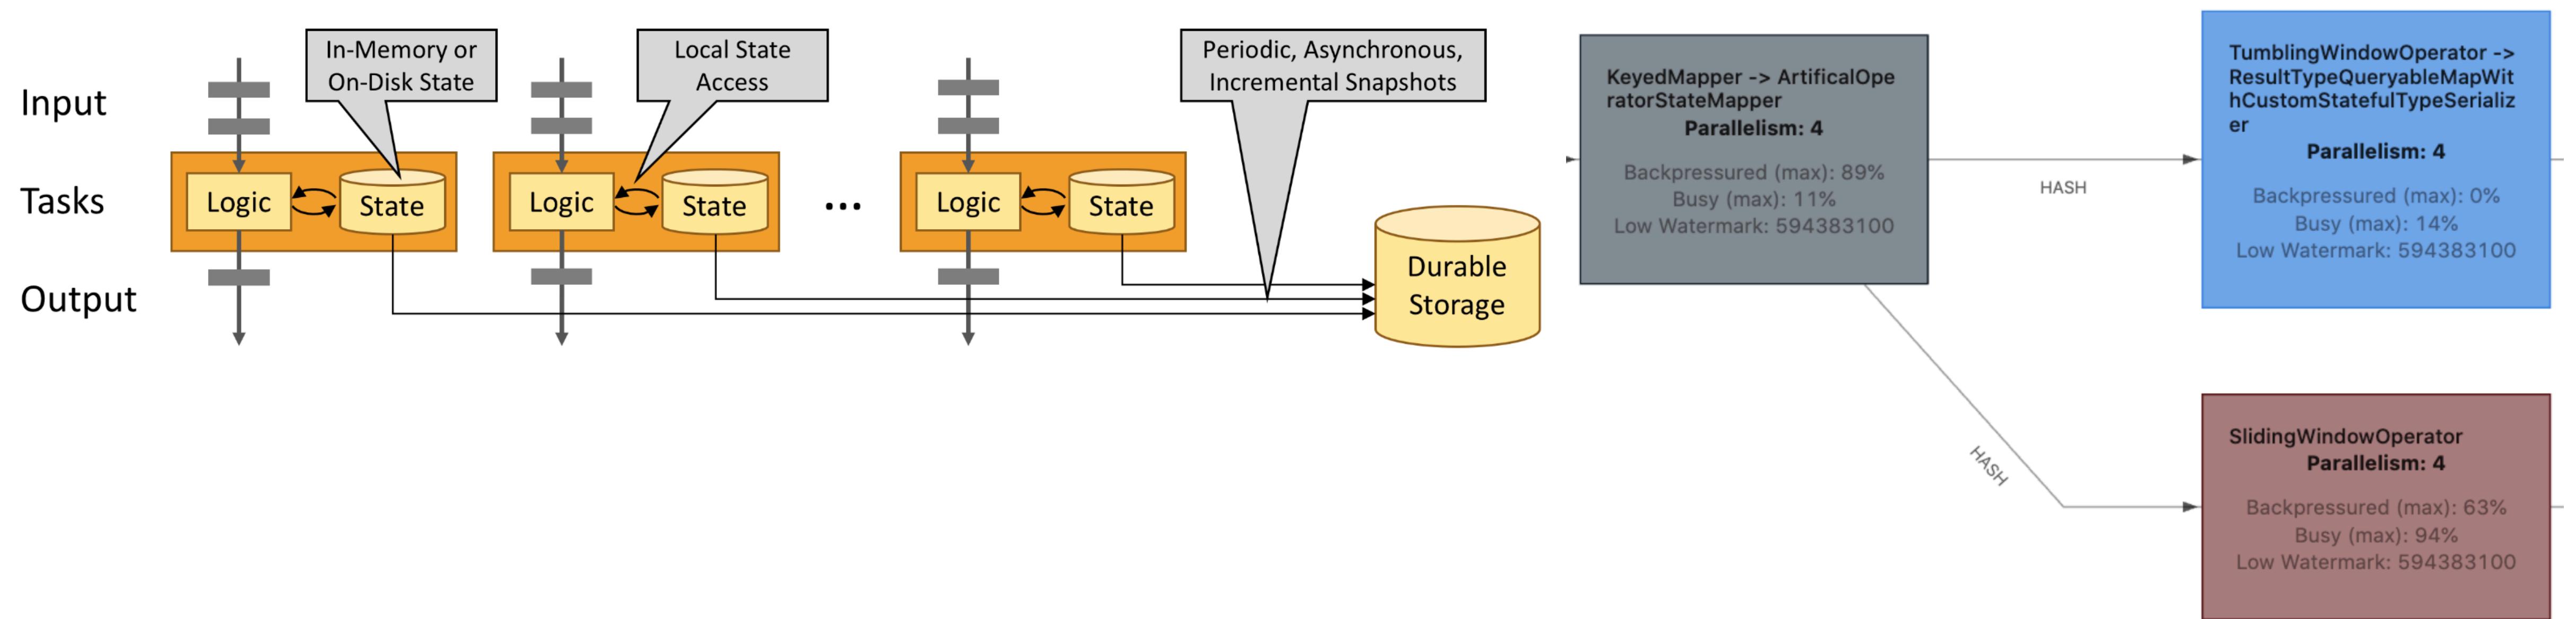
# Moving to the Cloud and Kubernetes

- Auto-adjustment for system consistency and resiliency
- Adaptability enabled by a modularized architecture



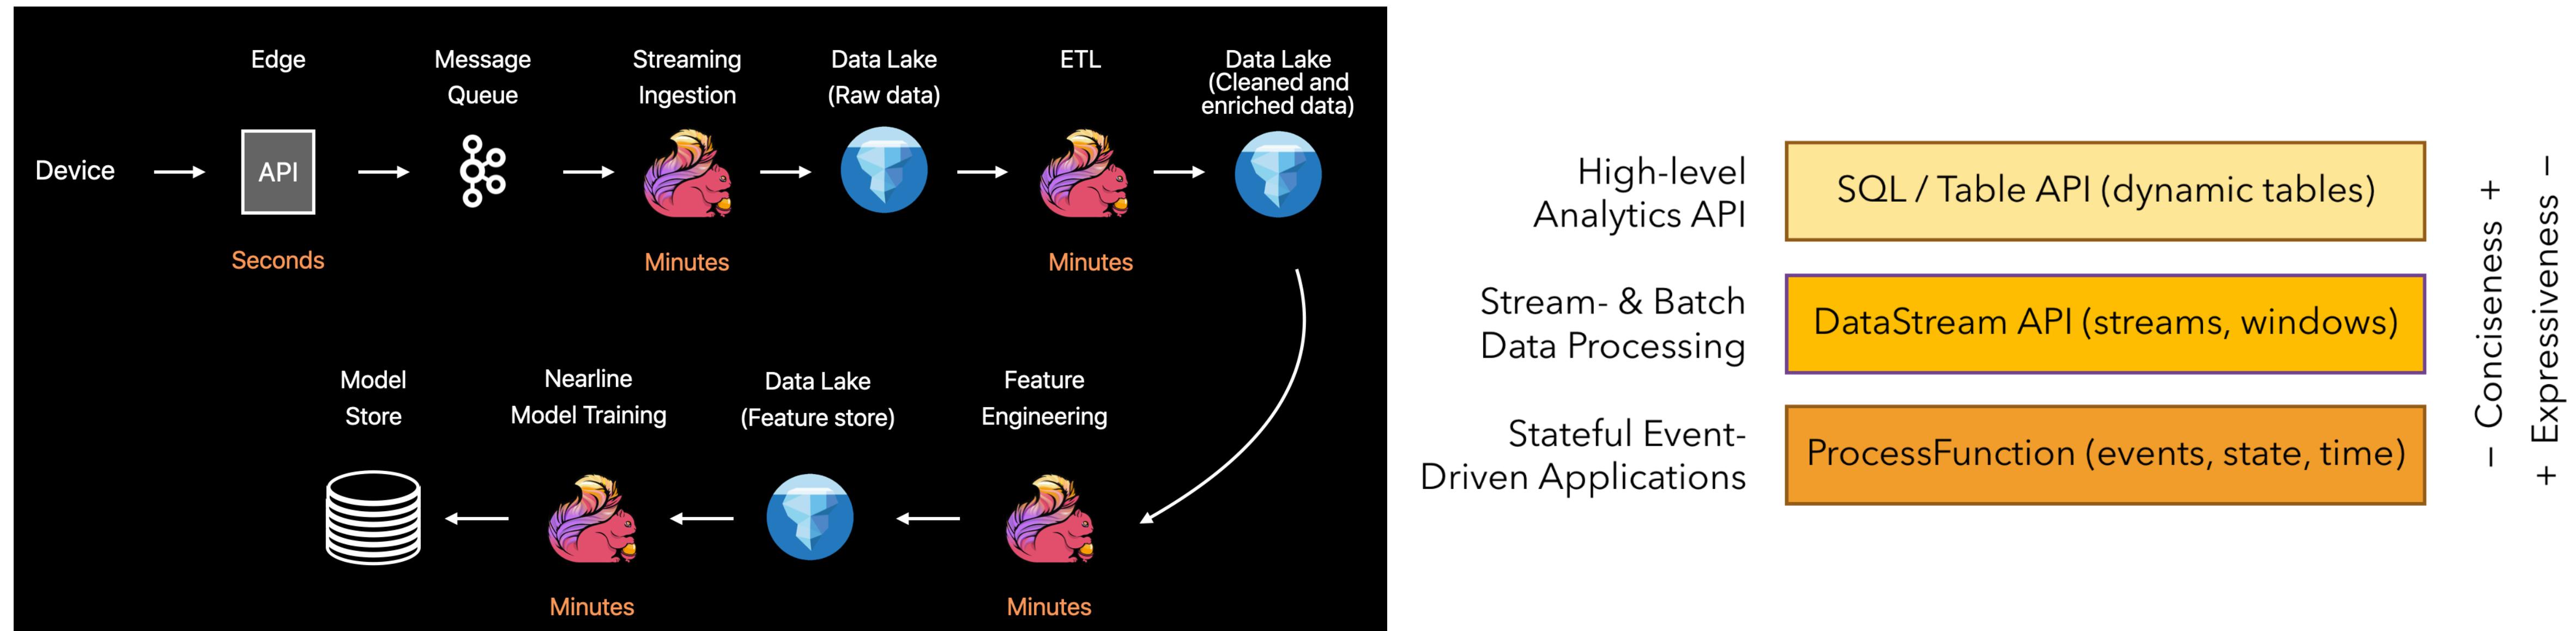
# Adopting Apache Flink

- In-memory pipelined execution with late data and back pressure handling
- Consistent snapshot to preserve accuracy and resilience of TB-state scale



# Adopting Apache Flink

- Built-in connectors for Kafka/Iceberg/... to simplify system integration
- Layered APIs with ProcessFunction and advanced windowing mechanism

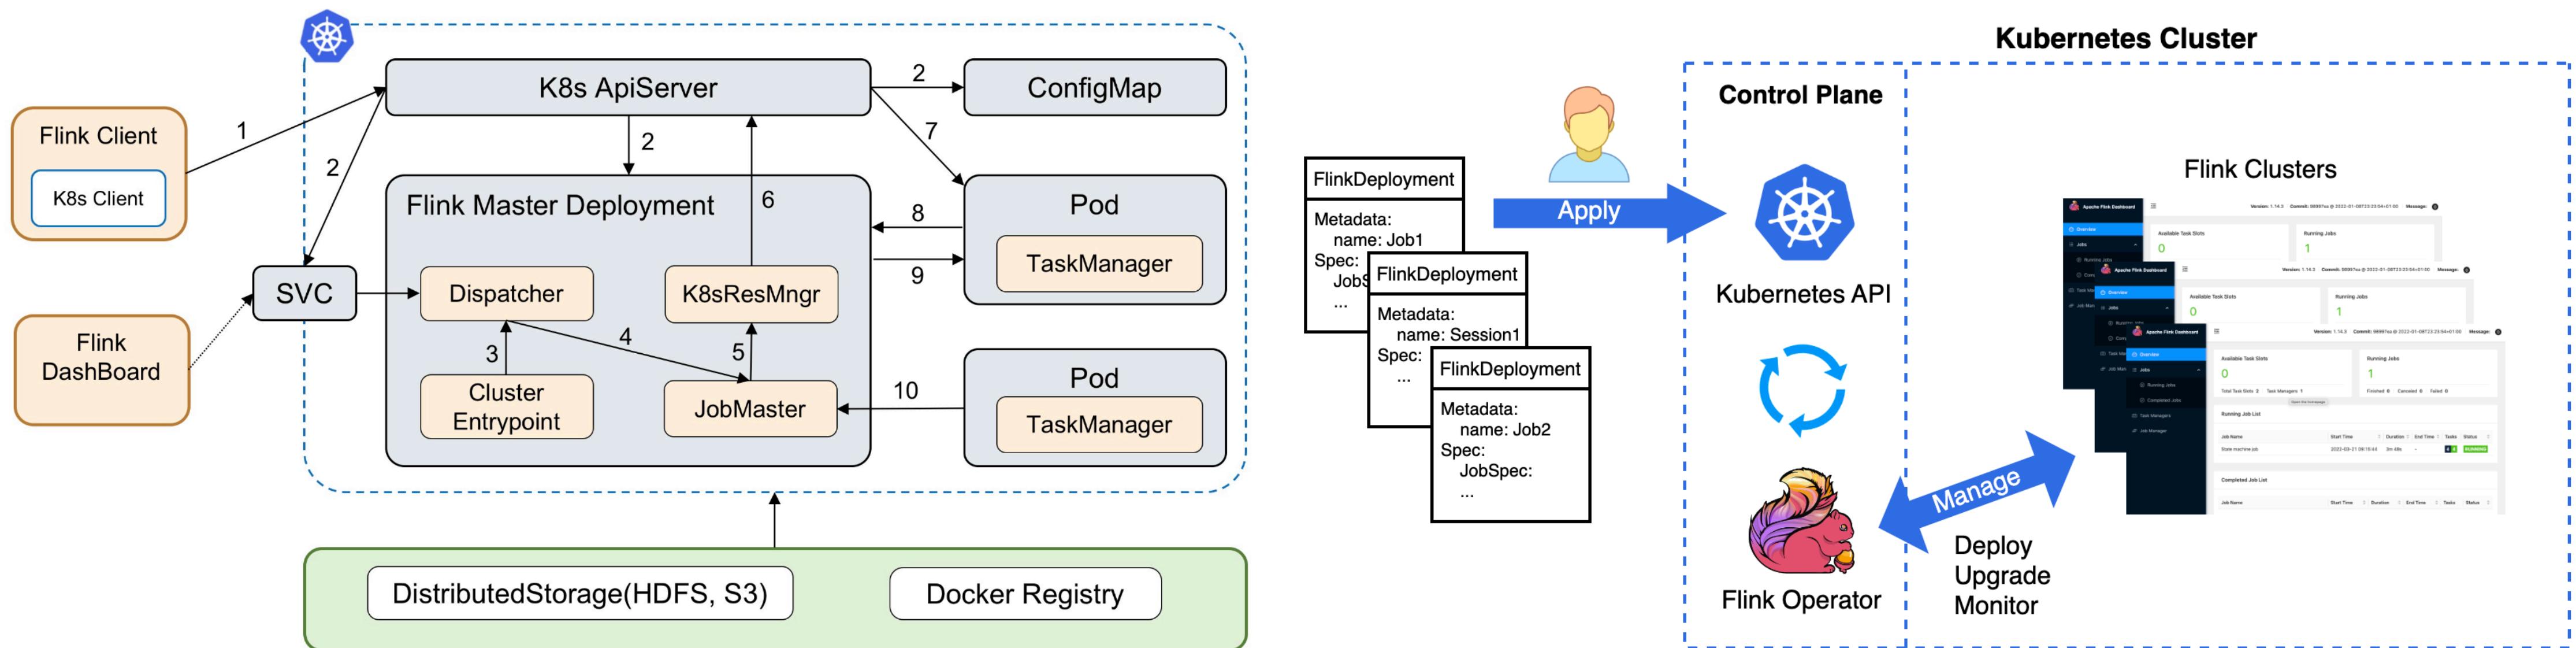


# Challenges at Production Scale

- Need an automated approach to deploy Flink in Kubernetes mode
- Miss a control plane to scale beyond a single Kubernetes cluster
- Self-onboard, observe, and troubleshoot without admins in the loop
- Bulk operate on a large amount of deployments by users and admins
- Automatically scale stateless and stateful apps per incoming traffic
- Manage resource availability and cluster upgrade in the cloud

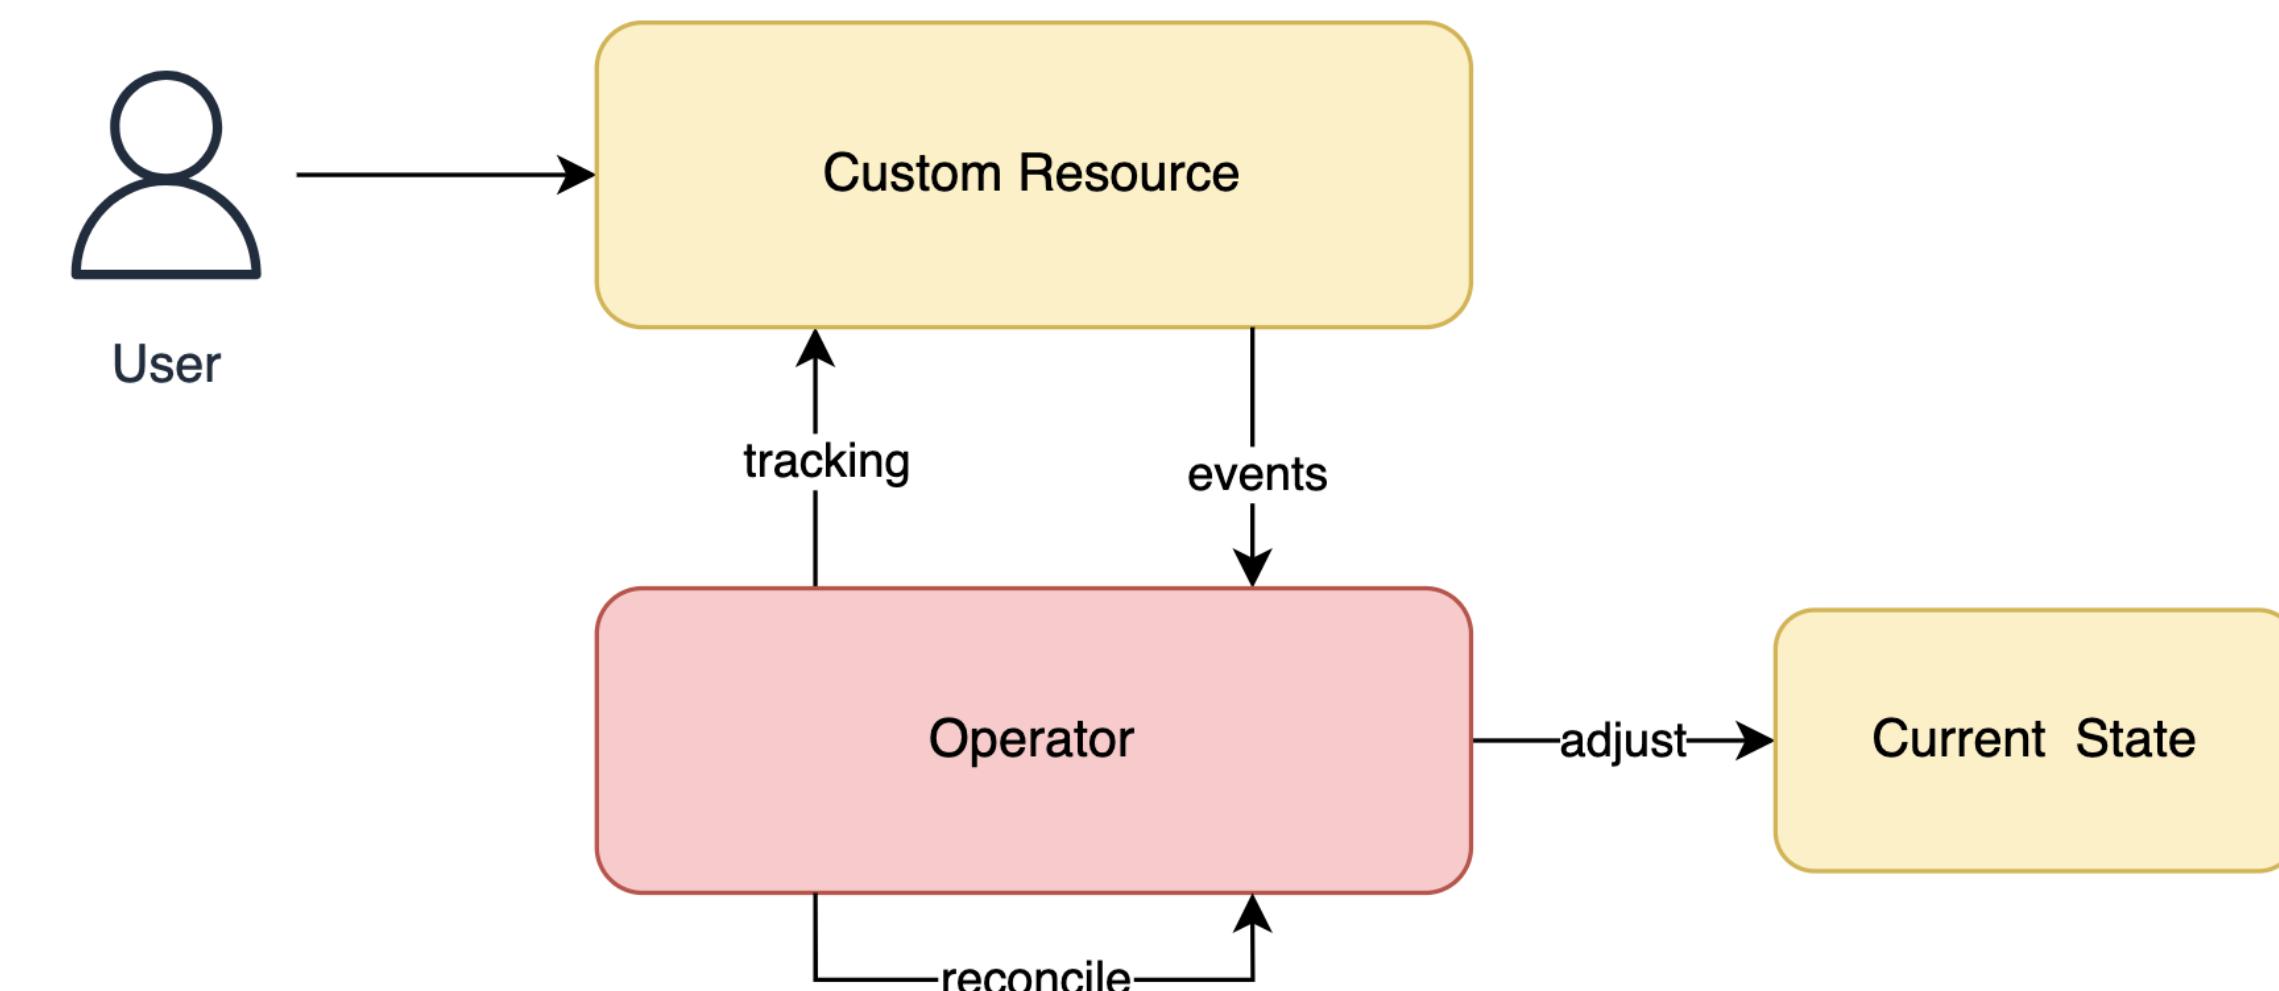
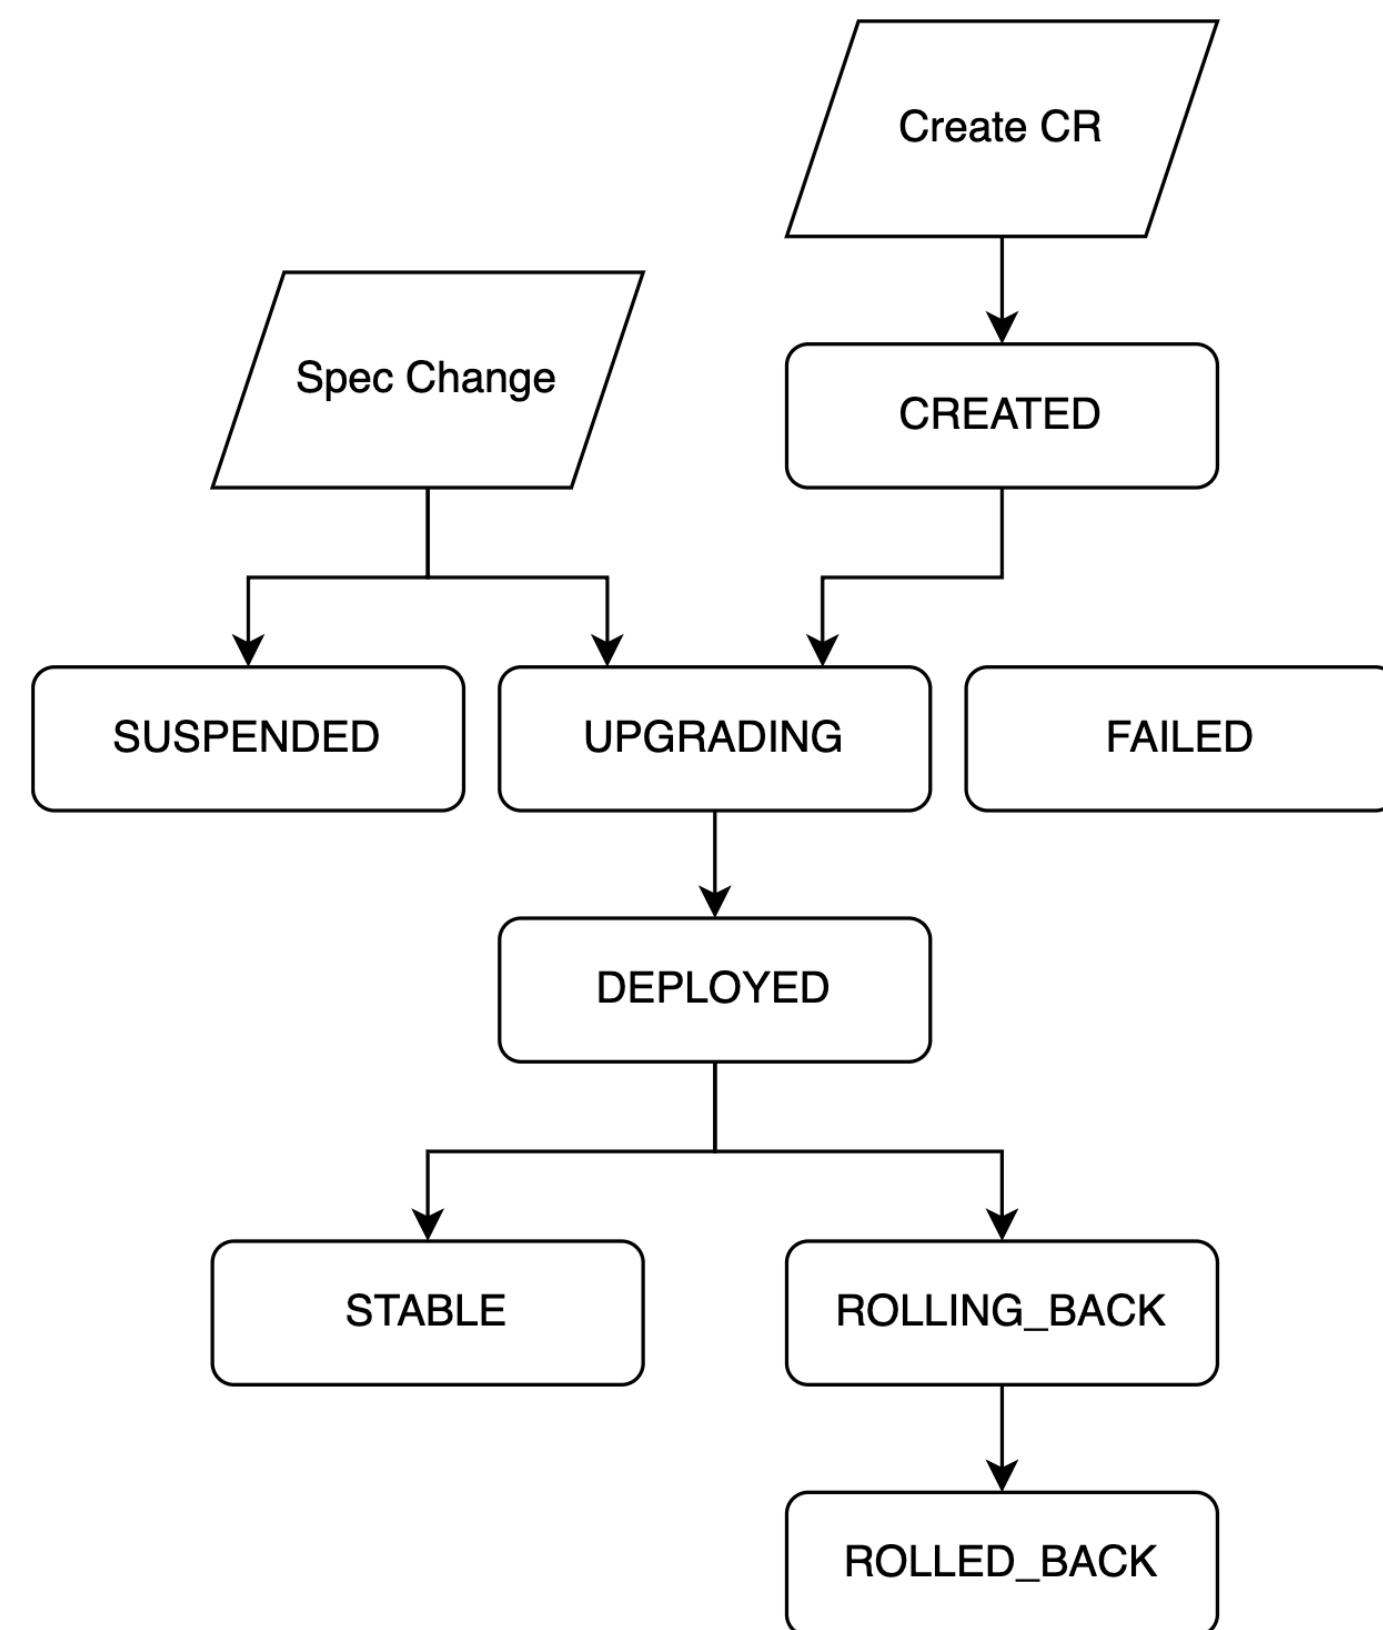
# Deploy Flink on Kubernetes

- Self-contained Flink native K8s support with embedded K8s client
- K8s Custom Resource and Operator to manage application lifecycle
- Automation and abstraction of application upgrade & failure recovery



# Deploy Flink on Kubernetes

- FLIP-212: Introduce Flink Kubernetes Operator
- Reconcile any required changes and execute upgrades

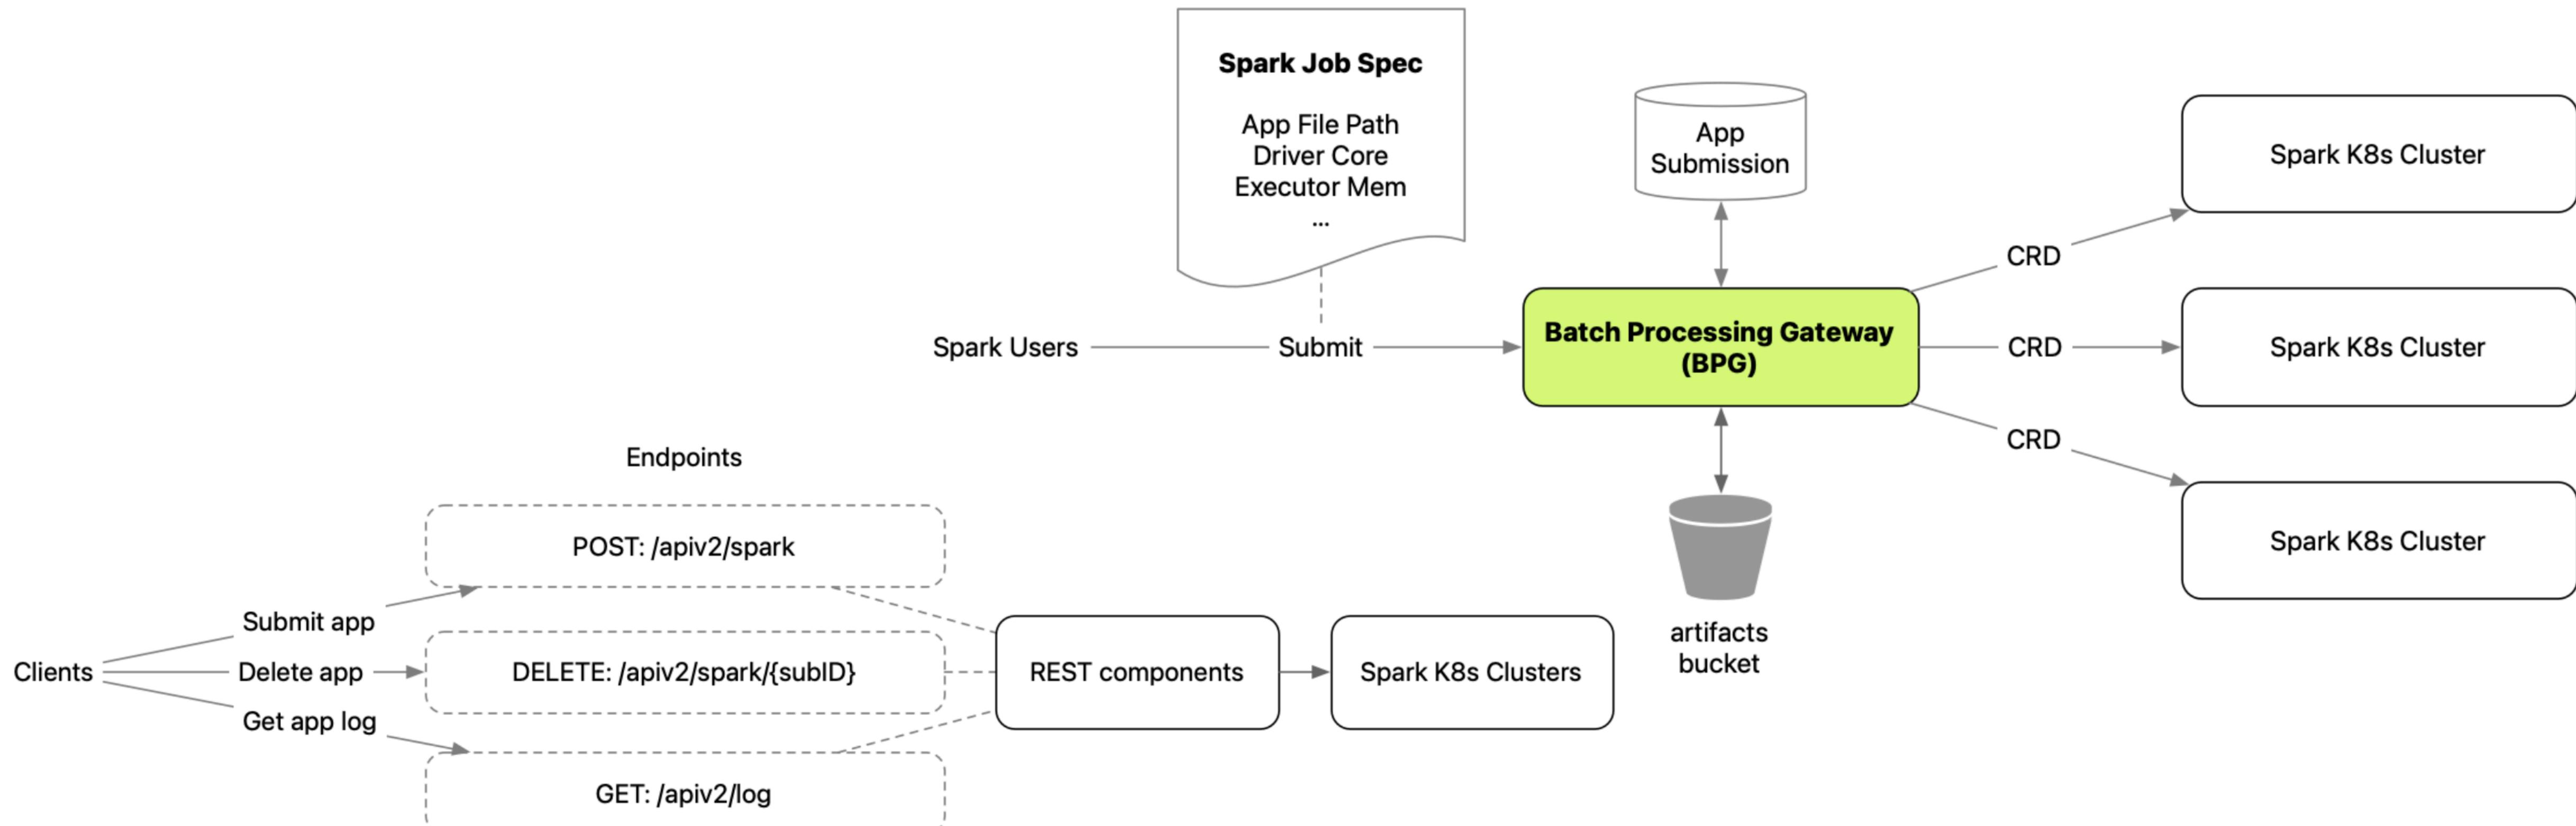


## What If ...

- Kubernetes cluster or Flink K8s Operator fails and goes down?
- User group need to map beyond one Kubernetes namespace?
- Authentication and authorization need to meet extra requirements?
- Deployment information need to be persisted for audit or recovery?
- Admins need to operate across multiple clusters and accounts?
- Other service or platform need to integrate for stream processing?

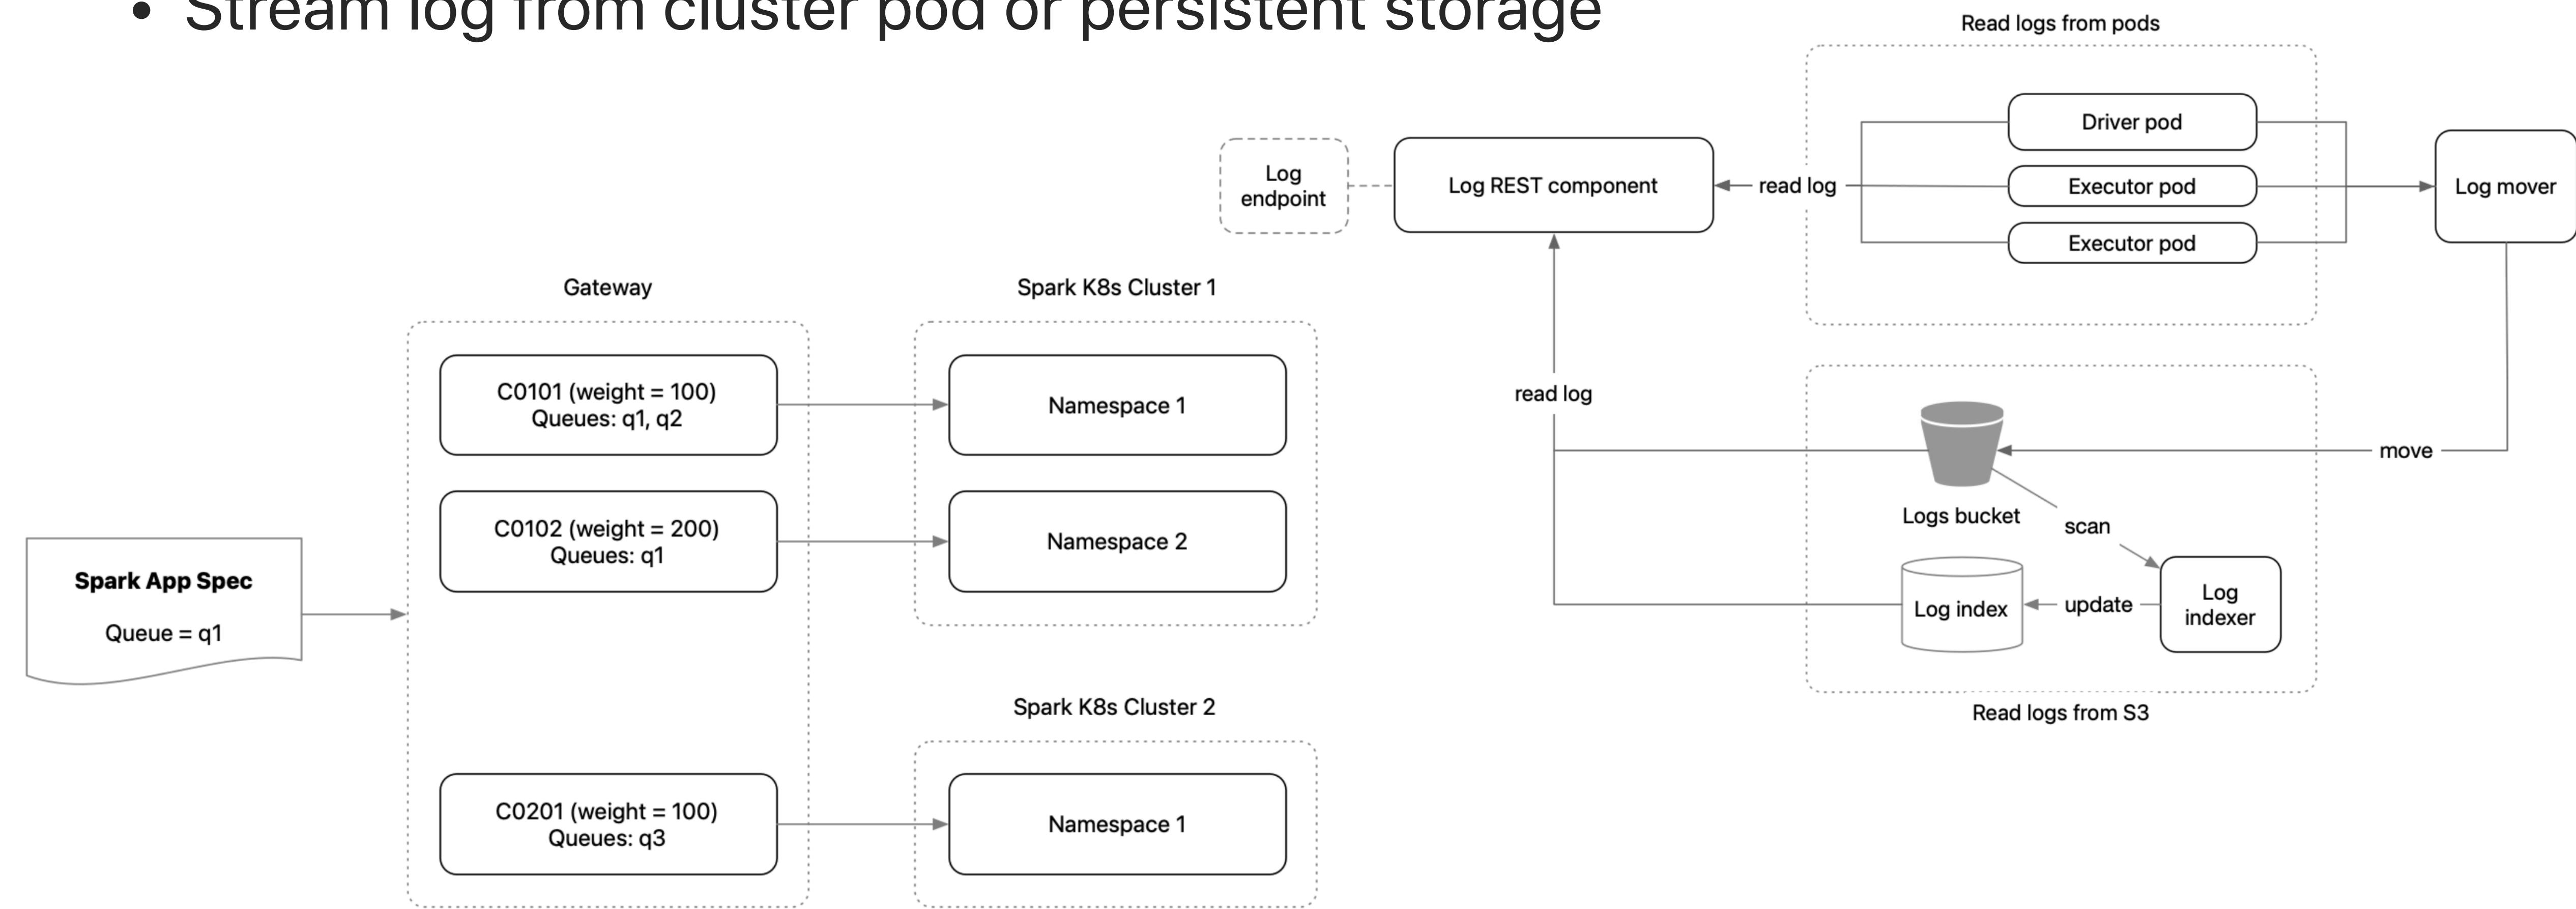
# Lessons from BPG

- Apple's batch processing gateway - REST APIs exposed to clients
- Translate user requests to CRDs supported by Spark K8s operator

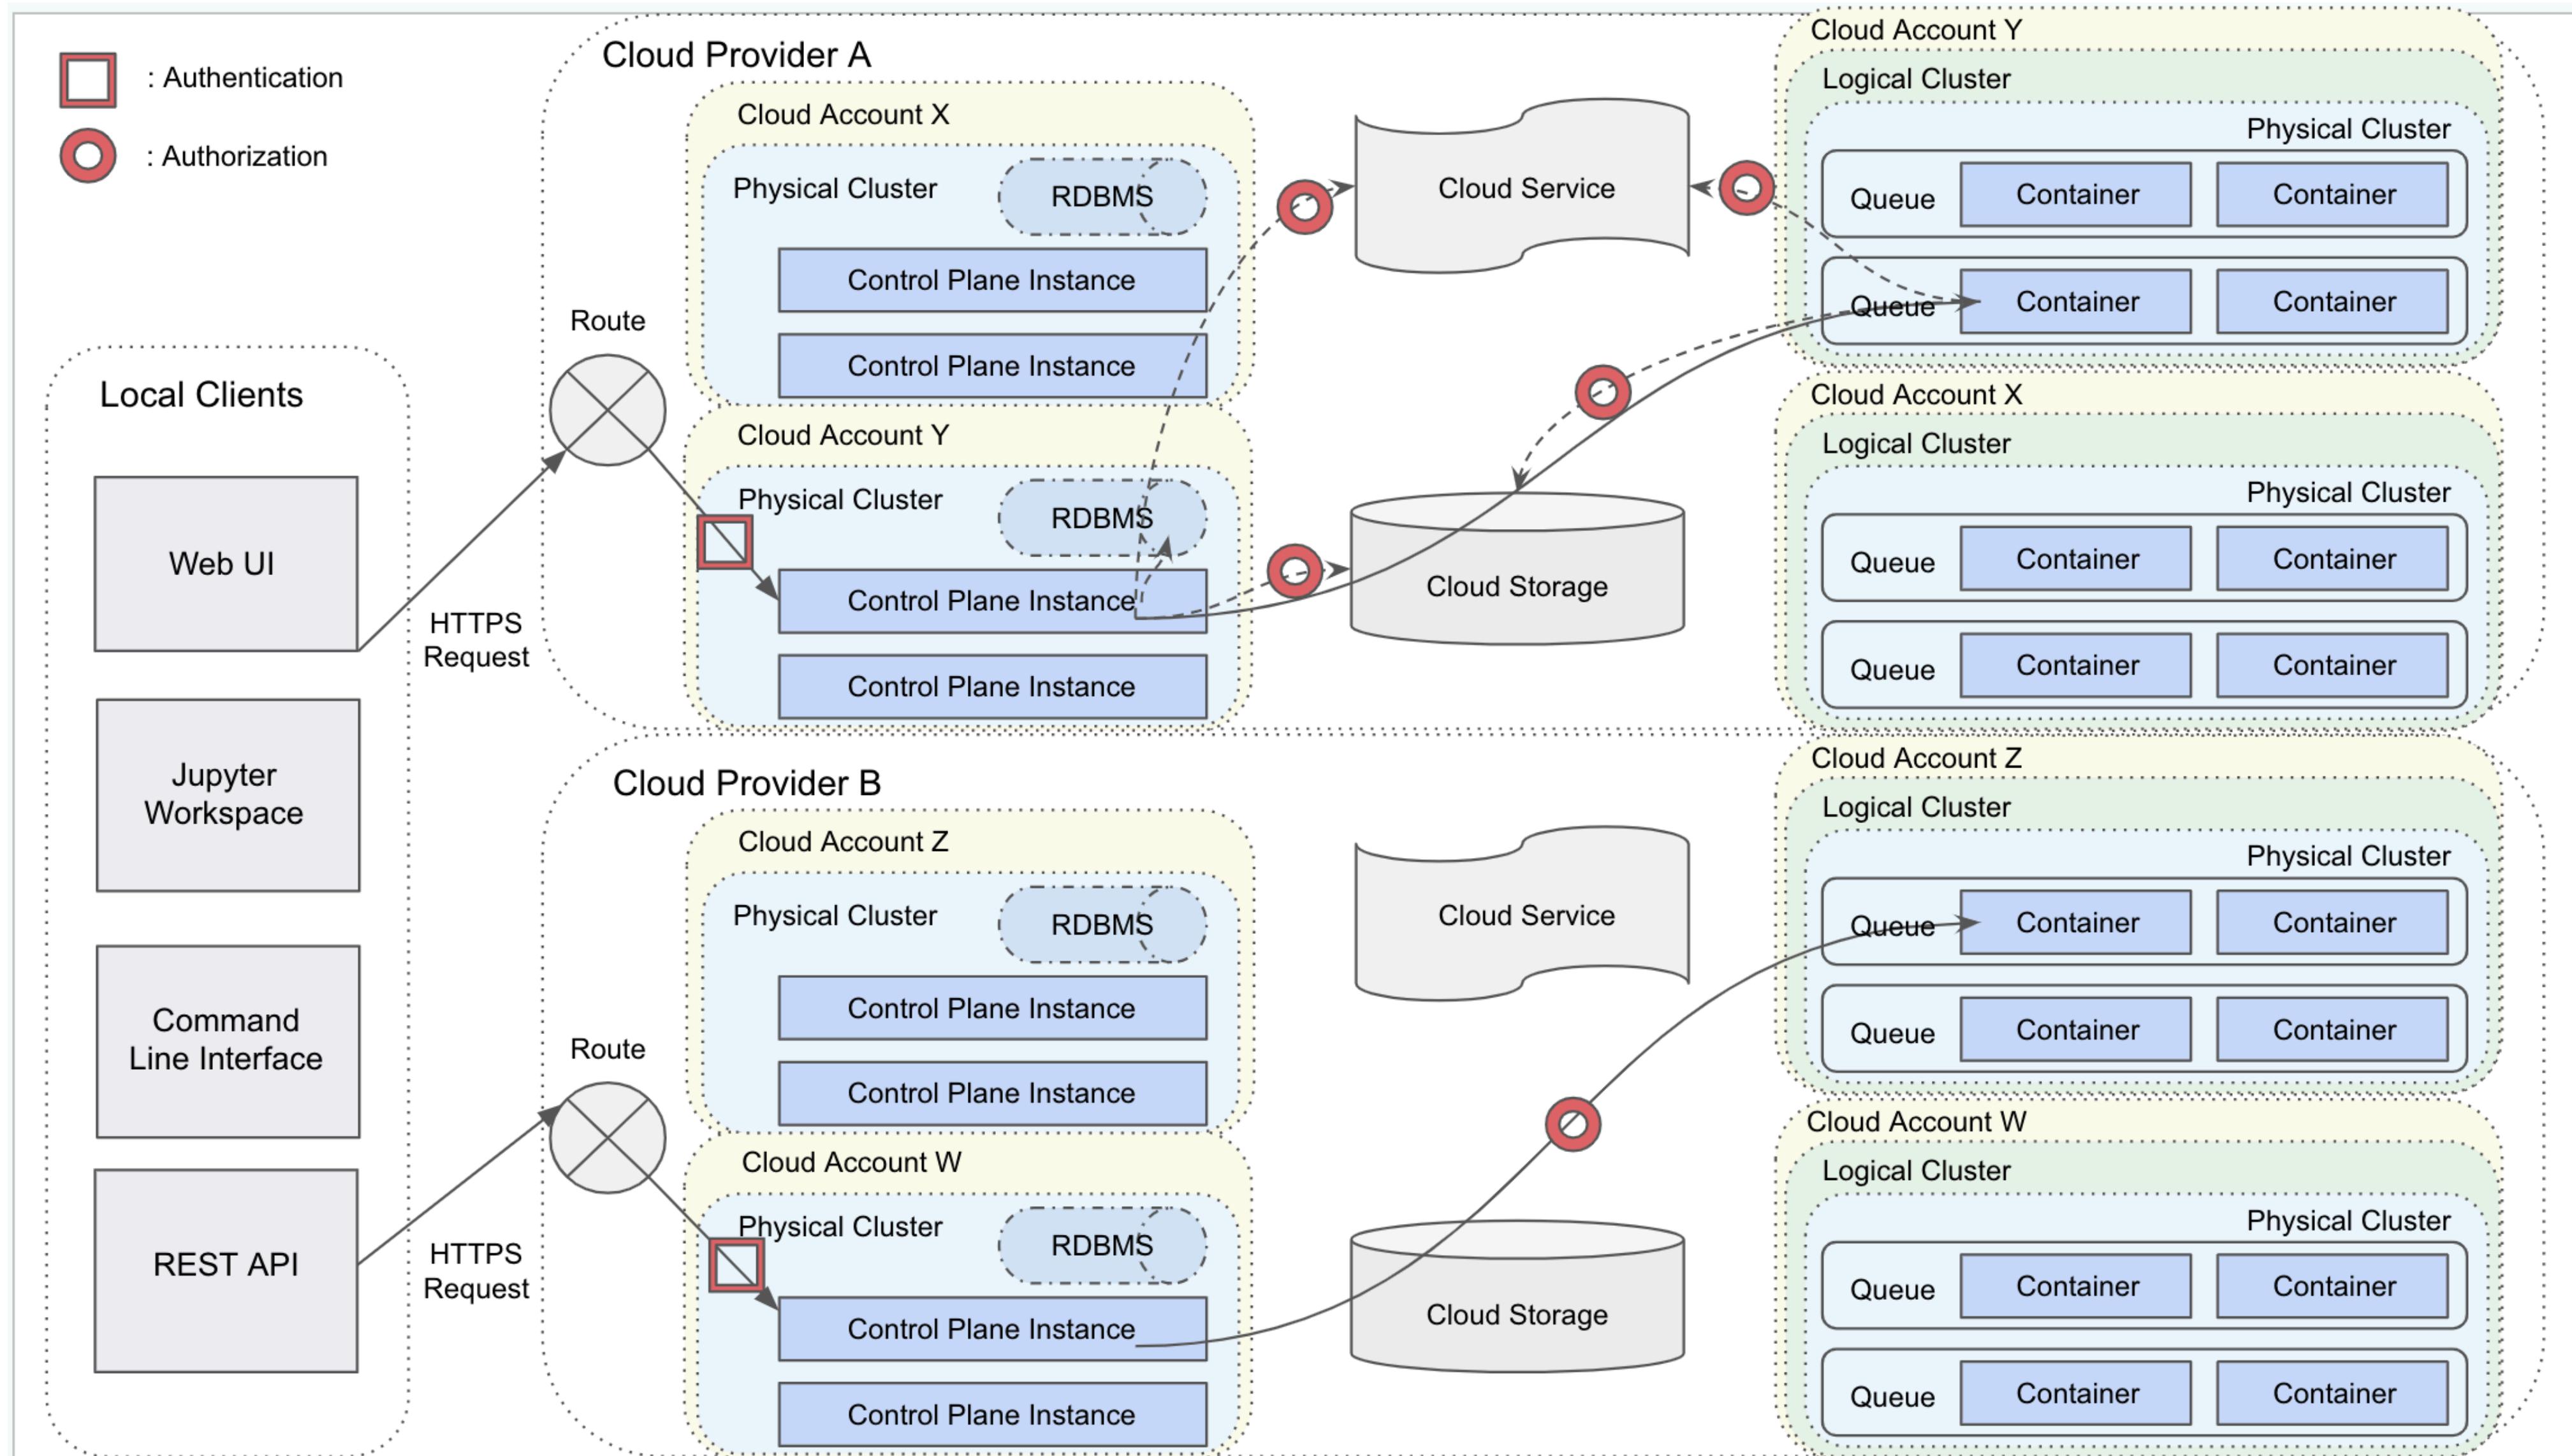


# Lessons from BPG

- Route requests based on queues and weights
- Stream log from cluster pod or persistent storage

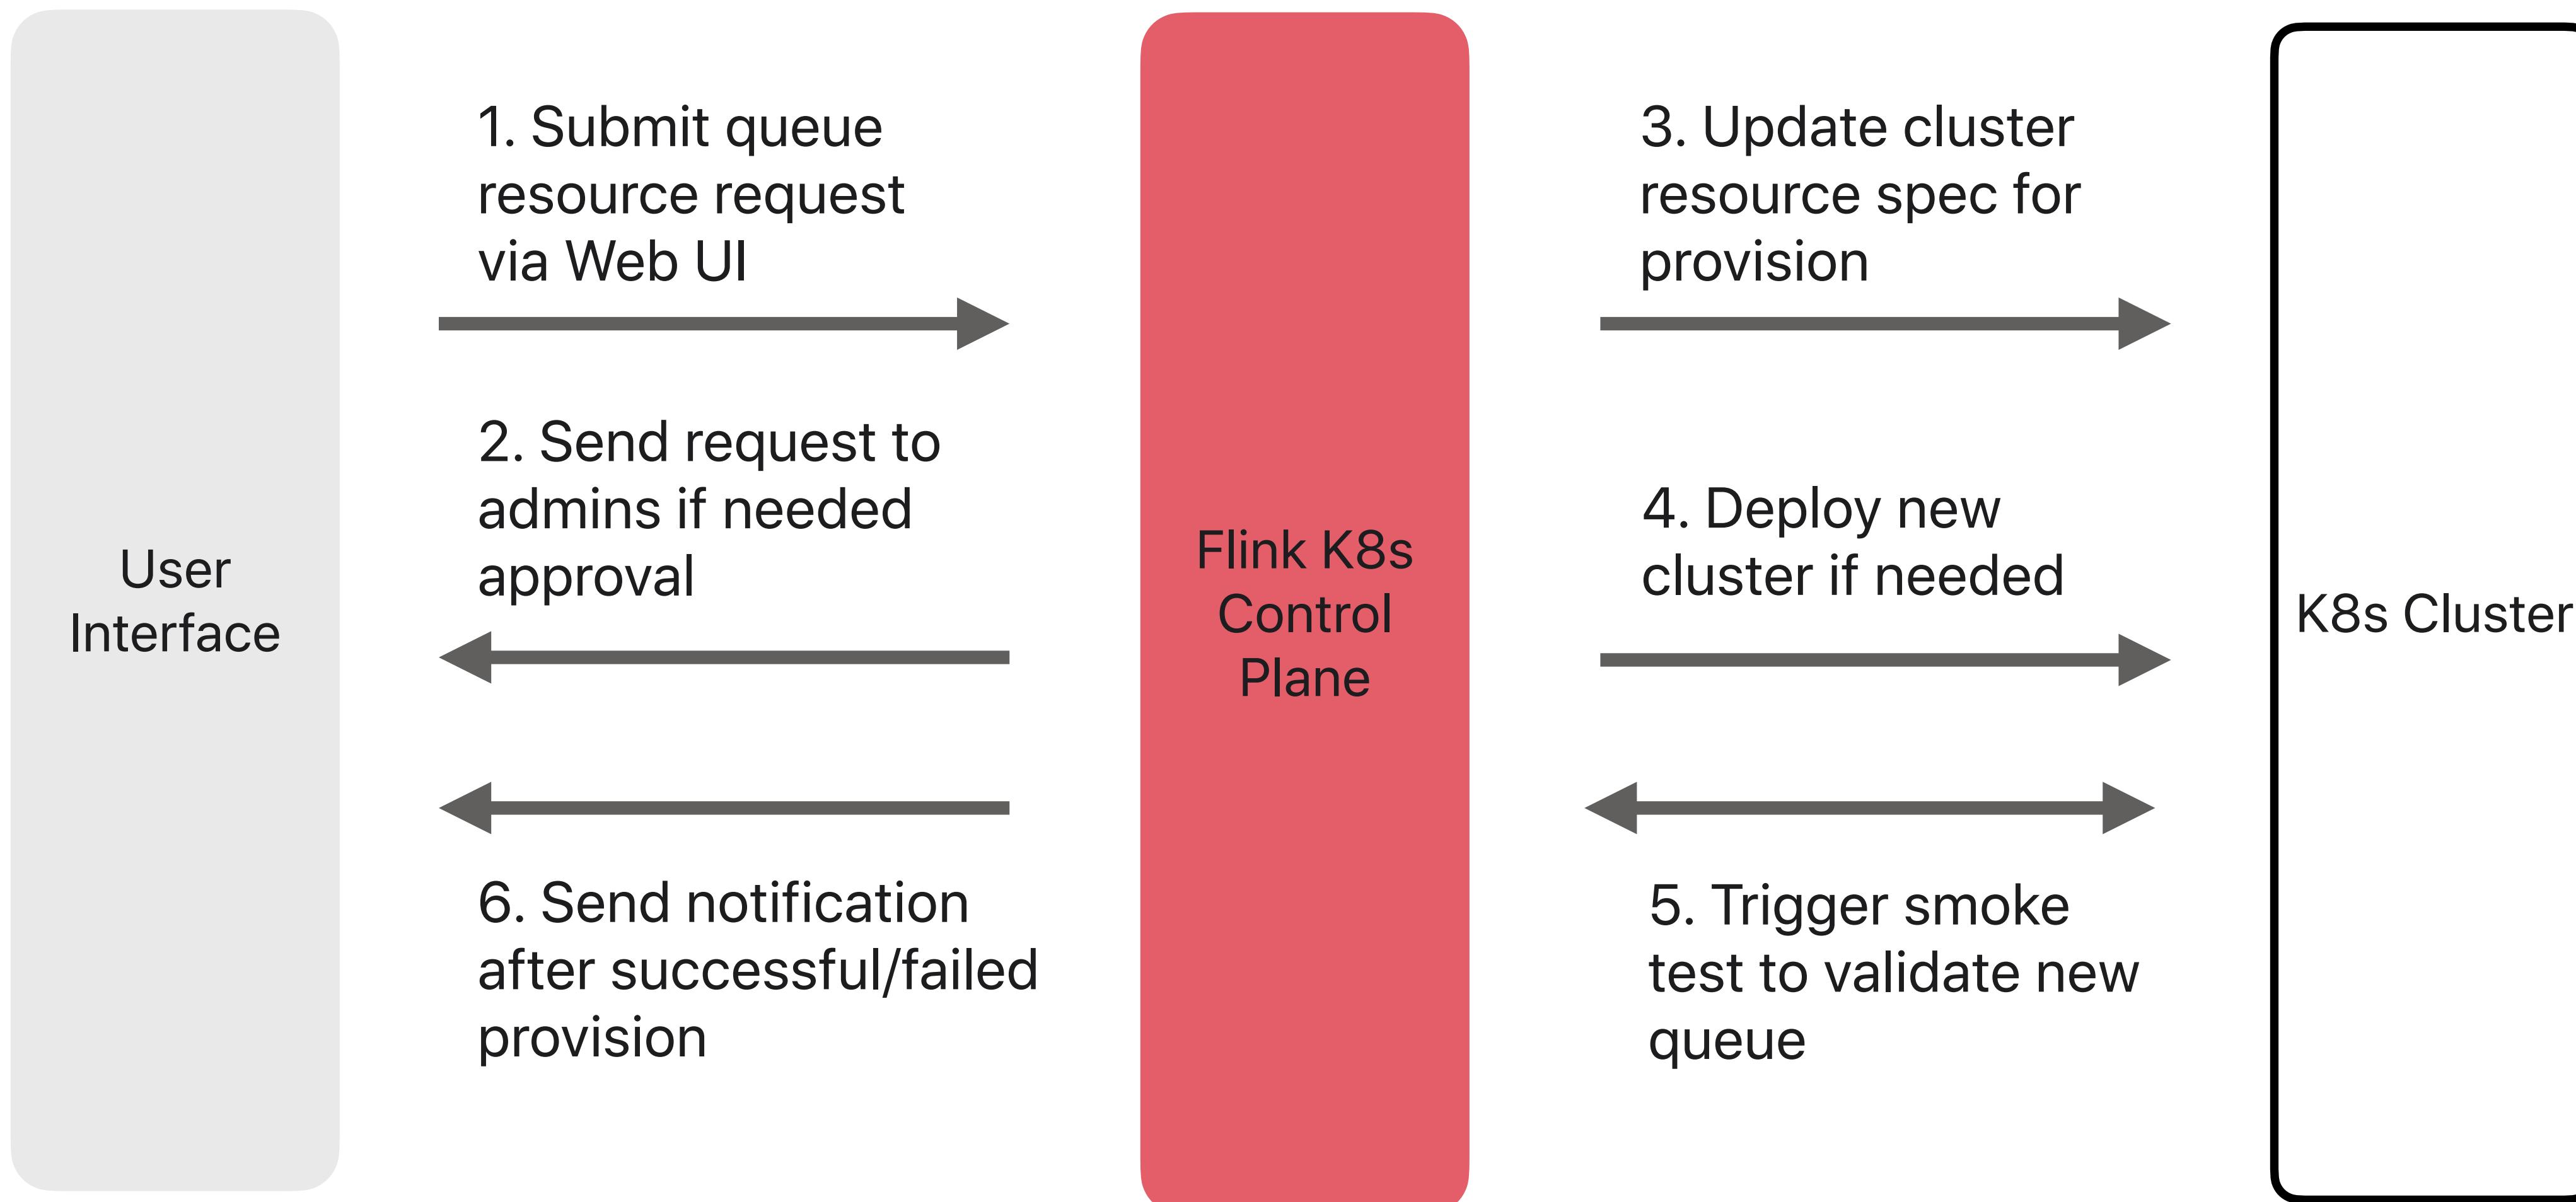


# Introduce Flink Control Plane



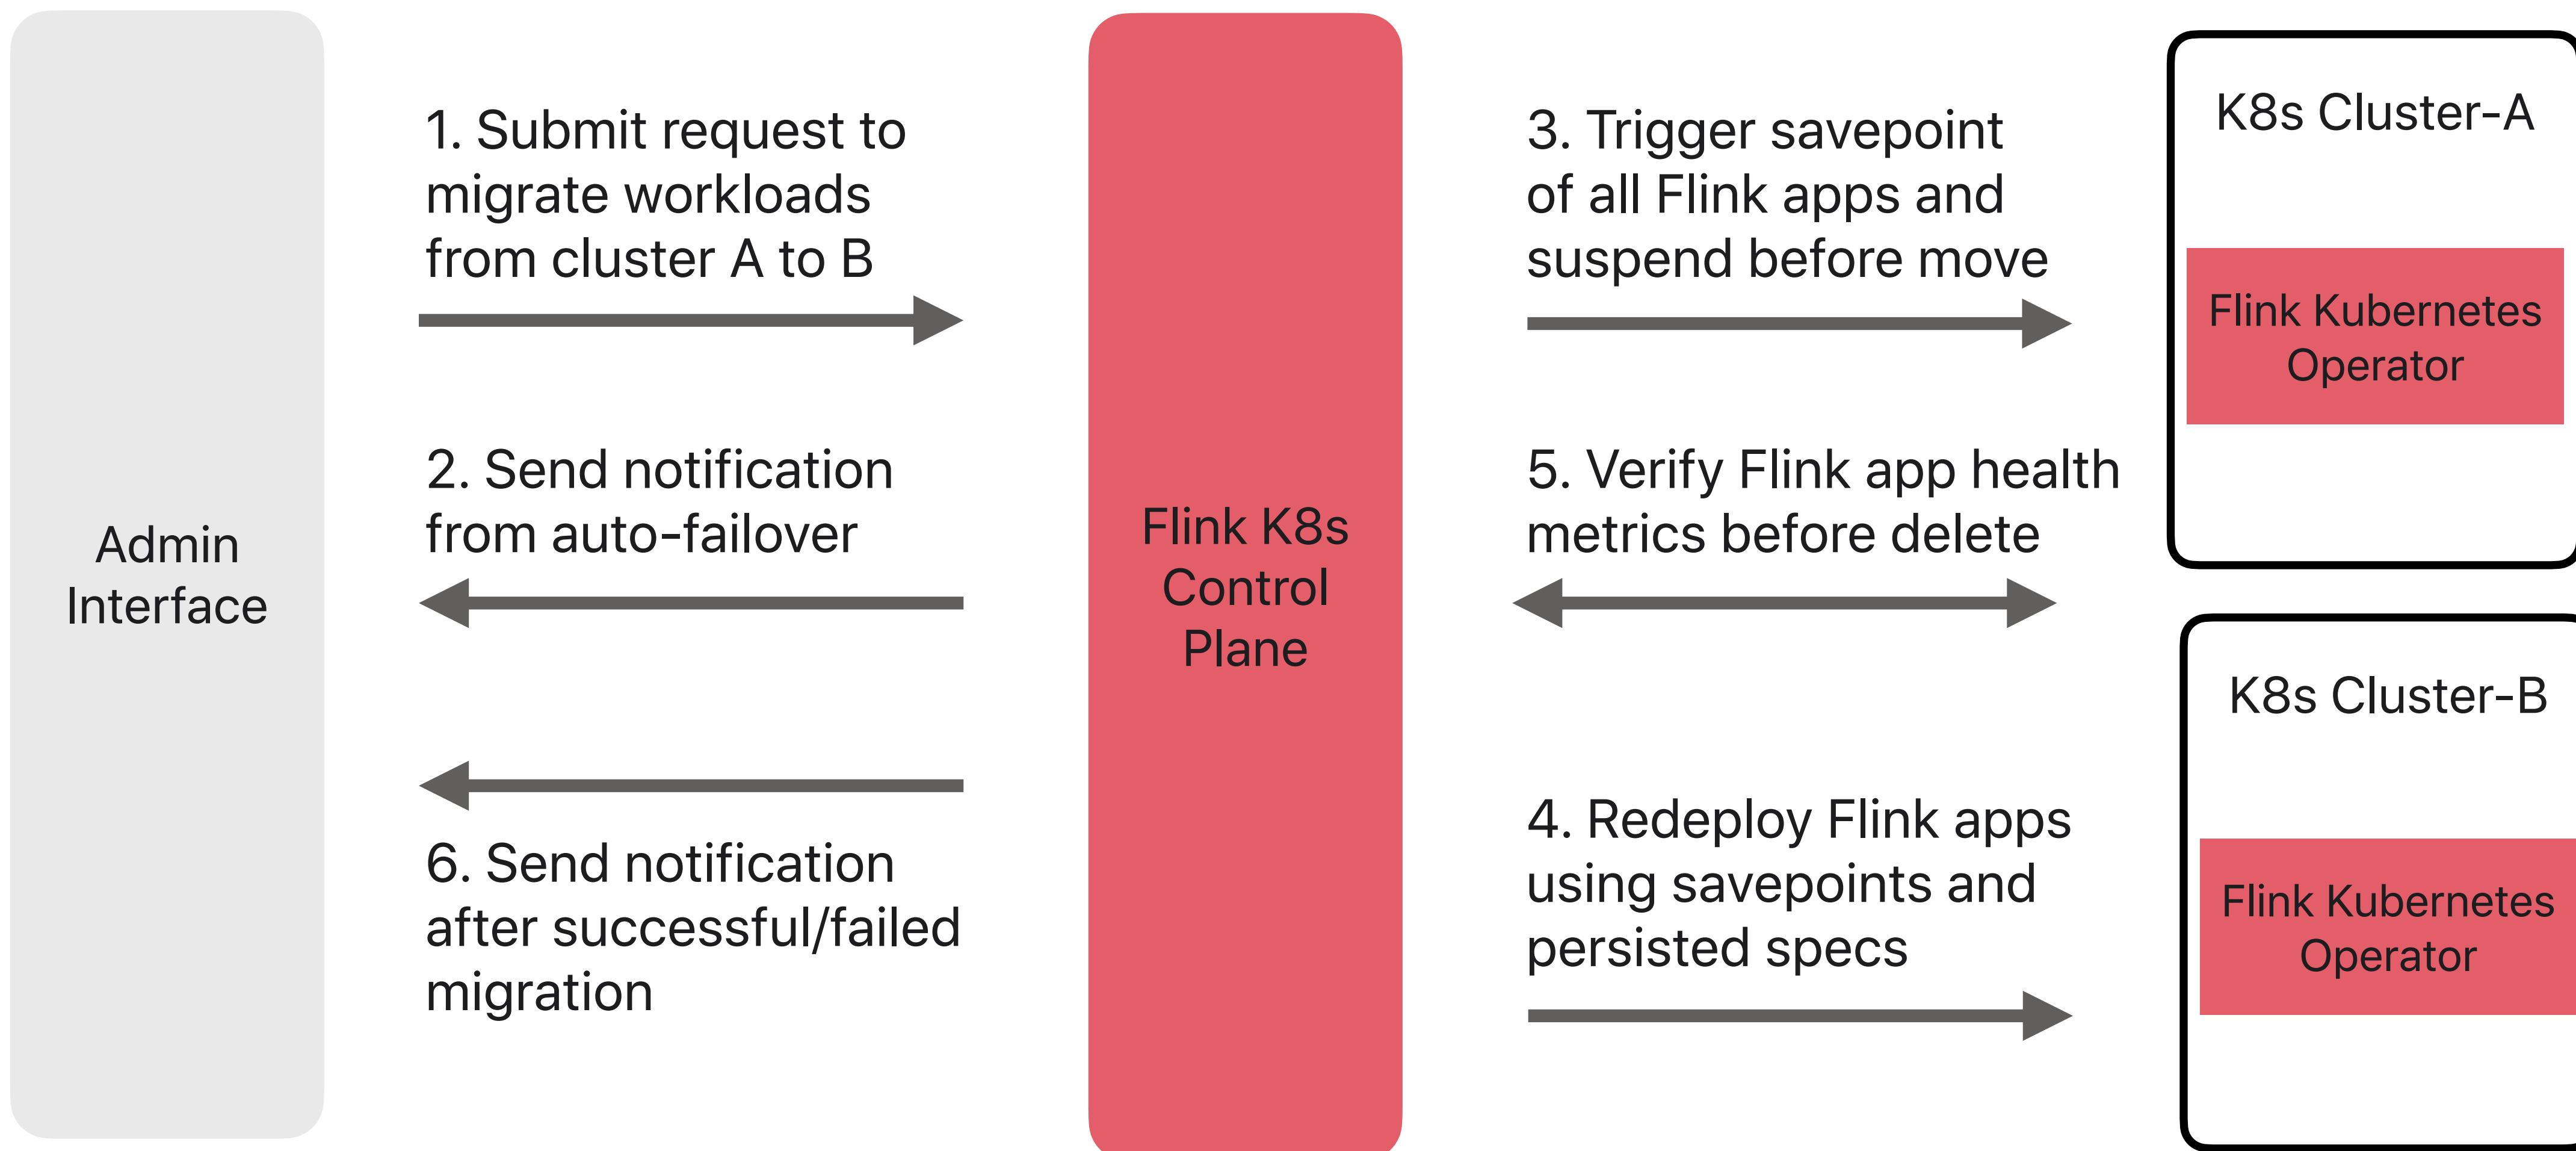
# Introduce Flink Control Plane

- Self-service user onboarding experience



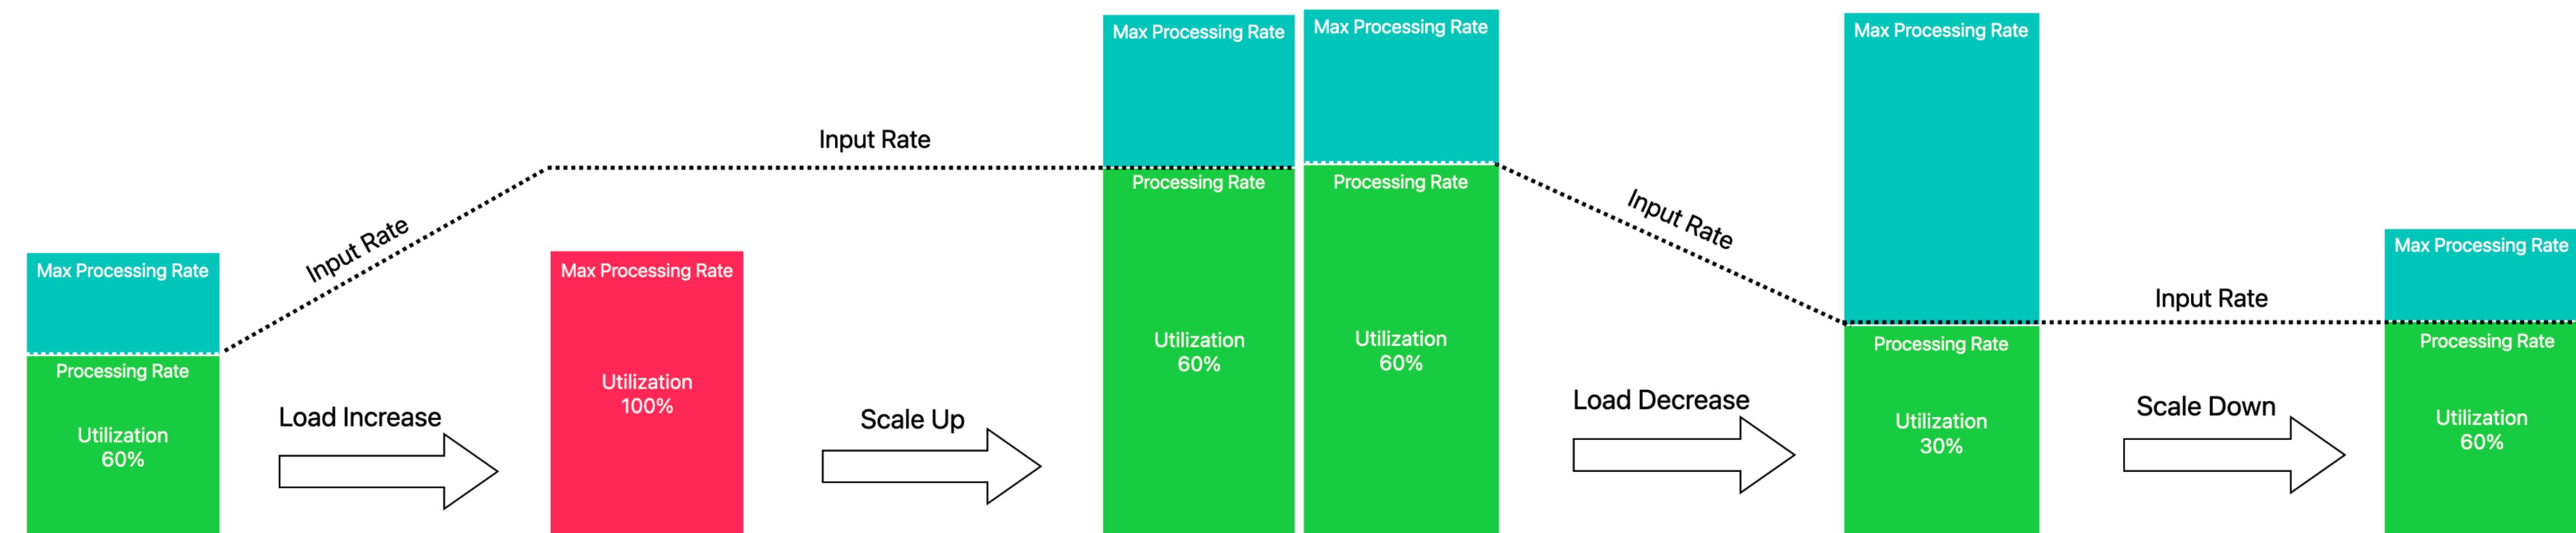
# Introduce Flink Control Plane

- Automated cluster migration and failover



# Autoscale Flink Application

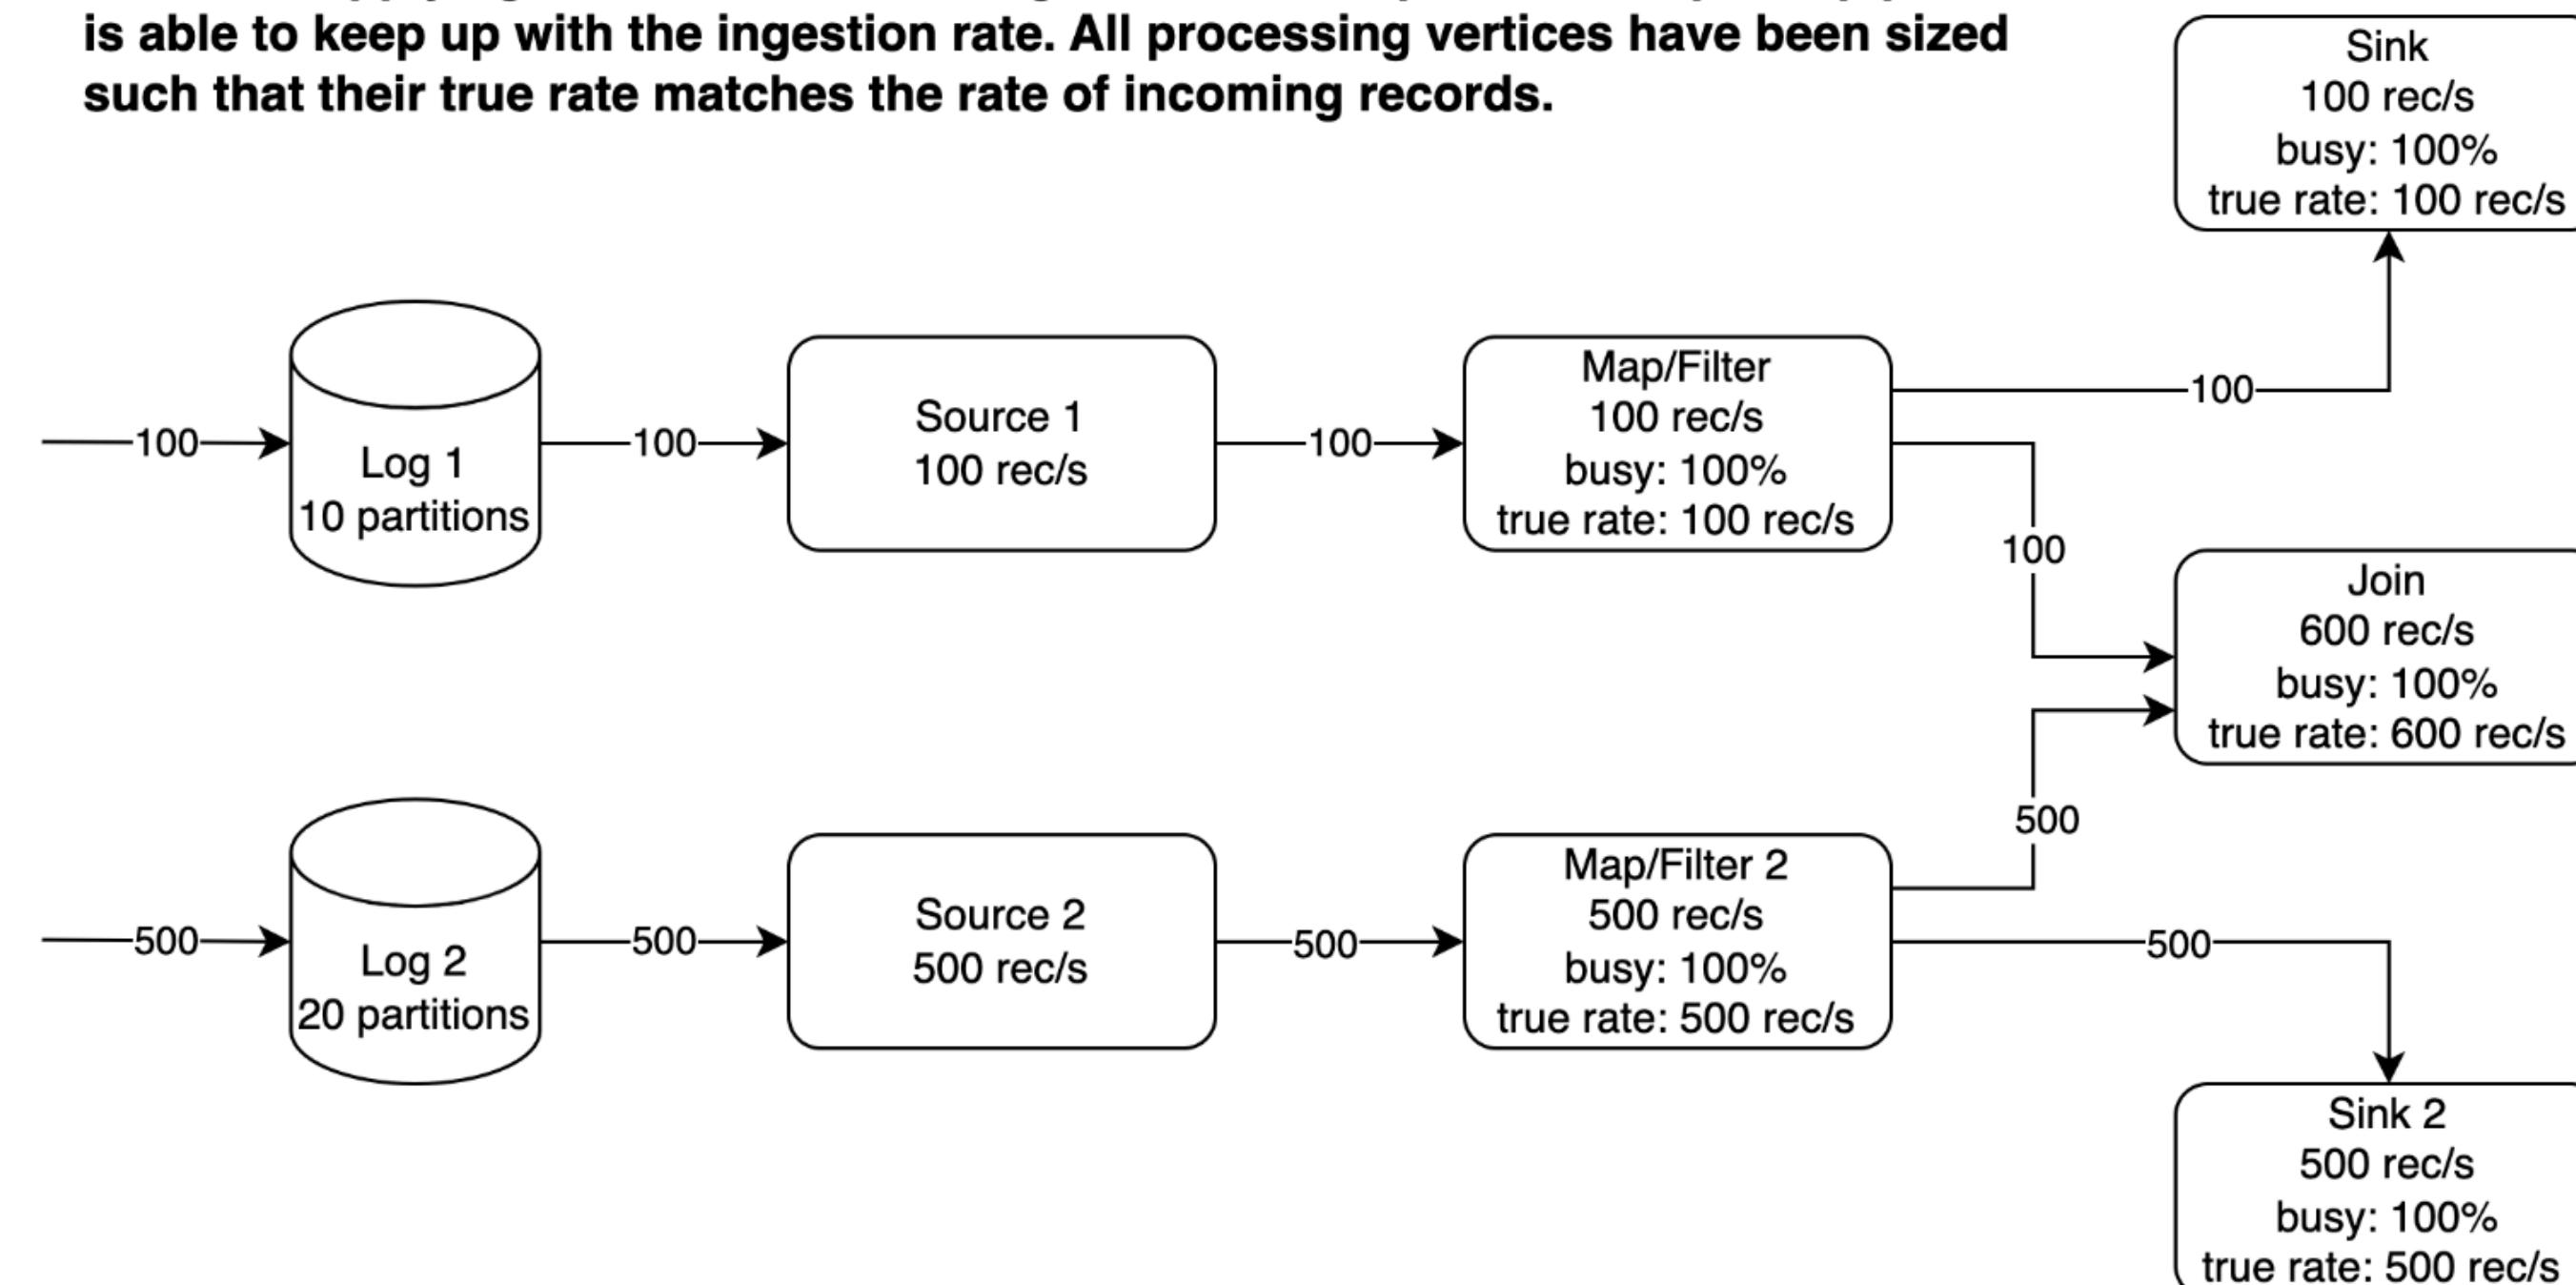
- FLIP-271: Autoscaling - when and how much to scale?
- Each scaling decision has an associated cost



# Autoscale Flink Application

- Predict parallelism change effects on downstream job vertices
- Flink Kubernetes Operator integration and deployment metrics

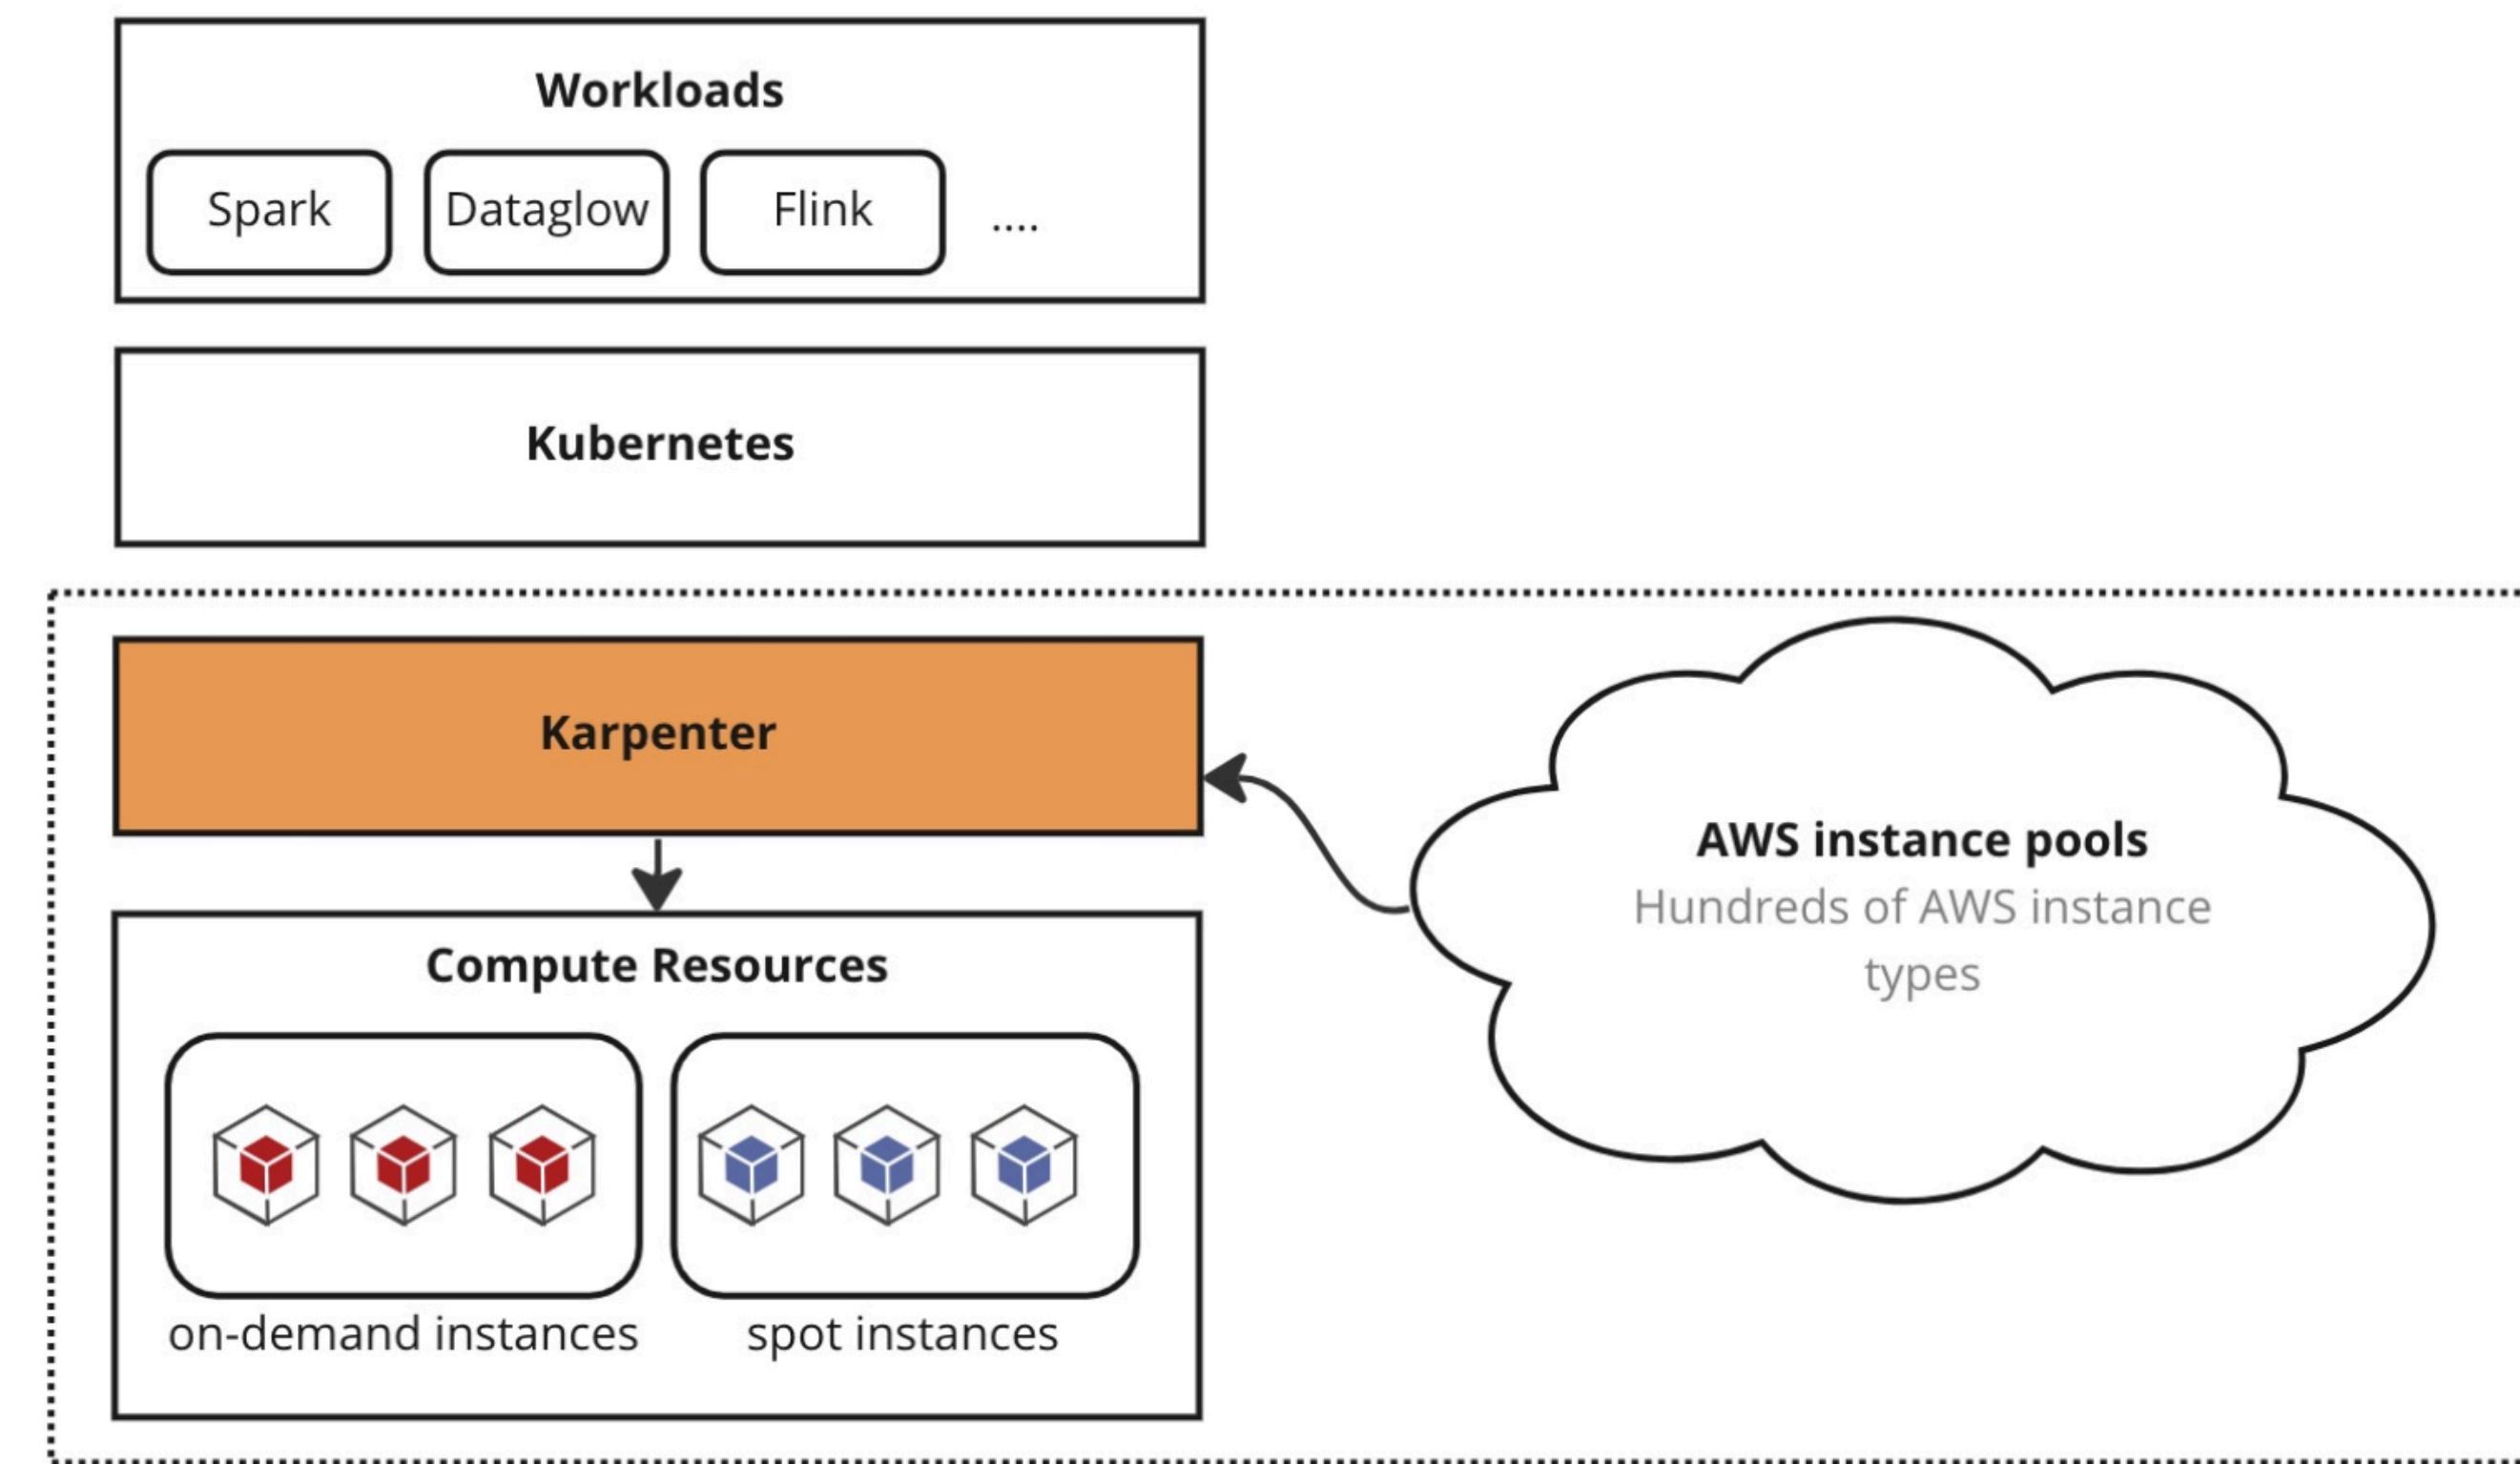
Based on applying the calculated scaling factors in the previous step, the pipeline is able to keep up with the ingestion rate. All processing vertices have been sized such that their true rate matches the rate of incoming records.



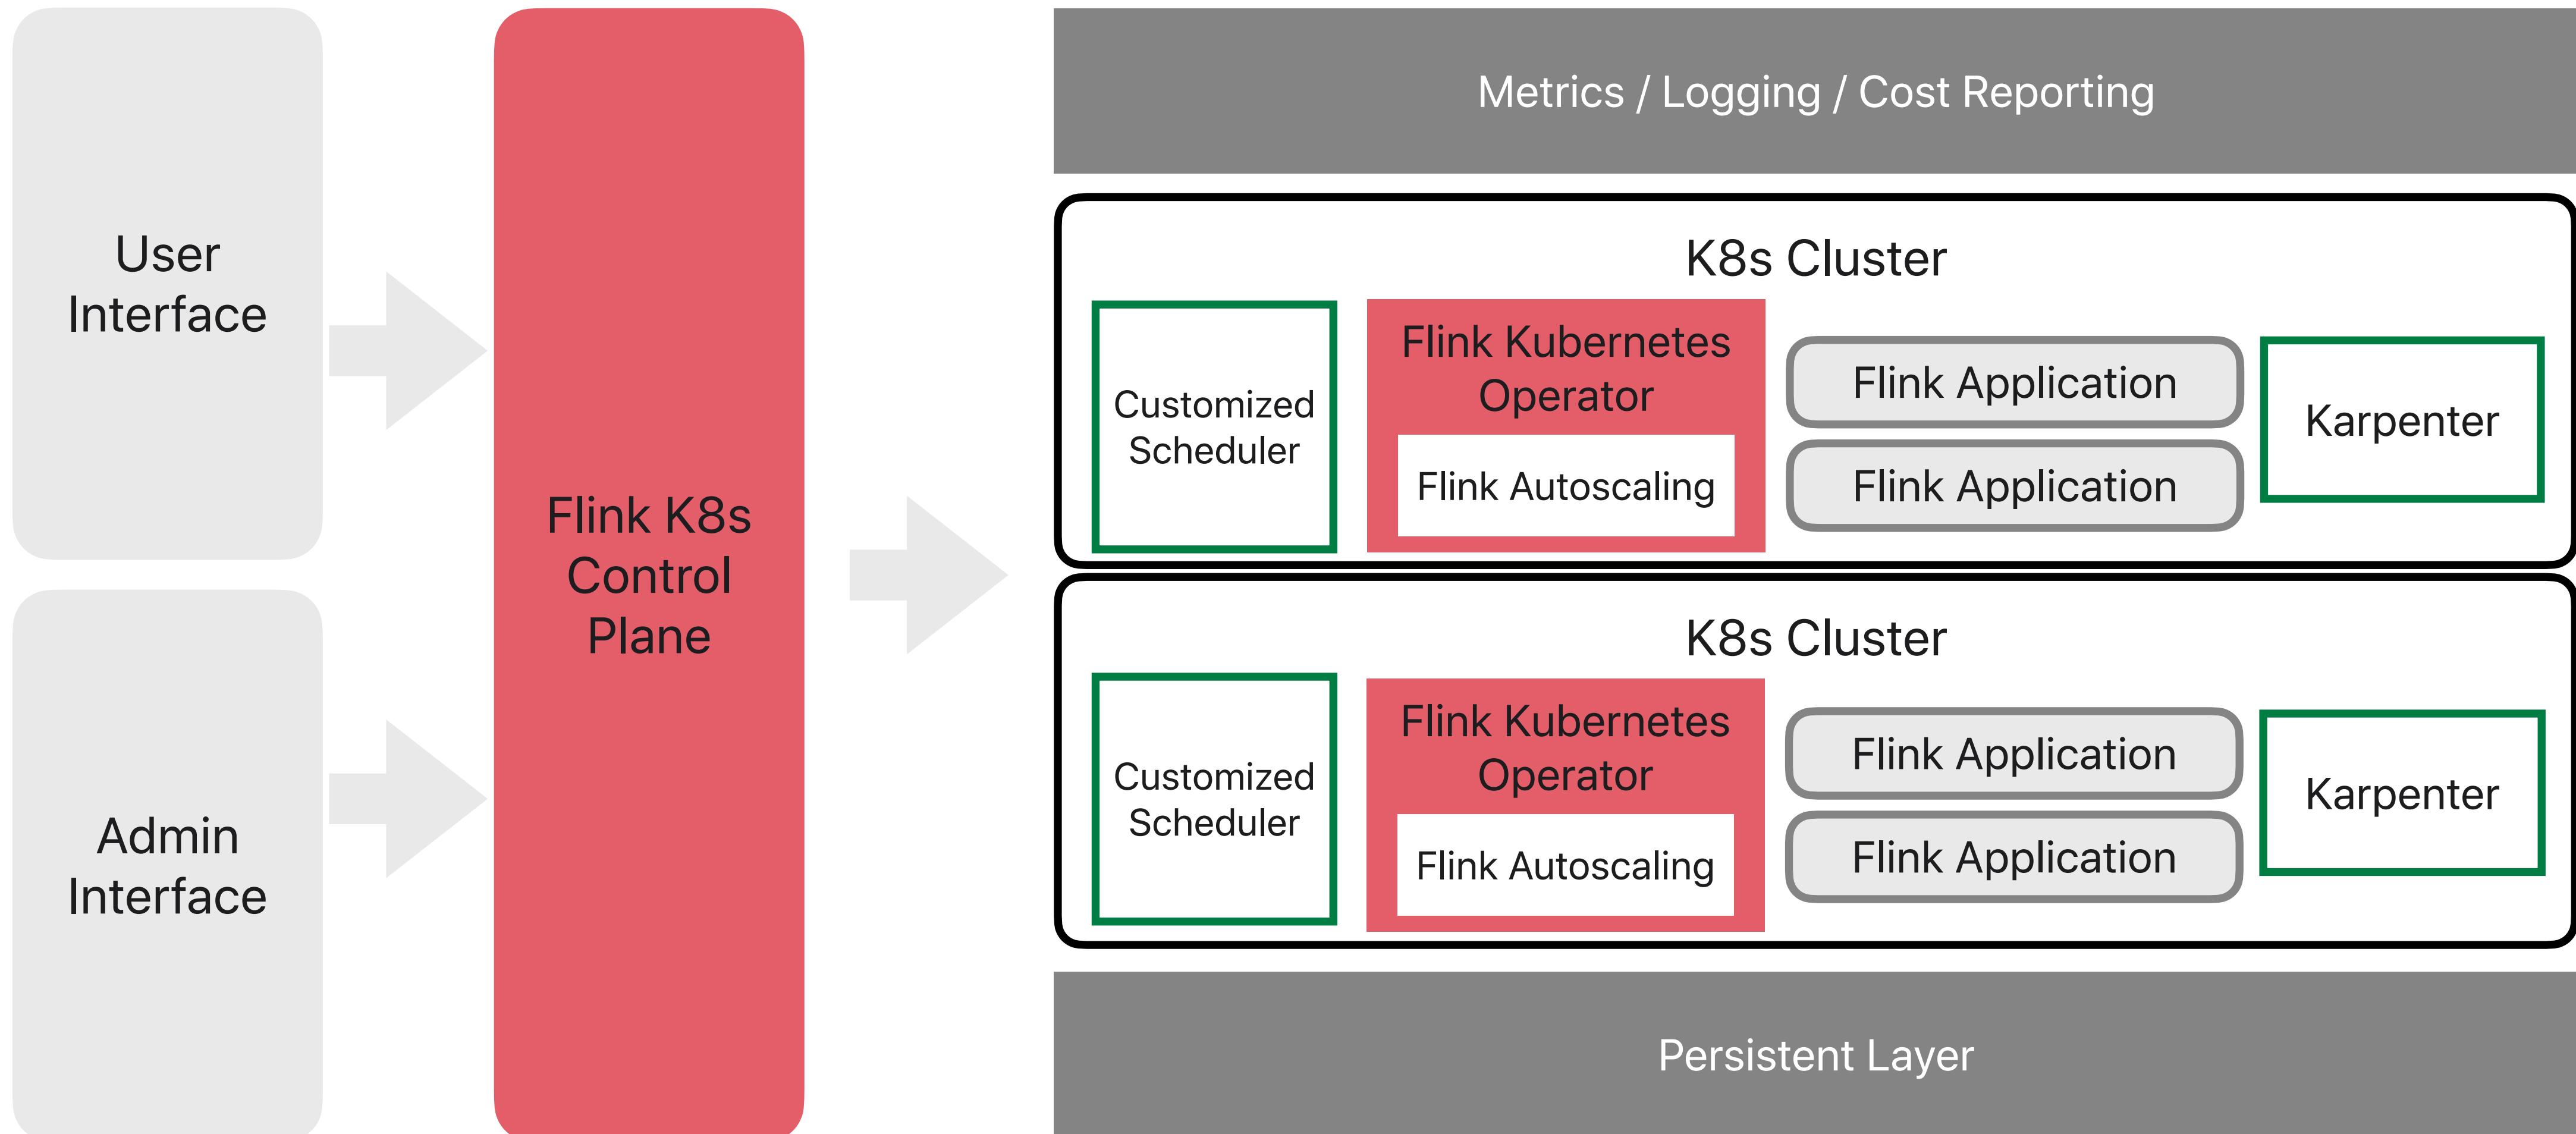
Scale 3x

# Stay Resource-Aware

- From Cluster Autoscaler to Karpenter
- Customized scheduler to improve bin-packing of sticky deployment



# Wrap up: Self-Service Platform





TM and © 2023 Apple Inc. All rights reserved.