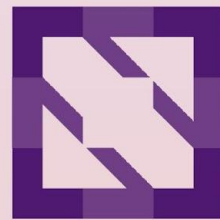




KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

Service Mesh Journey at DoorDash: The Good, the Bad, and the Ugly

Hochuen Wong
DoorDash



Hochuen Wong
Software Engineer, DoorDash

Doordash:

- Joined Core Infrastructure Org in 2020
- Compute and Traffic Infrastructure Teams

Prior to Doordash:


- Distributed systems for Machine Learning

Contact:

- <https://www.linkedin.com/in/hochuenw/>

Why Service Mesh?






DoorDash service graph in 2023

WHY SERVICE MESH?

Move to microservices

- Started [microservice journey](#) in 2019
- 50+ microservices by the end of Q1 2021
- Classic microservice challenges:
 - Observability and debugging
 - Service to service communication
 - Authentication and authorization
 - Standardization
 - ...



DoorDash service graph in 2023

WHY SERVICE MESH?

Move to microservices

- Started microservice journey in 2019
- 50+ microservices by the end of Q1 2021
- Classic microservice challenges:
 - Observability
 - Networking
 - Debugging
 - Authentication and authorization
 - Standardization
 - ...
- **Service Mesh journey began in Q2 2021**

Initial Design





Requirements

SCALABILITY

- Support the largest Kubernetes cluster with **2000+** nodes

FLEXIBILITY

- Support DoorDash Consul-based multi-cluster architecture

MATURITY

- Successful user stories
- Community

CONFIGURATION & DOCUMENTATION

- Easy to configure
- Documentation

FEATURES

- Observability
- Security
- Reliability
- Traffic Management



Requirements

SCALABILITY

- Support the largest Kubernetes cluster with 2000+ nodes

FLEXIBILITY

- Support DoorDash Consul-based multi-cluster architecture

MATURITY

- Successful user stories
- Community

CONFIGURATION & DOCUMENTATION

- Easy to configure
- Documentation

FEATURES

- Observability
- Security
- Reliability
- Traffic Management



Requirements

SCALABILITY

- Support the largest Kubernetes cluster with 2000+ nodes

FLEXIBILITY

- Support DoorDash Consul-based multi-cluster architecture

MATURITY

- Successful user stories
- Community

CONFIGURATION & DOCUMENTATION

- Easy to configure
- Documentation

FEATURES

- Observability
- Security
- Reliability
- Traffic Management



Requirements

SCALABILITY

- Support the largest Kubernetes cluster with 2000 nodes

FLEXIBILITY

- Support DoorDash Consul-based multi-cluster architecture

MATURITY

- Successful user stories
- Community

CONFIGURATION & DOCUMENTATION

- Easy to configure
- Documentation

FEATURES

- Observability
- Security
- Reliability
- Traffic Management



Requirements

SCALABILITY

- Support the largest Kubernetes cluster with 2000 nodes

FLEXIBILITY

- Support DoorDash Consul-based multi-cluster architecture

MATURITY

- Successful user stories
- Community

CONFIGURATION & DOCUMENTATION

- Easy to configure
- Documentation

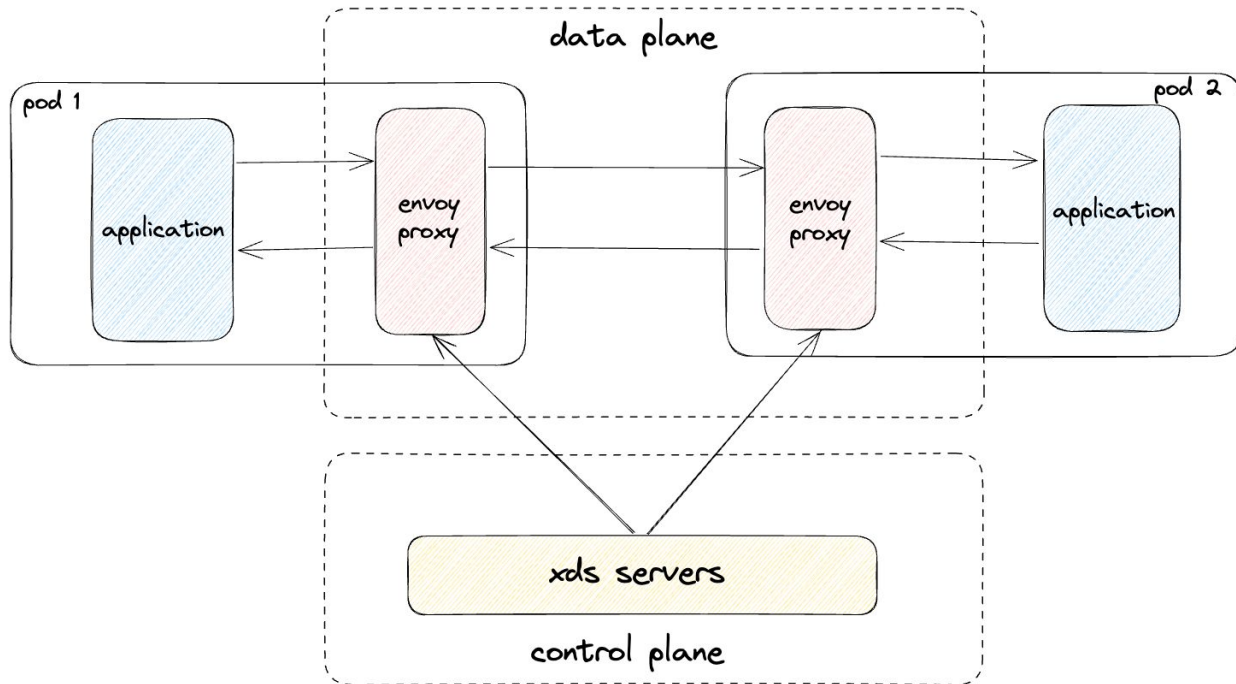
FEATURES

- Observability
- Security
- Reliability
- Traffic Management



Architecture

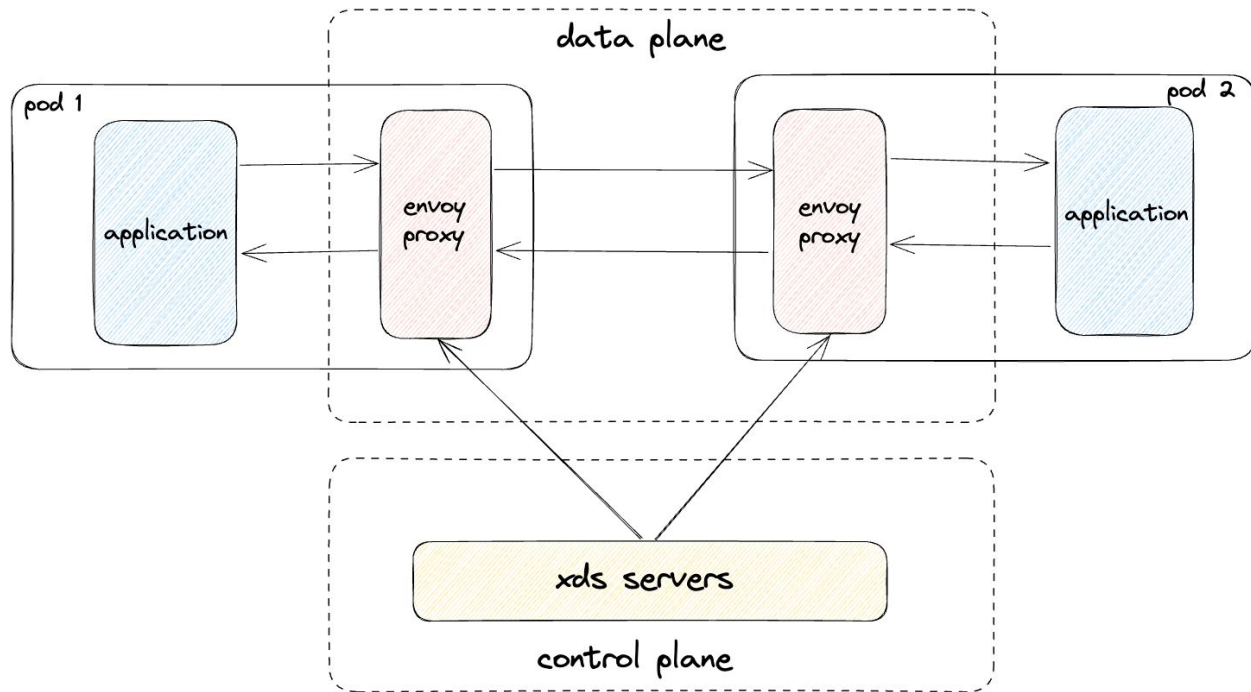
- K8s mutating webhook
- iptables
- Data Plane
 - Envoy sidecar container
 - HTTP1/HTTP2/gRPC
- Control plane
 - Custom version





Architecture

- Data Plane
 - Envoy sidecar container
- Control plane
 - Custom version
 - Consul for EDS
 - K8s custom resource for other resources
- Adoption
 - Onboard everyone
 - Minimum features



Sitewide Outage







SITE-WIDE OUTAGE

What happened

- CASCADING FAILURE (RCA)

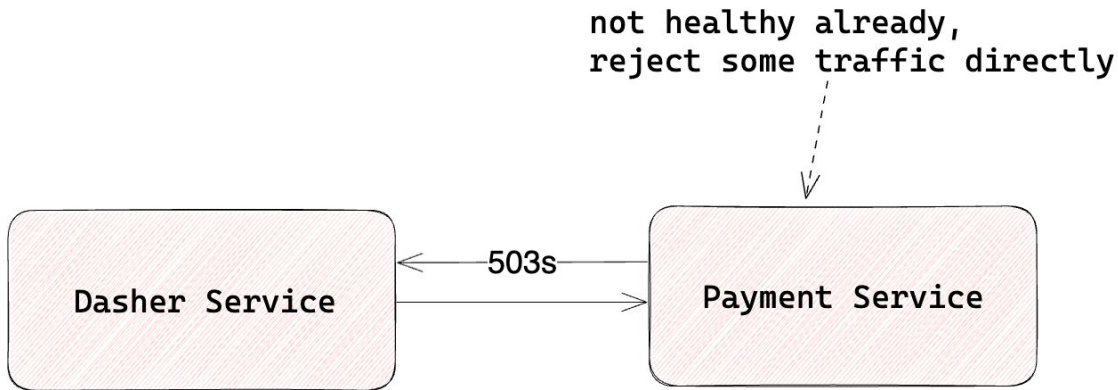




SITE-WIDE OUTAGE

What should have happened

- LOAD SHEDDING
- CIRCUIT BREAKING

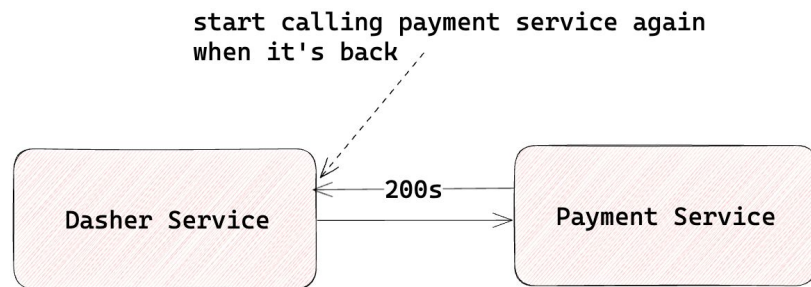
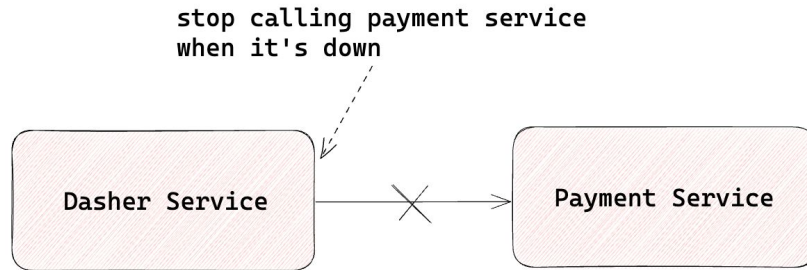




SITE-WIDE OUTAGE

What should have happened

- LOAD SHEDDING
- CIRCUIT BREAKING





Service Mesh Journey at DoorDash

SITE-WIDE OUTAGE

Code freeze.



SITE-WIDE OUTAGE

Code freeze.

A project called Service Mesh?

- Just reviewed the initial design doc
- No operational experience with running Envoy
- No control plane



SITE-WIDE OUTAGE

Code freeze.

A project called Service Mesh?

Shift Priority !



Design after the outage






Design after the outage

Implementation

- Envoy sidecar injector & iptables
- Configuration
 - File-based dynamic configurations
 - HTTP passthrough proxy
- Features
 - Adaptive concurrency (load shedding)
 - Outlier detection (circuit breaking)






INITIAL VERSION

A project called ~~Service Mesh~~ Envoy Sidecar

- Primitive configuration management
- Focus on the data plane and two reliability features



Onboarding First Customers



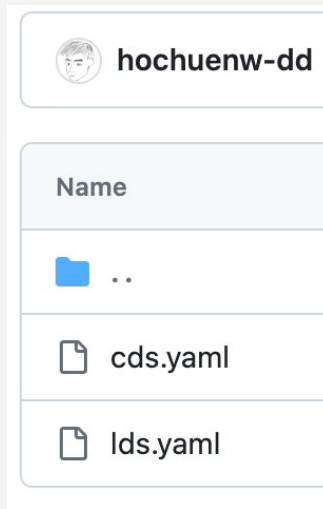


ONBOARDING FIRST CUSTOMERS

Steps to onboard one service

- label namespace
- label deployment
- create **RAW(!)** Envoy configurations (**1K+ lines**)

```
labels:  
  envoy-sidecar-injector.doordash.com/enabled: "true"
```






ONBOARDING FIRST CUSTOMERS

Rollout Strategy

- Canary deployment
- Two week bake-in period





ONBOARDING FIRST CUSTOMERS

Did we make the right decision?



LEARNINGS

- ✓ Dream Big, Start Small
- ✓ Solve real-world problems





General Availability!

WHAT WE HAVE IMPROVED SINCE ONBOARDING FIRST CUSTOMERS

- Configuration management
 - Templated Envoy configurations
- Observability
 - Common observability dashboard
 - Common alert module and runbook
- Operational experience
 - Support for more programming languages
 - Onboard more services



More features







More features

- Zone-aware Routing
 - Route egress traffic to its local zone
 - Ensure traffic is still balanced
 - Save data transfer cost
 - Reduce impact of AZ outage
 - Performance

before (everyone talks to everyone)



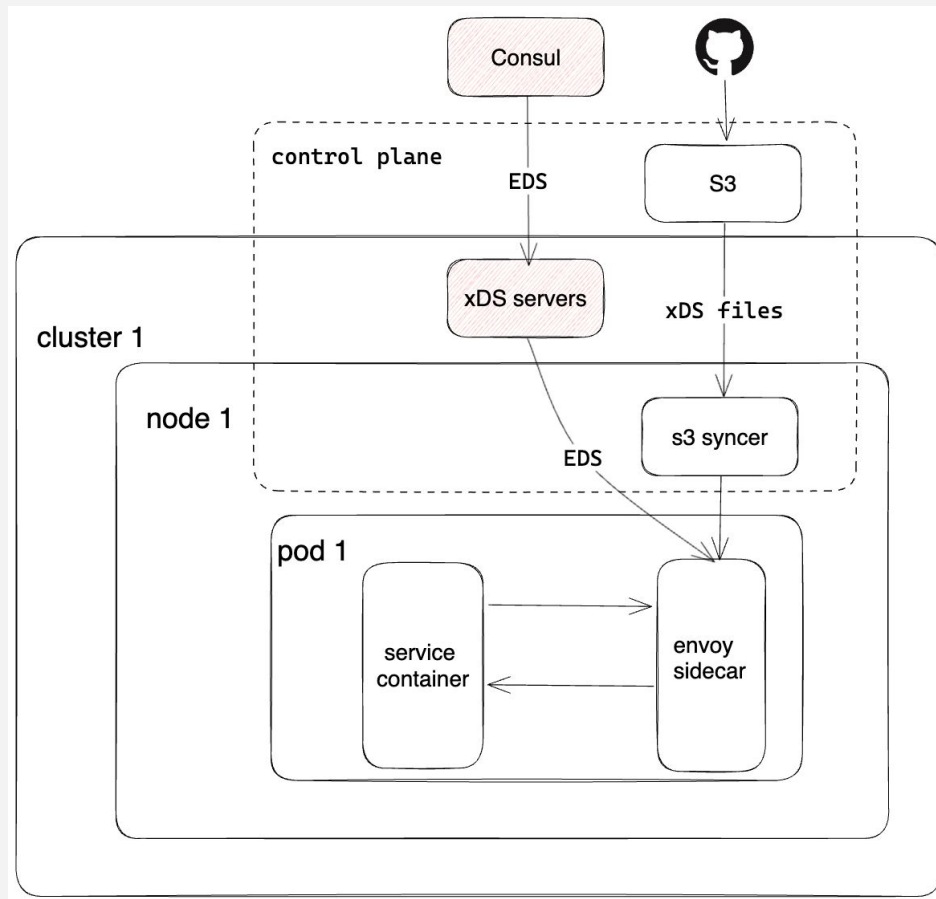
after (prefer local AZ)





More features

- Zone-aware Routing
 - Evolved configuration management systems
 - Introduced API-based dynamic configurations





More features

- There are many more feature request
 - Zone-aware Routing
 - Client-side load balancing
 - Sandbox traffic routing
 - Slow Start Mode
 - Traffic redirection
 - Global rate limiting
 - ...
- Deeper understanding of customer pain points



LEARNINGS

- ✓ Co-develop your solution with initial customers



Adopt 100 Services!



Four more years!



Q4 2022: Prepare for large scale onboarding






Prepare for large scale onboarding

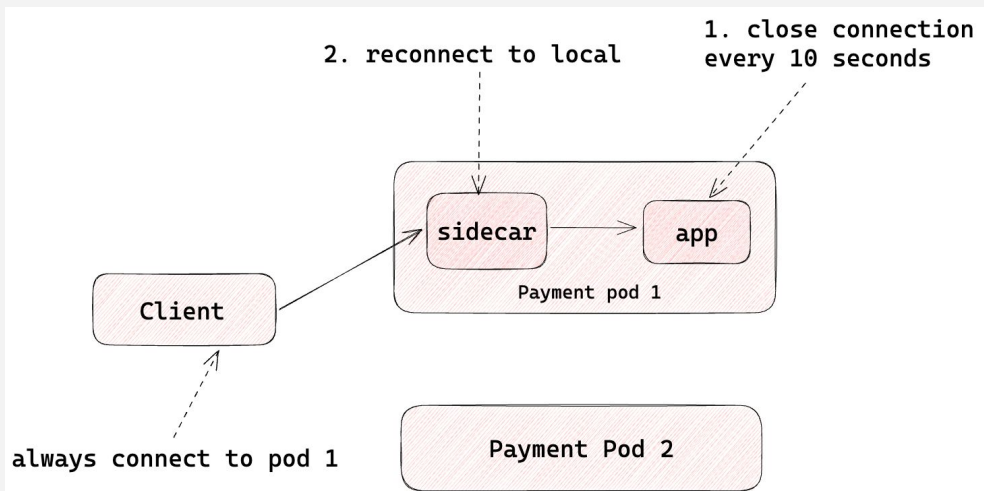
Unknowns discovered

- Unnoticed client behaviors became apparent
- First consumer payment service as an example

Before



After





Prepare for large scale onboarding

Unknowns discovered

- There are many more examples...
- But it slowed down as more services being onboarded



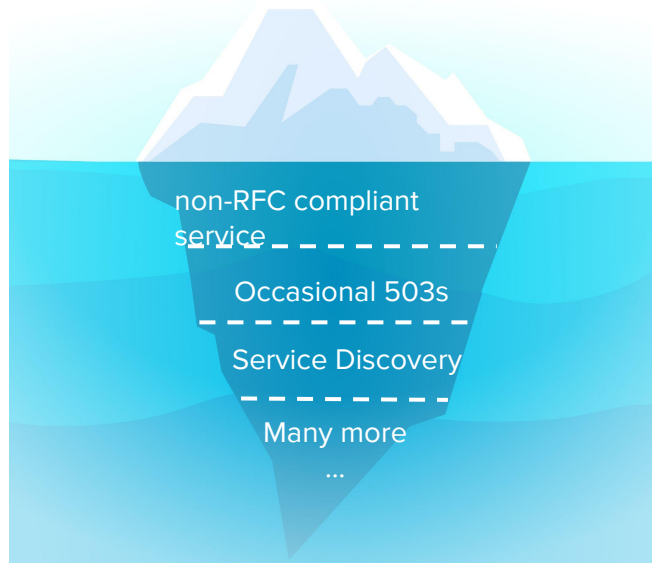
LEARNINGS

- ✓ Envoy sidecar is not transparent
- ✓ Expect the unexpected
- ✓ Take a bet at the right time



ACTION ITEMS

- ✓ Switch to use gradual rolling update approach





Prepare for large scale onboarding

Challenges in Developer Experience

- ✗ The manual onboarding process and Envoy resource tuning, led by developers, required a learning curve



LEARNINGS

- ✓ Decentralized onboarding and manual Envoy resource tuning doesn't scale



ACTION ITEMS

- ✓ Infra team owns onboarding and Envoy resource management
- ✓ Streamline the onboarding process
- ✓ Pre-generate all Envoy configs and labels



Prepare for large scale onboarding

Challenges in Developer Experience

- ✗ Caused confusion:
 - Onboarding status
 - Envoy configs
 - Architecture



ACTION ITEMS

- ✓ Invest in educating/enabling service owners



Prepare for large scale onboarding

Challenges in Developer Experience

✗ Complexity of observability features



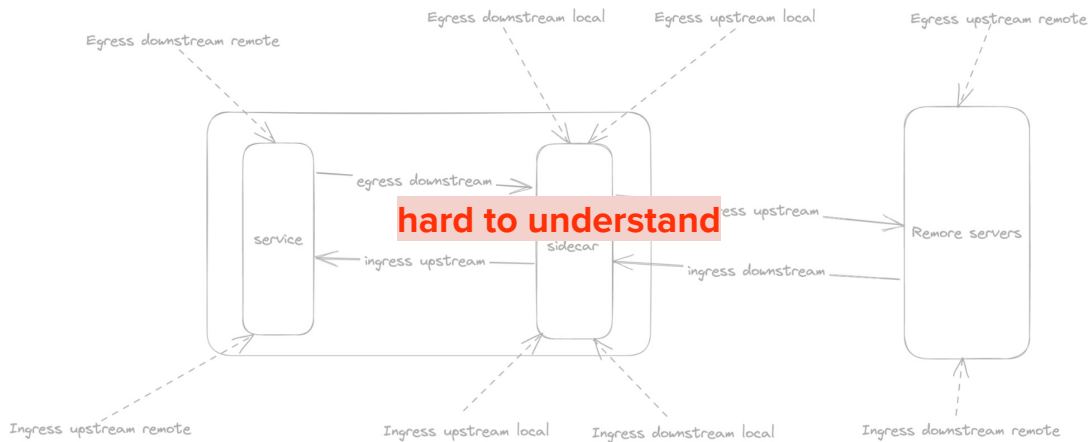
LEARNINGS

- ✓ Have a product mindset / customer obsessed



ACTION ITEMS

- ✓ Simplify the observability dashboards





Prepare for large scale onboarding

Challenges in Developer Experience

✗ Debugging challenges



ACTION ITEMS

- ✓ Availability SLOs
- ✓ Introducing distributed tracing with OpenTelemetry



Prepare for large scale onboarding

Challenges in Developer Experience

- ✗ Many features are only enabled when egress clusters are defined
- ✗ Building an accurate service graph is not trivial



ACTION ITEMS

- Get accurate service graph from tracing data
- Build a tool to generate egress configs





LARGE SCALE ONBOARDING

Large scale adoption

- Onboarding was much faster
- Uncovered a few unknowns
- Manageable additional maintenance responsibilities for the infrastructure team
- Features are widely adopted






Current State

WHERE WE'RE AT TODAY!

- 10 production Kubernetes clusters
- 500+ microservices
- 5 isolated mesh
- Multiple clusters with one mesh
- 10K pods
- 5M RPS






Current and future work

WHAT'S NEXT!

- Developer Velocity
 - Improve configuration management experience
 - Unify the experience of using observability features
- Efficiency
 - Compute
 - Metrics
- Traffic Simplification
 - Refactor multi-cluster service discovery infrastructure
- Feature enrichment





Recap!

BEFORE YOU START!

- ✓ Understand your requirements and use cases
- ✓ Co-develop your solution with initial customers

WHEN WORKING WITH ENVOY

- ✓ Sidecar is not always transparent
- ✓ Expect the unexpected



Recap (continued)!

WHEN ONBOARDING SERVICES!

- ✓ Test things gradually at the beginning
- ✓ Make well-informed bets (at the right time)
- ✓ Decentralized onboarding doesn't scale
- ✓ Streamline/Automate the onboarding process

WHEN DELIVERING PRODUCT TO THE REST OF THE TEAM

- ✓ Envoy metrics can overwhelm service owners
- ✓ Have a product mindset / customer obsessed
- ✓ Invest in training/enabling service owners
- ✓ Increase velocity through more user friendly solutions



Thank you



**Please scan the QR Code above
to leave feedback on this session**