



KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

The Cluster Killer Bug: Learning API Priority and Fairness the Hard Way

Eddie Zaneski

Hi I'm Eddie!



KubeCon



CloudNativeCon

North America 2023



- @eddiezane
- Denver, CO
- Climb big mountains
- Co-chair & tech lead for Kubernetes SIG-CLI
- Staff DevRel + OSS Engineer (hire me)

Flashback to Feb 1st 2023



KubeCon



CloudNativeCon

North America 2023



<https://events.linuxfoundation.org/cloudnativesecuritycon-north-america/>

Karaoke



KubeCon



CloudNativeCon

North America 2023



<https://giphy.com/gifs/YbXtbKoi2ZUOc>

Media Interview Tomorrow



KubeCon



CloudNativeCon

North America 2023

The New Stack logo, featuring the words 'THE NEW STACK' in a bold, sans-serif font. 'THE' and 'STACK' are in a light purple color, while 'NEW' is in white. The text is centered on a square background with a vertical gradient from purple at the top to pink at the bottom.

THE
NEW
STACK

<https://thenewstack.io/>

Last Minute Prep



KubeCon



CloudNativeCon

North America 2023

LEAVING A PARTY WITHOUT SAYING BYE.



<https://giphy.com/gifs/ShalitaGrant-sneaking-out-leaving-a-party-not-saying-bye-xlcXy1FF94YaYgxxaO>

Nothing Works



KubeCon



CloudNativeCon

North America 2023



<https://giphy.com/gifs/5Zesu5VPNGJIm>

Delete The Cluster & Go To Bed



KubeCon



CloudNativeCon

North America 2023



<https://giphy.com/gifs/this-is-fine-9M5jK4GXmD5o1irGrF>

Wake Up...



KubeCon



CloudNativeCon

North America 2023



<https://giphy.com/gifs/pnPfEgZi3InO>

Feb 2nd 2023



KubeCon



CloudNativeCon

North America 2023

- Push interview back
- Triage triage triage
- Here we go...

Flashback to Feb 1st 2023



KubeCon



CloudNativeCon

North America 2023



https://youtu.be/uPBUKazKjag?si=Rwj3l_erzPe7QzPs

Early Takeaways



KubeCon



CloudNativeCon

North America 2023

- Don't delete the cluster
- Start ruling out what you can
- Think out loud
- *Glance* at logs + metrics
- Your goal is to reproduce the issue, not understand it yet

The Controller



KubeCon



CloudNativeCon

North America 2023

- Built around the Sigstore Policy Controller
- Opened long lived watches to the API Server
- Knative
- Ability to run "agentless"

Controllers

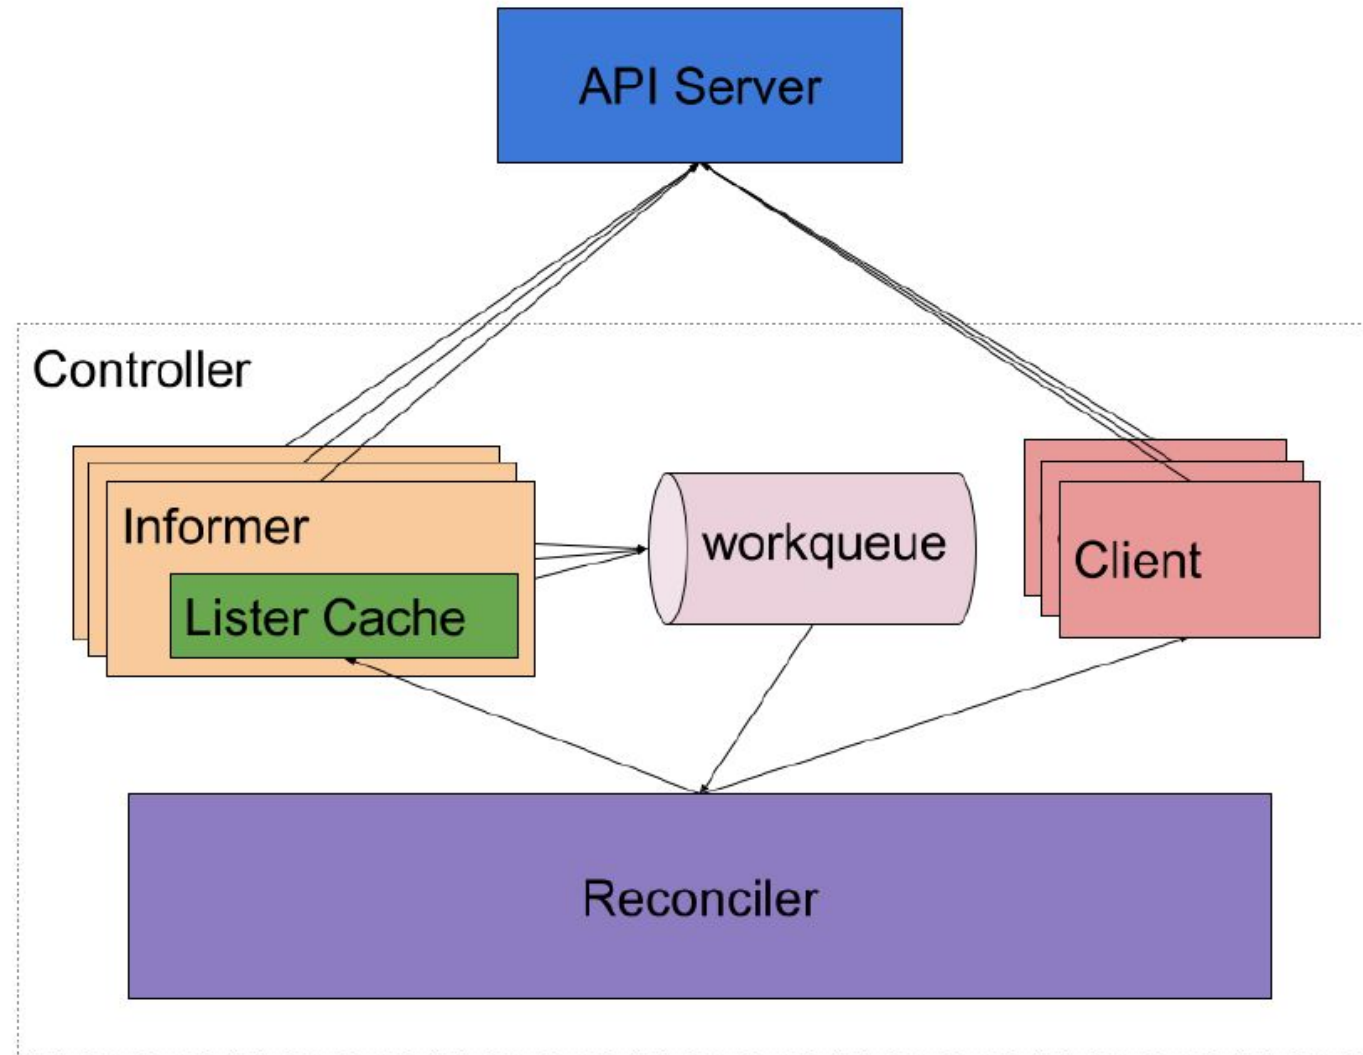


KubeCon



CloudNativeCon

North America 2023



<https://edu.chainguard.dev/chainguard/chainguard-enforce/concepts/gulfstream-overview/>

"Agentless"

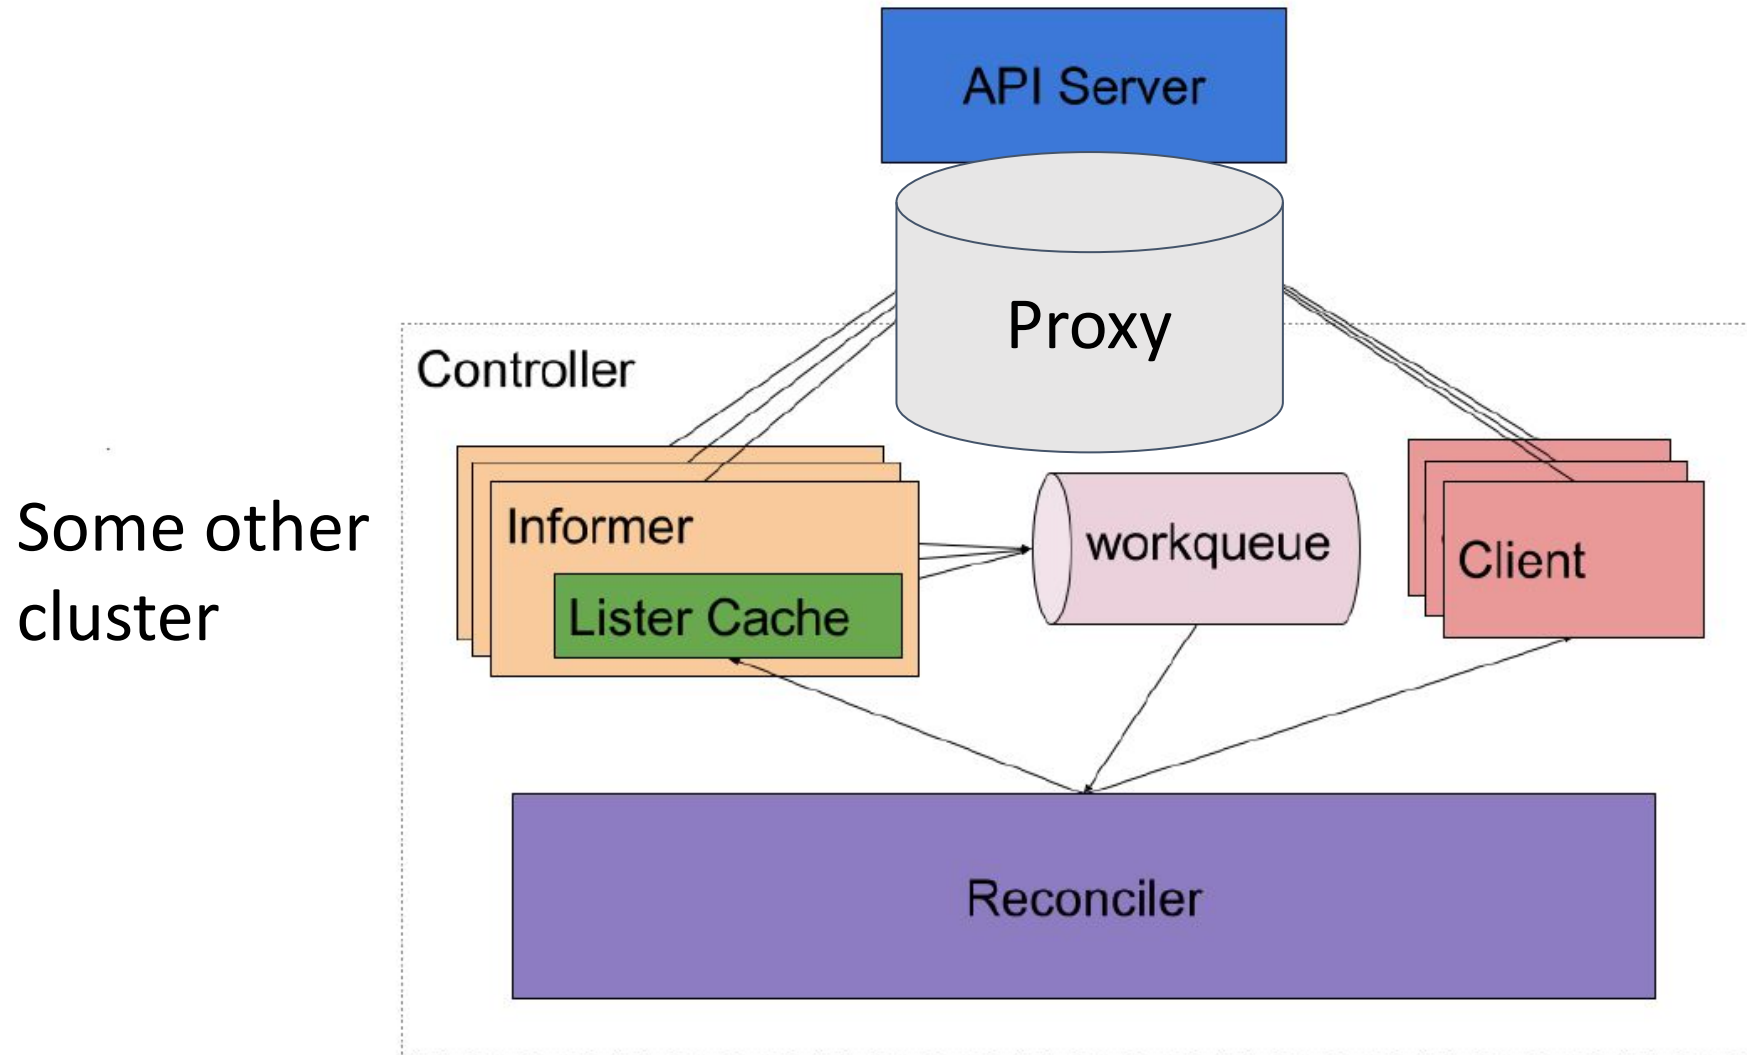


KubeCon



CloudNativeCon

North America 2023



What We Saw



KubeCon



CloudNativeCon

North America 2023

- The cluster appeared super dead
- Client requests were timing out
- GKE dashboard was dead
- Happened after 4-6 hours

The Cluster Killer Bug



KubeCon



CloudNativeCon

North America 2023

- Visibly only affected Kubernetes 1.25
- Only GKE
- Only "agentless"
- Only solution seemed to be deleting the cluster

WTF Is This?



KubeCon



CloudNativeCon

North America 2023

```
{
  "insertId": "twyb6f6sqrvstyfn",
  "jsonPayload": {
    "pid": "11",
    "message": "cachier (*unstructured.Unstructured): unexpected ListAndWatch error: failed to list
policy.sigstore.dev/v1beta1, Kind=ClusterImagePolicy: conversion webhook for
policy.sigstore.dev/v1alpha1, Kind=ClusterImagePolicy failed: the server was unable to return a
response in the time allotted, but may still be processing the request; reinitializing..."
  },
  "resource": {
    "type": "k8s_control_plane_component",
    "labels": {
      "component_name": "apiserver",
      "project_id": "eddiezane-chainguard",
      "cluster_name": "testing-zonal-again",
      "location": "us-central1-c",
      "component_location": "us-central1-c"
    }
  },
  "timestamp": "2023-02-16T16:02:59.841417Z",
  "severity": "ERROR",
  "logName": "projects/eddiezane-chainguard/logs/container.googleapis.com%2Fapiserver",
  "sourceLocation": {
    "file": "cachier.go",
    "line": "440"
  },
  "receiveTimestamp": "2023-02-16T16:03:03.472867173Z"
}
```

Or This??



KubeCon



CloudNativeCon

— North America 2023 —

[illegible]

API Server Memory



KubeCon



CloudNativeCon

North America 2023



First Theory



KubeCon



CloudNativeCon

North America 2023

Fixed: 22422 use singleflight to alleviate simultaneous calls to #112696

<> Code

Merged k8s-ci-robot merged 1 commit into `kubernetes:master` from `aimuz:fix-22422` on Oct 24, 2022

Conversation 41

Commits 1

Checks 0

Files changed 2

+118 -13



aimuz commented on Sep 23, 2022 • edited

Member

Signed-off-by: aimuz mr.aimuz@gmail.com

What type of PR is this?

/kind bug

What this PR does / why we need it:

Which issue(s) this PR fixes:

Fixes [#22422](#)

Special notes for your reviewer:

Does this PR introduce a user-facing change?

Fixed: #22422 Admission controllers can cause unnecessary significant load on apiserver



Additional documentation e.g., KEPs (Kubernetes Enhancement Proposals), usage docs, etc.:



Reviewers

aojea

liggitt

Assignees

liggitt

aojea

deads2k

Labels

approved cncf-cla: yes kind/bug lgdm

needs-priority needs-triage ok-to-test

release-note sig/scalability size/L

Projects

None yet

Milestone

v1.26

Development

Successfully merging this pull request may close these issues.

More Lessons



KubeCon



CloudNativeCon

North America 2023

- Have a baseline of what your logs + metrics look like
- Slow down
- Get creative and experiment

0xDEADBEEF



KubeCon



CloudNativeCon

North America 2023

Last 1 hour

foobar

×

Resource

Log name

Severity

Show query

1 foobar

2

Log fields

Log fields

Search fields and values

RESOURCE TYPE

Kubernetes Control Plane Component 2

Kubernetes Cluster 1

SEVERITY

Default 1

Error 1

Info 1

Query results 3 log entries

SEVERITY

TIMESTAMP ↑

MDT

SUMMARY

EDIT

Showing logs for last 1 hour from 4/6/23, 9:45 AM to 4/6/23, 10:45 AM.

Extend time by: 1 hour

Edit time

Find in results

Correlate by

Download

More actions

2023-04-06 10:43:19.313 MDT

"HTTP" verb="GET" URI="/api/foobar" latency="1m0.000926771s" userAgent="kubect1/v1.26.3 (linux/amd64) kubernetes/9e64410" audit-ID="23c8b52f-8df9-4d74-b9e7-e0b0a0ef116a" srcIP="75.166.125.163:37998" apf_pl="global-default" apf_fs="global-default" apf_iseats=1 apf_fseats=0 apf_additionallatency="0s" resp=504 statusStack=<goroutine 47652332 [running]:k8s.io/apiserver/pkg/server/httplog.(*respLogger).recordStatus(0xc015c57080, 0x0?)vendor/k8s.io/apiserver/pkg/server/httplog/httplog.go:324 +0x90k8s.io/apiserver/pkg/server/httplog.(*respLogger).WriteHeader(0xc015c57080, 0x1?)vendor/k8s.io/apiserver/pkg/server/httplog/httplog.go:306 +0x25k8s.io/apiserver/pkg/server/filters.(*baseTimeoutWriter).timeout(0xc01511a7e0, 0xc0152ca8c0)vendor/k8s.io/apiserver/pkg/server/filters/timeout.go:255 +0xd8k8s.io/apiserver/pkg/server/filters.(*timeoutHandler).ServeHTTP(0xc003434870, {0x56afe48, 0xc0195b7600}, 0xdf8475800?)vendor/k8s.io/apiserver/pkg/server/filters/timeout.go:146 +0x2e3k8s.io/apiserver/pkg/endpoints/filters.withRequestDeadline.func1({0x56afe48, 0xc0195b7600}, 0xc00e8b6900)vendor/k8s.io/apiserver/pkg/endpoints/filters/request_deadline.go:100 +0x362net/http.HandlerFunc.ServeHTTP(0x1?, {0x56afe48?, 0xc0195b7600?}, 0x1?)usr/local/go/src/net/http/server.go:2109 +0x2fk8s.io/apiserver/pkg/server/filters.withWaitGroup.func1({0x56afe48?, 0xc0195b7600}, 0xc00e8b6900)vendor/k8s.io/apiserver/pkg/server/filters/waitgroup.go:77 +0x766net/http.HandlerFunc.ServeHTTP(0x56b0fd8?, {0x56afe48?, 0xc0195b7600?}, 0x7f4ef9d3f5b8?)usr/local/go/src/net/http/server.go:2109 +0x2fk8s.io/apiserver/pkg/endpoints/filters.WithAuditAnnotations.func1({0x56afe48, 0xc0195b7600}, 0xc00e8b6800)vendor/k8s.io/apiserver/pkg/endpoints/filters/audit_annotations.go:36 +0x125net/http.HandlerFunc.ServeHTTP(0x56b0fd8?, {0x56afe48?, 0xc0195b7600?}, 0x566ae40?)usr/local/go/src/net/http/server.go:2109 +0x2fk8s.io/apiserver/pkg/endpoints/filters.WithWarningRecorder.func1({0x56afe48?, 0xc0195b7600}, 0xc00e8b6700)vendor/k8s.io/apiserver/pkg/endpoints/filters/warning.go:35 +0x18dnet/http.HandlerFunc.ServeHTTP(0x488efe0?, {0x56afe48?, 0xc0195b7600?}, 0xd?)usr/local/go/src/net/http/server.go:2109 +0x2fk8s.io/apiserver/pkg/endpoints/filters.WithCacheControl.func1({0x56afe48, 0xc0195b7600}, 0xc0195b75e0?)vendor/k8s.io/apiserver/pkg/endpoints/filters/cachecontrol.go:31 +0x126net/http.HandlerFunc.ServeHTTP(0x56b22a8?, {0x56afe48?, 0xc0195b7600?}, 0x...

But it's not dead?



KubeCon



CloudNativeCon

North America 2023



eddiezane 11:57 AM

And it's definitely not dead. I can ssh to a worker and reach it fine

```
gke-flow-schema-1-25-default-pool-d1a2791c-t06p ~ # kubectl --kubeconfig /var/lib/kubelet/kubeconfig get nodes
```

```
I0406 16:56:35.970985 564356 cache.go:40] re-using cached key and certificate
```

NAME	STATUS	ROLES	AGE	VERSION
gke-flow-schema-1-25-default-pool-d1a2791c-5s2j	Ready	<none>	2d	v1.25.7-gke.1000
gke-flow-schema-1-25-default-pool-d1a2791c-s8k4	Ready	<none>	2d	v1.25.7-gke.1000
gke-flow-schema-1-25-default-pool-d1a2791c-t06p	Ready	<none>	2d	v1.25.7-gke.1000

Bread Crumbs



KubeCon



CloudNativeCon

North America 2023

```
{
  "protoPayload": {
    "@type": "type.googleapis.com/google.cloud.audit.AuditLog",
    "authenticationInfo": {
      "principalEmail": "106836737900675175911"
    },
    "authorizationInfo": [
      {
        "granted": true,
        "permission": "io.k8s.admissionregistration.v1.mutatingwebhookconfigurations.update",
        "resource":
"admissionregistration.k8s.io/v1/mutatingwebhookconfigurations/enforcer.chainguard.dev"
      }
    ],
    "methodName": "io.k8s.admissionregistration.v1.mutatingwebhookconfigurations.update",
    "requestMetadata": {
      "callerIp": "34.132.193.40",
      "callerSuppliedUserAgent": "mcp/v0.0.0 (linux/amd64) kubernetes/$Format"
    },
    "resourceName":
"admissionregistration.k8s.io/v1/mutatingwebhookconfigurations/enforcer.chainguard.dev",
    "serviceName": "k8s.io",
    "status": {
      "code": 8,
      "message": "Too Many Requests"
    },
    "insertId": "545a61a0-0a7c-4071-877d-95736172c3d7",
    "resource": {
      "type": "k8s_cluster",
      "labels": {
        "project_id": "eddiezane-chainguard",
        "cluster_name": "prod-demo",
        "location": "us-central1-c"
      }
    },
    "timestamp": "2023-02-15T23:37:43.748760Z",
    "labels": {
      "apiserver.latency.k8s.io/response-write": "1.53µs",
      "apiserver.latency.k8s.io/total": "1m0.001364292s"
    },
    "logName": "projects/eddiezane-chainguard/logs/cloudaudit.googleapis.com%2Factivity",
    "operation": {
      "id": "545a61a0-0a7c-4071-877d-95736172c3d7",
      "producer": "k8s.io",
      "first": true,
      "last": true
    },
    "receiveTimestamp": "2023-02-15T23:37:45.222199510Z"
  }
}
```

API Server Tuning



KubeCon



CloudNativeCon

North America 2023

- `--max-requests-inflight`
- `--max-mutating-requests-inflight`

API Priority and Fairness



KubeCon



CloudNativeCon

North America 2023

The API Priority and Fairness feature (APF) is an alternative that improves upon aforementioned max-inflight limitations. APF classifies and isolates requests in a more fine-grained way. It also introduces a limited amount of queuing, so that no requests are rejected in cases of very brief bursts.

API Priority and Fairness



KubeCon



CloudNativeCon

North America 2023

Without APF enabled, overall concurrency in the API server is limited by the *kube-apiserver* flags *--max-requests-inflight* and *--max-mutating-requests-inflight*. With APF enabled, the concurrency limits defined by these flags are summed and then the sum is divided up among a configurable set of priority levels. Each incoming request is assigned to a single priority level, and each priority level will only dispatch as many concurrent requests as its particular limit allows.

API Priority and Fairness



KubeCon




CloudNativeCon

North America 2023

- Beta in 1.20
- Enabled by default
- How you control traffic in an overload situation
- "It doesn't matter until it does"

- `flowcontrol.apiserver.k8s.io/v1beta3`
- Similar to RBAC. The "who" and "what"
 - Subjects - Users, Service Accounts, Groups
 - (non)Resources - GVRs (Pods, Deployments), Paths (/metrics)
- Matching precedence
- Maps to a `PriorityLevelConfiguration`



```
apiVersion: flowcontrol.apiserver.k8s.io/v1beta2
kind: FlowSchema
metadata:
  name: demo
spec:
  distinguisherMethod:
    type: ByUser
  matchingPrecedence: 1000
  priorityLevelConfiguration:
    name: demo
  rules:
    - nonResourceRules:
        - nonResourceURLs:
            - '*'
          verbs:
            - '*'
      resourceRules:
        - apiGroups:
            - '*'
          clusterScope: true
          namespaces:
            - '*'
          resources:
            - '*'
          verbs:
            - '*'
    subjects:
      - serviceAccount:
          name: demo
          namespace: default
          kind: ServiceAccount
```

PriorityLevelConfiguration



KubeCon



CloudNativeCon

North America 2023

- "Limited" or "Exempt"
- "Queue" or "Reject"
- Knobs to turn for queueing



```
apiVersion: flowcontrol.apiserver.k8s.io/v1beta2
kind: PriorityLevelConfiguration
metadata:
  name: demo
spec:
  limited:
    assuredConcurrencyShares: 30
    limitResponse:
      queuing:
        handSize: 5
        queueLengthLimit: 20
        queues: 16
      type: Queue
    type: Limited
```


QueuingConfiguration



KubeCon



CloudNativeCon

North America 2023

- `assuredConcurrencyShares`
- `handSize`
- `queueLengthLimit`
- `queues`

QueuingConfiguration



KubeCon



CloudNativeCon

North America 2023

- Increasing *queues* reduces the rate of collisions between different flows, at the cost of increased memory usage. A value of 1 here effectively disables the fair-queuing logic, but still allows requests to be queued.
- Increasing *queueLengthLimit* allows larger bursts of traffic to be sustained without dropping any requests, at the cost of increased latency and memory usage
- Changing *handSize* allows you to adjust the probability of collisions between different flows and the overall concurrency available to a single flow in an overload situation.



- The **workload-low** priority level is for requests from any other service account, which will typically include all requests from controllers running in Pods.
- The **global-default** priority level handles all other traffic, e.g. interactive kubectl commands run by nonprivileged users.

Special Cases



KubeCon



CloudNativeCon

North America 2023

- Lists
- Watches
- `system:masters`

Health Checks



KubeCon



CloudNativeCon

North America 2023



```
apiVersion: flowcontrol.apiserver.k8s.io/v1beta2
kind: FlowSchema
metadata:
  name: health-for-strangers
spec:
  matchingPrecedence: 1000
  priorityLevelConfiguration:
    name: exempt
  rules:
    - nonResourceRules:
      - nonResourceURLs:
          - "/healthz"
          - "/livez"
          - "/readyz"
        verbs:
          - "*"
    subjects:
      - kind: Group
        group:
          name: "system:unauthenticated"
```

<https://kubernetes.io/docs/concepts/cluster-administration/flow-control/>

Levels



KubeCon



CloudNativeCon

North America 2023

```
cluster-killer-bug/repro k get prioritylevelconfigurations.flowcontrol.apiserver.k8s.io
```

NAME	TYPE	ASSURED	CONCURRENCY	SHARES	QUEUES	HANDSIZE	QUEUELENGTHLIMIT	AGE
catch-all	Limited	5			<none>	<none>	<none>	48m
exempt	Exempt	<none>			<none>	<none>	<none>	48m
global-default	Limited	20			128	6	50	48m
leader-election	Limited	10			16	4	50	48m
node-high	Limited	40			64	6	50	48m
system	Limited	30			64	6	50	48m
workload-high	Limited	40			128	6	50	48m
workload-low	Limited	100			128	6	50	48m

```
cluster-killer-bug/repro k get prioritylevelconfigurations.flowcontrol.apiserver.k8s.io
```

NAME	TYPE	ASSURED CONCURRENCY SHARES	QUEUES	HANDSIZE	QUEUE LENGTH LIMIT	AGE
catch-all	Limited	5	<none>	<none>	<none>	48m
exempt	Exempt	<none>	<none>	<none>	<none>	48m
global-default	Limited	20	128	6	50	48m
leader-election	Limited	10	16	4	50	48m
node-high	Limited	40	64	6	50	48m
system	Limited	30	64	6	50	48m
workload-high	Limited	40	128	6	50	48m
workload-low	Limited	100	128	6	50	48m

245

ACS = 245

--max-requests-inflight = 400 (default)

--max-mutating-requests-inflight = 200 (default)

workload-low = 100

$(400 + 200) / 245 * 100 = 244$ requests per second

Add another PriorityLevel with ACS 55

ACS = 300

$(400 + 200) / 300 * 100 = 200$ requests per second workload-low

$(400 + 200) / 300 * 55 = 110$ requests per second new level

Demo



KubeCon



CloudNativeCon

North America 2023

So What Was The Bug?



KubeCon



CloudNativeCon

North America 2023

- Off cluster traffic was getting bucketed in `global-default`
- Fix was to create APF resources
- When would you want to do this?

Hmm



KubeCon



CloudNativeCon

North America 2023

- Why not 1.26?

<https://kubernetes.io/docs/concepts/cluster-administration/flow-control/#suggested-configuration-objects>

Borrowing



KubeCon



CloudNativeCon

North America 2023

Add borrowing between priority levels in APF #113485

<> Code

Merged

k8s-ci-robot merged 8 commits into `kubernetes:master` from `MikeSpreitzer:apf-borrowing` on Nov 9, 2022

Conversation 85

Commits 8

Checks 0

Files changed 72

+2,226 -459



MikeSpreitzer commented on Oct 31, 2022 • edited

Member

What type of PR is this?

/kind feature

What this PR does / why we need it:

This PR implement the design change in [kubernetes/enhancements#3391](#) and [kubernetes/enhancements#3479](#). That is, this PR adds borrowing of concurrency between priority levels in the API Priority and Fairness feature.

Which issue(s) this PR fixes:

Fixes #

Special notes for your reviewer:

The API change part of this was already reviewed in [#112830](#).

The first six commits of this PR are the commits of that earlier PR, rebased onto a later revision of the `master` branch.

The "apiserver: update borrowing parameters for apf bootstrap objects" commit is a cherry-pick of [#113016](#).

Does this PR introduce a user-facing change?

```
Priority and Fairness has introduced a new feature called _borrowing_ that allows an API priority level to borrow a number of seats from other priority level(s). As a cluster operator, you can enable borrowing for a certain priority level configuration object via the two newly introduced fields 'lendablePercent', and 'borrowingLimitPercent' located under the '.spec.limited' field of the designated priority level. This PR adds the following metrics.
```

- `'apiserver_flowcontrol_nominal_limit_seats'`: Nominal number of execution seats configured for each priority level
- `'apiserver_flowcontrol_lower_limit_seats'`: Configured lower bound on number of execution seats available to each priority level
- `'apiserver_flowcontrol_upper_limit_seats'`: Configured upper bound on number of execution seats available to each priority level
- `'apiserver_flowcontrol_demand_seats'`: Observations, at the end of every nanosecond, of (the number of seats requested by) each priority level
- `'apiserver_flowcontrol_demand_seats_high_watermark'`: High watermark, over last adjustment period, of demand seats
- `'apiserver_flowcontrol_demand_seats_average'`: Time-weighted average, over last adjustment period, of demand seats
- `'apiserver_flowcontrol_demand_seats_stddev'`: Time-weighted standard deviation, over last adjustment period, of demand seats
- `'apiserver_flowcontrol_demand_seats_smoothed'`: Smoothed seat demands
- `'apiserver_flowcontrol_target_seats'`: Seat allocation targets
- `'apiserver_flowcontrol_seat_fair_frac'`: Fair fraction of server's concurrency to allocate to each priority level
- `'apiserver_flowcontrol_current_limit_seats'`: current derived number of execution seats available to each priority level

The possibility of borrowing means that the old metric `apiserver_flowcontrol_request_concurrency_limit` can no longer be used to determine the number of execution seats available to each priority level.

Reviewers

lavalamp

logicalhan

ligitt

wojtek-t

dgrisonnet

andrewsykim

deads2k

tkashem

Assignees

lavalamp

wojtek-t

Labels

approved

area/apiserver

area/code-generation

area/test

cncf-cla: yes

kind/api-change

kind/feature

lgdm

needs-priority

release-note

sig/api-machinery

sig/testing

size/XXL

triage/accepted

Projects

API Reviews

Status: API review completed, 1.26

+2 more

Milestone

v1.26

Changes in 1.26



KubeCon



CloudNativeCon

North America 2023

- ``borrowingLimitPercent``
- ``lendablePercent``
- ``nominalConcurrencyShares``
- $$\text{NominalCL}(i) = \text{ceil}(\text{ServerCL} * \text{NCS}(i) / \text{sum_ncs})$$
$$\text{sum_ncs} = \text{sum}[\text{limited priority level } k] \text{NCS}(k)$$
- Bigger numbers mean a larger nominal concurrency limit, at the expense of every other Limited priority level.



- Why did the GKE Dashboard die?

Ohhhhh



KubeCon



CloudNativeCon

North America 2023

```
{
  insertId: "qwss6d19ybb1rqjz"
  jsonPayload: {
    message:
      "HTTP verb=LIST URI=/api/v1/nodes?resourceVersion=0 latency=944.552µs userAgent=GoogleContainerEngine audit-ID=d57291af-86a3-4f2f-ac57-135dca9c5c25"
      srcIP=127.0.0.1:38710 apf_pl=global-default apf_fs=global-default apf_iseats=1 apf_fseats=0 apf_additionalLatency=0s apf_execution_time=534.969µs resp=200"
    pid: "11"
  }
  logName: "projects/eddiezane-chainguard/logs/container.googleapis.com%2Fapiserver"
  receiveTimestamp: "2023-04-18T13:55:29.464435182Z"
  resource: {
    labels: {
      cluster_name: "flow-schema-2-electric-boogaloo"
      component_location: "us-central1-c"
      component_name: "apiserver"
      location: "us-central1-c"
      project_id: "eddiezane-chainguard"
    }
    type: "k8s_control_plane_component"
  }
  severity: "INFO"
  sourceLocation: {
    file: "httplog.go"
    line: "131"
  }
  timestamp: "2023-04-18T13:55:26.457251Z"
}
```

Hide log summary

Collapse nested fields

Copy to clipboard

Copy link

Similar er

Sidequest



KubeCon



CloudNativeCon

North America 2023

```
▼ {
  insertId: "097cdaa2-c853-4d13-b9a8-9c6dc44882e7"
  ▶ labels: {2}
  logName: "projects/eddiezane-chainguard/logs/cloudaudit.googleapis.com%2Fdata_access"
  ▶ operation: {3}
  ▼ protoPayload: {
    @type: "type.googleapis.com/google.cloud.audit.AuditLog"
    ▼ authenticationInfo: {
      principalEmail: "106836737900675175911"
    }
    ▶ authorizationInfo: [1]
    methodName: "io.k8s.batch.v1.cronjobs.watch"
    ▶ requestMetadata: {2}
    resourceName: "batch/v1/cronjobs"
    serviceName: "k8s.io"
    ▶ status: {1}
  }
  receiveTimestamp: "2023-04-06T19:09:40.415993919Z"
  ▶ resource: {2}
  timestamp: "2023-04-06T19:09:36.659037Z"
}
```

More More Lessons



KubeCon



CloudNativeCon

North America 2023

- Being clever means you need to do more research
- Need people to test new things and give feedback
- Read the docs and release notes
- Get involved upstream!!

To Be Continued?



KubeCon



CloudNativeCon

North America 2023

Thursday, April 13th ▾



eddiezane 10:09 AM

Calling this one solved. I'll send over the write up when I finish it but it was an APF issue. Our watches were clogging up the global-default bucket. Creating our own FlowSchema and queue is the fix. It worked on 1.26 because of that bucket sharing change.

Thanks for all your help on this one!

Future kubecon talk

| [Learning API Priority and Fairness the Hard Way](#)



liggitt 10:18 AM

mm

I'm still really surprised you were getting request timeout issues, not P&F 429 responses that doesn't seem right

Shoutouts



KubeCon



CloudNativeCon

North America 2023

- Jordan Liggitt (Google)
- Mike Danese (Google)
- Billy Lynch (Chainguard)

Thank You!



KubeCon



CloudNativeCon

North America 2023

- Questions?
- @eddiezane

Please scan the QR Code ->
to leave feedback on this session

