



KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

Turbocharging Compilation Times: Kubernetes Orchestration with Distributed Compilation Tools

Diogo Guerra, formerly @ CERN

Who



KubeCon



CloudNativeCon


North America 2023

Diogo Guerra

Former Computing Engineer, Kubernetes team @ CERN

www.linkedin.com/in/diogoguerraist

dy090.guerra@gmail.com






KubeCon



CloudNativeCon

———— North America 2023 ————



Once Upon a Time...



KubeCon



CloudNativeCon

North America 2023




Source: <https://www.reddit.com/media?url=https%3A%2F%2Fi.redd.it%2F38iqur4116x41.jpg>

- What makes compilation time-consuming?
- When to go distributed? Are there ready-to-use solutions available?
- Can we use Kubernetes to scale a distributed compilation service on-demand?

CERN develops and maintains many applications:

- ROOT - framework to analyse and visualize large amounts of data
- CVMFS - read-only file system for software distribution

These solutions are usually big and resource-consuming to compile.



CERN Root homepage

Teams have constrained resources:

- Cost of infrastructure for the team/organization
 - Expensive to set up, maintain and run
- Users' default is to opt for the easiest solution



CERN CVMFS homepage

Compute challenges:

- Small teams have lower budget for resources
- Big computer node (64+ threads) or distributed compilation required

Organization challenges:

- Solution easily integratable with existing compile definitions/workflows
- No 'standard' way offered to the community
- Use of public machines in shared Hypervisors for Heavy Workloads

Landscape of Some Distributed Compile Tools



KubeCon



CloudNativeCon

North America 2023

Criteria:

- Compatibility
- Ease of configuration
- Does it work with Kubernetes

Explored tools:

- distcc
- icecc
- sccache
- gg



Source: <https://twitter.com/AmyfStuart/status/735108951138766849/photo/1>

Setup - distcc

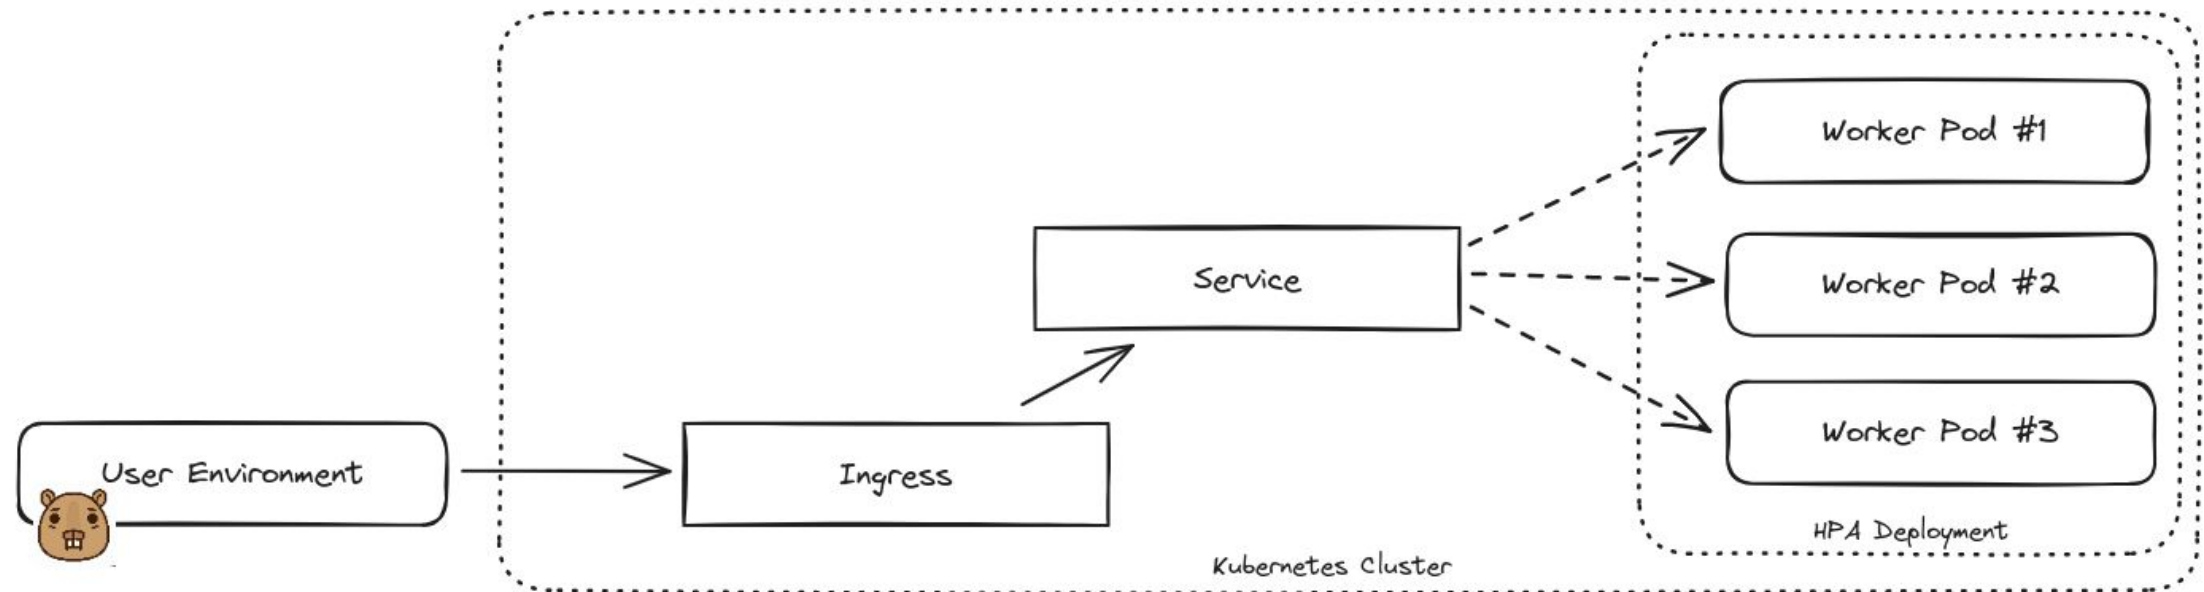


KubeCon



CloudNativeCon

North America 2023



Key Takeaways - distcc



KubeCon



CloudNativeCon

North America 2023

Pros:

- Most known and widely used tool

Cons:

- distcc is static
- Kubernetes service will assign jobs to random workers
- The client controls the whole compile job

Setup - icecream

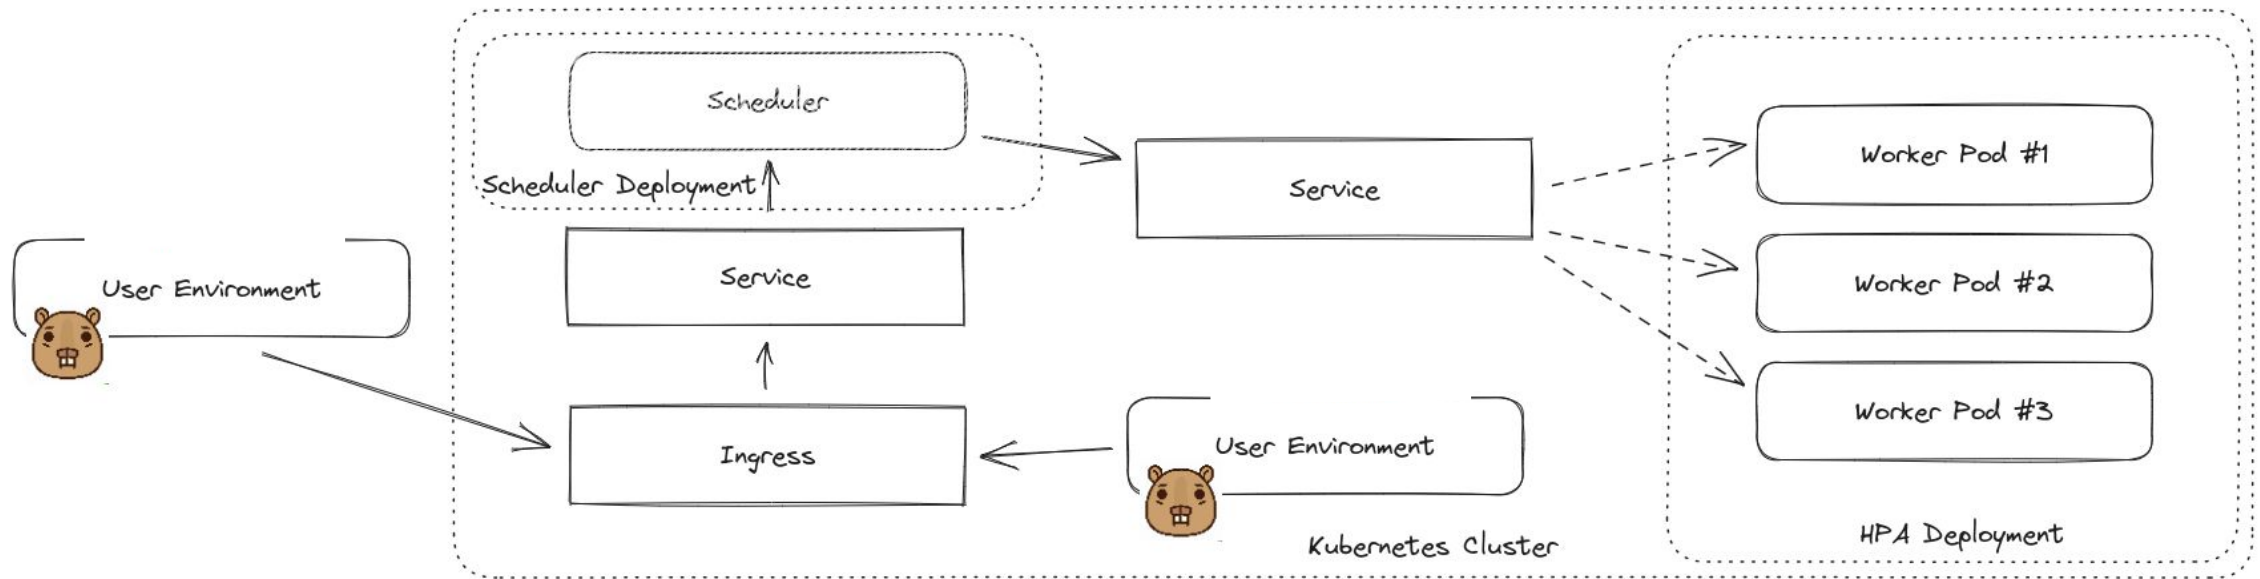


KubeCon



CloudNativeCon

North America 2023



Testing - icecream




KubeCon




CloudNativeCon

North America 2023

8 CPU - Scheduler & 2 CPU - Workers



8 CPU - Scheduler & 1 CPU - Workers



Key Takeaways - icecream



KubeCon



CloudNativeCon

North America 2023

Pros:

- Icecc supports worker discovery
- Failed jobs fallback to local computer

Cons:

- Requirement of setting up dependencies
- Scaling is a problem



KubeCon



CloudNativeCon

North America 2023

No real stable solution...



Source: <https://images.app.goo.gl/SEr8n38eKqioXQ9s8>

Setup - sccache

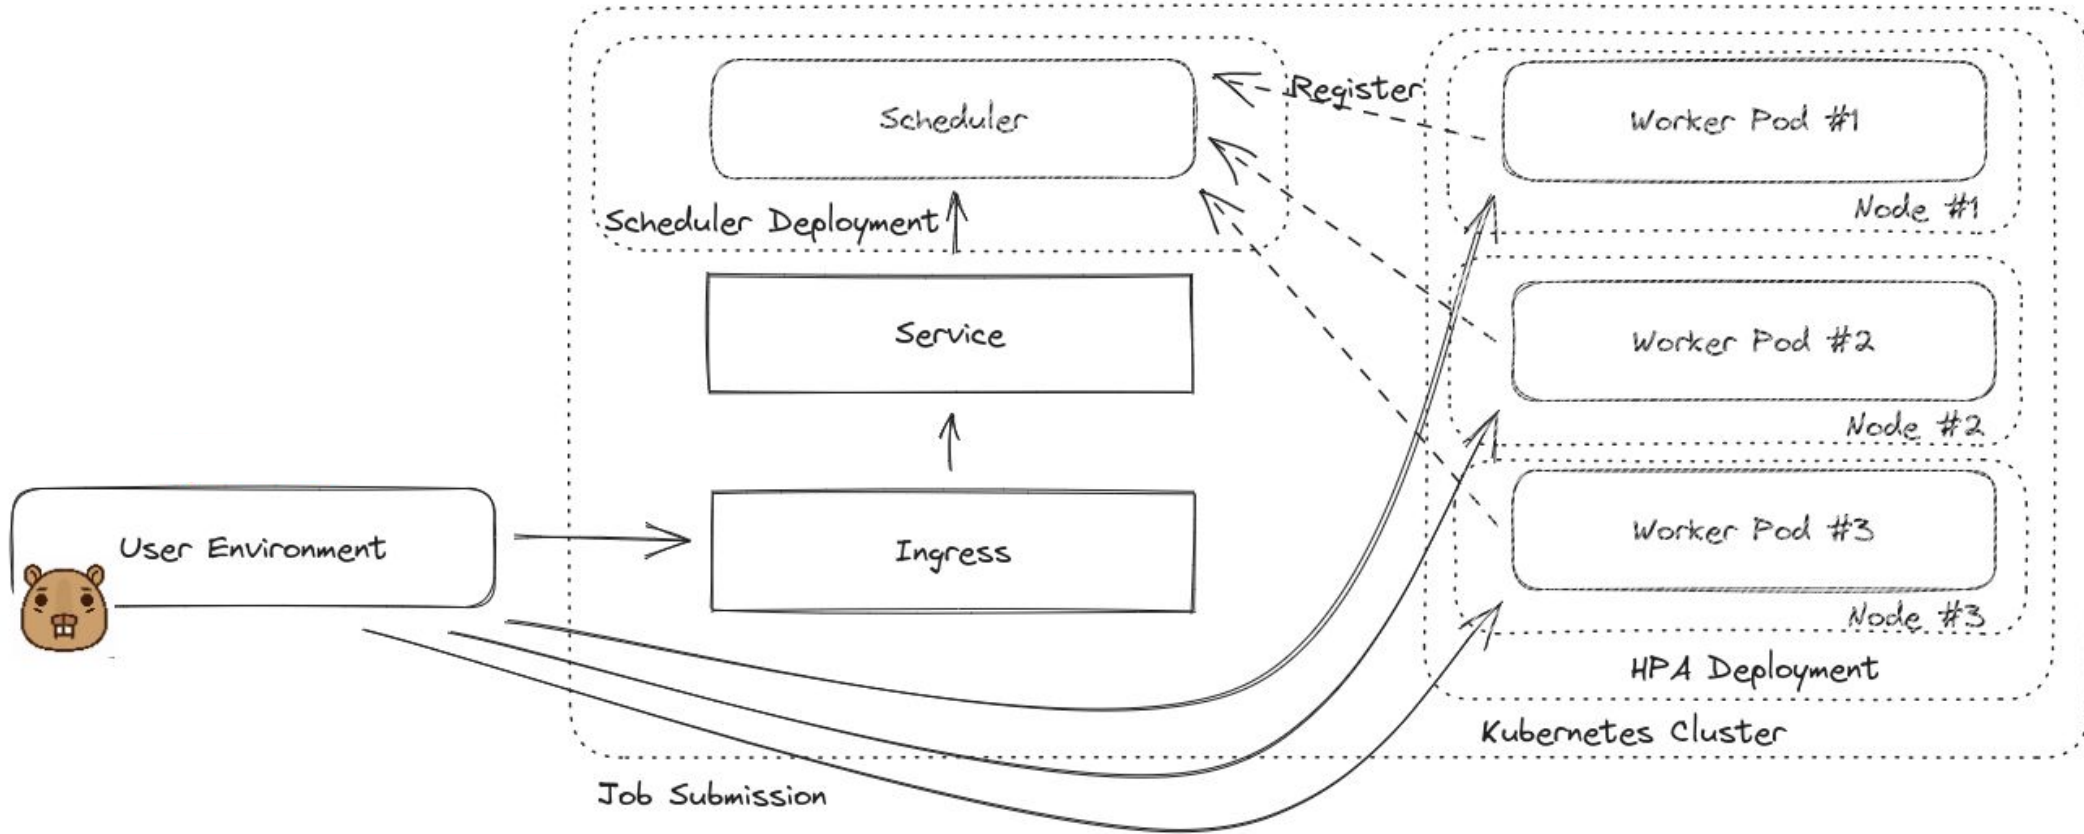


KubeCon



CloudNativeCon

North America 2023



Testing - sccache

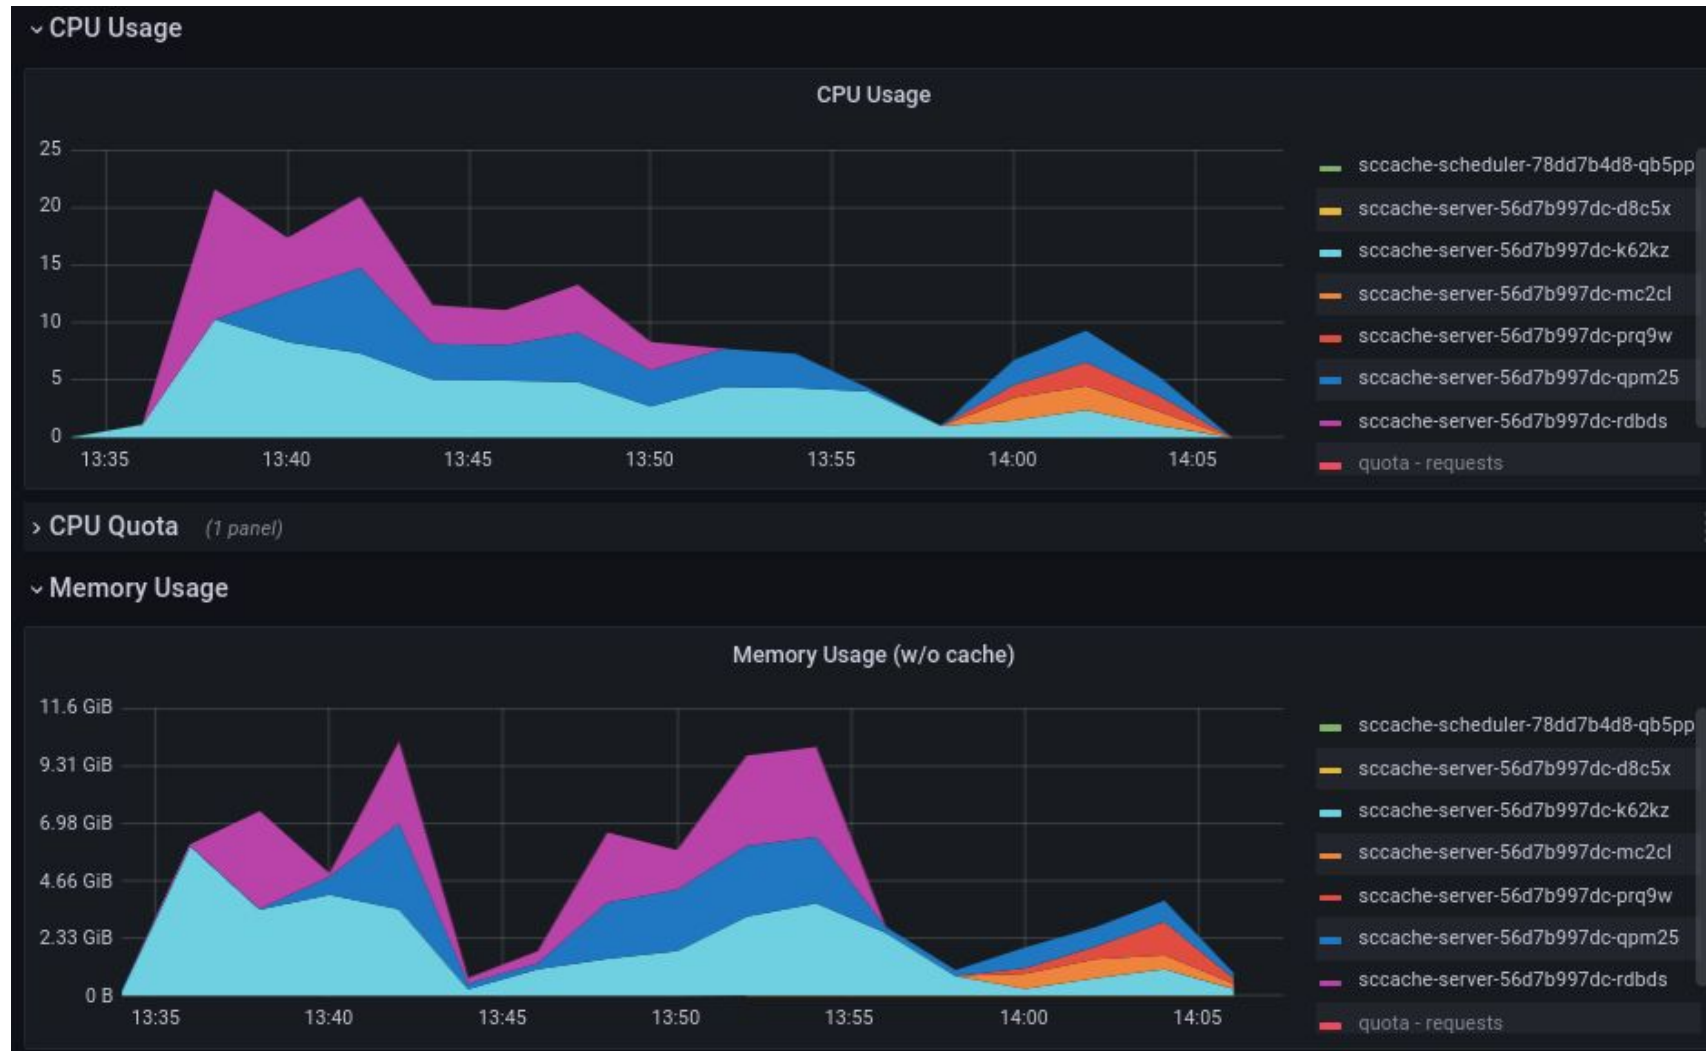


KubeCon



CloudNativeCon

North America 2023



Key Takeaways - sccache



KubeCon



CloudNativeCon

North America 2023

Pros:

- No dependency management needed
- Compatible with rust projects

Cons:

- Client needs access to worker node
- sccache uses bubblewrap
 - Access to namespaces & privileges
- Job limit not supported



KubeCon



CloudNativeCon

North America 2023

Can we go 'Cloud Native?'



Source: https://www.reddit.com/r/capybara/comments/vmt4ax/capybara_on_vacation/

Compilation Database



KubeCon



CloudNativeCon

North America 2023

- Many Compilation Tools create a compilation database.

```
dtomasgu@r2d2:~/ws/github.com/root/_build$ cat compile_commands.json | jq length
2604
dtomasgu@r2d2:~/ws/github.com/root/_build$ cat compile_commands.json | jq '.[0]'
{
  "directory": "/home/dtomasgu/ws/github.com/root/_build/interpreter/llvm-project/llvm/lib/Demangle",
  "command": "/usr/bin/c++ -D_GNU_SOURCE -D__STDC_CONSTANT_MACROS -D__STDC_FORMAT_MACROS -D__STDC_LIMIT_MACROS -I/
home/dtomasgu/ws/github.com/root/_build/interpreter/llvm-project/llvm/lib/Demangle -I/home/dtomasgu/ws/github.com/
root/interpreter/llvm-project/llvm/lib/Demangle -I/home/dtomasgu/ws/github.com/root/_build/interpreter/llvm-projec
t/llvm/include -I/home/dtomasgu/ws/github.com/root/interpreter/llvm-project/llvm/include -std=c++17 -Wno-implicit
-fallthrough -Wno-noexcept-type -pipe -Wshadow -Wall -W -Woverloaded-virtual -fsigned-char -pthread -fno-semantic
-interposition -fPIC -fno-semantic-interposition -fvisibility-inlines-hidden -Werror=date-time -w -ffunction-sect
ions -fdata-sections -O3 -DNDEBUG -fvisibility=hidden -fvisibility-inlines-hidden -fno-exceptions -std=c++17 -o C
MakeFiles/LLVMDemangle.dir/Demangle.cpp.o -c /home/dtomasgu/ws/github.com/root/interpreter/llvm-project/llvm/lib/D
emangle/Demangle.cpp",
  "file": "/home/dtomasgu/ws/github.com/root/interpreter/llvm-project/llvm/lib/Demangle/Demangle.cpp"
}
```



KubeCon



CloudNativeCon

North America 2023

We were having lunch...



Source: https://www.youtube.com/watch?v=UZeaR2N4l8c&ab_channel=%E9%95%B7%E5%B4%8E%E3%83%90%E3%82%A4%E3%82%AA%E3%83%91%E3%83%BC%E3%82%AF%E5%85%AC%E5%BC%8F

Setup - gg

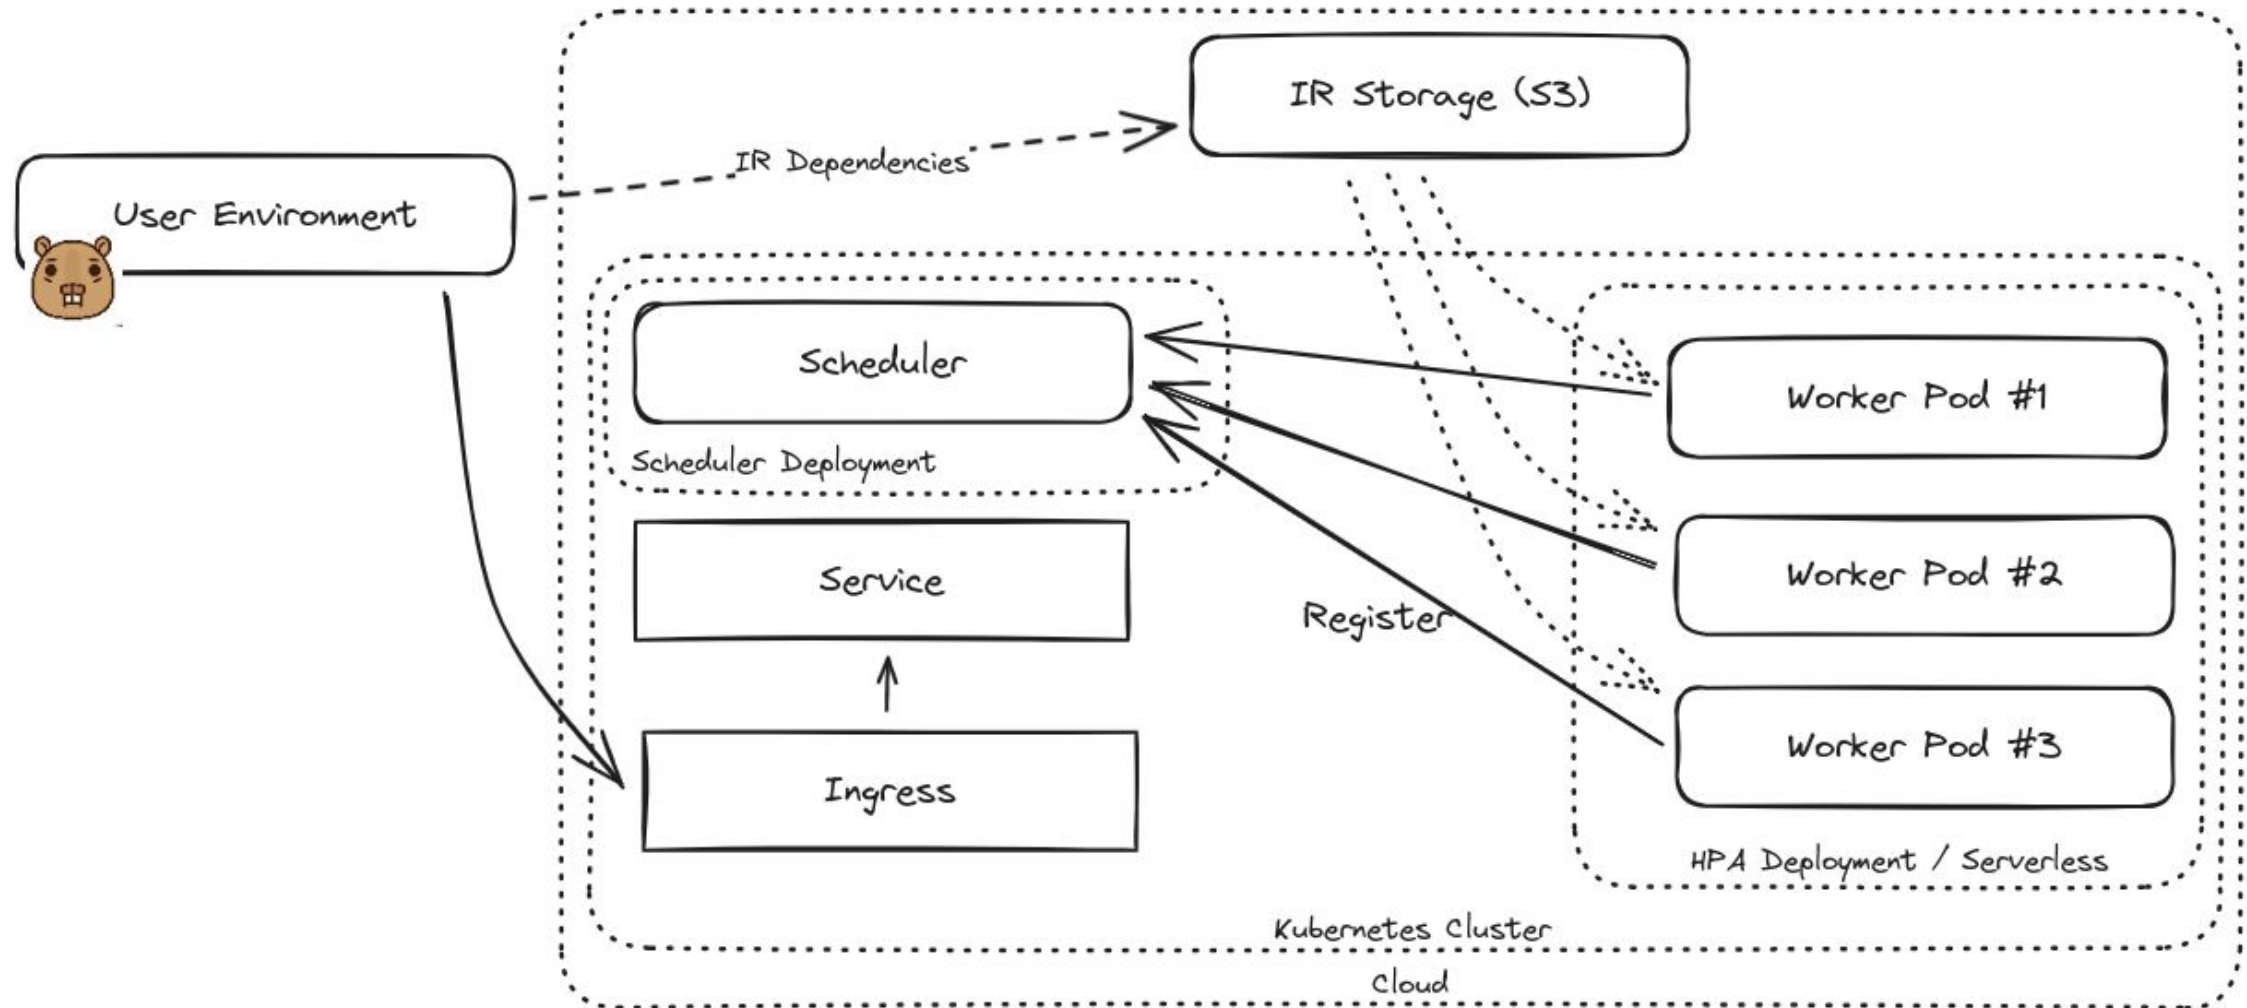


KubeCon



CloudNativeCon

North America 2023



Testing - gg



KubeCon



CloudNativeCon

North America 2023

```
| preprocessed: TjSTQgdT0ChDax2RE59j2juf9AgINFM2VDw7rLUkAhDE0000bdb1
| compiled: Tosbia_zkQxh0Lx3aE4f0DodublMM838kqFfAYkIAX500000066e
| assembled: T_XKtAIx6RQ0JvDKixD9VjGSoFSUnAHhZh51sU9AI200000064c
| output: CMakeFiles/LLVMX86AsmParser.dir/X86AsmParser.cpp.o
[100%] Linking CXX static library ../../../../libLLVMX86AsmParser.a
[100%] Built target LLVMX86AsmParser
[100%] Building CXX object tools/llvm-ar/CMakeFiles/llvm-ar.dir/llvm-ar.cpp.o
| generating model for g++
| input: /home/dtomasgu/ws/github.com/gg/examples/llvm/llvm-project/llvm/tools/llvm-ar/llvm
-ar.cpp
| generating make dependencies file... done.
| preprocessed: TxcAt4A8LueqSrDLI03TP492eiurVs3nXePUTZ99Qdzs00009c69
| compiled: Tmkfo2vfDNrYEVvgD8Nu4dDR7Qriu21N1f7ci2gBAfz8000005bb
| assembled: TDcdQsS0NZ.nDrrrZYWP6kfA4_g0ZHAJbZPFFGJq08pc0000058b
| output: CMakeFiles/llvm-ar.dir/llvm-ar.cpp.o
[100%] Linking CXX executable ../../bin/llvm-ar
| generating model for g++
| linked: TnQwFE8RniDfhV.w8CvDL0An_WqKCUaLiW0dcLB.NlKk00002d90
| output: ../../bin/llvm-ar
[100%] Built target llvm-ar
→ Loading the thunks... done (116 ms).
↗ Uploading 1715 files (94.8 MiB)... done (2975 ms).

in queue: 6      λ (1000): 388    done: 841    remaining: 850 | cost: ~$0.02
```

Key Takeaways - gg



KubeCon



CloudNativeCon

North America 2023

Pros:

- Fastest or equal speed
- Fits very well in the 'Cloud Native' ideology
- Scalable and flexible
- No dependency management needed

Cons:

- Recent project (2019) with no active users/development
- Some lack of supporting documentation
- Integration work required for more complex projects

Conclusions



KubeCon



CloudNativeCon


North America 2023

- Bigger nodes that accommodate the workload are ALWAYS preferable
- Distributed compilations are useful when:
 - User compile machine is not powerful enough
 - Infrastructure is intended to be shared across an organization
- Legacy distributed compilation tools do not work well in modern architectures
- New paradigm to take advantage of the dynamic presented by cloud native environments

SCCACHE Helm



StanfordSNR/gg src and Paper



Special thanks



KubeCon




CloudNativeCon

North America 2023



Source: <https://www.reddit.com/media?url=https%3A%2F%2Fi.redd.it%2Fxfkifu5yylqx41.jpg>



Diana Gaponcic




KubeCon



CloudNativeCon

———— North America 2023 ————



Thank you!