



KubeCon



CloudNativeCon

North America 2023





KubeCon



CloudNativeCon

North America 2023

Learn How to Build an eBPF CNI Plugin from Scratch

Adam Sayah, Solo.io

Speaker



KubeCon



CloudNativeCon

North America 2023

Adam Sayah

Solo.io

_asayah

linkedin.com/in/adamsayah



Agenda



KubeCon



CloudNativeCon

North America 2023

- **What is a CNI ?**
- **Exercise 1: Create a CNI**
- **Exercise 2: eBPF Basics**
- **Exercise 3 : eBPF Maps / Monitoring**
- **Exercise 4: eBPF for Security.**
- **Conclusion**

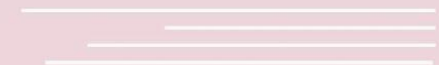


KubeCon



CloudNativeCon

———— North America 2023 ————



What is a CNI ?

What is a CNI?

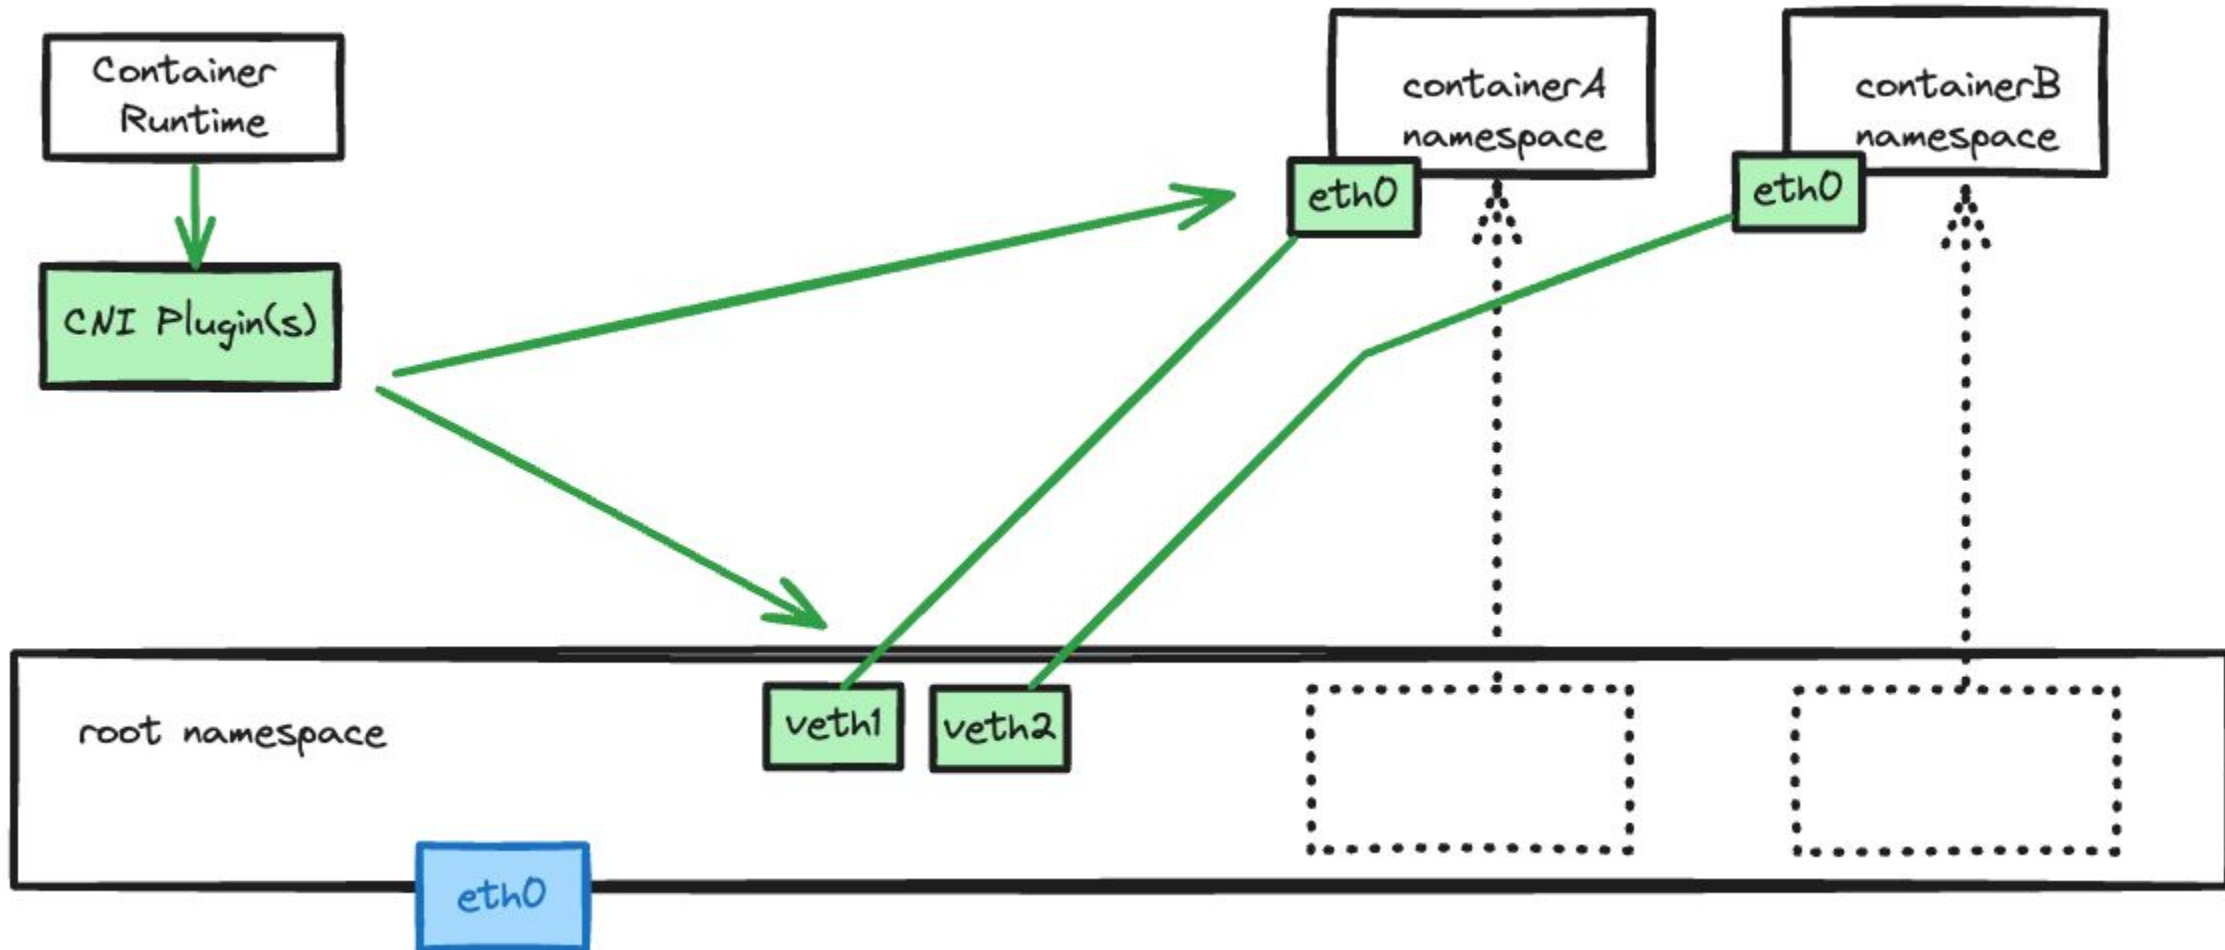



KubeCon



CloudNativeCon

North America 2023





CNI

<https://www.cni.dev>



- CNI Specification
- Implementations
- Lib

CNI - Spec

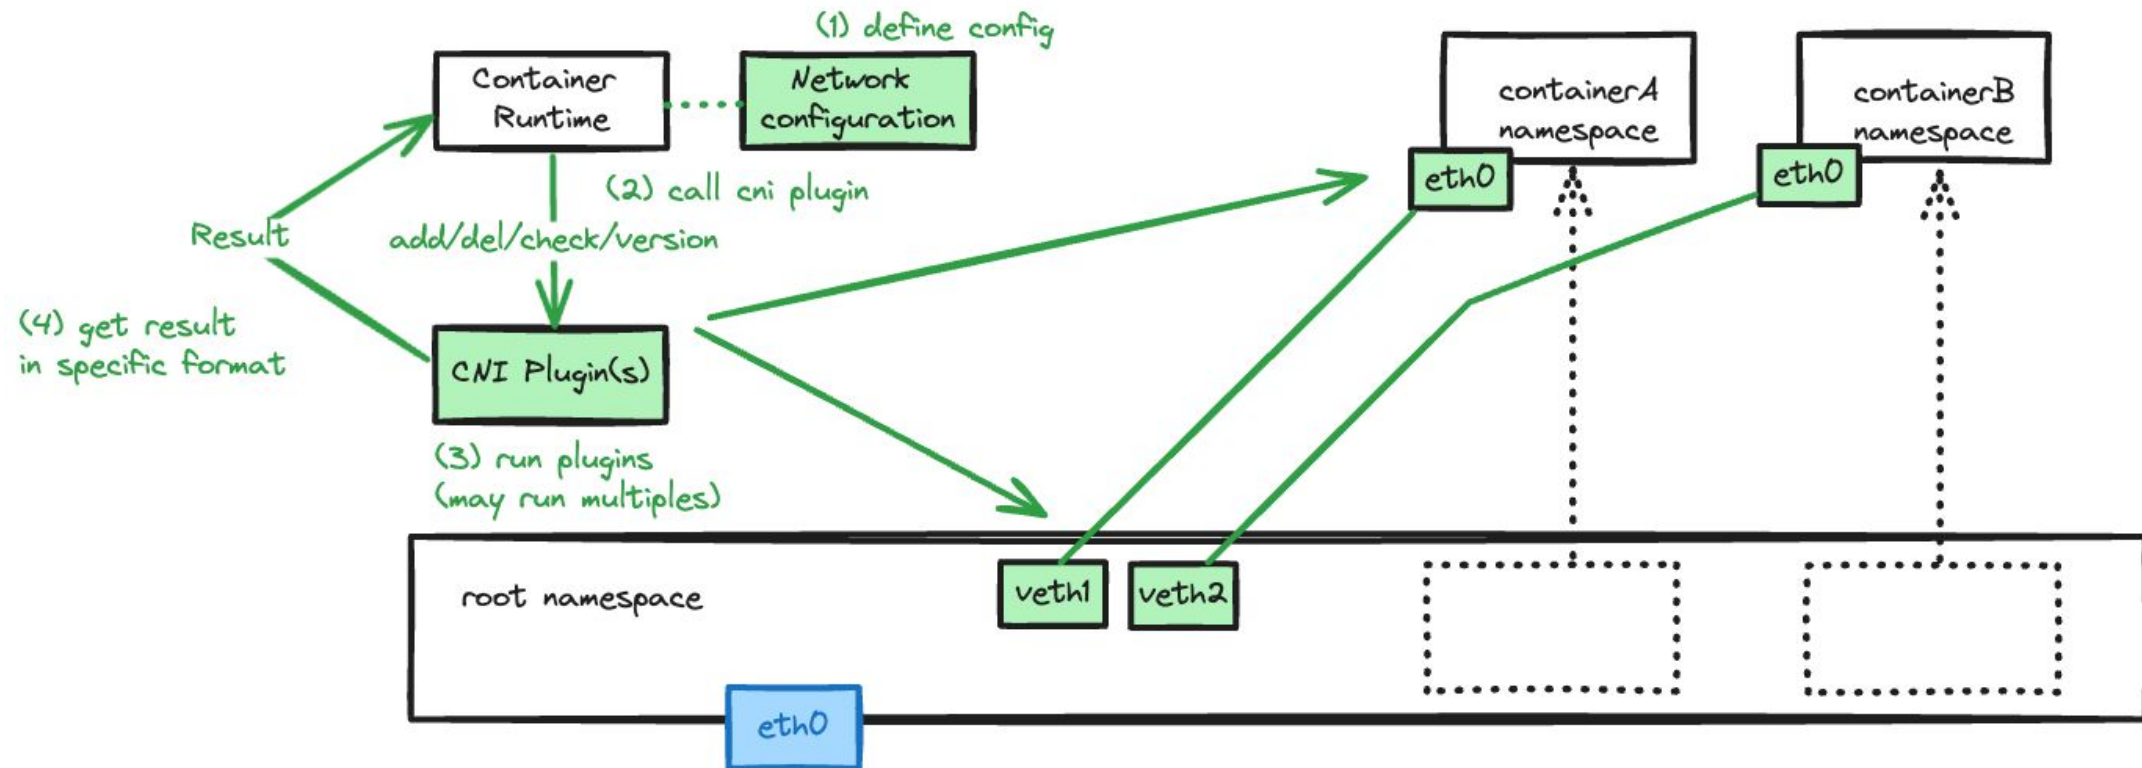


KubeCon



CloudNativeCon

North America 2023



CNI - Configuration



KubeCon



CloudNativeCon

North America 2023



CNI - Result



KubeCon



CloudNativeCon

North America 2023

```
{
  "cniVersion": "1.0.0",
  "name": "dbnet",
  "type": "bridge",
  "bridge": "cni0",
  "keyA": ["some more", "plugin specific", "configuration"],
  "ipam": {
    "type": "host-local",
    "subnet": "10.1.0.0/16",
    "gateway": "10.1.0.1"
  },
  "dns": {
    "nameservers": [ "10.1.0.1" ]
  }
}
```

CNI - Execution

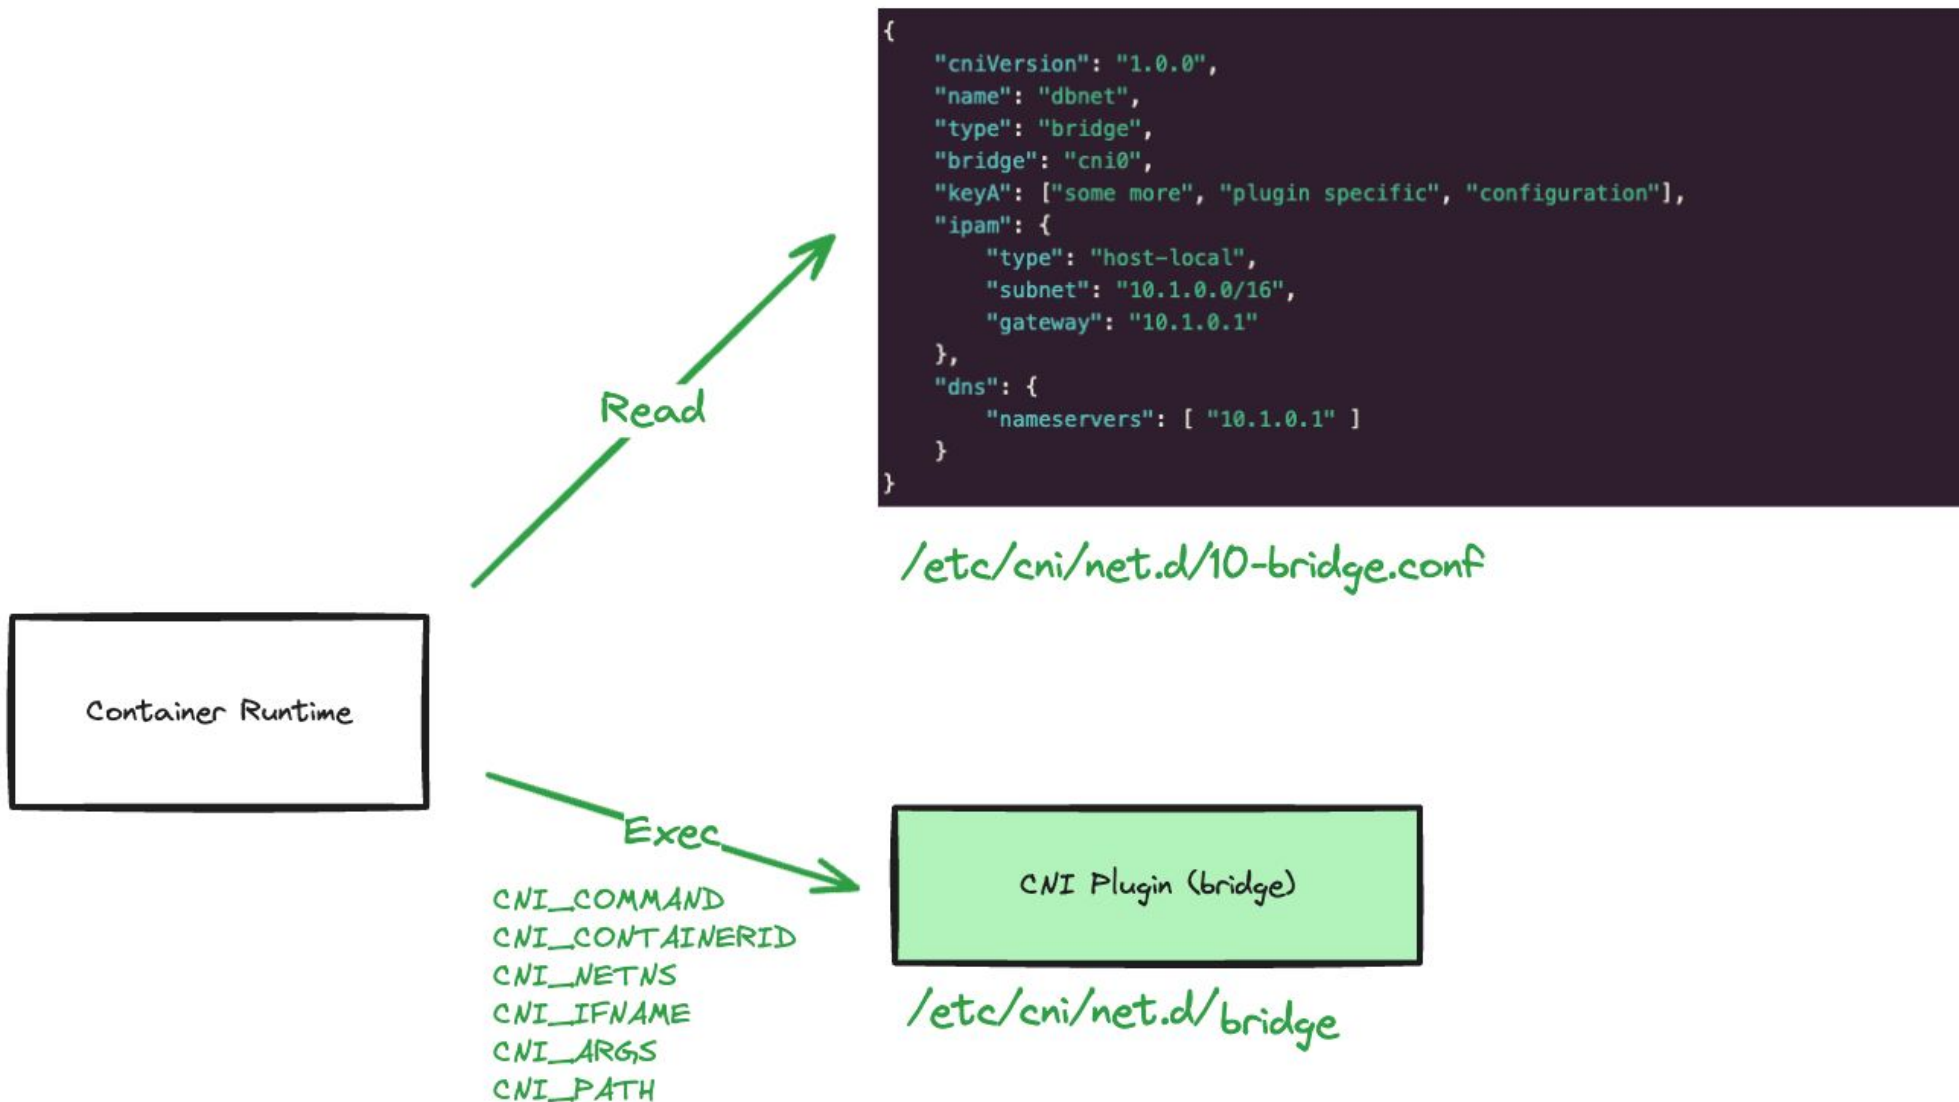


KubeCon



CloudNativeCon

North America 2023



CNI - Parameters



KubeCon



CloudNativeCon

North America 2023

- **CNI_COMMAND**: indicates the desired operation; **ADD**, **DEL**, **CHECK**, **GC**, Or **VERSION**.
- **CNI_CONTAINERID**: Container ID. A unique plaintext identifier for a container, allocated by the runtime.
- **CNI_NETNS**: A reference to the container's "isolation domain". If using network namespaces, then a path to the network namespace (e.g. `/run/netns/[nsname]`)
- **CNI_IFNAME**: Name of the interface to create inside the container
- **CNI_ARGS**: Extra arguments passed in by the user at invocation time. Alphanumeric key-value pairs separated by semicolons; for example, "FOO=BAR;ABC=123"
- **CNI_PATH**: List of paths to search for CNI plugin executables.

Source: <https://github.com/containernetworking/cni/blob/main/SPEC.md#parameters>

CNI - Multiple Plugins




KubeCon



CloudNativeCon

North America 2023






KubeCon



CloudNativeCon

North America 2023



Your First CNI Plugin

Exercise 1 - Bridge Plugin

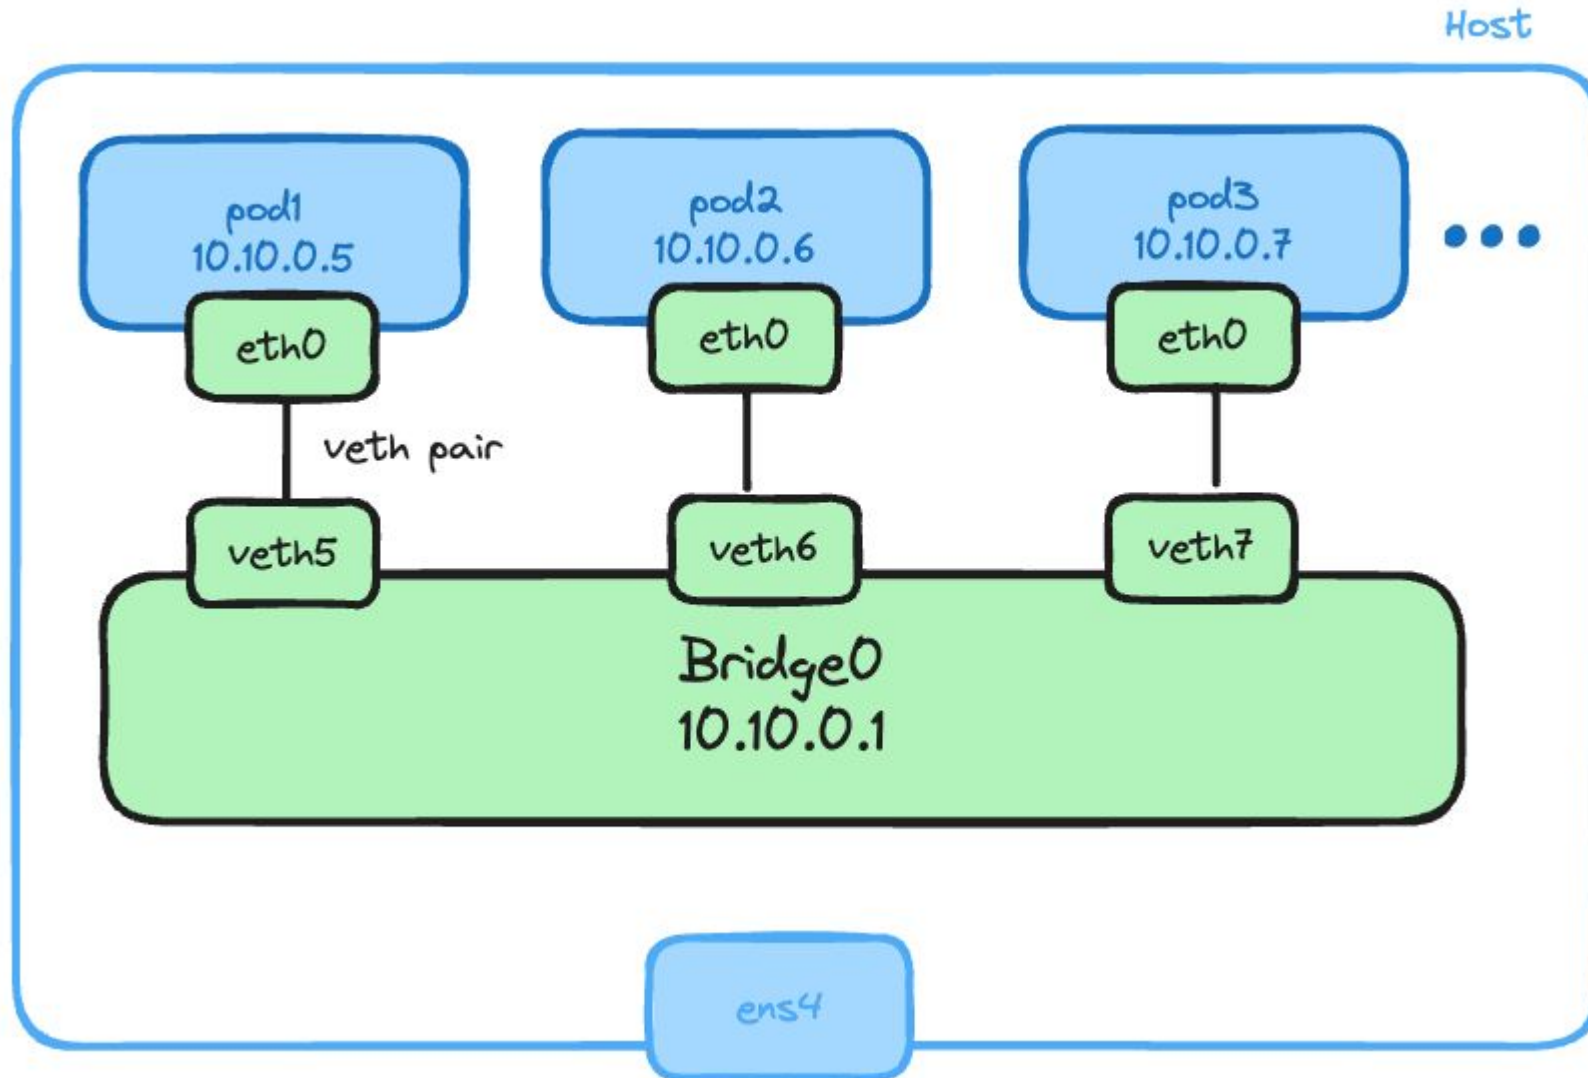


KubeCon



CloudNativeCon

North America 2023





bit.ly/ebpfcni

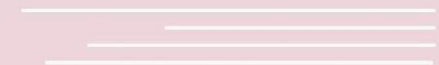


KubeCon



CloudNativeCon

———— North America 2023 ————



eBPF

What is eBPF



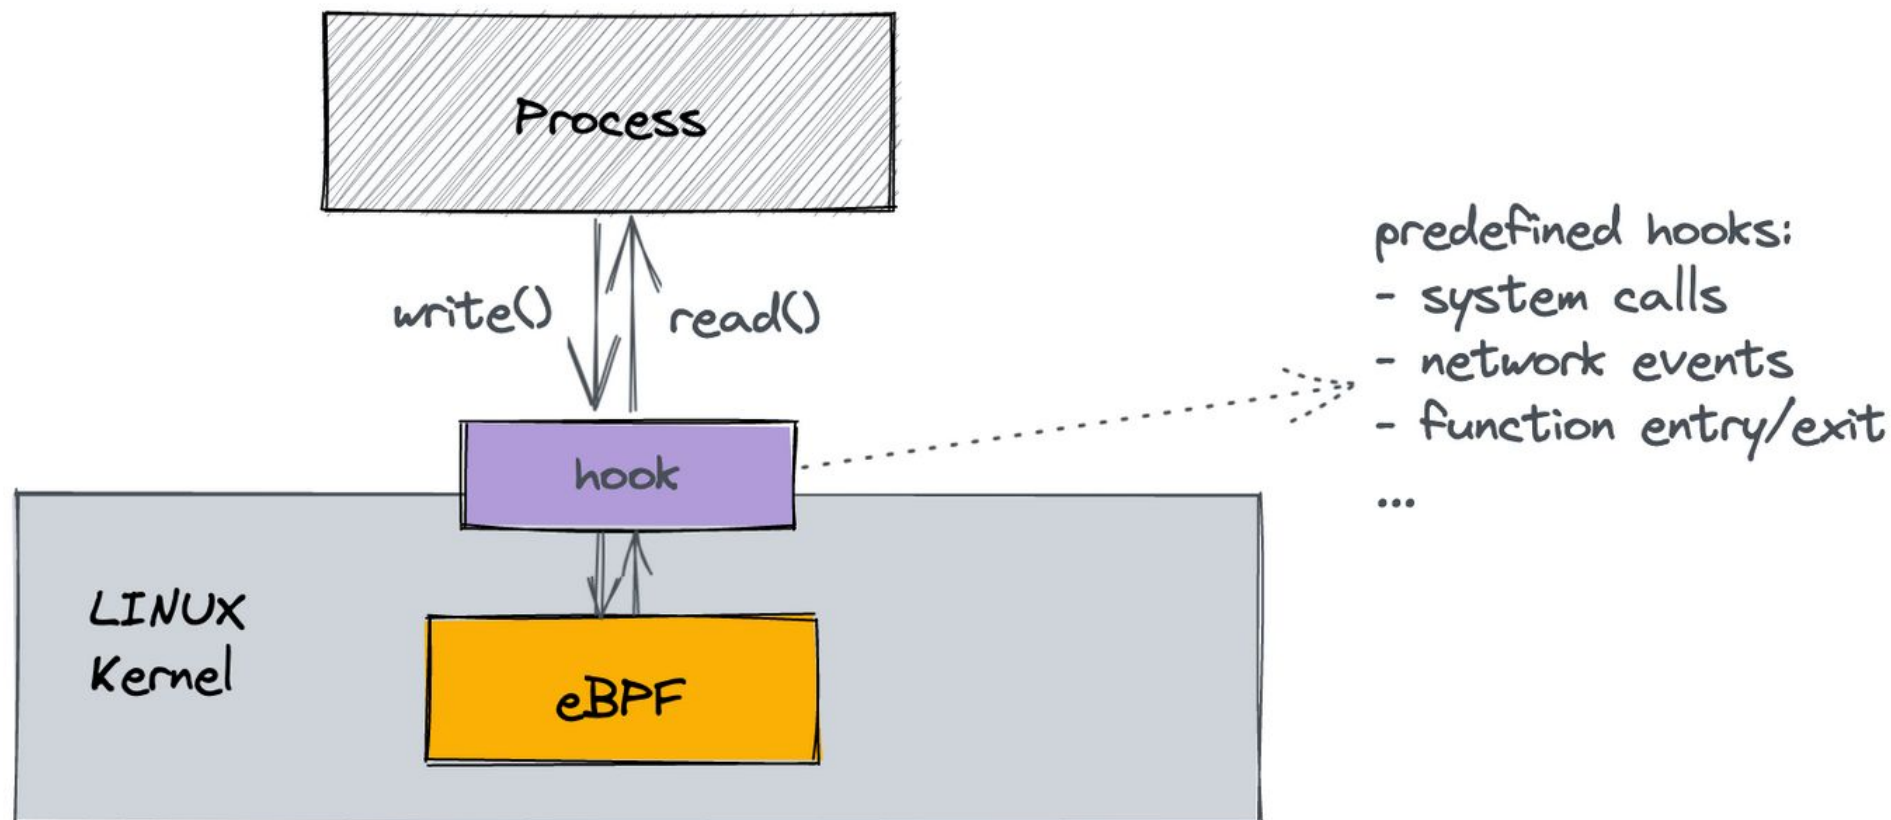
KubeCon



CloudNativeCon

North America 2023

- Linux technology which enables users to run custom programs “sandboxed” in the kernel
- **extended Berkeley Packet Filter**, evolution of “classic BPF” – think tcpdump
 - ‘eBPF’ & ‘BPF’ will be used interchangeably
- Event-based – programs are attached to “hook points” that are triggered by certain events
 - E.g. ‘kprobe’ type programs are attached to kernel functions and are then executed when that function is called
- BPF programs are verified to be “safe” – won’t crash the kernel, guaranteed to return (no infinite loops), can only access specific sections of memory, etc.



xBPF Hooks

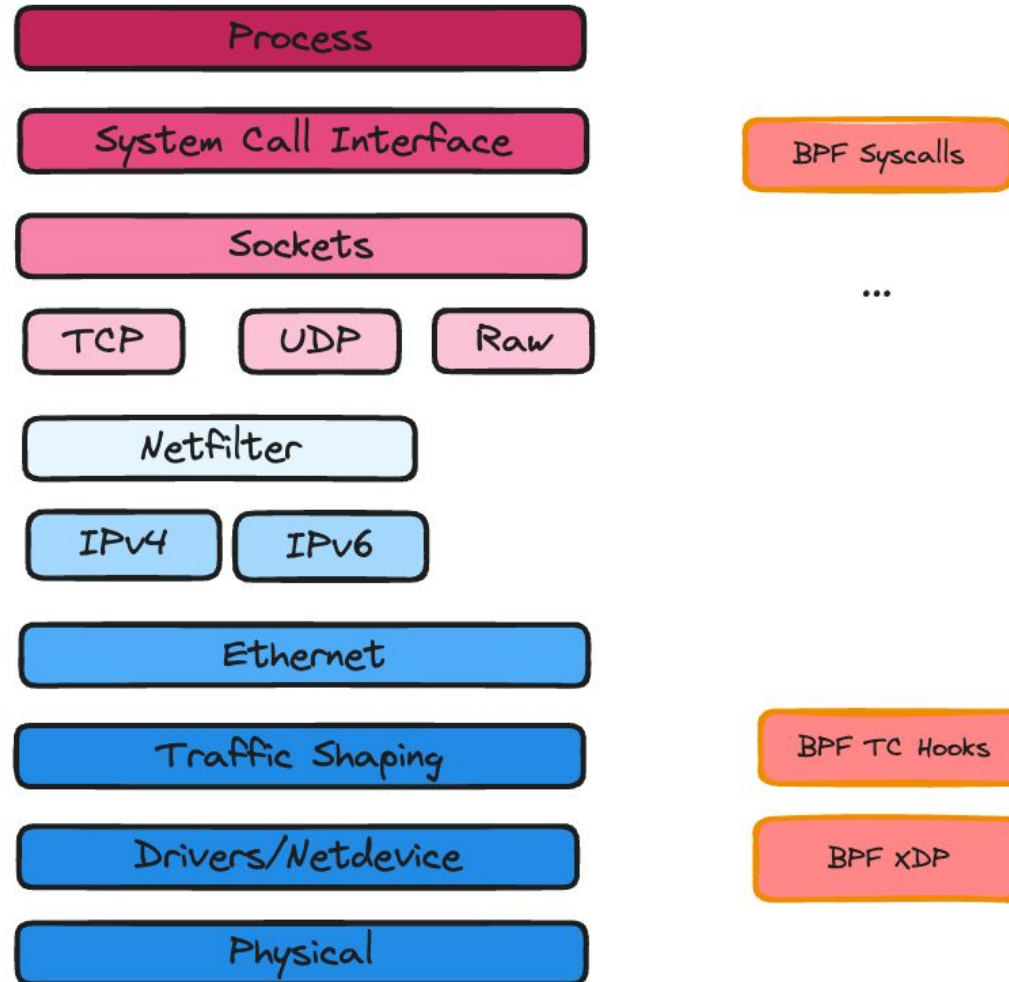


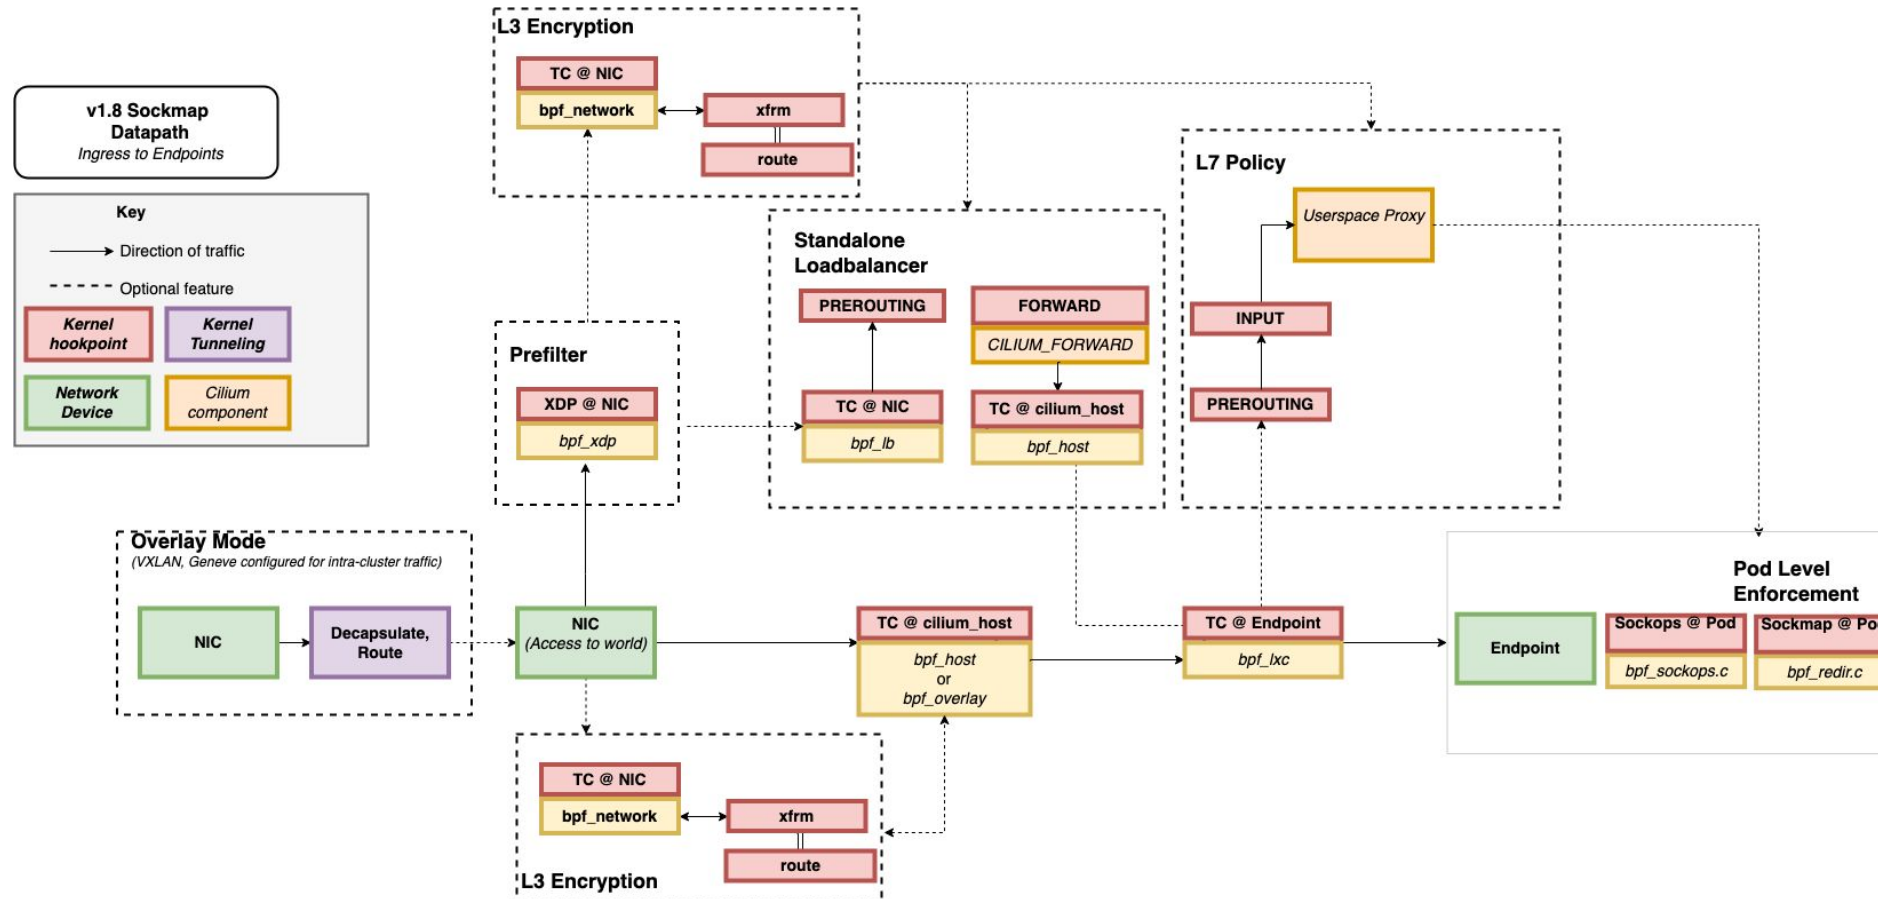
KubeCon



CloudNativeCon

North America 2023





source: <https://docs.cilium.io/en/stable/concepts/ebpf/lifeofapacket/#ingress-to-endpoint>

eBPF in Cilium




KubeCon



CloudNativeCon

North America 2023

Available CPU capacity under forwarding (higher is better)



source: <https://cilium.io/blog/2020/06/22/cilium-18/>

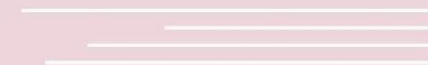


KubeCon



CloudNativeCon

———— North America 2023 ————



Your First eBPF program

Exercise 2 - eBPF with our CNI plugin

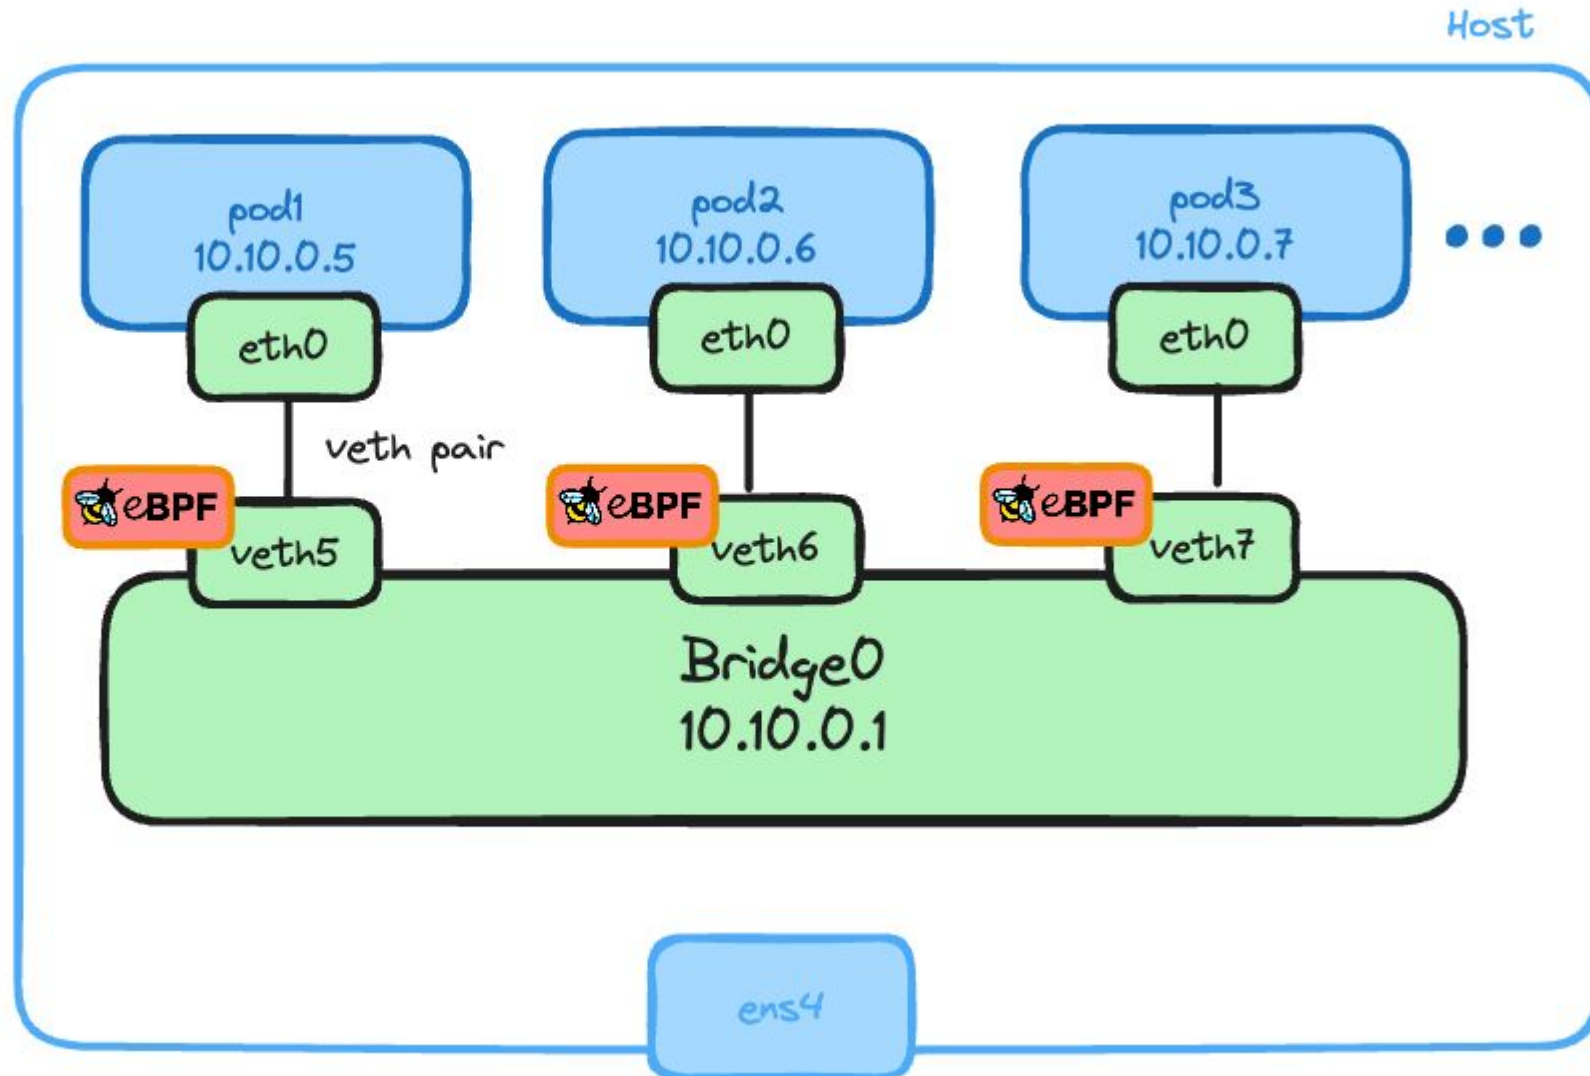


KubeCon



CloudNativeCon

North America 2023



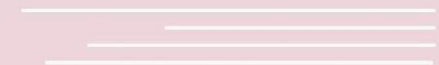


KubeCon



CloudNativeCon

———— North America 2023 ————



eBPF for Monitoring

Why eBPF for Observability




KubeCon



CloudNativeCon

North America 2023

- BPF programs can be attached to almost any kernel function, making it possible to extract data or metrics for observing the state of your systems
- Specific hookpoints (e.g. 'kprobes' and 'tracepoints') and infrastructure exist to enable BPF-based observability, performance monitoring, and tracing
- Due to eBPF's efficiency, you can process raw events as they happen in the kernel, enabling visibility that isn't possible with non-eBPF based solutions



eBPF Maps

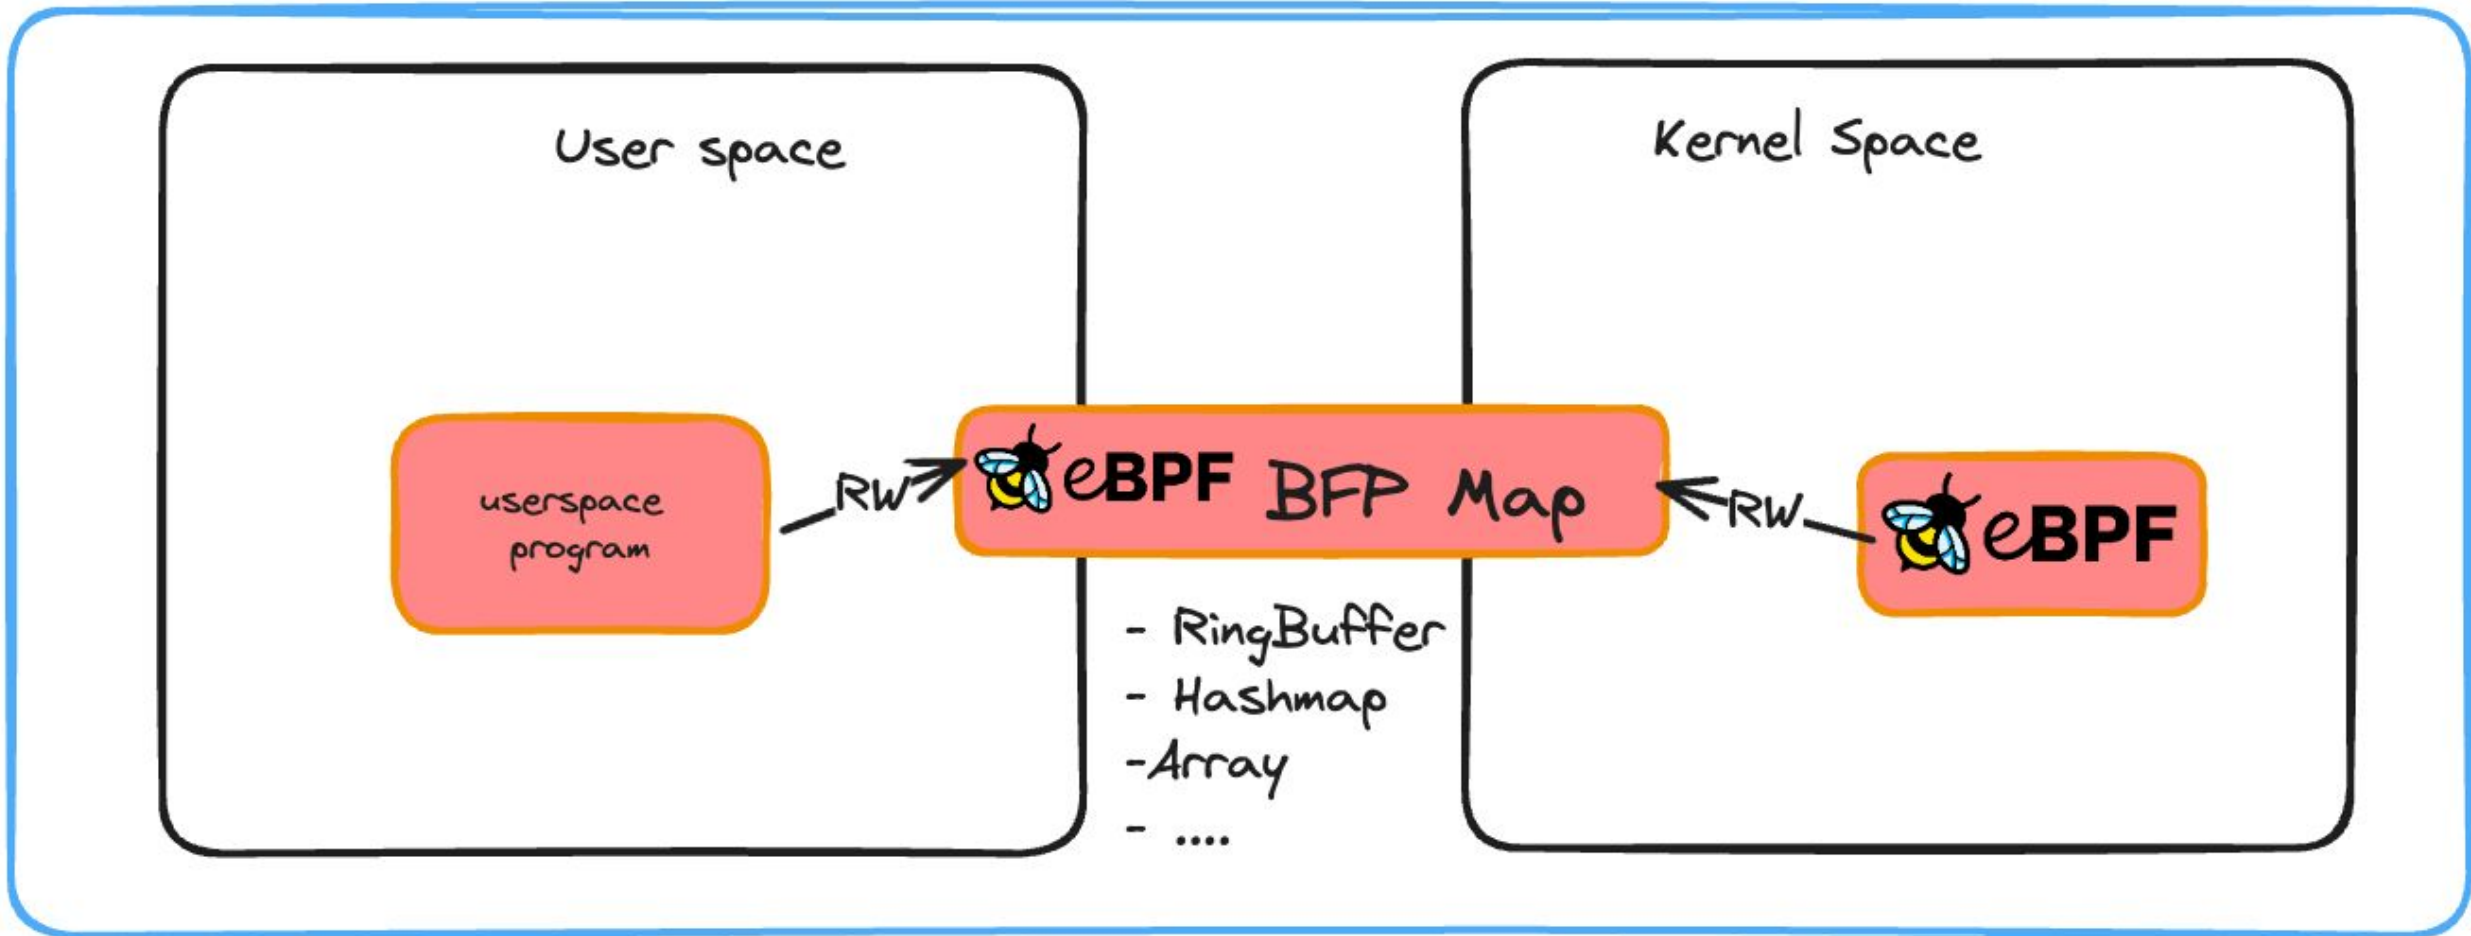


KubeCon



CloudNativeCon

North America 2023



Exercise 3 - Monitoring



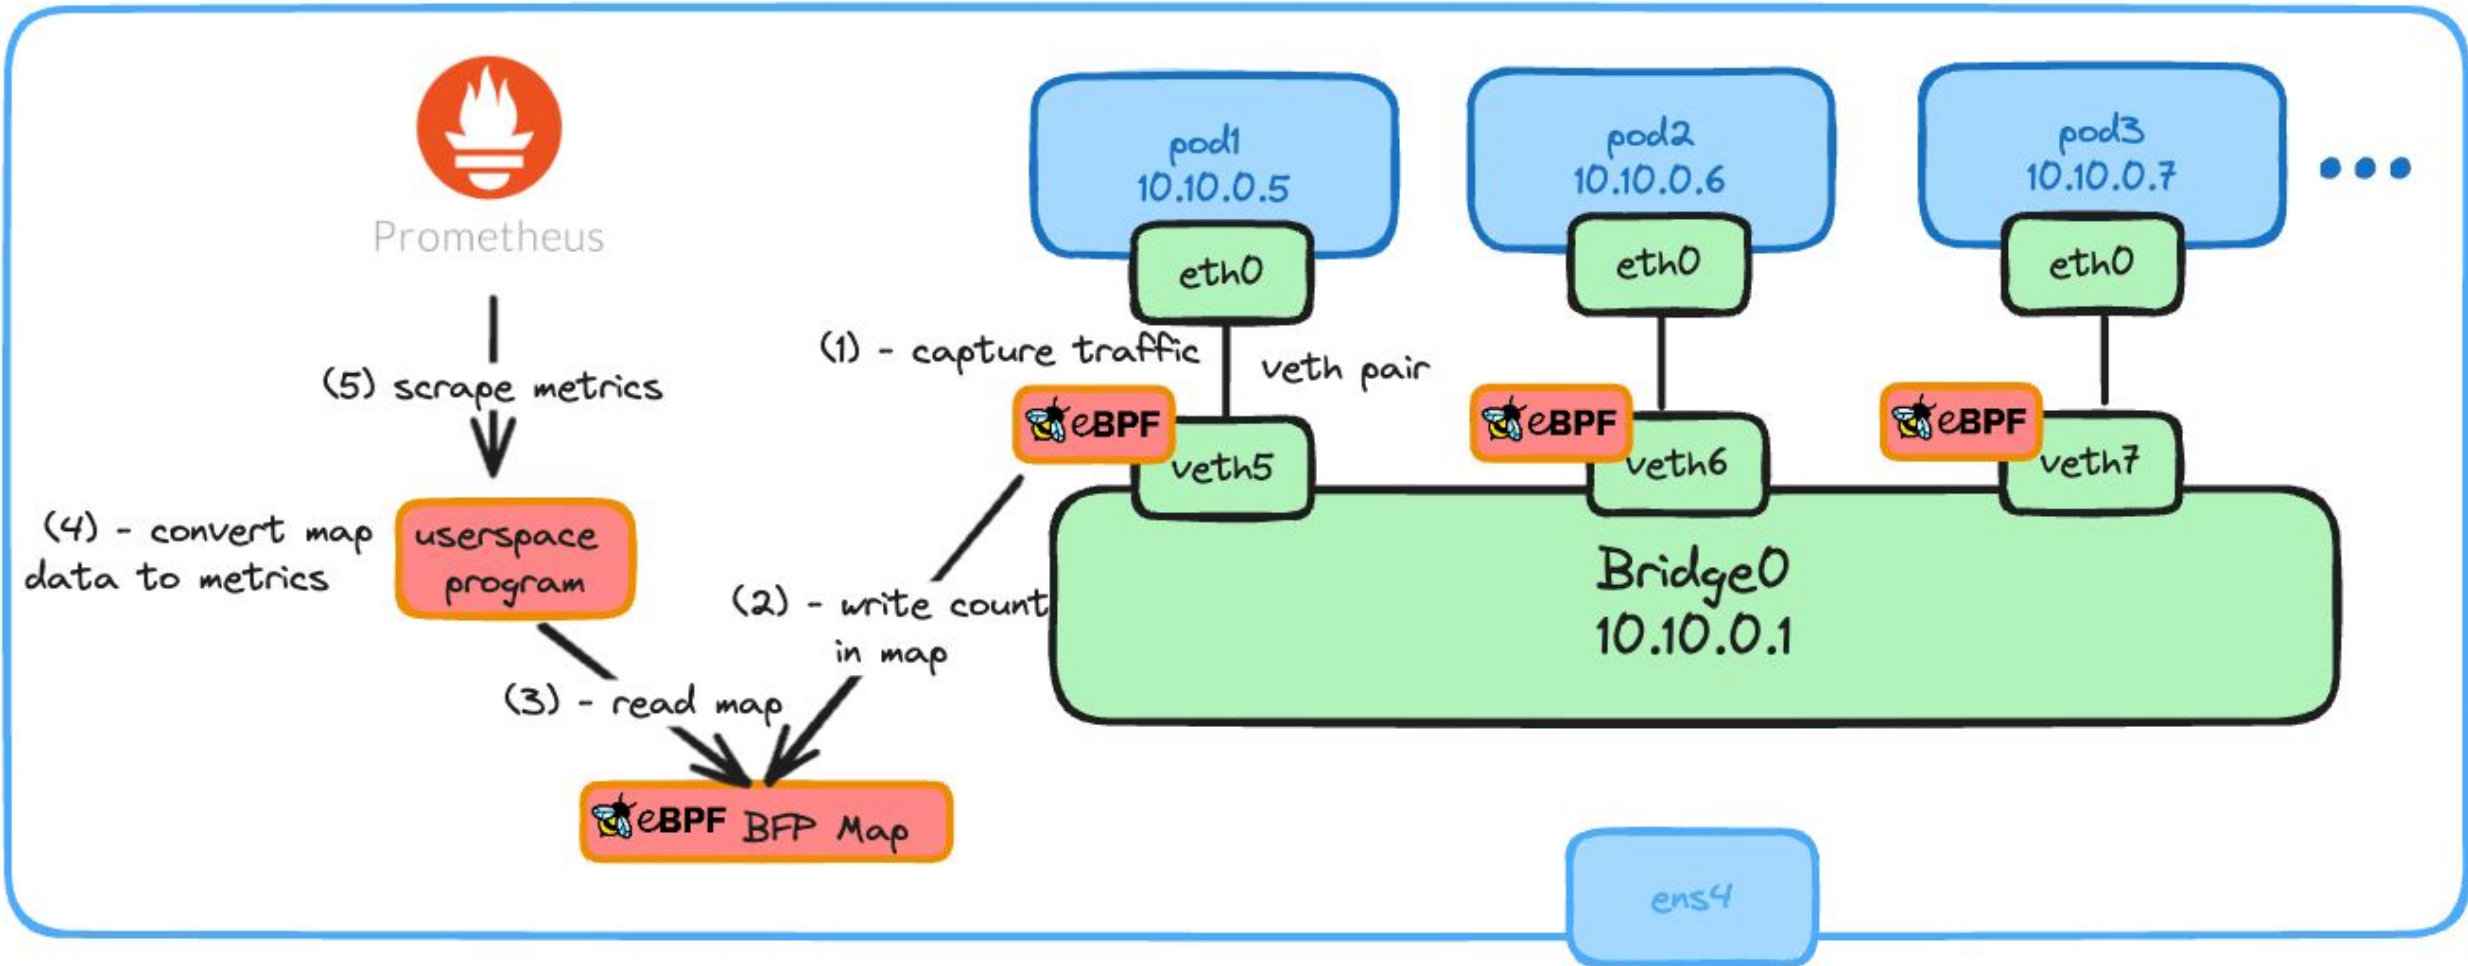
KubeCon



CloudNativeCon

North America 2023

Host



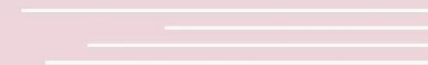


KubeCon



CloudNativeCon

North America 2023



eBPF for Security

Why eBPF for Pod Networking



KubeCon




CloudNativeCon

North America 2023


- **Networking**

- As classic BPF was designed for efficiently filtering packets (e.g. tcpdump), eBPF naturally lends itself well to numerous applications in the networking space
- eBPF enables programmatic processing of packets at several locations in the Linux networking stack



- **Security**

- Since eBPF programs can be attached to low-level system events, it is a great fit for codifying security policies or monitoring sensitive operations
- Additionally, as previously mentioned, since eBPF programs can be attached to various network operations, it is a logical fit for enforcing network security





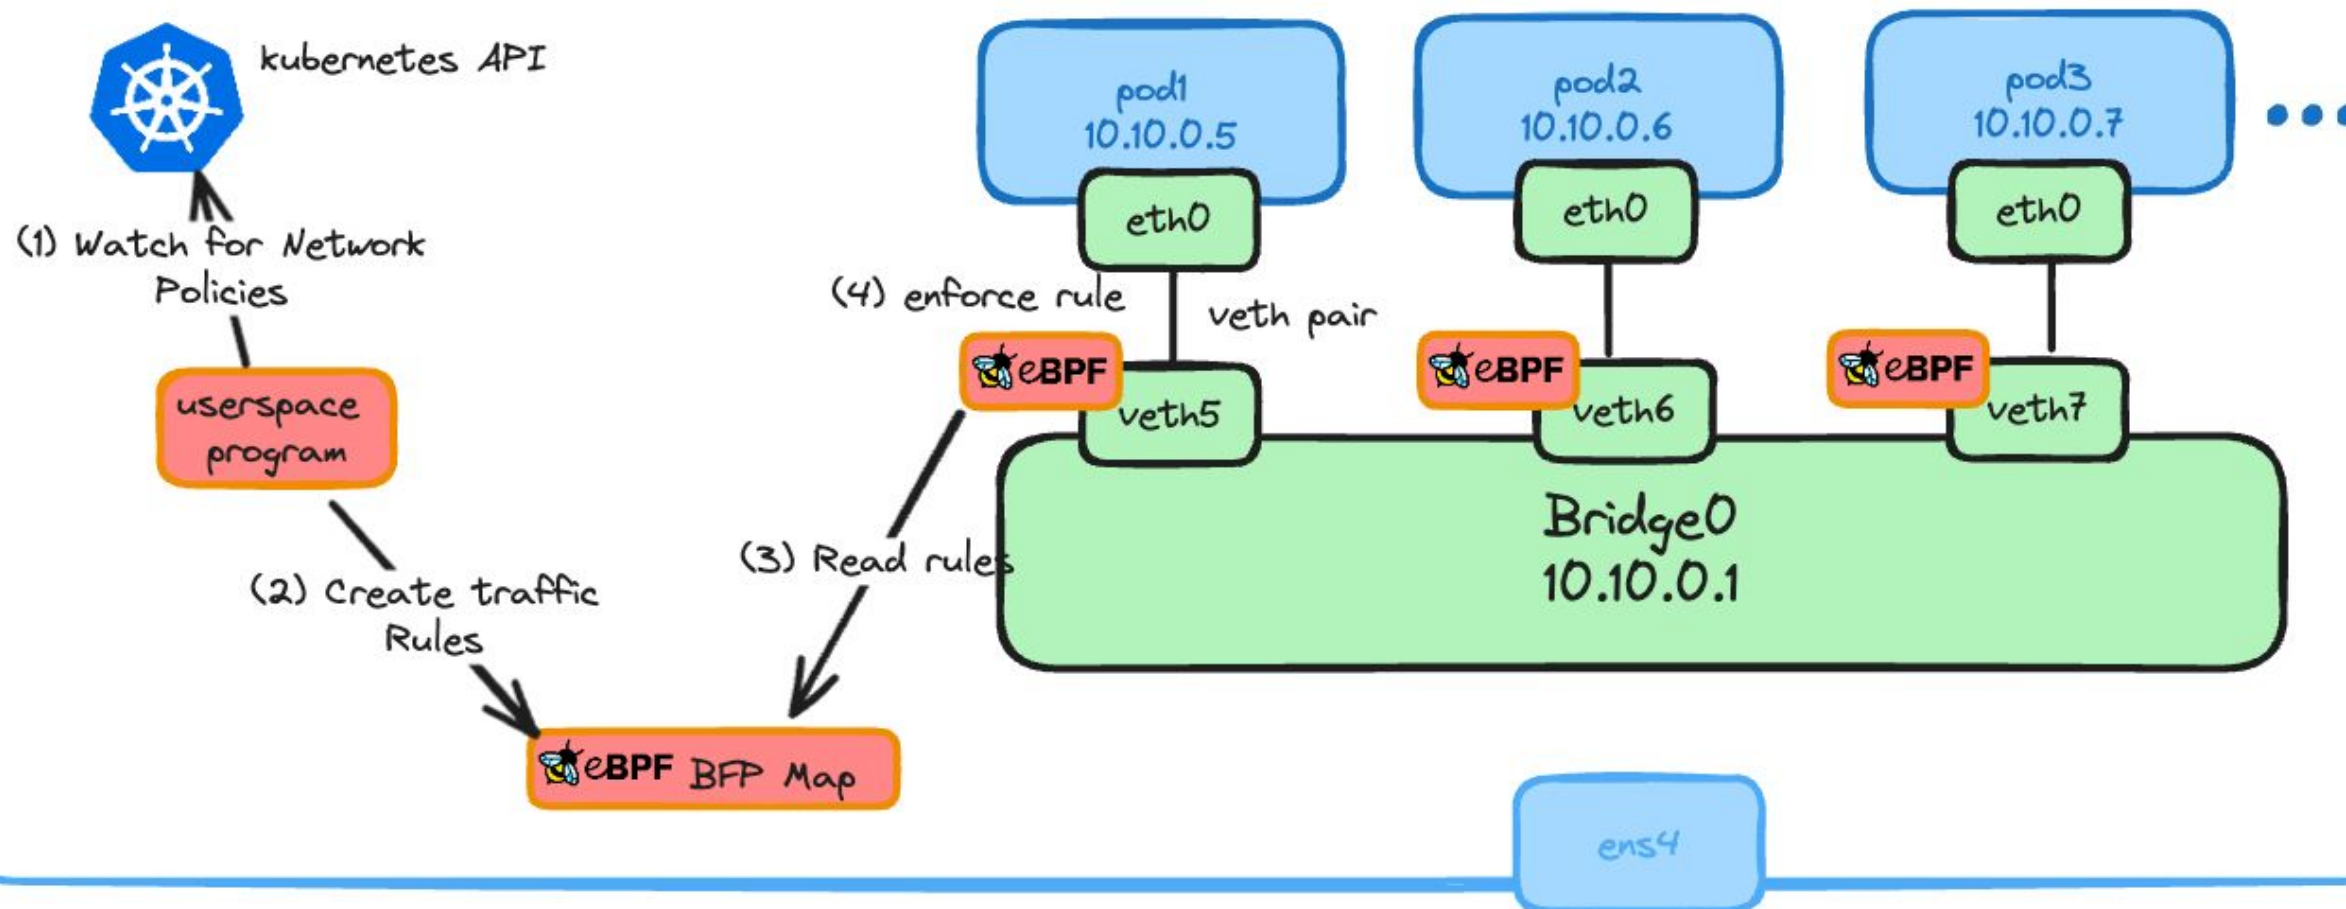
KubeCon



CloudNativeCon

North America 2023

Host



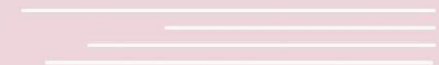


KubeCon



CloudNativeCon

———— North America 2023 ————



Conclusion



- <https://github.com/asayah/ebpf-cni-from-scratch-kubecon-na-2023>
- <https://github.com/libbpf/libbpf-bootstrap>
- https://github.com/cloudflare/ebpf_exporter
- <https://github.com/cilium/ebpf>
- <https://github.com/container networking/plugins/>



PromCon
North America 2021



**Please scan the QR Code above
to leave feedback on this session**